CertyIQ

CertyIQ

Premium exam material

Get certification quickly with the CertyIQ Premium exam material.
Everything you need to prepare, learn & pass your certification exam easily. Lifetime free updates
First attempt guaranteed success.

https://www.CertyIQ.com

About CertyIQ

We here at CertyIQ eventually got enough of the industry's greedy exam paid for. Our team of IT professionals comes with years of experience in the IT industry Prior to training CertIQ we worked in test areas where we observed the horrors of the paywall exam preparation system.

The misuse of the preparation system has left our team disillusioned. And for that reason, we decided it was time to make a difference. We had to make In this way, CertyIQ was created to provide quality materials without stealing from everyday people who are trying to make a living.

Doubt Support

We have developed a very scalable solution using which we are able to solve 400+ doubts every single day with an average rating of 4.8 out of 5.

https://www.certyiq.com

Mail us on - certyiqofficial@gmail.com

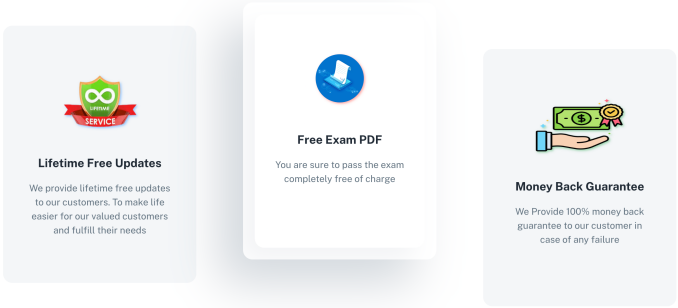

John

October 19, 2022

★★★★★

Thanks you so much for your help. I scored 972 in my exam today. More than 90% were from your PDFs!

Dana

September 04, 2022

★★★★★

Thanks a lot for this updated AZ-900 Q&A. I just passed my exam and got 974, I followed both of your Az-900 videos and the 6 PDF, the PDFs are very much valid, all answers are correct. Could you please create a similar video/PDF for DP900, your content/PDF's is really awesome. The team did a really good job. Thank You 😊.

Ahamed Shibly

2 months ago

★★★★★

Customer support is realy fast and helpful, I just finished my exam and this video along with the 6 PDF helped me pass! Definitely recommend getting the PDFs. Thank you!

October 22, 2022

★★★★★

Passed my exam today with 891 marks. Out of 52 questions, 51 were from certyiq PDFs including Contoso case study.
Thank You certyiq team!

Henry Rome

2 months ago

★★★★★

These questions are real and 100 % valid. Thank you so much for your efforts, also your 4 PDFs are awesome, I passed the DP900 exam on 1 Sept. With 968 marks. Thanks a lot, buddy!

Esmaria

2 months ago

★★★★★

Simple easy to understand explanations. To anyone out there wanting to write AZ900, I highly recommend 6 PDF's.Thank you so much, appreciate all your hard work in having such great content. Passed my exam Today - 3 September with 942 score.

Microsoft

(AB-100)

Agentic AI Business Solutions Architect

Total: **124 Questions**

Link: https://certyiq.com/papers/microsoft/ab-100

Question: 1 CertyIQ

HOTSPOT -
Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Background -
Fabrikam, Inc., is a global consumer goods company that is undergoing a digital transformation initiative to migrate its entire infrastructure to the Microsoft cloud. As a key element of this cloud migration, the company will implement Microsoft Dynamics 365 Sales, moving away from the current on-premises proprietary technologies used by its business-to-business (B2B) sales team.
As part of the cloud migration, Fabrikam will adopt an AI-first approach to its business solutions and implement AI solutions, wherever possible, to streamline operations.
Problem Statements -
Fabrikam's infrastructure currently relies on various on-premises systems that require sales executives to use corporate computers with physical keyboards to access business information during customer interactions. Mobile phones cannot be used for these purposes, as the systems depend on keyboard input. As a result, the sales executives spend a lot of time using keyboards to search for data on several disparate systems and file servers, rather than focusing on the customers. This affects the customer experience.
Fabrikam stakeholders are concerned that users will be hesitant to adopt AI. If the AI initiatives are NOT adopted, cost savings will never be realized. Additionally, funding for future AI initiatives will depend on demonstrating an increase in AI adoption month over month. As the AI agent initiative for the sales team will be the first for Fabrikam, the rapid adoption of the agent is a high priority.
Planned Initiatives -
General -
Fabrikam management has prioritized AI-driven projects to improve efficiency, customer engagement, and responsible AI adoption. The current application infrastructure is on-premises and must be migrated to the cloud to support the adoption of these technologies.
Infrastructure Migration -
Fabrikam plans to migrate from its current on-premises infrastructure to a completely cloud-based topology; this will include user authentication, the security framework, and, primarily, the adoption of the services by end users. All the data from the different systems will be consolidated into a single data source - a common data model that will use a Microsoft Dataverse environment as a single source of truth (SSOT) for the sales team.
Sales Cycle Enablement -
To achieve the company's objectives, Fabrikam intends to implement the following strategies to enhance the sales cycle:
Use low-code development to create a single AI agent that has Dataverse as its core component.
Ensure that sales managers can access unanswered correspondence from prospects and intervene as appropriate.
Replace the previous proprietary software with Dynamics 365 Sales to track sales cycles and customer interactions.
Have the sales executives use Dynamics 365 Sales to track interactions for open opportunities and send follow-up communications to prospects.
Have the sales executives use handsfree headsets to interact with an AI agent when they have questions about internal policies or customer data.
Requirements -
Infrastructure Migration -
Fabrikam has identified the following infrastructure migration requirements:
Azure must be used for all future infrastructure workloads.
The company must follow Microsoft-recommended methodologies for infrastructure migration to the cloud.
Any created AI agents must have their return on investment (ROI) calculated to ensure that the solution will save the company money.

Sales Cycle Enablement -
Fabrikam has identified the following requirements for sales cycle enablement:
The final AI agent must follow Microsoft recommendations for a conversational user experience.
A designated checklist must be reviewed to ensure that the AI agent follows Microsoft deployment recommendations for a compliant solution.
Detailed telemetry must be logged for the first created AI agent to help troubleshoot and optimize the agent during the initial AI agent adoption process.
Unexpected AI agent actions must end in an escalation to a live representative. For example, a sales executive must be rerouted to a representative if the agent cannot answer a question after two failed attempts.
The return on investment (ROI) of switching from the current process to the future process is required for stakeholder sign off.
The sales team must use Dynamics 365 Sales to correspond with prospects more quickly and efficiently than currently.
Sales managers must report on the adoption of the AI agent to key Fabrikam stakeholders on a monthly basis.
Any sensitive information, such as user IDs and names, shared via the AI agent must be tracked for future auditing.
Which framework should you use to meet the AI agent requirements for the sales cycle enablement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

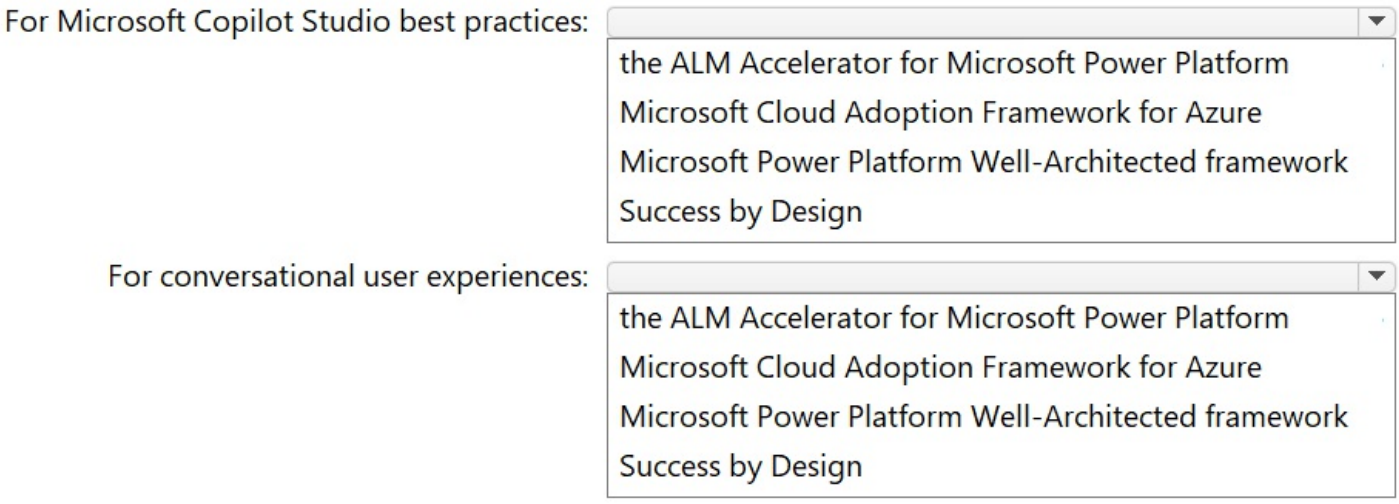

Answer:

Answer Area

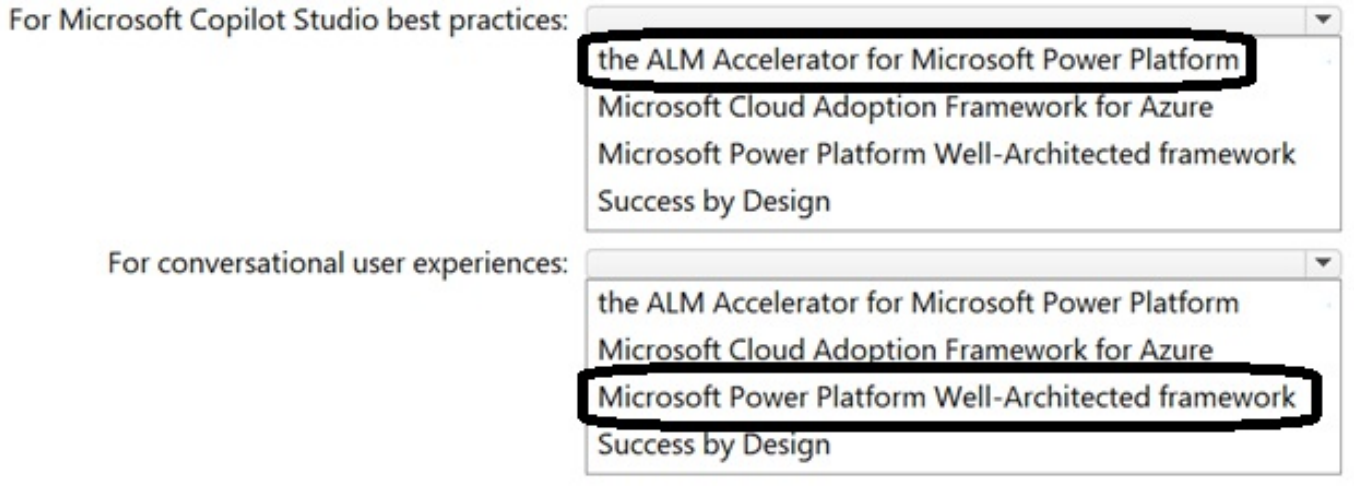

Explanation:

the ALM Accelerator for Microsoft Power Platform.

ALM = Application Lifecycle Management

It manages the entire lifecycle of apps and copilots, including:

Development

Version control

Testing

Deployment

Updates

For conversational user experiences

Microsoft Power Platform Well-Architected framework.

What this framework does

The Power Platform Well-Architected Framework provides guidance on designing solutions with:

Performance

Security

Reliability

Operational excellence

User experience

Why the other answer are incorrect:

Microsoft Cloud Adoption Framework for Azure

This framework is for:

Azure migration

Governance

Cloud strategy

Enterprise cloud adoption

It is not specifically about Copilot Studio development or ALM.

Microsoft Power Platform Well-Architected framework

This framework focuses on:

Reliability

Security

Performance

Operational excellence

It is not primarily an ALM or deployment framework.

Success by Design.

This methodology is intended for:

Dynamics 365 implementations

Enterprise project guidance

Solution architecture reviews

It is not specifically a Copilot Studio best-practice framework.

ALM Accelerator for Microsoft Power Platform.

This is for:

Deployment pipelines

Source control

Environment management

It does not provide guidance on conversational UX design.

Microsoft Cloud Adoption Framework for Azure.

This is an Azure governance and migration framework, not a conversational design framework.

Success by Design.

This provides implementation guidance for Dynamics 365 projects but does not focus on chatbot or copilot conversational experiences.

Question: 2 CertyIQ

Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Background -
Fabrikam, Inc., is a global consumer goods company that is undergoing a digital transformation initiative to migrate its entire infrastructure to the Microsoft cloud. As a key element of this cloud migration, the company will implement Microsoft Dynamics 365 Sales, moving away from the current on-premises proprietary technologies used by its business-to-business (B2B) sales team.
As part of the cloud migration, Fabrikam will adopt an AI-first approach to its business solutions and implement AI solutions, wherever possible, to streamline operations.
Problem Statements -
Fabrikam's infrastructure currently relies on various on-premises systems that require sales executives to use corporate computers with physical keyboards to access business information during customer interactions. Mobile phones cannot be used for these purposes, as the systems depend on keyboard input. As a result, the sales executives spend a lot of time using keyboards to search for data on several disparate systems and file servers, rather than focusing on the customers. This affects the customer experience.

Fabrikam stakeholders are concerned that users will be hesitant to adopt AI. If the AI initiatives are NOT adopted, cost savings will never be realized. Additionally, funding for future AI initiatives will depend on demonstrating an increase in AI adoption month over month. As the AI agent initiative for the sales team will be the first for Fabrikam, the rapid adoption of the agent is a high priority.
Planned Initiatives -
General -
Fabrikam management has prioritized AI-driven projects to improve efficiency, customer engagement, and responsible AI adoption. The current application infrastructure is on-premises and must be migrated to the cloud to support the adoption of these technologies.
Infrastructure Migration -
Fabrikam plans to migrate from its current on-premises infrastructure to a completely cloud-based topology; this will include user authentication, the security framework, and, primarily, the adoption of the services by end users. All the data from the different systems will be consolidated into a single data source - a common data model that will use a Microsoft Dataverse environment as a single source of truth (SSOT) for the sales team.
Sales Cycle Enablement -
To achieve the company's objectives, Fabrikam intends to implement the following strategies to enhance the sales cycle:
Use low-code development to create a single AI agent that has Dataverse as its core component.
Ensure that sales managers can access unanswered correspondence from prospects and intervene as appropriate.
Replace the previous proprietary software with Dynamics 365 Sales to track sales cycles and customer interactions.
Have the sales executives use Dynamics 365 Sales to track interactions for open opportunities and send follow-up communications to prospects.
Have the sales executives use handsfree headsets to interact with an AI agent when they have questions about internal policies or customer data.
Requirements -
Infrastructure Migration -
Fabrikam has identified the following infrastructure migration requirements:
Azure must be used for all future infrastructure workloads.
The company must follow Microsoft-recommended methodologies for infrastructure migration to the cloud.
Any created AI agents must have their return on investment (ROI) calculated to ensure that the solution will save the company money.
Sales Cycle Enablement -
Fabrikam has identified the following requirements for sales cycle enablement:
The final AI agent must follow Microsoft recommendations for a conversational user experience.
A designated checklist must be reviewed to ensure that the AI agent follows Microsoft deployment recommendations for a compliant solution.
Detailed telemetry must be logged for the first created AI agent to help troubleshoot and optimize the agent during the initial AI agent adoption process.
Unexpected AI agent actions must end in an escalation to a live representative. For example, a sales executive must be rerouted to a representative if the agent cannot answer a question after two failed attempts.
The return on investment (ROI) of switching from the current process to the future process is required for stakeholder sign off.
The sales team must use Dynamics 365 Sales to correspond with prospects more quickly and efficiently than currently.
Sales managers must report on the adoption of the AI agent to key Fabrikam stakeholders on a monthly basis.
Any sensitive information, such as user IDs and names, shared via the AI agent must be tracked for future auditing.
Which framework should you use for the infrastructure migration?

A. Microsoft Cloud Adoption Framework for Azure
B. Success by Design
C. Microsoft Power Platform Center of Excellence (CoE)
D. Microsoft Power Platform Project Setup Wizard

Answer: A

Explanation:

Why optionA –Microsoft Cloud Adoption Framework for Azure – is the correct choice

The Cloud Adoption Framework (CAF) provides a **Microsoft-recommended end-to-end methodology** for moving on-premises workloads to Azure, covering assessment, solution design, migration, validation, and

operations.
It aligns with Fabrikam's requirement to "follow Microsoft-recommended methodologies for infrastructure migration to the cloud" and includes guidance on **ROI calculation, security, and telemetry**, which are explicitly part of the migration objectives.
CAF's **migration planning workstreams** map directly to the stated needs: consolidating data into a single Dataverse source, establishing a unified authentication and security model, and ensuring audit-ready telemetry for AI agents.
The framework also integrates with other Microsoft services (Azure Migrate, Azure Arc, Power Platform, Dynamics365) needed for the downstream AI-agent and sales-automation components.

Why the other options are less suitable

Success by Design (B) – This is a go-to-market and consumption model for Power Platform that focuses on licensing and value realization, not a migration methodology. It does not address the low-level infrastructure lift-and-shift or the detailed migration checklist required by Fabrikam.
Microsoft Power Platform Center of Excellence (CoE) (C) – A CoE provides governance, best-practice guidance, and community support for building and scaling low-code solutions, but it is not a migration framework. It cannot replace the structured, Azure-centric migration process mandated for moving all workloads to the cloud.
Power Platform Project Setup Wizard (D) – This wizard creates a new Power Platform project scaffolding (dataverse, Power Apps, Power Automate templates). It is a **development-level tool** for prototyping solutions, not a comprehensive migration framework that covers network, identity, security, and data-migration phases required for the overall infrastructure transformation.

Conclusion

Only the **Microsoft Cloud Adoption Framework for Azure** supplies the full, Microsoft-endorsed migration pathway that satisfies Fabrikam's migration, security, ROI, and telemetry requirements, making it the optimal choice for the infrastructure migration question.

References

Microsoft Cloud Adoption Framework (CAF) – Overview and migration guidance:
https://learn.microsoft.com/en-us/azure/architecture/cloud-adoption-framework/
CAF Migration best practices – Detailed steps for assessing, planning, and executing cloud migrations:
https://learn.microsoft.com/en-us/azure/architecture/cloud-adoption-framework/migrate

Question: 3 CertyIQ

A company uses Microsoft Dynamics 365 Sales to manage leads that are stored in a Microsoft Dataverse table named Lead and use non-standard terminology and custom columns.
You need to configure business terms in the Lead table so that Microsoft Copilot controls can summarize the leads efficiently. The solution must minimize administrative effort.
How should you configure the business terms?

Under Microsoft Power Platform and Copilot Studio architectural standards for generative AI summaries and semantic grounding, Copilot relies on a business glossary / business terms mapping to understand Dataverse schema elements.

A. Combine all the fields into one custom field.
B. Map the field display names as business terms.
C. Add the schema names as business terms.
D. Create new business terms for each field.

When customizing Copilot summaries or generative responses over Dataverse tables (such as Lead), administrators can associate Dataverse columns with standardized business terms.

Answer: B

Explanation: Mapping existing display names to business terms allows Copilot's natural language understanding (NLU) model to interpret custom or localized fields using familiar, human-readable labels already defined on the entity, avoiding redundant manual entry and minimizing administrative effort.

Business terms in Dynamics365 must map to standard field names that Copilot can reference – by mapping the display names of the custom columns to business terms, the platform creates readable, searchable labels that Copilot's summarisation engine recognizes without additional configuration.
OptionB (Map the field display names as business terms) directly achieves this mapping, requiring only a one-time entry per term and no changes to the underlying schema. This minimizes admin effort while keeping the terminology aligned with the data model that Copilot consumes.
OptionA (Combine all fields into one custom field) loses granularity, prevents Copilot from distinguishing different data points, and is not supported for business-term mapping.
OptionC (Add the schema names as business terms) introduces technical, internal identifiers that are not user-friendly; Copilot expects human-readable terms, so summarisation would be harder to understand and maintain.
OptionD (Create new business terms for each field) duplicates effort and does not tie the terms to the actual field metadata, leading to unsynchronised terminology and higher maintenance overhead.

Conclusion: Mapping display names (OptionB) aligns terminology with Copilot's expectations, maximises summarisation accuracy, and requires the least administrative overhead.

References

Configure business terms in Dynamics365 Sales: https://learn.microsoft.com/en-us/dynamics365/sales-enterprise/business-terms
How Copilot uses business terms for summarisation: https://learn.microsoft.com/en-us/dynamics365/sales-enterprise/copilot-summarisation-terms-guide

Question: 4 CertyIQ

DRAG DROP -
You are designing two Microsoft Copilot Studio agents named Agent1 and Agent2. Each agent must meet the following requirements:
Each agent must use a standard model.
Each agent must NOT use generative orchestration.
Agent1 must support simple and short phrases for a given topic.
Agent2 must integrate with Microsoft Dynamics 365 Contact Center voice channel.
You need to recommend language models for the agents.
What should you recommend for each agent? To answer, drag the appropriate language models to the correct agents. Each language model may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

Language models

- Azure Language in Foundry Tools
- Azure OpenAI
- Conversational language understanding (CLU)
- Natural language understanding (NLU)
- Natural language understanding + (NLU+)

Answer Area

	Language model
Agent1:	Natural language understanding (NLU)
Agent2:	Natural language understanding + (NLU+)

Answer:

Under Microsoft Copilot Studio architecture and model selection standards (specifically non-generative, standard intent trigger routing):

Natural language understanding (NLU): The classic, standard intent-recognition model in Copilot Studio. It is optimized for shorter, straightforward trigger phrases (5–10 words) explicitly defined within topic triggers.

Natural language understanding + (NLU+): An enhanced deep-learning model built for more complex intent recognition, extended conversational variance, and specific multimodal channel capabilities. Crucially, NLU+ is required for telephony/voice interactions (such as integration with the Dynamics 365 Contact Center voice channel/Omnichannel voice) to properly process speech-to-text transcriptions with ambient noise, filler words, and varied grammatical structures without relying on generative orchestration.

Language models

- Azure Language in Foundry Tools
- Azure OpenAI
- Conversational language understanding (CLU)
- Natural language understanding (NLU)
- Natural language understanding + (NLU+)

Answer Area

	Language model
Agent1:	Natural language understanding (NLU)
Agent2:	Natural language understanding + (NLU+)

Explanation:

Agent1 – **Natural Language Understanding (NLU)**

What NLU is

Natural Language Understanding (NLU) is a capability that allows a system to:

Understand user intent

Extract entities

Process natural language inputs

Incorrect Choices

Natural language understanding (NLU): This refers to the original legacy NLU engine. While it supports simple phrases, it is treated as a deprecated or secondary option compared to the updated CLU model for standard intent mapping.

Natural language understanding + (NLU+): This model is over-engineered for Agent 1. It is designed for multi-turn entity extraction and advanced telephony, not just simple, short phrases.

Azure OpenAI & Azure Language in Foundry Tools: Both of these require advanced generative or custom orchestration workflows, violating the prompt's constraint that the agent must not use generative orchestration.

Agent2 – **Natural Language Understanding + (NLU+)**

NLU+ is an enhanced language model capability that extends NLU with additional intelligence such as:

More advanced reasoning

Context awareness

Generative AI responses

Improved conversation handling

It can support:

Complex multi-step conversations

Dynamic responses

Integration with knowledge sources.

Incorrect Choices

Natural language understanding (NLU) / Conversational language understanding (CLU): Neither of these

basic intent-classification models possesses the specialized acoustic processing, advanced slot filling, or conversational memory necessary to handle complex voice and telephony scenarios smoothly.

Azure OpenAI & Azure Language in Foundry Tools: These utilize generative AI capabilities that bypass the classic model requirements and directly violate the constraint against generative orchestration.

Question: 5 CertyIQ

A company uses Microsoft Dynamics 365 finance and operations apps.
The company plans to use Microsoft Copilot in-app help and guidance to generate responses for internal business processes.
You need to add an additional knowledge source for the business processes. The solution must NOT add new topics to the Copilot agent for the finance and operations apps.
Which knowledge source should you add?

A. Microsoft Dataverse
B. a public website
C. Azure AI Search
D. a file upload

According to Microsoft Learn documentation ("Add knowledge to generative help and guidance with Copilot" under Dynamics 365 finance and operations apps architecture), administrators can extend the out-of-the-box in-app help Copilot agent by uploading files directly (e.g., DOCX, PDF, RTF) to the Knowledge section in Microsoft Copilot Studio.

When uploading document files directly via the Upload files feature, Copilot automatically indexes and grounds responses across internal SOPs and business process manuals.

Crucially, Microsoft specifies that linking to certain external sources or enterprise repositories (such as custom SharePoint configurations or complex search feeds) requires authoring or customizing a topic in the Copilot for finance and operations apps agent, whereas file upload works directly within the existing knowledge configuration without creating any new topics.

Answer: D

Explanation:

File upload is the supported knowledge-source type for in-app Copilot guidance in Dynamics365 Finance & Operations.
It can be added directly from the Copilot studio without creating new “topics” in the agent’s navigation model.
The uploaded document (e.g., a PDF or Excel guide) is indexed locally and made available to the Copilot engine, enabling it to retrieve context-relevant passages during user interaction.
This approach meets the requirement of adding a knowledge source while not altering the existing topic structure.
Why the other options are unsuitable
Microsoft Dataverse – Used for structured data storage and custom entity design, not for delivering ad-hoc guidance content to the Copilot UI. Adding content here would require custom plugin development and would not be consumed as a “knowledge source” by the in-app guidance engine.
A public website – External URLs are not ingestible by the Copilot knowledge-source service; they would need a separate web-crawling or Search index, which adds unnecessary complexity and does not satisfy the “no new topics” constraint.
Azure AI Search – Although capable of semantic search, it introduces an external indexing pipeline and topic creation steps (e.g., defining index schema, mapping fields). This exceeds the minimal, out-of-the-box solution required for the scenario.

Conclusion – The only option that directly satisfies the requirement of adding a knowledge source without creating new topics is **D. a file upload**.

References

Add a knowledge source to Copilot in Dynamics365 –https://learn.microsoft.com/en-us/dynamics365/customer-engagement/copilot/knowledge-source
Upload and manage files as knowledge sources –https://learn.microsoft.com/en-us/dynamics365/customer-engagement/copilot/upload-files

Question: 6 CertyIQ

A company has an AI business solution.
You need to extend the solution so that Microsoft 365 Copilot can invoke external logic hosted in Azure services.
What should you include in the solution?

Under the Microsoft 365 Copilot extensibility architecture and the Architect agentic AI business solutions learning path, Microsoft 365 Copilot connects to external external systems, logic, and APIs using Power Platform connectors (specifically certified or custom connectors).

A. Microsoft Copilot Studio skills
B. Microsoft Power Platform connectors
C. custom engine agents

When external business logic is hosted in Azure (e.g., deployed in an Azure Function, API Management, or Azure Container App), the established architectural pattern to expose these capabilities to Microsoft 365 Copilot is to wrap the OpenAPI definition (Swagger) in a Power Platform connector (or Power Platform connector plugin/action).

Answer: B

Explanation:

This integration pattern provides built-in enterprise authentication (Microsoft Entra ID, OAuth2, API key), data loss prevention (DLP) policy enforcement, and parameter mapping, allowing Copilot's orchestrator to dynamically call external Azure logic at runtime.

Correct choice –Microsoft Power Platform connectors
Power Platform connectors provide a standardized, bi-directional integration layer that enables Azure-hosted logic (e.g., Azure Functions, Azure Logic Apps, Azure API Management) to be exposed as a connector. This allows Microsoft365Copilot to call external services directly, leveraging the connector's authentication, schema mapping, and invocation capabilities. The integration is managed through the Power Platform connector catalog, supports custom OpenAPI definitions, and can be securely shared across tenant boundaries, making it the recommended mechanism for exposing Azure services to Copilot.
Why optionA (Microsoft Copilot Studio skills) is less suitable
Copilot Studio skills are primarily used for building conversational agents within Copilot Studio and are not designed to expose arbitrary Azure APIs to Microsoft365Copilot. Skills operate within the confines of the studio environment and lack the necessary connector infrastructure for secure, governed, cross-service calls. Consequently, they cannot serve as a robust integration point for external Azure logic.
Why optionC (custom engine agents) is less suitable
Custom engine agents refer to bespoke AI engine implementations that run outside the Microsoft 365 ecosystem. They do not provide a native, supported bridge to Microsoft365Copilot and generally require custom middleware that must be manually provisioned, monitored, and maintained. This approach bypasses the standardized connector model, resulting in higher operational overhead, limited governance, and potential compliance risks.

References

Power Platform connectors documentation: https://learn.microsoft.com/en-us/power-platform/connectors/
Extend Copilot with custom connectors: https://learn.microsoft.com/en-us/copilot-extensibility/connectors/

Question: 7 CertyIQ

HOTSPOT -
You need to design a shared prompt library that will be used across multiple business units. The solution must meet the following requirements:
Ensure consistent AI responses with reusable formats.
Support governance and version control.
Minimize administrative effort.
Minimize ongoing costs.
What should you recommend for each requirement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Standardized Prompt Templates: To ensure consistency across models and users while enabling reusable output formats (such as structured JSON, standardized markdown tables, or uniform tone), organizations establish parameterized, standardized prompt templates. Templates define the persona, explicit constraints, input dynamic parameters, and expected output structure.

Version Control & Governance via Git: Prompts deployed across multiple business units are treated as code artifacts ("Prompts-as-Code"). Storing prompt templates in a centralized Git repository (such as Azure DevOps or GitHub) provides built-in branching, pull requests, peer reviews, commit history, rollback mechanisms, and audit logging with minimal ongoing administrative overhead and no additional tooling costs.

Answer Area

Ensure consistent AI responses:

- Delegate department-specific prompt templates.
- Define standardized prompt templates.
- Maintain a prompt history.

Support governance and version control:

- Define standardized prompt templates.
- Store prompts in a Git repository.
- Categorize prompts by business function.

Answer:

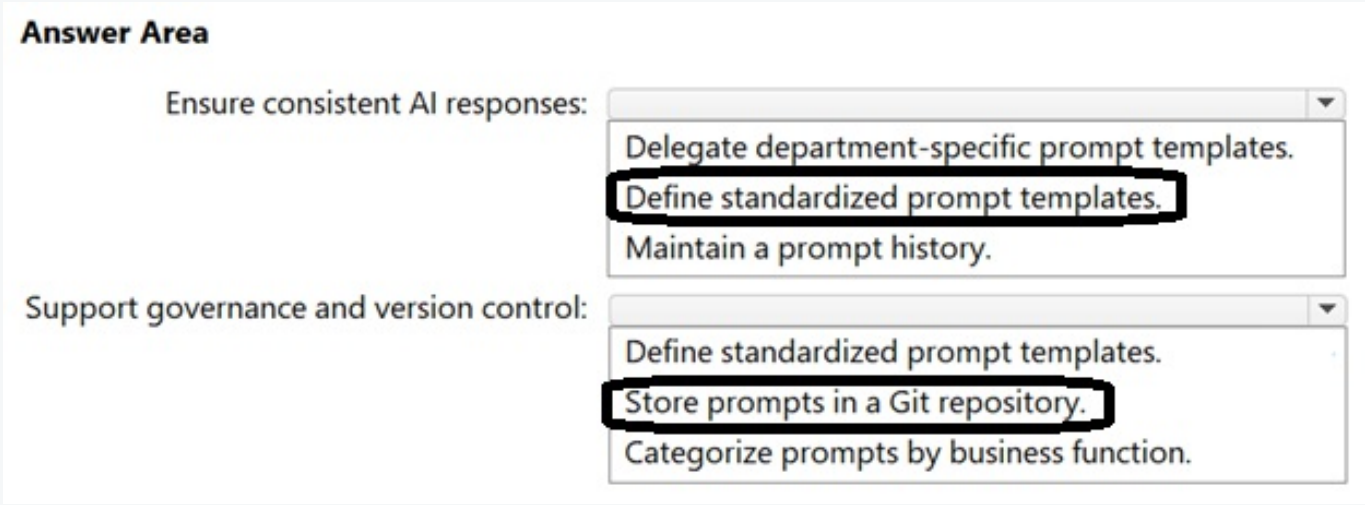

Explanation:

Define standardized prompt templates.

Prompt templates ensure that AI models receive structured and consistent instructions.

Benefits include:

Consistent responses across teams

Reduced prompt variability

Easier maintenance and updates

Better response quality

Support governance and version control.

Incorrect answers

Delegate department-specific prompt templates

Why it is incorrect: Allowing different departments to create their own isolated prompt versions leads to fragmentation. Without centralized control, the AI responses across the organization will drift apart and

become highly inconsistent rather than standardized.

Maintain a prompt history,

Why it is incorrect: A prompt history tracks what was sent to the AI model in the past (logging/auditing). While valuable for security and debugging, a historical log does not actively enforce or guide the system to generate uniform, predictable responses for new queries.

Store prompts in a Git repository.

A Git repository supports governance by enabling:

Version control

Change tracking

Collaboration

Rollback to previous versions

Auditability.

Incorrect answers:

Define standardized prompt templates.

Why it is incorrect: Standardizing the structure of a prompt ensures content consistency (the first requirement), but it does not provide a mechanism to track code-level changes, manage release cycles, or roll back to historical versions if a prompt fails.

Categorize prompts by business function.

Why it is incorrect: Organizing or labeling prompts by department (e.g., Marketing, HR) is a folder/metadata management practice. It does not offer technical governance capabilities like access auditing, change approval workflows, or explicit version tracking.

Question: 8 CertyIQ

DRAG DROP -
A company has a Microsoft Foundry project that uses a single agent and a single prompt to complete a series of tasks.
The agent encounters the following issues:
It frequently produces incomplete results.
It struggles with domain-specific reasoning.
Agent response times are remarkably slow.
You need to recommend a solution to improve the overall performance and accuracy of the agent.
What should you include in the recommendation? To answer, drag the appropriate actions to the correct requirements. Each action may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

Performance / Latency Bottlenecks: When a monolithic single agent uses a massive single prompt loaded with instructions for a complex series of tasks, the system experiences high token load, context saturation, and execution serialization, leading to incomplete execution and slow response times. Decomposing the workload into a multi-agent architecture (using specialized worker agents orchestrated by a lead or coordinator agent) enables task distribution, parallel execution, and smaller, faster prompt contexts, directly improving response times and throughput.

Accuracy & Domain Reasoning: Generative models often hallucinate or fail at domain-specific reasoning when relying solely on parametric weights. Adding a grounding data source (Retrieval-Augmented Generation / RAG via enterprise data indices, Azure AI Search, or structured document stores) supplies the model with precise, authoritative domain facts and context, eliminating hallucinations and directly improving answer accuracy.

Actions	Answer Area	
Add a grounding data source.		Action
Add a prebuilt connector.	To improve performance:	Move to a multi-agent architecture.
Move to a multi-agent architecture.	To improve accuracy:	Add a grounding data source.
Upgrade to a larger generative AI model.		

Answer:

Actions	Answer Area	
Add a grounding data source.		Action
Add a prebuilt connector.	To improve performance:	Move to a multi-agent architecture.
Move to a multi-agent architecture.	To improve accuracy:	Add a grounding data source.
Upgrade to a larger generative AI model.		

Explanation:

Move to a multi-agent architecture.

What performance means in AI systems

Performance refers to:

Speed of responses

Ability to handle complex workflows

Scalability for large workloads

Why a multi-agent architecture improves performance

A multi-agent architecture splits responsibilities across multiple specialized AI agents.

Add a grounding data source.

Grounding connects the AI model to reliable external data sources.

Instead of relying only on the model's training data, the AI can reference:

Company documents

Databases

Knowledge bases

Internal systems

This approach is commonly implemented using RAG (Retrieval Augmented Generation).

Why the other answer are incorrect:

Upgrade to a larger generative AI model

Why it is incorrect: While a larger model (such as moving from GPT-4o-mini to a full frontier model) has better

native reasoning capabilities, it typically increases latency and inference times. In a complex workflow bottlenecked by a single overloaded agent, upgrading to a larger model will often slow down response times further rather than improve performance.

Add a grounding data source.

Why it is incorrect: Adding a data source for retrieval-augmented generation (RAG) introduces an additional search and extraction step before the LLM can generate a response. This extra overhead adds to the execution path and increases latency, making it an ineffective tool for purely improving speed/throughput.

Add a prebuilt connector.

Why it is incorrect: Connectors simply provide integration endpoints to external services (like Outlook or SharePoint). They do not optimize internal token limits, task transitions, or parallelization bottlenecks.

Move to a multi-agent architecture.

Why it is incorrect: While multi-agent division of labor avoids task overloading, it does not inject missing specialized internal domain data or proprietary company information into a model's weights. If the agent doesn't have access to the specific source documents, restructuring its code into multiple agents will just mean multiple agents hallucinating or missing the context.

Upgrade to a larger generative AI model.

Why it is incorrect: A larger generic model possesses more generalized world knowledge but is still completely blind to an organization's internal, proprietary domain information. It will not fix specific gaps in domain-specific reasoning without proper external grounding data.

Add a prebuilt connector.

Why it is incorrect: Connectors only establish the bridge or API tunnel to a service; they do not automatically handle indexing, semantic vector searching, or proper context formatting needed for the LLM to yield high-accuracy answers.

Question: 9 CertyIQ

A financial services company uses Microsoft Dynamics 365 Finance.
Currently, the company's support staff manually reviews customer transaction histories to detect potential fraud cases before escalating the cases.
You need to recommend an automation solution for the review process. The solution must ensure that escalations reach a human analyst for final decision making. What should you recommend?

A. Deploy an autonomous agent that closes non-fraud cases automatically.
B. Use Microsoft 365 Copilot in Word to automatically finalize fraud detection policies.
C. Configure a task agent to generate fraud risk scores for the human analyst to review.
D. Export the data to a data lake for analysis in Microsoft Power BI.

Microsoft's agentic AI architectural guidelines prescribe assistive task agents to perform labor-intensive preliminary data collection, pattern recognition, and risk-scoring tasks, while reserving critical gatekeeping and final adjudication for human experts.

Answer: C

Explanation: Configuring a task agent to analyze transaction history, generate an objective fraud risk score, and surface findings directly to a human analyst fulfills the automation objective while preserving the required governance guardrail.

Why optionC is the best choice

Task agents in Dynamics365Finance can ingest transaction data, apply fraud-detection models, and output a

risk score.
The generated score is presented to the **human analyst** for final review and escalation decision, satisfying the requirement that "escalations reach a human analyst".
This approach leverages built-in **business process automation** (e.g., PowerAutomate/Flow) to pre-filter records, reducing manual effort while keeping a human in the decision loop.
It integrates natively with Dynamics365Finance security roles and audit trails, ensuring compliance and traceability.

Why the other options are unsuitable

A. Deploy an autonomous agent that closes non-fraud cases automatically – This removes the human analyst from the escalation path, violating the explicit requirement that a human must perform the final decision.

B. Use Microsoft365Copilot in Word to automatically finalize fraud detection policies – Copilot is a productivity assistant, not a workflow automation tool for risk scoring; it does not connect to Finance data or enforce human-in-the-loop escalation.

D. Export the data to a data lake for analysis in Microsoft PowerBI – Moving data to PowerBI creates a reporting layer but does not provide an automated **escalation mechanism** or guarantee that analysts receive the information in a timely, decision-oriented format within the Finance environment.

References

Task agents and AI-driven business process automation in Dynamics365 Finance: https://learn.microsoft.com/en-us/dynamics365/finance/audit-management/automated-risk-scoring
Human-in-the-loop escalation patterns with PowerAutomate: https://learn.microsoft.com/en-us/power-automate/automated-flows-human-in-the-loop

Question: 10 CertyIQ

A company plans to deploy a Microsoft Copilot Studio agent that will analyze historical business data to predict customer behavior.
The data is currently stored in an Azure SQL database, flat files, APIs, and logs.
You need to organize the data into a format that can be used as a knowledge source in Copilot Studio.
What should you include in the solution?

A. Azure AI Search
B. Azure Data Lake Storage
C. Azure Cosmos DB
D. Azure Translator in Foundry Tools

Microsoft Copilot Studio does not natively query or ground against dispersed raw endpoints, database tables, API payloads, and application logs in real time as a direct conversational knowledge source.

To make fragmented structured, semi-structured, and unstructured business data usable for Retrieval-Augmented Generation (RAG), the data must be ingested, chunked, vectorized, and organized into a unified searchable index.

Azure AI Search is the purpose-built PaaS information retrieval engine designed to pull data from diverse storage sources (via indexers and data pipelines) and serve as a native, enterprise-grade knowledge source in Copilot Studio.

Answer: A

Explanation:

Justification

Azure AI Search is purpose-built for ingestion, indexing, and semantic search over heterogeneous data sources — including structured relational tables, files, APIs, and logs — making it the ideal knowledge-store layer for a Copilot Studio agent that must retrieve contextual business data in real time. Its built-in support for relational connections, vector indexing, and AI-enhanced query pipelines enables low-latency retrieval of historical business data needed for predictive modelling.

Azure Data Lake Storage offers massive, retained raw storage but lacks native indexing, full-text search, or query optimization needed for efficient knowledge-source retrieval; it would require additional services (e.g., Synapse, Event-Based Ingest) to make the data searchable.
Azure Cosmos DB provides globally replicated NoSQL data with low-latency reads, yet it is optimized for transactional workloads rather than comprehensive document-level search and analytics; without integrated search capabilities it would not satisfy the requirement for a unified knowledge source.
Azure Translator in Foundry Tools is focused on language translation and offers no data storage, indexing, or retrieval features; therefore it cannot serve as a knowledge-source repository.

Conclusion – Azure AI Search directly addresses the need for scalable, indexed, and semantically searchable data across multiple storage formats, aligning perfectly with the knowledge-source requirements of a Copilot Studio agent.

References

Azure AI Search documentation: https://learn.microsoft.com/en-us/azure/search/
Copilot Studio knowledge-source integration guide: https://learn.microsoft.com/en-us/copilot-studio/knowledge-sources/

Question: 11 CertyIQ

A retail company plans to deploy Microsoft Copilot Studio agents to support:
Microsoft Dynamics 365 Commerce scenarios.
A Microsoft Power Apps inventory management solution.
You need to recommend a solution to organize product catalog data as a consistent source for multiple AI systems.
What should you recommend?

A. Let each agent scrape product details from Microsoft SharePoint Online libraries.
B. Store the product catalog data in a separate custom table for each agent.
C. Configure prompts to pull product details from the PDFs of external vendors.
D. Centralize the product catalog data in Microsoft Dataverse and expose the data to both agents.

Structured enterprise master data (such as product catalogs, inventory items, and pricing tiers) shared across first-party Dynamics 365 applications (via Dual-write / virtual tables) and Power Apps must follow a centralized single-source-of-truth pattern.

Answer: D

Explanation: Centralizing the catalog in Dataverse ensures semantic consistency, eliminates data duplication, and enables multiple downstream AI agents to ground their reasoning and execute actions against identical, synchronized records.

Why OptionD is the correct recommendation

Single source of truth – Centralizing product catalog data in **Microsoft Dataverse** guarantees that every system consumes the exact same, up-to-date records, eliminating synchronization drift between agents and Power Apps.
Built-in governance & security – Dataverse enforces row-level security, field-level permissions, and data-type validation, which are essential for e-commerce compliance and for protecting proprietary product information.
Native connectivity – Both **Copilot Studio** and **Power Apps** provide first-class connectors to Dataverse, allowing the agents to query, update, and cache catalog items without custom APIs or scraping logic.
Scalability & performance – Dataverse scales horizontally, supports rich query capabilities (filter, order, lookup), and can be indexed for fast retrieval — critical for handling large retail catalogs.

Why the other options are less suitable

A – Scrape from SharePoint

Relies on unstructured file content and manual extraction scripts.
Introduces latency, potential data loss, and lacks schema enforcement.
Increases operational overhead and breaks governance contracts.

B – Separate custom tables per agent

Creates data duplication, increasing storage cost and risk of inconsistency.
Requires duplicated maintenance (schema changes must be applied repeatedly).
Introduces synchronization challenges when product information evolves.

C – Prompt-based extraction from vendor PDFs

Depends on static PDFs that are often outdated or mismatched in format.
Harvesting data via natural-language prompts is brittle and not deterministic.
Limits support for complex catalog attributes (e.g., hierarchical categories, multimedia assets).

Conclusion
Centralizing the catalog in **Dataverse** delivers a governed, performant, and future-proof data foundation that both Copilot agents and Power Apps can consume reliably.

References

Copilot Studio documentation – Connecting to Dataverse: https://learn.microsoft.com/en-us/power-copilot-studio/connectors/dataverse
Power Apps & Power Automate – Using Dataverse as a data source: https://learn.microsoft.com/en-us/power-apps/maker/data-platform/dataverse-intro

These resources provide detailed guidance on configuring agents to read and write data from a unified Dataverse environment.

Question: 12 CertyIQ

A company has a portfolio of AI initiatives at different stages of development.
You need to recommend a structured approach to evaluating the return on AI investment (ROAI) across all the initiatives. The solution must balance immediate results with long-term values and strategic innovations.
What should you include in the recommendation?

A. a simple cost and benefit analysis
B. a horizon-based framework
C. the internal rate of return (IRR) function
D. a prioritization grid

Answer: B

Explanation:

Justification

Horizon-based framework (Correct option)

Aligns initiatives with distinct time-frames (short-term, medium-term, long-term) to capture both immediate efficiencies and future strategic innovation.

Microsoft prescribes a horizon-based framework (derived from the Three Horizons model) to structure and evaluate AI investments across distinct time-to-value phases:

Horizon 1 (Immediate / Core): Optimizes existing business processes with rapid time-to-value (e.g., standard Copilot prompt templates, basic automated workflows).

Horizon 2 (Medium-term / Emerging): Extends current capabilities into new business models or specialized autonomous task agents.

Horizon 3 (Long-term / Transformational): High-impact, multi-agent systems and strategic R&D innovations.

Enables weighting of financial returns, risk exposure, and capability development across horizons, facilitating a balanced portfolio view.
Supports dynamic re-evaluation as models mature, ensuring that early-stage experiments are not over-prioritized and that high-impact long-term investments receive appropriate attention.
Provides a repeatable structure that can be documented in governance artefacts (e.g., investment charter, ROI dashboard).

Why the other options are less suitable

Simple cost-benefit analysis – focuses only on direct expenses and revenue; it ignores timing, uncertainty, and strategic fit across multi-year horizons.
Internal Rate of Return (IRR) function – useful for single-project financial modelling but does not address portfolio-level alignment or the need to balance near-term results with future innovation.
Prioritization grid – helpful for ranking projects but lacks a formal horizon dimension; without explicit time-frame weighting it can bias selections toward low-risk, low-impact initiatives.

Therefore, a **horizon-based framework** best satisfies the requirement to "balance immediate results with long-term values and strategic innovations" across a heterogeneous AI portfolio.

References

Microsoft AI Business School – Delivering AI ROI: A strategic guide (2023) – https://learn.microsoft.com/en-us/training/modules/ai-roi-strategy/
Microsoft Azure Architecture Center – AI investment planning and ROI framework (2022) – https://learn.microsoft.com/en-us/azure/architecture/framework/ai-innovation/roi

Question: 13 CertyIQ

You need to recommend a Microsoft Power Platform business solution that consolidates data from multiple internal and external data sources. The solution must meet the following requirements:
Provide the data as a centralized source for multiple AI systems, including Microsoft Copilot Studio agents, Dynamics 365 applications, and external AI models.
Support built-in data classification and protection policies.
Provide data for grounding and analytics.
What should you include in the recommendation?

A. Microsoft Dataverse
B. Azure Data Lake Storage
C. a Microsoft Power BI semantic model
D. Azure Cosmos DB

Integration & Centralization: Dataverse can consolidate data from internal and external sources using dataflows, Power Platform connectors, Azure Synapse Link, and Virtual Tables, acting as a single source of truth across Copilot Studio agents, Dynamics 365 first-party applications, and external models via REST/OData APIs.

Governance & Protection: Dataverse integrates directly with Microsoft Purview for automated data classification, sensitivity labeling, column-level security, and Data Loss Prevention (DLP) policies

Answer: A

Grounding & Analytics: Dataverse natively integrates with Microsoft Copilot Studio for zero-copy knowledge grounding and with Microsoft Fabric/Power BI (via Azure Synapse Link / Fabric Link) for enterprise analytics.

Explanation:

Justification

Microsoft Dataverse is the native, purpose-built data platform of the Power Platform. It offers a single, governed repository that can ingest and harmonise data from on-premises, cloud, and external SaaS sources, making it ideal for consolidating disparate inputs into one spot for AI consumption.
It includes **built-in data classification, sensitivity labeling, and security policies** (via Azure Information Protection integration), satisfying the requirement for automatic protection of regulated datasets.

Dataverse supports **direct connectivity from Copilot Studio agents, Dynamics365 applications, and external AI models** through out-of-the-box connectors and the Power Platform's AI Builder capabilities, enabling it to serve as the grounding layer for analytics and AI-driven workflows.
While still capable of powering analytics, Dataverse maintains **strong data-model fidelity and query performance** needed for downstream reporting without sacrificing governance.

Why the other options are less suitable

Azure Data Lake Storage provides massive raw storage but lacks native, out-of-the-box classification and protection policies; additional services (e.g., Azure Purview) must be layered on to meet the governance requirement.
Power BI semantic model is optimized for visual analytics consumption and does not natively host or expose the data to multiple AI agents or external models; it also offers limited built-in classification.
Azure Cosmos DB delivers a globally distributed NoSQL database with high scalability, yet it requires custom development to implement data-policy controls and does not integrate directly with Power Platform's native AI services for unified grounding.

Conclusion: The only solution that simultaneously fulfills data centralisation, built-in classification/protection, and seamless AI grounding across Copilot Studio, Dynamics365, and external models is **Microsoft Dataverse**.

References

Microsoft Dataverse documentation: https://learn.microsoft.com/en-us/power-platform/data-platform/
Data loss prevention & sensitivity labeling in Dataverse: https://learn.microsoft.com/en-us/power-platform/admin/sensitivity-labeling-dataverse

Question: 14 CertyIQ

A company plans to deploy an AI-based customer service app that will autonomously manage interactions, escalate complex cases, and learn from historical ticket data.
You need to perform a return on AI investment (ROAI) analysis of the app deployment. The solution must ensure that the analysis is accurate.
What should you do first?

To ensure the analysis is accurate from the start, architects must first capture the complete cost reality across the entire lifecycle: initial design/development, data preparation/fine-tuning, deployment infrastructure, runtime token/API consumption, licensing (Dynamics 365, Copilot Studio, Azure services), and ongoing operational maintenance.

A. Establish the AI performance metrics.
B. Conduct an AI market benchmarking study.
C. Model the customer experience.
D. Identify and quantify all the development, deployment, and operating costs.

Answer: D

Attempting to project returns, model business benefits, or establish downstream evaluation benchmarks without an accurate accounting of total investment costs leads to distorted ROI calculations and financial underestimation.

Explanation:

Why optionD is the correct first step

ROI (or ROAI) accuracy depends on a complete cost model. Before any performance metrics or market data can be evaluated, you must capture every expense associated with the solution — development, integration, infrastructure, licensing, ongoing maintenance, and training.
Quantified costs provide the denominator for any return calculation. Without a baseline financial footprint, estimates of cost-savings, productivity gains, or revenue uplift cannot be translated into a reliable return figure.
All subsequent analyses (metrics, benchmarking, CX modeling) rely on this cost foundation. Having the cost data enables you to align AI performance indicators with business outcomes and to assess whether the

projected benefits outweigh the investment.

Why the other options are less suitable as the first action

A – Establish the AI performance metrics: Metrics are essential for measuring success, but they only have meaning when you know the financial context in which they will be applied. Selecting metrics prematurely can lead to mis-aligned KPIs.
B – Conduct an AI market benchmarking study: Benchmarks are useful for setting realistic expectations, yet they do not address the organization-specific investment required. Benchmarking without knowing your own cost structure risks irrelevant comparisons.
C – Model the customer experience: While CX modeling informs the value proposition, its quantitative impact must be tied to cost and revenue forecasts derived from a full cost model; otherwise the financial justification remains speculative.

Key takeaway:

Start by **identifying and quantifying all development, deployment, and operating costs** (OptionD). This creates a solid financial baseline that makes every later analysis — metric definition, benchmarking, and CX modeling — accurate, comparable, and defensible in a ROAI assessment.

References

Microsoft AI Business Cases & ROI Toolkit – Guide on building ROI models for AI solutions:
https://learn.microsoft.com/en-us/azure/ai-services/ai-roi
ROI Calculator for Azure AI Projects – Interactive calculator to estimate costs and benefits:
https://azure.microsoft.com/en-us/pricing/calculator/?service=ai-generator

Question: 15 CertyIQ

DRAG DROP
-
You are designing end-to-end test scenarios for a business solution that uses Microsoft Dynamics 365 Sales and Dynamics 365 Finance.
You need to ensure that the business solution meets the following test requirements:
•Properly exchanges data between the Dynamics 365 apps
•Aligns with defined user workflows and business processes
Which type of testing should you use for each requirement? To answer, drag the appropriate testing types to the correct requirements. Each testing type may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

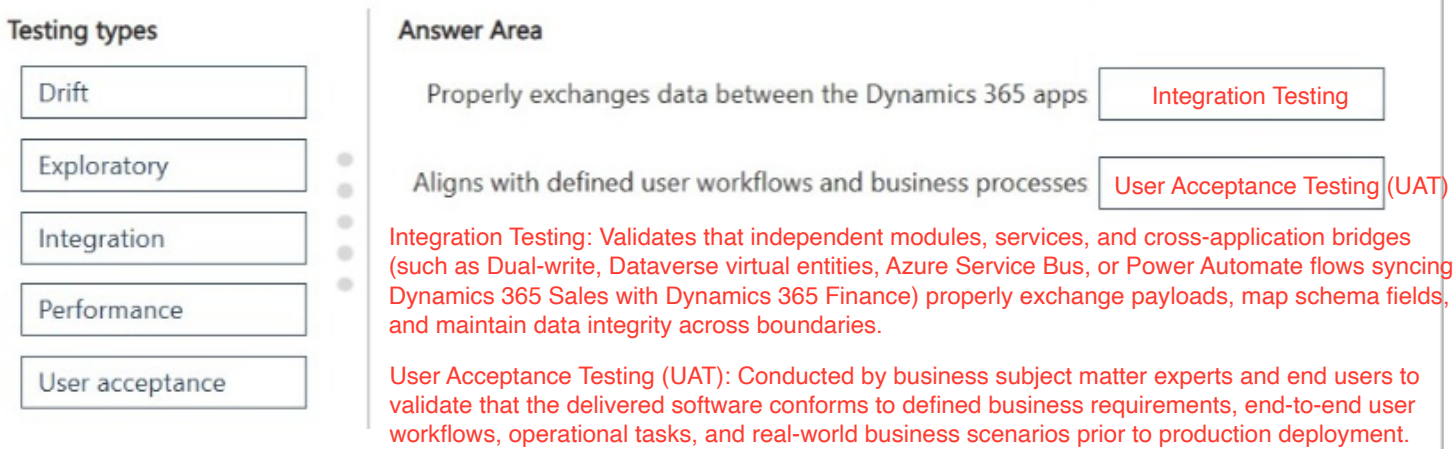

Integration Testing: Validates that independent modules, services, and cross-application bridges (such as Dual-write, Dataverse virtual entities, Azure Service Bus, or Power Automate flows syncing Dynamics 365 Sales with Dynamics 365 Finance) properly exchange payloads, map schema fields, and maintain data integrity across boundaries.

User Acceptance Testing (UAT): Conducted by business subject matter experts and end users to validate that the delivered software conforms to defined business requirements, end-to-end user workflows, operational tasks, and real-world business scenarios prior to production deployment.

Answer:

Explanation:

Integration Testing: This type of testing evaluates how effectively different applications or components pass information back and forth. It ensures that data mappings, APIs, dual-write connections, or batch synchronization protocols work reliably between Dynamics 365 Sales and Dynamics 365 Finance.

User Acceptance :This phase acts as the final verification layer performed by end-users or business analysts. It evaluates whether the configured system behaves as expected according to real-world day-to-day business processes, security roles, and user workflow journeys.

Why Other Options Are Incorrect:

Drift: Used in AI/ML scenarios to monitor if data patterns or model predictions decay over time; it is not a software functional testing type.

Exploratory: Involves unstructured, manual test sessions to discover unexpected bugs; it does not explicitly target predefined, formal workflow alignment.

Performance: Checks the technical limits, volume handling, load behavior, and response times of the software environment, rather than the correctness of functional data flows.

Question: 16 CertyIQ

A company has a Microsoft 365 tenant in Canada and multiple Microsoft Power Platform environments in Canada and the United States.
The company plans to deploy a Microsoft Copilot Studio agent to the Canadian environment that will use:
•Microsoft Dataverse data stored in Canada
•A connector that connects to an Azure OpenAI instance in the United States
You need to ensure that the agent adheres to data residency and data movement policies before being deployed.
What should you do?

A. Ensure that the data processed by Azure OpenAI is stored in the United States.
B. From the Microsoft Purview portal, validate the Data loss prevention settings.
C. Migrate the tenant to the United States.
D. Ensure that cross-region data movement is enabled for the Canadian environment and connector dependencies.

When a Copilot Studio agent or Power Platform solution hosted in one sovereign data boundary (such as Canada) communicates with an external endpoint or generative AI model located in another boundary (such as the United States), it triggers cross-border data transmission.

Answer: D

Explanation: Microsoft Power Platform enforces geographic data boundary safeguards: generative AI features and cross-geographic model routing are restricted by default unless an administrator explicitly enables the "Move data across regions" (cross-region data movement) toggle in the Power Platform admin center (PPAC) for that specific environment.

Enabling cross-region data movement provides the necessary administrative authorization for prompts, grounded Dataverse payloads, and connector telemetry to cross geographic boundaries to reach the US-hosted Azure OpenAI

Technical Justification for Correct Answer D

Scenario Recap: A company with a Microsoft 365 tenant in Canada and multiple Power Platform environments across Canada and the U.S. plans to deploy a Microsoft Copilot Studio agent in the Canadian environment. The agent will utilize:

Microsoft Dataverse data stored in Canada
A connector to an Azure OpenAI instance in the United States

Requirement: Ensure the agent adheres to data residency and data movement policies.

Why D is the Correct Answer:

Data Residency Compliance: The primary data (in Dataverse) is already stored in Canada, complying with the implied Canadian data residency requirement for this part of the setup.
Cross-Region Data Movement Necessity: Since the connector reaches out to an Azure OpenAI instance in the U.S., **cross-region data movement is inevitable** for the processing part of the workflow. Enabling this feature ensures transparency and deliberate allowance of this movement, aligning with data movement policies.

Why Other Options are Less Suitable:

A. Ensure data processed by Azure OpenAI is stored in the U.S.:
Ineffective for Data Residency of Source Data: This does not address the Canadian storage requirement for the Dataverse data.
Misinterpretation of Requirement: The focus is on ensuring compliant data movement, not altering the storage location of processed data.
B. Validate Data Loss Prevention (DLP) settings from Microsoft Purview portal:
Relevant but Insufficient: While DLP settings are crucial for protecting data, this action does **not directly address the cross-border data movement policy compliance** required for this specific deployment scenario.
Complementary Measure: Recommended as part of a broader security strategy but not the solution to the question's primary concern.
C. Migrate the tenant to the United States:
Overly Invasive and Unnecessary: This would **not solve the cross-region data movement issue** (since the Power Platform environments are already in both regions) and could **complicate existing data residency compliance** for Canadian data.

Correct Action (D):

Ensure cross-region data movement is enabled for the Canadian environment and connector dependencies.
Directly Addresses the Requirement: Allows for the acknowledged and compliant movement of data between Canada and the U.S. for the specified workflow.
Minimal Invasion with Maximum Compliance: Does not require changing the storage locations of existing data or migrating the entire tenant.

References

Microsoft Documentation: Data residency in Microsoft 365
Microsoft Documentation: Understand data movement in Power Platform

Question: 17 CertyIQ

HOTSPOT
-
A company has a Microsoft Copilot Studio agent for customer support.
You are reviewing and validating the following prompts:
•A prompt that has instructions to "help the customer as best you can"

•A prompt that helps retrieve product information from a knowledge base

You need to ensure that the agent delivers consistent and accurate responses.

What should you do for each prompt? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

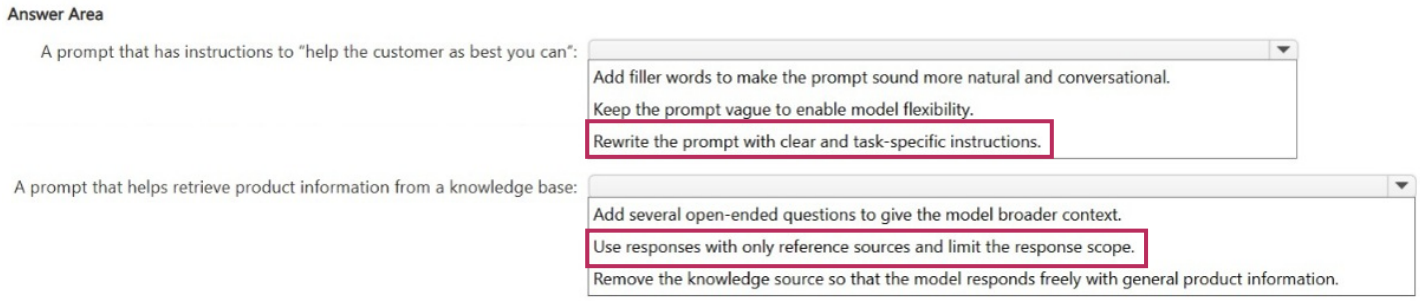

Answer:

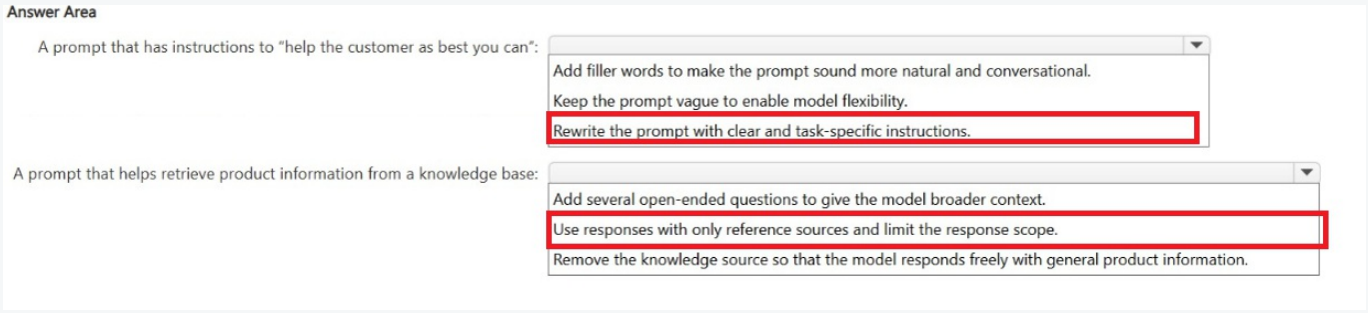

Explanation:

A prompt that has instructions to "help the customer as best you can"

Correct Choice

Rewrite the prompt with clear and task-specific instructions.
Why it is correct: Broad phrases like "help the customer as best you can" are too vague for Large Language Models (LLMs). Effective prompts require explicit goals, boundaries, formatting steps, and rules (e.g., specifying tone, what data to look up, or how to handle edge cases) to produce reliable business outputs.

Incorrect Choices

Keep the prompt vague to enable model flexibility.
Why it is incorrect: Vague prompts do not grant helpful flexibility; instead, they drastically increase the risk of hallucinations, erratic formatting, and unpredictable or unhelpful customer service answers.
Add filler words to make the prompt sound more natural and conversational.
Why it is incorrect: LLMs do not need human-like conversational courtesy inside system prompts. Adding conversational fluff or filler words consumes unnecessary tokens and degrades instruction clarity.

A prompt that helps retrieve product information from a knowledge base

Correct Choice

Use responses with only reference sources and limit the response scope.
Why it is correct: When an AI agent accesses an internal knowledge base (Retrieval-Augmented Generation / RAG), it must be heavily constrained. Forcing the model to rely only on the retrieved documents prevents it from hallucinating fake facts or using its general pre-training data to answer proprietary company questions.

Incorrect Choices

Add several open-ended questions to give the model broader context.
Why it is incorrect: Open-ended questions introduce ambiguity and widen the model's generation scope, which directly contradicts the goal of retrieving specific, accurate facts from a structured knowledge repository.
Remove the knowledge source so that the model responds freely with general product information.
Why it is incorrect: Removing the knowledge base deprives the model of your specific company documents. The model will fail completely or invent fake product features using its outdated, public training weights.

Question: 18 CertyIQ

HOTSPOT
-
You are designing a testing solution for Microsoft Copilot Studio agents.
You need to validate prompt engineering best practices to ensure that the agents generate accurate and contextually relevant responses.
Which prompt validation techniques and metrics should you include in the solution? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

Prompt validation techniques:

must execute test suites utilizing varied phrasing, alternate sentence constructs, and synonyms representing natural user variance.

- Exclude domain-specific terminology from the prompts.
- Use prompts that have varied phrasing.
- Use only simple, one-word prompts.

Metrics:

- The number of words generated per response
- Response relevance and accuracy
- The response generation time

When evaluating generative answers and prompt quality, architects track qualitative alignment metrics—specifically response relevance (how well the output addresses the query) and accuracy / groundedness (freedom from hallucination and fidelity to underlying reference data).

Answer:

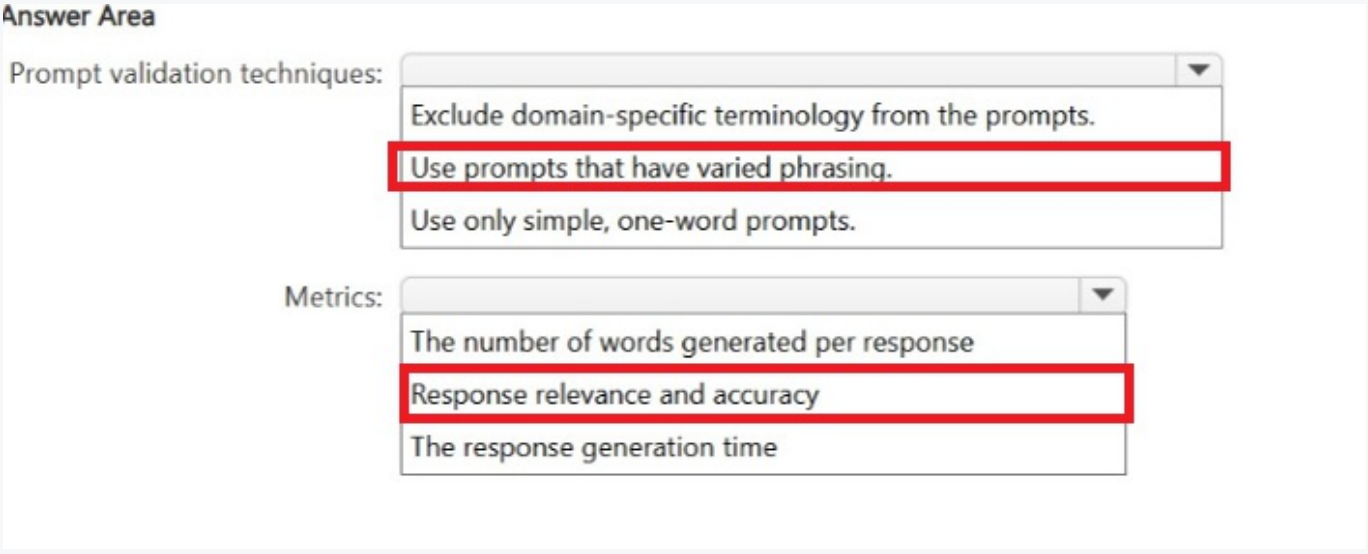

Explanation:

1. Prompt validation techniques

Use prompts that have varied phrasing: Validating prompts requires testing them against diverse linguistic styles, synonym adjustments, and semantic variations. This ensures that the system remains robust and behaves predictably regardless of slight changes in how a user asks a question.

Incorrect answers:

Exclude domain-specific terminology: This is incorrect because prompt validation must specifically test the model's ability to accurately handle specialized industry jargon and company terms it will face in production.

Use only simple, one-word prompts: This is incorrect because complex enterprise tasks cannot be evaluated with single words; prompts must mimic real-world usage patterns to be validated effectively.

2. Metrics

Response relevance and accuracy: These are core quality metrics (often evaluated using ground-truth datasets or LLM-as-a-judge frameworks) that directly measure whether the model's output satisfies the user's intent without hallucinating incorrect facts.

Incorrect answers:

The number of words generated per response: This is a trivial length metric that does not indicate whether an answer is helpful, accurate, or aligned with business rules.

The response generation time: This measures systemic latency and hardware performance rather than the linguistic quality, safety, or factual correctness driven by the prompt engineering design.

Question: 19 CertyIQ

A company has two Microsoft Power Platform environments named Dev1 and Prod1. A Microsoft Copilot Studio agent named Agent1 is built into a solution in the Dev1 environment.
You plan to deploy Agent1 to Prod1.
You need to make Agent1 available to the users in Prod1. The solution must minimize administrative effort.
What should you do?

Solutions are the core mechanism for implementing ALM in Power Platform and Microsoft Copilot Studio.

A. Share Agent1 with the users in Prod1.

B. Export the solution as an unmanaged solution and import the solution into Prod1.

C. Export the solution as a managed solution and import the solution into Prod1.

D. Create a new Copilot Studio agent in Prod1 by replicating the configuration of Agent1.

When distributing components to non-development downstream environments (test, UAT, and production like Prod1), Microsoft ALM architecture mandates deploying solutions as managed.

Answer: C

Explanation: Managed solutions protect production assets against unauthorized direct customizations, ensure clean servicing, support atomic updates/upgrades, and allow predictable rollbacks, achieving production availability with minimal administrative maintenance.

Technical Justification for Correct Answer: C

To deploy Agent1 from the Dev1 environment to Prod1 in Microsoft Power Platform while minimizing administrative effort, the most suitable approach is to **export the solution as a managed solution and import it into Prod1 (Option C)**. Here's why this option stands out and why others are less suitable:

Why Option C is Correct:

Managed Solution: Exporting as a managed solution allows for the encapsulation of Agent1 and its dependencies within a single, self-contained package. This ensures that the deployment to Prod1 is straightforward and less prone to errors due to missing components.
Minimizes Administrative Effort: The managed solution approach streamlines the deployment process. Once imported into Prod1, the solution is ready for use with minimal to no additional configuration required, assuming the environments are similarly configured.
Security and Isolation: Managed solutions provide an added layer of security by not allowing direct modification of the solution's components once imported, which is beneficial in a production environment like Prod1.

Why Other Options are Less Suitable:

A. Share Agent1 with the users in Prod1:
Environment Boundary: Sharing across environments (Dev1 to Prod1) isn't directly supported in the same way as sharing within an environment. This approach would not actually deploy the agent to Prod1 but rather attempt to access it across environments, which is not feasible for making the agent fully available in Prod1.
Security and Access Control: Could potentially bypass intended security and access controls designed to keep production environments isolated.
B. Export the solution as an unmanaged solution and import into Prod1:
Administrative Effort: Unmanaged solutions require more manual configuration post-import, as components are not encapsulated in the same way as managed solutions. This increases administrative effort, contrary to the requirement.
Dependency Management: There's a higher risk of dependency issues or missing components, as the import doesn't ensure all dependencies are properly migrated.
D. Create a new Copilot Studio agent in Prod1 by replicating the configuration of Agent1:
Administrative Effort: Manually replicating the configuration maximizes administrative effort, which is the opposite of the requirement. This approach is error-prone and time-consuming.
Consistency: Ensuring consistency between Agent1 in Dev1 and the new agent in Prod1 would require meticulous documentation and replication, increasing the likelihood of discrepancies. While data replication tools could assist, for a single agent, exporting/importing is more efficient.

Conclusion

Exporting the solution as a managed solution and importing it into Prod1 (Option C) is the most efficient and administratively light approach to make Agent1 available in the production environment while ensuring consistency, security, and minimal post-import configuration.

References

(For further reading on managed solutions and Power Platform deployment strategies)

Microsoft Documentation: Managed and Unmanaged Solutions in Power Apps
https://docs.microsoft.com/en-us/powerapps/developer/data-platform/solutions-overview#managed-and-unmanaged-solutions
Microsoft Learn: Deploying Custom Business Apps in Power Apps
https://learn.microsoft.com/en-us/learn/modules/deploy-custom-business-apps-power-apps/ (Specifically, the section on Solution Deployment Strategies)

Question: 20 CertyIQ

A company has a Microsoft Power Platform environment that contains Microsoft Dataverse data.
You create a Microsoft Copilot Studio agent named Agent1 that processes the Dataverse data.
You discover that Agent1 fails to return relevant or accurate results.
You need to improve the quality and reliability of data grounding.
What should you do?

If the grounding source contains noisy, contradictory, duplicate, incomplete, or corrupted records, the agent will inevitably produce irrelevant or inaccurate responses regardless of model capabilities.

A. Retrain Agent1.

B. Verify and cleanse the Dataverse data.

C. Use an adaptive card in Agent1.

D. Add example user inputs to the training data of Agent1.

Microsoft architectural guidelines state that the first and most critical prerequisite for high-fidelity data grounding is data preparation and hygiene—verifying data accuracy, eliminating duplicates, completing missing values, and establishing clear semantic field descriptions within Dataverse before attempting prompt or model adjustments.

Answer: B

Explanation:

Technical Justification for Correct Answer: B

Why B (Verify and cleanse the Dataverse data) is the best option:

Improving the quality and reliability of data grounding in the context of Microsoft Copilot Studio (Agent1) processing Microsoft Dataverse data hinges on ensuring the foundational data's accuracy and relevance.

Data Grounding Fundamentals: Data grounding refers to the process of ensuring AI models (like Agent1) operate on accurate, relevant, and consistent data to produce reliable outputs. Poor data quality directly impacts the model's performance.
Impact of Dataverse Data Quality on Agent1: Since Agent1's failures are attributed to non-relevant or inaccurate results, it suggests the data it is processing might be flawed. Verification and cleansing of the Dataverse data address the root cause by ensuring the data is accurate, complete, and consistent, thereby enhancing the grounding.
Precedence Over Other Actions: While retraining (A), using adaptive cards (C), or adding example user inputs (D) can be beneficial, these actions are more effective when the base data is of high quality. For instance, retraining with poor data would only reinforce inaccurate patterns.

Why Other Options are Less Suitable:

A. Retrain Agent1: Retraining without addressing the data quality issue risks reinforcing existing biases or inaccuracies present in the Dataverse data, leading to no significant improvement in result relevance or accuracy.
C. Use an Adaptive Card in Agent1: Adaptive cards enhance user interaction and feedback display but do not directly impact the data quality or the accuracy of the data grounding process. This is a UI/UX improvement rather than a data quality fix.
D. Add Example User Inputs to the Training Data of Agent1: While beneficial for handling user query variations, this does not correct the underlying issue of poor data quality in Dataverse, which is the primary cause of Agent1's failures.

Conclusion: Verifying and cleansing the Dataverse data (Option B) is the most effective first step to improve the quality and reliability of data grounding for Agent1, as it tackles the root cause of the problem.

References

1. **Microsoft Documentation - Data Quality in Dataverse**: https://learn.microsoft.com/en-us/powerapps/maker/data-platform/data-quality
2. **Microsoft Copilot Studio Documentation - Improving Model Accuracy**: https://learn.microsoft.com/en-us/microsoft-365/content-intelligence/copilot-studio/improve-model-accuracy?view=o365-worldwide

Question: 21

CertyIQ

A company plans to deploy a Microsoft Copilot Studio agent to enhance customer support.
The company stores customer data across ServiceNow, Microsoft Dynamics 365 Finance, Dynamics 365 Supply Chain Management, and Excel files in SharePoint Online.
You need to recommend a solution to ensure that the agent can deliver accurate and timely responses.
What should you recommend?

When an agent must retrieve live, transactional customer data across disparate operational systems (third-party platforms like ServiceNow, enterprise ERPs like Dynamics 365 Finance and Supply Chain Management, and Microsoft 365 services like SharePoint), the standard extensibility mechanism in Microsoft Copilot Studio is Power Platform connectors.

A. Implement a model router for query handling.
B. Create custom prompts.
C. Implement Microsoft Power Platform connectors.
D. Enable incremental indexing in Azure AI Search.

Connectors expose read and write actions via REST APIs and OpenAPI definitions, enabling the agent's orchestration engine to dynamically formulate parameters and call external endpoints at runtime.

Answer: C

Explanation:

Technical Justification for Recommended Solution

Recommended Solution: C. Implement Microsoft Power Platform connectors

Why C is the best option: The company's customer data is dispersed across multiple platforms, including ServiceNow, Microsoft Dynamics 365 (Finance and Supply Chain Management), and Excel files in SharePoint Online. To enable the Microsoft Copilot Studio agent to deliver accurate and timely responses, it must seamlessly integrate with these diverse data sources. Implementing Microsoft Power Platform connectors is the most suitable approach because:

Unified Integration Platform: Power Platform (including Power Apps, Power Automate, and Power BI) offers pre-built connectors for all mentioned platforms (ServiceNow, Dynamics 365, SharePoint Online), facilitating unified access to customer data.
Real-time Data Accessibility: Connectors provide real-time or near-real-time data synchronization, ensuring the agent's responses are based on the most current information.
Scalability and Maintainability: Leveraging Power Platform connectors simplifies the integration architecture, making it easier to manage and scale as the company's data ecosystem evolves.

Why other options are less suitable:

A. Implement a model router for query handling:
Insufficient Addressal of Data Source Diversity: A model router primarily manages query distribution among AI models, not integrating disparate data sources.
Additional Complexity Without Solving the Core Issue: Would add a layer of complexity without ensuring the agent has unified access to customer data.
B. Create custom prompts:
Does Not Address Underlying Data Integration Need: Custom prompts enhance user input formatting but do not resolve the lack of unified data access for the agent.
Limited Impact on Response Accuracy Due to Potential Data Silos: Prompts won't improve accuracy if the agent can't access all relevant data.
D. Enable incremental indexing in Azure AI Search:
Partial Solution: While helpful for keeping search indexes up-to-date, it assumes data is already centralized or connected in a manner that Azure AI Search can effectively index, which is not the case.
Does Not Solve the Fundamental Integration Challenge: The primary issue of integrating multiple, distinct data sources remains unaddressed.

References

Microsoft Power Platform Connectors Documentation
Microsoft Copilot Studio Integration Scenarios

Question: 22 CertyIQ

A manufacturing company wants to deploy an agent that will automate supplier invoice processing.
You are designing a solution to evaluate the financial implications of the deployment. The company is especially concerned about budget overruns.
You need to ensure that the solution considers the total cost of ownership (TCO), the expected savings from using automation, and whether to extend the existing AI capabilities.
What should you include in the design?

A. a break-even analysis only
B. adopting prebuilt agents to reduce the deployment time
C. training a custom model
D. a return on AI investment (ROAI) analysis

Answer: D

Explanation:

Justification for Correct Answer: D

The correct answer, **D. a return on AI investment (ROAI) analysis**, is the most comprehensive approach to address the manufacturing company's concerns when deploying an AI agent for automating supplier invoice processing. Here's why D is the best choice and why the other options are less suitable:

Why D (ROAI analysis) is the best choice:

Comprehensive Financial Evaluation: An ROAI analysis provides a thorough assessment of the return on investment (ROI) specifically for AI initiatives, aligning perfectly with the company's request to evaluate financial implications, including budget overruns.
Total Cost of Ownership (TCO) Consideration: ROAI inherently includes calculating TCO, which encompasses not just the initial deployment costs but also ongoing expenses, maintenance, and potential upgrades.
Expected Savings Evaluation: It quantitatively assesses the anticipated savings from automation, directly addressing the company's interest in understanding the financial benefits of the AI solution.
Extension of Existing AI Capabilities: An ROAI analysis can also inform decisions on whether extending current AI capabilities is financially viable, considering scalability and integration costs against potential additional benefits.
Why Other Options are Less Suitable:

A. A Break-Even Analysis Only:

Too Narrow: While useful, break-even analysis only tells you when the investment will pay for itself, not the overall ROI or long-term financial health of the project.
Lacks Comprehensive Insight: Doesn't fully address TCO, detailed savings analysis, or the viability of extending AI capabilities.

B. Adopting Prebuilt Agents to Reduce Deployment Time:

Implementation Strategy, Not Financial Justification: This option focuses on the deployment method rather than providing a financial justification or evaluation of the investment.
Financial Implications Not Assessed: Doesn't directly address the company's concerns about budget overruns, TCO, or expected savings.

C. Training a Custom Model:

Technical Approach, Not Financial Evaluation: Similar to B, this outlines a technical strategy (custom vs. prebuilt) without offering a framework for financial analysis or justification.
Ignores Financial Evaluation Need: Fails to address the primary requirement of evaluating the financial implications and ROAI.
Conclusion:

Given the manufacturing company's specific concerns about financial implications, budget overruns, expected savings, and the potential extension of AI capabilities, **D. a return on AI investment (ROAI) analysis** is the most appropriate choice. It uniquely combines a comprehensive financial evaluation with the specific needs of AI investment analysis.

References:

(For further study on ROAI and TCO in the context of AI deployments)

1. **Microsoft Azure - Total Cost of Ownership Calculator**https://azure.microsoft.com/en-us/pricing/tco/calculator/Useful for understanding and calculating TCO in cloud and AI deployments.
2. **IEEE - Return on Investment (ROI) Analysis for AI Projects**https://ieeexplore.ieee.org/document/9316984 (Subscription or Purchase Required)Provides insights into conducting ROI analyses specifically tailored for AI projects, akin to ROAI.

Question: 23 CertyIQ

A company has a Microsoft Power Platform solution that contains the following components:
•Microsoft Dataverse tables
•A Microsoft Power BI workspace named WS1
•A canvas app named App1 that uses Dataverse
•A Power BI semantic model that connects to Dataverse by using DirectQuery
You plan to use generative AI to provide answers to queries based on a subset of corporate data.
You need to ensure that the data is available as a grounding data source for AI systems.
What should you do?

A Power BI semantic model built on DirectQuery already encapsulates the curated schema, relationships, business measures, and the intended "subset of corporate data" from Dataverse.

A. Populate a Dataverse table.
B. Share WS1.
C. Endorse the semantic model.
D. Export the semantic model.

To make Power BI semantic models discoverable, reusable, and qualified as authoritative grounding data sources for generative AI and Copilot consumption across an organization, Microsoft prescribes endorsing the semantic model (promoting or certifying it).

Answer: C

Explanation: Endorsement signals to enterprise AI agents, Copilot orchestrators, and organizational data consumers that the semantic definitions, calculations, and data boundaries have been verified by business owners as trusted corporate sources of truth.

Technical Justification for Correct Answer: C - Endorse the semantic model

Why C is the best option: The goal is to make the data available as a grounding data source for AI systems, which implies ensuring the data is properly structured, accessible, and validated for AI model training or query answering. Endorsing the semantic model (Option C) is the most suitable choice because:

Semantic Model Endorsement in Power BI signifies that the model is approved for use across the organization, ensuring consistency and reliability - crucial for training AI models.
The semantic model already **connects to Dataverse via DirectQuery**, meaning it's designed to leverage the up-to-date data from Dataverse without needing additional data migration or synchronization efforts.
Endorsing it **does not require additional data storage or sharing configurations**, keeping the solution efficient and secure.

AI systems can then be integrated to query or learn from this endorsed model, benefiting from its validated and organization-approved structure.

Why other options are less suitable:

A. Populate a Dataverse table:
Redundancy: Data is already in Dataverse tables, which the semantic model draws from.
Additional Maintenance: Would introduce another layer of data synchronization to manage.
No Direct AI Benefit: Populating another table doesn't directly enhance AI's access to validated data structures.
B. Share WS1 (Power BI Workspace):
Security and Access Control: Sharing the entire workspace might grant unnecessary access to other components or sensitive data within WS1.
Lack of Specificity: Sharing the workspace doesn't specifically highlight or prepare the data for AI system consumption.
D. Export the semantic model:
Static Data Issue: Exporting would provide a static snapshot, not a dynamic, up-to-date data source for AI systems.
Integration Complexity: AI systems would need to be configured to ingest this exported model, introducing additional integration challenges.

Conclusion: Endorsing the semantic model (C) is the most efficient, secure, and effective way to make the data available for AI systems, leveraging the existing DirectQuery connection to Dataverse and ensuring the data model is organization-approved for broader, including AI, usage.

References:

1. **Microsoft Documentation - Endorse a model in Power BI**: https://learn.microsoft.com/en-us/power-bi/admin/service-admin-endorse-models
2. **Power BI Semantic Models and DirectQuery**: https://learn.microsoft.com/en-us/power-bi/service-dataset-overview#semantic-models-and-directquery

Question: 24 CertyIQ

DRAG DROP
-
A company plans to implement an AI business solution for a consumer goods company.
You need to create agents that meet the following requirements:
•Orchestrate the sales order fulfillment and shipping of goods to customers.
•Analyze historical data and trends to replenish stock.
Which type of agent should you use for each requirement? To answer, drag the appropriate agent types to the correct requirements. Each agent type may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

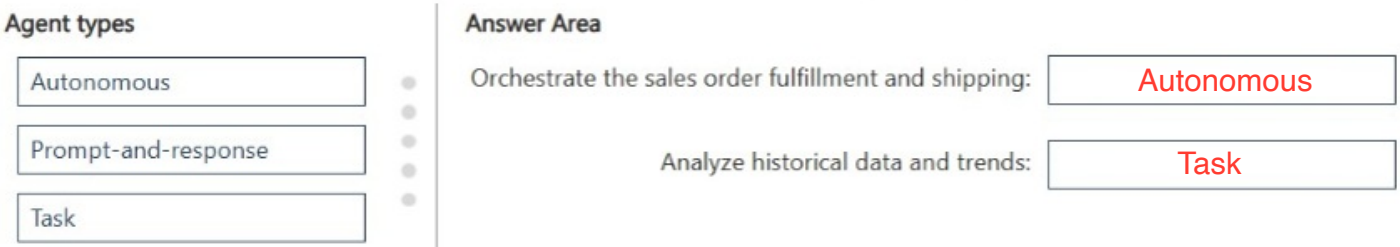

Answer: Autonomous Agents: Designed to coordinate end-to-end operational lifecycles without requiring continuous, turn-by-turn user steering. Orchestrating order fulfillment and shipping spans long-running asynchronous workflows, dynamic multi-system actions

Task Agents: Bounded, domain-specific AI agents engineered to execute a targeted analytical or computational workflow (such as evaluating historical consumption velocity, seasonal variance, and reorder point calculations to generate stock replenishment recommendations).

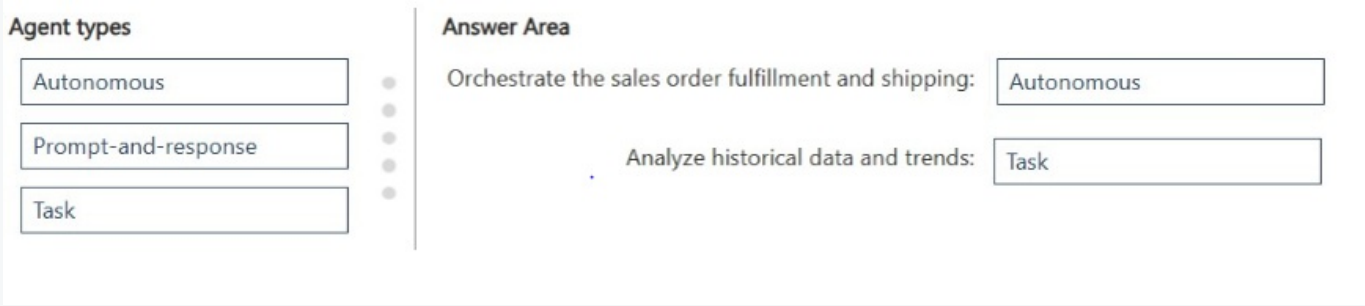

Explanation:

Autonomous Agent: This agent type is designed to operate independently toward multi-step objectives, manage end-to-end workflows (like order processing or dispatch), execute sequential decisions, and automatically adapt or escalate whenever system states change.

Task Agent: This agent type is designed for bounded, well-defined analytical, computational, or specific operational tasks. It acts on demand or structured flows to perform specialized sub-routines like processing historical analytics, generating specific data reports, or forecasting replenishment parameters.

Why the other answer are incorrect:

Prompt-and-response

Why it is incorrect: This agent type is designed only for basic, single-turn information retrieval (like a traditional chatbot answering FAQs). It lacks the capability to chain multiple steps together or actively update backend systems to fulfill a complex, multi-system operational workflow like shipping.

Question: 25 CertyIQ

A company has an AI agent that automates the review of customer feedback stored in a cloud database.
You plan to generate monthly reports from the agent's output to provide insights into customer sentiment and guide product development and marketing.
You need to ensure that the data ingested by the agent is clean and suitable for the intended use.
What should you do to prepare the data?

A. Create a workflow in Microsoft Power Automate.
B. Identify and address biased data.
C. Create an agent flow in Microsoft Copilot Studio.
D. Sort the database by customer last name.

Preparing the data specifically entails establishing data hygiene: identifying, cleansing, and addressing biased or unrepresentative records so that the agent's synthesized outputs reflect objective, balanced customer reality rather than distorted feedback loops.

Answer: B

Explanation:

Technical Justification for Correct Answer: B

Why B is the Correct Option: Identifying and addressing biased data is crucial to ensure the AI agent's output is reliable and trustworthy for generating insights into customer sentiment. Biased data can lead to skewed analysis, resulting in misguided product development and marketing strategies. By addressing biases, the company can increase the fairness and accuracy of the feedback analysis, directly impacting the quality of the monthly reports. This step is fundamental in data preparation for AI-driven analytics, as it ensures the data's integrity and suitability for the intended use.

Why Other Options are Less Suitable:

A. Create a workflow in Microsoft Power Automate: While Power Automate is useful for automating business processes, it does not directly address the issue of data cleanliness or bias. This tool would be more relevant for streamlining the report generation process after the data has been prepared.
C. Create an agent flow in Microsoft Copilot Studio: Copilot Studio is designed for building conversational AI experiences. Creating an agent flow here does not pertain to the preparation of existing customer feedback data for analysis, but rather to building interactive AI interfaces.
D. Sort the database by customer last name: Sorting the database by customer last name is an organizational task that does not impact the cleanliness, bias, or analytical suitability of the data. It's a basic data management task unrelated to preparing data for AI-driven sentiment analysis.

Key Technical Rationale for B:

Data Integrity: Addressing biased data ensures the integrity of the insights generated.
AI Model Reliability: Feeding clean, unbiased data to an AI model enhances its reliability and the accuracy of its outputs.
Business Impact: Accurate sentiment analysis directly influences the effectiveness of product development and marketing strategies.

References:

Microsoft Learn - Ensure data quality for machine learning: Provides insights into the importance of data quality in AI and ML projects.
MIT Sloan - The Dangerous Bias of Artificial Intelligence: Highlights the risks of biased data in AI systems and the need for addressing it.

Question: 26 CertyIQ

A company is designing a Microsoft Power Platform solution to reduce the manual steps of a business process by deploying an existing AI model.
You need to calculate the return on AI investment (ROAI) by identifying the metadata and telemetry of the solution.
What should you use?

The Business value toolkit (part of Microsoft's Power Platform Center of Excellence / CoE guidance) is the purpose-built suite of sample applications, processes, and Dataverse schemas designed specifically to track solution metadata, capture telemetry/survey feedback, and quantify savings using an integrated value calculator.

A. Microsoft Power Platform admin center
B. Success by Design
C. the Business value toolkit
D. Microsoft Cloud Adoption Framework for Azure

Answer: C

It enables architects to calculate ROAI and total financial impact by correlating solution usage metrics directly with organizational KPIs and labor cost savings.

Explanation:

Technical Justification for Correct Option (C)

Why C (The Business Value Toolkit) is the Best Option: The Business Value Toolkit is specifically designed to help organizations quantify the value of their Microsoft Power Platform investments, including calculating return on investment (ROI) or in this case, return on AI investment (ROAI). This toolkit provides templates, calculators, and methodologies to assess the business value of digital transformation initiatives, such as integrating an existing AI model into a Power Platform solution. For calculating ROAI by identifying metadata and telemetry, the Business Value Toolkit offers the most relevant tools and frameworks to quantify efficiencies gained from automating manual steps through AI integration.

Key Capabilities Making C Suitable:

ROI/ROAI Calculation Templates: Pre-built templates facilitate the calculation of return on investment, adaptable for AI-specific investments.
Metadata and Telemetry Integration: Guides on how to leverage metadata (e.g., process efficiency metrics) and telemetry (e.g., usage and performance data) to inform ROAI calculations.
Alignment with Power Platform: Designed with Microsoft ecosystem investments in mind, ensuring a seamless fit for Power Platform solutions.

Why Other Options are Less Suitable:

A. Microsoft Power Platform Admin Center:

Primary Focus: Administration, Monitoring, and Configuration of Power Platform environments.
Limitation: While it provides some analytics and usage metrics, it lacks the specific ROI/ROAI calculation frameworks and templates focused on business value quantification for AI integrations.

B. Success by Design:

Primary Focus: A methodology for ensuring successful adoption and implementation of Microsoft products, focusing on planning, design, and deployment.
Limitation: More about the upfront planning and less about the post-implementation evaluation of ROI/ROAI for specific technology integrations like AI.

D. Microsoft Cloud Adoption Framework for Azure:

Primary Focus: Comprehensive framework for cloud adoption, including governance, security, and migration strategies, tailored for Azure.
Limitation: Too broad in scope, not specifically tailored for calculating ROAI within a Power Platform and AI model integration context. Its focus is more on the cloud adoption journey than on the business value calculation of specific solutions.

Question: 27 CertyIQ

HOTSPOT
-
A company has a Microsoft Power Platform environment.
You need to build two agents named Agent1 and Agent2. The solution must meet the following requirements:
•Agent1 must be extendable by using the Semantic Kernel and must connect to multiple business apps and APIs.
•Agent2 must connect directly to data stored in Microsoft Dataverse and must be embeddable in a Microsoft Power Apps canvas app.
What should you use to build each agent? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Microsoft Foundry (Azure AI Foundry): Designed for code-first, developer-driven enterprise AI and multi-agent solutions. It provides deep integration with SDKs and orchestration frameworks like Semantic Kernel, AutoGen, and LangChain, enabling pro-code developers to author plugins, define complex native memories, and interface programmatically with heterogeneous enterprise APIs and custom microservices.

Microsoft Copilot Studio: The low-code conversational AI studio native to Power Platform. Copilot Studio agents natively ground directly against Microsoft Dataverse tables with zero-copy architecture and can be embedded seamlessly into Microsoft Power Apps canvas apps via the native Copilot canvas control or modern component container.

Answer Area

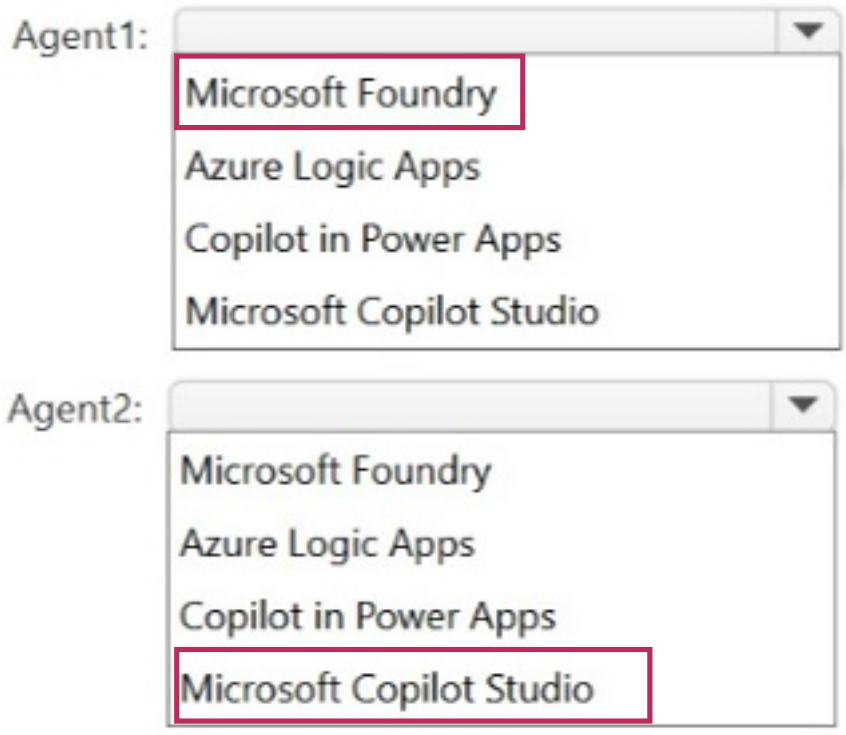

Answer:

Answer Area

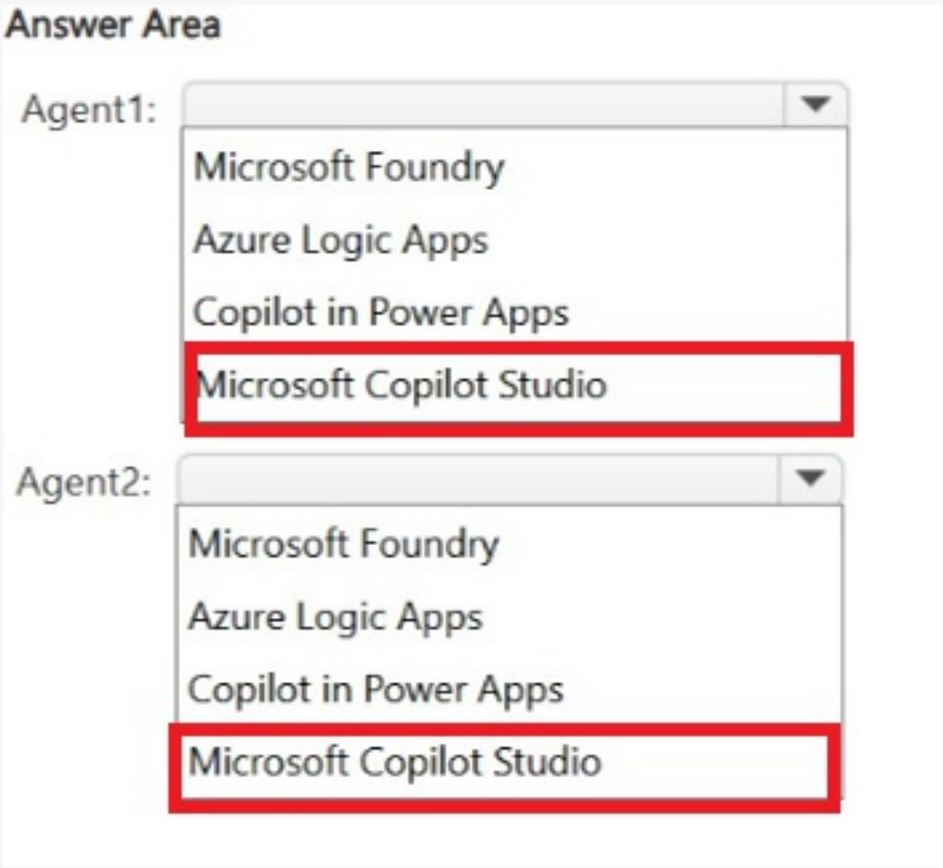

Explanation:

Microsoft Copilot Studio (Correct for Both): This is the unified enterprise platform used to build, customize, and orchestrate conversational and agentic AI. It natively handles standard custom chatbots (using intent mapping like CLU for Agent1) as well as advanced telephony and voice-channel configurations (using speech-optimized engines like NLU+ for Agent2).

Why Other Options Are Incorrect

Microsoft Foundry: This is a specialized development environment primarily focused on lower-level AI models, custom foundational tooling, or research workflows rather than ready-to-deploy low-code/no-code business agent orchestration.

Azure Logic Apps: This is a workflow automation tool designed to connect apps, data, and systems via APIs (integration/ETL). While it can be called by an agent to execute an action, it cannot serve as the core conversational agent hosting platform itself.

Copilot in Power Apps: This refers specifically to the embedded AI assistant features inside canvas or model-driven apps that help end-users build apps or query data tables. It is not an environment for building standalone, omni-channel customer service bots.

Question: 28 CertyIQ

A company has an Azure environment that supports multiple business units.
The company plans to implement an AI solution that will perform sentiment analysis on customer product reviews.
You need to evaluate the potential cost of the solution to support return on AI investment (ROAI) analysis.
What should you use?

A. Cost Management + Billing
B. Microsoft Fabric SKU Estimator
C. Total Cost of Ownership (TCO) Calculator
D. Azure Reservations

Tricky: Azure Pricing Calculator

Answer: C

Explanation:

Technical Justification for Correct Answer: C (Total Cost of Ownership (TCO) Calculator)

Why C (Total Cost of Ownership (TCO) Calculator) is the Best Option: The Total Cost of Ownership (TCO) Calculator is the most suitable choice for evaluating the potential cost of an AI solution for sentiment analysis in this scenario. This tool provides a comprehensive estimate of all costs associated with owning and operating the proposed AI solution on Azure, including initial setup, ongoing expenses (such as compute, storage, and potentially Azure Cognitive Services for AI functionalities), maintenance, and any other related expenditures. For an ROI analysis on AI investments (ROAI), understanding the total cost is crucial to accurately assess the return against the investment.

Key Reasons:

Comprehensive Cost Assessment: TCO Calculator evaluates all aspects of cost, which is essential for a thorough ROAI analysis.
Azure AI Service Integration: It can account for the specific costs of Azure services that will be used for the AI solution (e.g., Azure Machine Learning, Cognitive Services).
Long-term Cost Projection: Enables planning based on projected long-term costs, aligning with the strategic evaluation of ROAI.

Why Other Options are Less Suitable:

A. Cost Management + Billing:
Limited Scope: Primarily focused on tracking and managing current spending rather than projecting total ownership costs for a new solution.
Lack of Predictive Capability: Not designed for estimating future costs of a proposed deployment.

B. Microsoft Fabric SKU Estimator:
Specific Use Case Limitation: Assuming "Microsoft Fabric" refers to a product/service with a specific SKU estimator, this tool would likely be too narrow in scope, focusing on estimating costs for a particular product or service rather than a custom AI solution.
Insufficient for AI Solution: May not fully capture the dynamic costs associated with developing and deploying an AI-driven sentiment analysis solution.
D. Azure Reservations:
Cost Optimization Tool: While useful for reducing costs by committing to resource usage over a period, it does not estimate the total cost of ownership for a new solution.
Post-Implementation Focus: More beneficial after the solution's cost structure is understood and reservations can be strategically purchased.

Correct Answer Justification in Brief:

Best for ROAI Analysis: TCO Calculator for its comprehensive and predictive cost analysis capabilities.
Less Suitable Due to: Limited scope (A), narrow focus (B), and post-implementation optimization (D).

References

(For further reading on the recommended tool and Azure cost management in general)

1. **Microsoft Total Cost of Ownership (TCO) Calculator**
2. URL: https://www.microsoft.com/en-us/microsoft-365/calculator
3. Note: While the link provided is broad, navigating through can lead to TCO calculators relevant to Azure and AI services. For a more direct experience, consider searching "Azure TCO Calculator" on Microsoft's official website.
4. **Azure Pricing and Cost Management Documentation**
5. URL: https://azure.microsoft.com/en-us/pricing/
6. Relevance: Comprehensive resource for understanding Azure pricing models, including tools for cost estimation and management.

Question: 29 CertyIQ

HOTSPOT
-
You need to recommend a Microsoft Power Platform solution for customer support. The solution must include AI capabilities in Microsoft Power Automate and must meet the following requirements:
•Use a centralized workspace for AI models.
•Generate short overviews from large amounts of unstructured text, such as case notes or transcripts, without requiring additional training or coding.
What should you include in the recommendation for each requirement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Centralized Workspace: In Microsoft's unified AI architecture, a Microsoft Foundry hub (Azure AI Foundry hub) serves as the top-level resource and centralized collaborative workspace. It governs access, security boundaries, shared connections, compute, and model deployments across multiple projects, providing a single control plane for managing foundational models and custom fine-tuned assets across enterprise workloads.

Prebuilt Prompt for Summarization: In Power Automate and AI Builder, Microsoft provides out-of-the-box prebuilt prompts (specifically the Summarize text prebuilt prompt). Because it utilizes large language models out-of-the-box to generate concise overviews from arbitrary unstructured text (case notes, support transcripts) without requiring prompt authoring, custom dataset training, or code, it directly addresses the second requirement.

Answer Area

Use a centralized workspace:

- An Microsoft Foundry hub
- Azure OpenAI Foundry
- Microsoft Copilot Studio
- Microsoft Dataverse

Generate short overviews:

- An AI Builder prebuilt model
- An AI Builder prebuilt prompt
- Azure OpenAI
- GitHub Copilot
- Microsoft Copilot in Power Automate

Answer:

Answer Area

Use a centralized workspace:

- An Microsoft Foundry hub
- Azure OpenAI Foundry
- Microsoft Copilot Studio
- Microsoft Dataverse

Generate short overviews:

- An AI Builder prebuilt model
- An AI Builder prebuilt prompt
- Azure OpenAI
- GitHub Copilot
- Microsoft Copilot in Power Automate

Explanation:

Use a centralized workspace

Microsoft Dataverse.

Why it is correct: Inside the Microsoft Power Platform, the AI Hub runs entirely on top of Dataverse. It is the mandatory data foundation used to store, secure, govern, track versions, and share custom prompt templates and AI Builder configurations across all business applications.

Incorrect Choices

Microsoft Copilot Studio: This is an orchestration platform designed for building standalone AI bots and conversational agents, not a system-level backend metadata repository for general application storage.

Azure OpenAI Foundry / An Microsoft Foundry hub: These are code-first, developer-centric spaces intended for fine-tuning baseline models and testing LLM applications. They do not act as the low-code storage backend for Power Apps or Power Automate solutions.

Generate short overviews

An AI Builder prebuilt prompt.

Why it is correct: Microsoft provides a native, pre-configured AI Builder template specifically named "Summarize text". Choosing a prebuilt prompt lets an application instantly consume massive paragraphs of text and return concise bulleted summaries with zero coding or instruction writing.

Incorrect Choices

An AI Builder prebuilt model: Prebuilt models are designed for structured cognitive perception (e.g., receipt scanning, sentiment extraction, object detection). Generating summaries via prompt text falls explicitly under generative prompts, not structured classification models.

Azure OpenAI: While capable, calling a custom model via Azure OpenAI directly requires manual API management, deployment endpoints, and manual prompt creation, adding unnecessary configuration overhead.

GitHub Copilot: This is a developer-focused assistant integrated inside code editors to write scripts or functions, rather than a citizen-developer tool meant for handling automated data pipelines.

Microsoft Copilot in Power Automate: This refers to the auxiliary sidebar chat that helps a developer build or troubleshoot their automation flow logic, rather than a run-time tool used to clean and summarize business datasets.

Question: 30 CertyIQ

HOTSPOT
-
A company uses Microsoft Dynamics 365 Finance for accounts payable and customer debt recovery.
You are designing an AI finance process that meets the following requirements:
•Provides AI-driven details to help staff identify overdue vendor invoices and outstanding balances
•Helps staff reduce how long it takes to review overdue invoices and payment history
You need to recommend which Microsoft Copilot features to include in the design.
What should you recommend for each requirement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

AI Summaries with Copilot (Vendor Summary): For Accounts Payable (AP), hovering over or opening a vendor record leverages AI summaries with Copilot. This feature synthesizes transactions, active holds, open purchase orders, payment statuses, and highlights overdue vendor invoices alongside outstanding balances without manual subledger navigation.

Collections coordinator summary: For Customer Debt Recovery (Accounts Receivable / Credit and Collections), the Collections coordinator summary workspace feature displays AI-generated summaries powered by Azure OpenAI Service. It evaluates customer payment history over the past year, aging buckets, and open invoices over the past six months to drastically cut the review time needed to initiate collection activities or draft AI reminder letters.

Answer Area

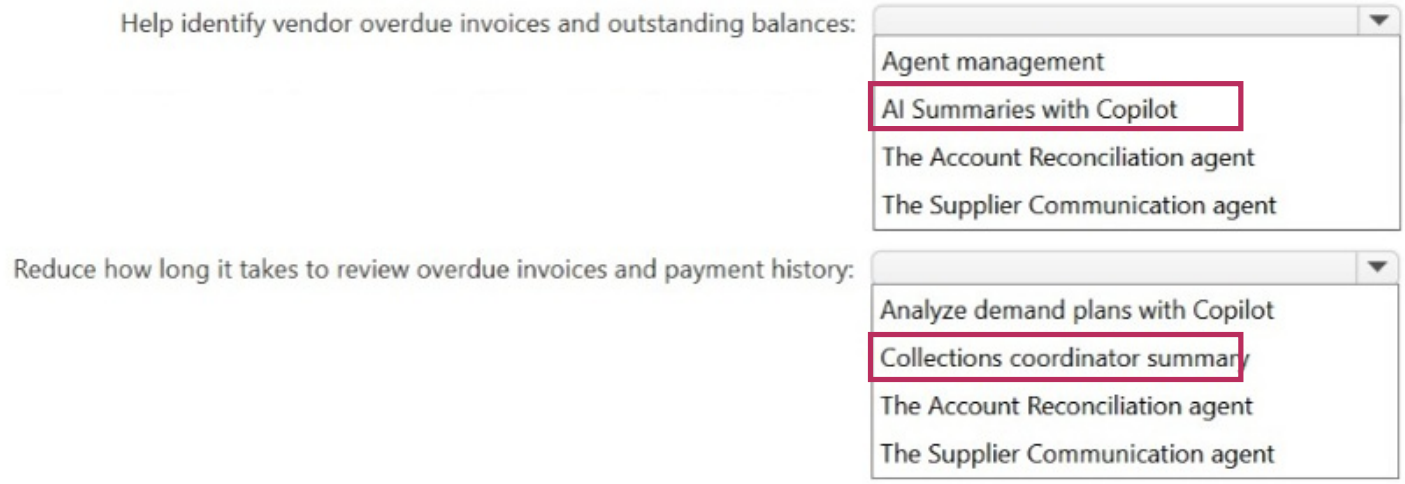

Answer:

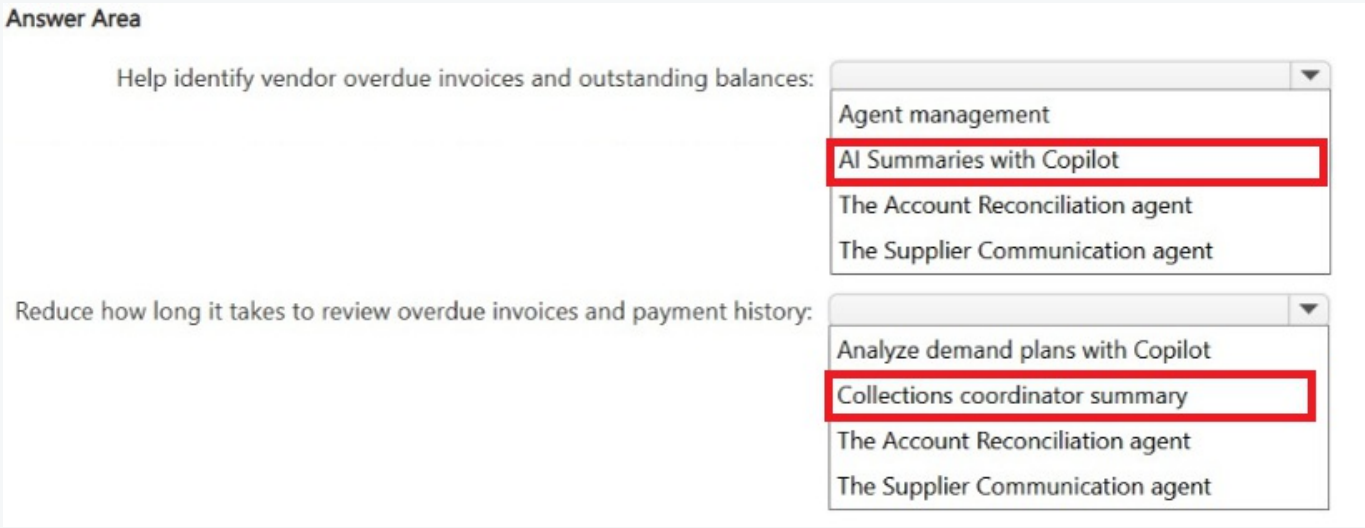

Explanation:

1. Help identify vendor overdue invoices and outstanding balances

AI Summaries with Copilot

In the context of accounts payable and general vendor monitoring, the AI summaries feature provides clear overview summaries directly inside records. This enables accounting staff to rapidly spot vendor overdue items and check outstanding balances across transactional data without manually expanding line elements.

2. Reduce how long it takes to review overdue invoices and payment history

Collections coordinator summary:

This workspace tool is specifically tailored for customer debt recovery. It uses automated workspace intelligence to evaluate customer records over specified periods (e.g., past 6 months of invoices and 1 year of payments), synthesizing an analytical overview to dramatically slash review time and even draft corresponding reminder communications.

Why Other Options Are Incorrect

Agent management: This is an administrative framework within feature management used by system administrators to install or activate specific functional agents (like enabling telemetry hooks); it has no direct

role in surfacing user-facing transactional analytics.

The Account Reconciliation agent: This agent focuses on automated general ledger matching against subledgers (e.g., bank statements, cash-to-ledger accounts) at fiscal closing periods to handle discrepancies, rather than aggregating vendor/customer communication overviews.

The Supplier Communication agent: This specific agent acts on procurement pipelines to track unconfirmed purchase orders or manage incoming email change requests regarding delivery delays, rather than reviewing accounts receivable collection tracking or customer ledger payment histories.

Analyze demand plans with Copilot: This is a production planning utility within the Dynamics 365 Supply Chain Management workspace intended to run mathematical calculations over future material requirements, completely unrelated to customer collection operations.

Question: 31 CertyIQ

Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Overview -
Contoso, Ltd. is a high-tech manufacturing company that uses Microsoft Dynamics 365 Finance. Dynamics 365 Supply Chain Management, and Dynamics 365 Commerce for its North American operations. The company designs and develops innovative products that have many patents and proprietary technologies. The patents and engineering designs are closely guarded secrets.
Contoso executives want to integrate and adopt AI solutions to help scale the company in preparation for an anticipated period of rapid growth.
The company has multiple legal entities and Azure subscriptions that will be used in the adopted AI solutions.
Requirements -
AI Adoption -
The following executives will have specific responsibilities in the overall AI adoption:
•Chief Technology Officer (CTO): Select one Dynamics 365 Finance, Dynamics 365 Supply Chain Management or Dynamics 365 Commerce prebuilt AI agent and one custom Microsoft Copilot Studio AI agent to prioritize and deploy during the initial AI adoption phase.
•Chief Information Officer (CIO): Ensure that appropriate security labels are assigned to the data used by the AI agents.
•Chief Financial Officer (CFO): Analyze the return on investment (ROI) for the AI agents being deployed.
•Chief Information Security Officer (CISO): Discover and inventory AI resources for auditing.
•Chief Executive Officer (CEO): Ensure that all solutions adhere to industry-standard responsible AI practices.
All AI initiatives and agents will have a detailed business use case, a defined audience profile, and an estimated ROI that will compare the cost savings of the current process against the estimated costs of using the new AI solutions.
The company's research and development (R&D) department already has a custom Model Context Protocol (MCP) server that contains comprehensive product specifications and compliance data.
Prebuilt AI Agent -
The CTO has NOT yet selected which prebuilt AI agent to use in Dynamics 365 Supply Chain Management. The CTO wants to view available agent templates to identify which agent will add the most business value.

Depending on which high-priority AI agents are identified, its agent capabilities must be previewed in a discovery meeting with the relevant business operation stakeholders.
Custom AI Agent -
Contoso has identified the following custom AI agent requirements:
•The custom AI agent will use data from Dynamics 365 Supply Chain Management to answer questions for the manufacturing team as a low-code solution.
•The custom AI agent will be accessible from within Microsoft Teams.
•The custom AI agent must be designed to eventually connect to other agents that can be selected based on their description.
•The topics used in the custom AI agent will be selected based NOT on a trigger phrase, but on a description of the purpose of the query, to make the interactions more conversational.
•The custom AI agent must be able to answer questions about product specifications by using existing technologies. The product specifications are maintained by the R&D department.
•The custom AI agent must be integrated with and accessible from Dynamics 365 Supply Chain Management.
•The custom AI agent must be able to use Dynamics 365 Supply Chain Management business logic that is stored outside of the application.
Analysis, Reporting, and Troubleshooting
Contoso has identified the following analysis, reporting, and troubleshooting requirements:
•The CISO will audit all the AI solutions monthly for compliance and security.
•The CFO will analyze all the AI solutions quarterly to compare the estimated ROI against actual measured efficiencies and adoption. The CFO will use the Copilot Studio agent usage estimator to perform this analysis.
•The CISO wants to identify how much sensitive data was accessed for a given AI agent run and who accessed the data. Too much sensitive data accessed by a single user might indicate a high security risk.
•The CTO wants to track user feedback on the quality of the AI agent responses during user interactions with the agents. Consistently poor feedback will trigger an escalated reengineering discussion.
•The CEO wants a quarterly assessment of all the required metrics for their specific responsibilities. The tools used for the assessments must be Microsoft-recommended and must verify reliability, interpretability, fairness, and compliance.
•The CFO wants to identify how many interactions with the AI agents are abandoned on a given day as compared to resolved conversations. Too many abandoned sessions might indicate that Copilot Studio credits are being used inefficiently by end users.
Which two components in the custom AI agent design should the CFO evaluate in the quarterly agent analysis? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. the GPT models used for the agent
B. the average characters in a chat message
C. the agent orchestration method
D. the average session time per agent

Answer: BD

Explanation:

Technical Justification for Correct Answer: BD

Question Analysis: The CFO needs to evaluate two components of the custom AI agent design in the quarterly agent analysis to assess efficiency and ROI. The correct answer, **BD**, is justified below, along with explanations for why other options are less suitable.

Correct Answer Justification: BD

B. the average characters in a chat message:
Relevance to Efficiency and ROI: Analyzing the average characters in a chat message helps the CFO understand the complexity and brevity of interactions. Shorter messages might indicate efficiency, while excessively long messages could suggest unnecessary complexity or inefficiency in user interactions, potentially wasting Copilot Studio credits.
Alignment with CFO's Responsibility: Directly relates to evaluating the efficient use of Copilot Studio credits, as mentioned in the requirements.

D. the average session time per agent:
Relevance to Efficiency and ROI: Average session time is a key metric for assessing the efficiency of AI agent interactions. Shorter session times typically indicate more efficient resolutions, directly impacting the ROI analysis by showing how quickly users can resolve queries.
Alignment with CFO's Responsibility: Crucial for comparing estimated ROI against actual efficiencies, as shorter sessions can imply lower operational costs and higher user satisfaction.

Why Other Options are Less Suitable

A. the GPT models used for the agent:
Suitability: While important for overall agent performance, the specific GPT models used are more relevant to technical performance and user feedback (CTO's responsibility) than to the CFO's focus on ROI and efficiency metrics.
C. the agent orchestration method:
Suitability: Orchestration methods are vital for scalability and integration (more aligned with CTO or CIO responsibilities) but do not directly provide the ROI and efficiency metrics the CFO needs for quarterly analyses.

Conclusion

Options **B** and **D** are the most relevant for the CFO's quarterly analysis because they provide direct insights into the efficiency and effectiveness of the custom AI agent's interactions, closely aligning with ROI analysis and the efficient use of resources (Copilot Studio credits).

References

Microsoft Dynamics 365 & AI Integration
Microsoft Copilot Studio Documentation

Question: 32 CertyIQ

Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Overview -
Contoso, Ltd. is a high-tech manufacturing company that uses Microsoft Dynamics 365 Finance. Dynamics 365 Supply Chain Management, and Dynamics 365 Commerce for its North American operations. The company designs and develops innovative products that have many patents and proprietary technologies. The patents and engineering designs are closely guarded secrets.
Contoso executives want to integrate and adopt AI solutions to help scale the company in preparation for an anticipated period of rapid growth.
The company has multiple legal entities and Azure subscriptions that will be used in the adopted AI solutions.
Requirements -
AI Adoption -

The following executives will have specific responsibilities in the overall AI adoption:
•Chief Technology Officer (CTO): Select one Dynamics 365 Finance, Dynamics 365 Supply Chain Management or Dynamics 365 Commerce prebuilt AI agent and one custom Microsoft Copilot Studio AI agent to prioritize and deploy during the initial AI adoption phase.
•Chief Information Officer (CIO): Ensure that appropriate security labels are assigned to the data used by the AI agents.
•Chief Financial Officer (CFO): Analyze the return on investment (ROI) for the AI agents being deployed.
•Chief Information Security Officer (CISO): Discover and inventory AI resources for auditing.
•Chief Executive Officer (CEO): Ensure that all solutions adhere to industry-standard responsible AI practices.
All AI initiatives and agents will have a detailed business use case, a defined audience profile, and an estimated ROI that will compare the cost savings of the current process against the estimated costs of using the new AI solutions.
The company's research and development (R&D) department already has a custom Model Context Protocol (MCP) server that contains comprehensive product specifications and compliance data.
Prebuilt AI Agent -
The CTO has NOT yet selected which prebuilt AI agent to use in Dynamics 365 Supply Chain Management. The CTO wants to view available agent templates to identify which agent will add the most business value.
Depending on which high-priority AI agents are identified, its agent capabilities must be previewed in a discovery meeting with the relevant business operation stakeholders.
Custom AI Agent -
Contoso has identified the following custom AI agent requirements:
•The custom AI agent will use data from Dynamics 365 Supply Chain Management to answer questions for the manufacturing team as a low-code solution.
•The custom AI agent will be accessible from within Microsoft Teams.
•The custom AI agent must be designed to eventually connect to other agents that can be selected based on their description.
•The topics used in the custom AI agent will be selected based NOT on a trigger phrase, but on a description of the purpose of the query, to make the interactions more conversational.
•The custom AI agent must be able to answer questions about product specifications by using existing technologies. The product specifications are maintained by the R&D department.
•The custom AI agent must be integrated with and accessible from Dynamics 365 Supply Chain Management.
•The custom AI agent must be able to use Dynamics 365 Supply Chain Management business logic that is stored outside of the application.
Analysis, Reporting, and Troubleshooting
Contoso has identified the following analysis, reporting, and troubleshooting requirements:
•The CISO will audit all the AI solutions monthly for compliance and security.
•The CFO will analyze all the AI solutions quarterly to compare the estimated ROI against actual measured efficiencies and adoption. The CFO will use the Copilot Studio agent usage estimator to perform this analysis.
•The CISO wants to identify how much sensitive data was accessed for a given AI agent run and who accessed the data. Too much sensitive data accessed by a single user might indicate a high security risk.
•The CTO wants to track user feedback on the quality of the AI agent responses during user interactions with the agents. Consistently poor feedback will trigger an escalated reengineering discussion.
•The CEO wants a quarterly assessment of all the required metrics for their specific responsibilities. The tools used for the assessments must be Microsoft-recommended and must verify reliability, interpretability, fairness, and compliance.
•The CFO wants to identify how many interactions with the AI agents are abandoned on a given day as compared to resolved conversations. Too many abandoned sessions might indicate that Copilot Studio credits are being used inefficiently by end users.
What should you configure for the custom AI agent?

A. AI-assisted evaluators
B. classic orchestration
C. generative orchestration
D. Azure OpenAI reasoning models

Answer: C

Explanation:

Technical Justification for Correct Answer: C (Generative Orchestration)

Why C (Generative Orchestration) is the Best Option:

Conversational Interaction Requirement: The custom AI agent needs to facilitate conversational interactions based on the "description of the purpose of the query," not trigger phrases. Generative Orchestration in Microsoft Copilot Studio is designed to handle more open-ended, natural language inputs, making it suitable for this conversational requirement.
Integration with Dynamics 365 Supply Chain Management: Generative Orchestration can integrate with external data sources and applications, such as Dynamics 365 Supply Chain Management, to fetch and utilize business logic stored outside the application, aligning with the requirement for the custom AI agent to leverage this data.
Low-Code Solution and Microsoft Teams Accessibility: While not exclusively tied to generative orchestration, Microsoft's emphasis on low-code solutions and integration with Microsoft Teams for various AI and automation tools suggests that generative orchestration would be compatible with these requirements, given its part of the broader Copilot Studio ecosystem which prioritizes ease of use and integration.
Future Connectivity with Other Agents: The flexible nature of generative orchestration makes it more adaptable for future integrations with other agents based on their descriptions, supporting the company's evolving AI landscape.

Why Other Options are Less Suitable:

A. AI-assisted Evaluators:
Primarily used for assessment and evaluation tasks, not for building conversational AI agents that integrate with external business logic.
Does not align with the requirement for open conversational interactions.
B. Classic Orchestration:
Better suited for predefined, rule-based workflows rather than the open-ended, conversational interactions required.
Less flexible for integrating dynamic external data sources and business logic compared to generative orchestration.
D. Azure OpenAI Reasoning Models:
While powerful for specific AI tasks, the question emphasizes the use of Microsoft Copilot Studio for the custom AI agent, making this option less relevant.
Does not directly address the integration and conversational requirements specified in the context of Copilot Studio.

References

1. **Microsoft Copilot Studio Documentation** - https://learn.microsoft.com/en-us/copilot/studio/ (Explore sections on Generative Orchestration and Custom AI Agents)
2. **Dynamics 365 and AI Integration** - https://docs.microsoft.com/en-us/dynamics365/release-waves/ (Review AI and Machine Learning enhancements in recent release waves)

Question: 33 CertyIQ

A company has an AI agent that automates the review of customer feedback stored in a cloud database.
You plan to generate monthly reports from the agent's output to provide insights into customer sentiment and guide product development and marketing.
You need to ensure that the data ingested by the agent is clean and suitable for the intended use.
What should you do to prepare the data?

A. Ensure that the size of the database does not exceed 100 GB.
B. Translate the data into a single language.
C. Identify and address biased data.
D. Sort the database by customer last name.

When raw customer feedback is ingested to extract sentiment and drive strategic product and marketing decisions, the primary risk to data suitability and reliability is data bias (e.g., non-representative sampling, duplicate submission loops, or skewed demographic feedback).

Answer: C

Explanation:

Technical Justification for Correct Answer: C

Why C is the Best Option: Identifying and addressing biased data (Option C) is crucial for preparing the data ingested by the AI agent. Biased data can lead to skewed analysis, resulting in misleading customer sentiment insights. This, in turn, can misguide product development and marketing strategies. By addressing biases, the company ensures the AI agent's output is reliable and representative of the true customer feedback spectrum. This step is fundamental in data preprocessing for machine learning and analytics tasks, as it directly impacts the accuracy and fairness of the outcomes.

Why Other Options are Less Suitable:

A. Ensure that the size of the database does not exceed 100 GB: While managing database size is important for performance, the specified size limit (100 GB) does not directly relate to data cleanliness or suitability for analyzing customer sentiment. Modern cloud databases can handle much larger sizes efficiently, and this approach does not address the quality or bias of the data.
B. Translate the data into a single language: Although unifying the language can facilitate analysis, this option assumes a monolingual target audience or development/marketing strategy. It does not ensure data cleanliness and might even introduce errors or lose nuances in translation. Moreover, the primary concern for data suitability in this context is not linguistic uniformity but the accuracy and bias of the feedback data.
D. Sort the database by customer last name: Sorting by customer last name is an organizational task that does not contribute to ensuring the data's cleanliness, suitability, or quality for analyzing customer sentiment. It's more related to data presentation or specific querying needs rather than preprocessing for AI-driven insights.

Key Technical Points Supporting C:

Data Quality: Addressing biases ensures higher data quality, crucial for reliable AI outputs.
Fairness and Accuracy: Biased data correction leads to fairer, more accurate sentiment analysis.
Strategic Impact: Unbiased insights better inform product and marketing strategies.

References

Microsoft Learn - Data Preparation for Machine Learning: Emphasizes the importance of data cleaning and preprocessing, including addressing biases, for reliable ML model outputs.
Microsoft Azure - Data Quality and Integrity: Discusses strategies for ensuring data integrity, which includes identifying and correcting biased data for accurate analytics.

Question: 34

CertyIQ

A company has an Azure environment that supports multiple business units.
The company plans to implement an AI solution that will perform sentiment analysis on customer product reviews.
You need to evaluate the potential cost of the solution to support return on AI investment (ROAI) analysis.
What should you use?

The Azure pricing calculator is Microsoft's dedicated pre-deployment financial modeling tool. It allows architects to configure target service tiers, estimate transaction/API call volumes, model prompt/completion token consumption, and project monthly and annual operating run-rates required to substantiate the cost baseline of an ROAI analysis.

A. Total Cost of Ownership (TCO) Calculator
B. Azure Reservations
C. Azure pricing calculator
D. Azure Monitor

Answer: C

Explanation:

Technical Justification for Choosing Azure Pricing Calculator (C)

Why C (Azure Pricing Calculator) is the best option: The Azure Pricing Calculator is the most suitable tool for evaluating the potential cost of an AI solution in Azure, specifically for sentiment analysis on customer product reviews. This tool allows for a granular,Estimated cost analysis by:

Specifying the exact Azure services required for the AI solution (e.g., Azure Machine Learning, Cognitive Services for Text Analysis, Storage for reviews data).
Configuring the anticipated usage patterns (e.g., compute hours, number of text analysis API calls, storage size).
Factoring in the company's existing Azure environment and potential discounts (e.g., reserved instances, enterprise agreements).
Providing a detailed, projected cost breakdown to inform the Return on AI Investment (ROAI) analysis.

Why other options are less suitable:

A. Total Cost of Ownership (TCO) Calculator: While useful for comparing cloud vs. on-premises costs and understanding long-term ownership costs, the TCO Calculator is not as precise for estimating the specific, upfront costs of a new Azure-based AI solution as the Azure Pricing Calculator.
B. Azure Reservations: This is a cost-saving mechanism for committed resource usage over 1 or 3 years. While it can reduce costs, it does not serve as a tool for estimating the initial costs of implementing an AI solution.
D. Azure Monitor: Primarily used for monitoring and optimizing resource usage post-deployment, Azure Monitor does not provide pre-implementation cost estimation capabilities for new solutions.

Correct Answer Justification Summary:

Precision for Specific Service Costs: C (Azure Pricing Calculator) offers the most accurate cost estimation for the chosen AI services.
Relevance to Pre-Implementation Analysis: C is designed for planning and estimation, aligning perfectly with ROAI analysis needs.
Irrelevance of Other Tools for This Specific Task: A, B, and D serve different purposes not aligned with the question's requirements.

References:

Azure Pricing Calculator - Official Microsoft resource for estimating Azure service costs.
Understanding Azure Pricing for AI Workloads - Microsoft Documentation on optimizing AI workloads costs in Azure.

Question: 35

CertyIQ

A company has an Azure environment that supports multiple business units.
The company plans to implement an AI solution that will perform sentiment analysis on customer product reviews.
You need to evaluate the potential cost of the solution to support return on AI investment (ROAI) analysis.
What should you use?

A. Azure savings plans
B. Microsoft Fabric SKU Estimator
C. Cost Management + Billing

D. Azure Monitor

Answer: C

Explanation:

Technical Justification for Choosing C: Cost Management + Billing

Why C is the Correct Option:

Direct Relevance to Cost Evaluation: Cost Management + Billing is a Azure-native service specifically designed for tracking, analyzing, and optimizing cloud expenditures. It provides detailed cost breakdowns by resource, tags (e.g., by business unit or project), and subscription, making it ideal for evaluating the potential cost of a new AI solution.
Support for ROAI Analysis: To perform a return on AI investment (ROAI) analysis, accurate cost measurement of the AI solution is crucial. Cost Management + Billing offers the granularity needed to isolate costs associated with the sentiment analysis AI solution, facilitating a precise ROAI calculation.
Integration with Azure Services: Since the AI solution will be implemented on Azure, Cost Management + Billing seamlessly integrates with other Azure services (e.g., Azure Cognitive Services for AI functionalities), ensuring comprehensive and accurate cost aggregation.
Why Other Options are Less Suitable:

A. Azure Savings Plans:

Purpose: Optimizing costs for consistent, predictable workloads by committing to spend a certain amount over a year.
Limitation: Not designed for evaluating or estimating costs of new, unspecified workloads like a yet-to-be-implemented AI solution.
B. Microsoft Fabric SKU Estimator:

Note: There's no widely recognized or documented "Microsoft Fabric SKU Estimator" as of my last update. Assuming it's a misnomer or less relevant tool, it's unclear how it would directly apply to estimating costs for Azure-based AI solutions without more context.
General Limitation (if assuming a similar tool): SKU estimators are helpful for specific product licensing costs but may not provide a holistic view of cloud service costs for an AI solution.
D. Azure Monitor:

Purpose: Monitoring and troubleshooting application performance, resource utilization, and logging.
Limitation: While Azure Monitor can provide insights into resource utilization, it is not primarily designed for financial cost analysis or estimation, lacking the direct cost management features needed for ROAI analysis.
Conclusion

Cost Management + Billing is the most appropriate choice for evaluating the potential cost of the AI sentiment analysis solution due to its direct relevance to cost analysis, support for ROAI calculations, and seamless integration with Azure services.

References

1. **Azure Cost Management + Billing Documentation**: https://docs.microsoft.com/en-us/azure/cost-management-billing/
2. **Understanding Azure Pricing for AI Workloads**: https://azure.microsoft.com/en-us/pricing/guide/#ai-machine-learning

Question: 36 CertyIQ

A company has an Azure environment that supports multiple business units.
The company plans to implement an AI solution that will perform sentiment analysis on customer product reviews.
You need to evaluate the potential cost of the solution to support return on AI investment (ROAI) analysis.
What should you use?

A. Azure Reservations
B. Microsoft Fabric SKU Estimator
C. Anomaly Detection in Azure Cost Management
D. Azure pricing calculator

Answer: D

Explanation:

Technical Justification for Choosing D: Azure Pricing Calculator

To evaluate the potential cost of an AI solution for sentiment analysis on customer product reviews in an Azure environment, the most suitable option is **D: Azure Pricing Calculator**. Here's why:

Accuracy and Directness for New Deployments: The Azure Pricing Calculator is the most direct tool for estimating costs for new Azure deployments, including AI solutions. It allows for the specification of exact services (e.g., Azure Cognitive Services for Sentiment Analysis, underlying compute resources, storage for reviews, etc.), providing a detailed, itemized cost estimate. This is crucial for a ROAI analysis where precision in cost projection is key.

Why Other Options are Less Suitable:

A. Azure Reservations: While Azure Reservations can offer cost savings for committed resource usage over a year or three years, this option is more about optimizing existing costs rather than estimating the upfront and total cost of a new solution. It doesn't provide a baseline cost estimate for a new AI deployment.
B. Microsoft Fabric SKU Estimator: This tool is primarily used for estimating costs related to migrating applications to Microsoft Fabric (now more commonly referred to in the context of Azure Service Fabric or generally for estimating costs of scalable, microservice-based architectures). It's less directly applicable for quickly estimating the cost of a specific AI service deployment like sentiment analysis.
C. Anomaly Detection in Azure Cost Management: Anomaly Detection is useful for identifying unexpected spending patterns in an **already deployed** environment. It does not serve to estimate the potential cost of a solution prior to its deployment, making it unsuitable for pre-implementation ROAI analysis.

Correct Choice Justification Summary (D)

Direct Application: Suitable for estimating new deployments.
Precision: Allows for service-specific cost estimation (critical for AI solutions with varied service dependencies).
Relevance to ROAI: Provides the necessary cost data for return on investment analysis prior to deployment.

References

1. **Azure Pricing Calculator** - For detailed cost estimation of Azure services: https://azure.com/pricing/calculator/
2. **Azure Cost Management + Billing Documentation** - For understanding cost analysis and optimization tools, including reservations and anomaly detection: https://learn.microsoft.com/en-us/azure/cost-management-billing/

Question: 37 CertyIQ

A company has an Azure environment that supports multiple business units.
The company plans to implement an AI solution that will perform sentiment analysis on customer product reviews.
You need to evaluate the potential cost of the solution to support return on AI investment (ROAI) analysis.
What should you use?

A. Anomaly Detection in Azure Cost Management
B. Azure pricing calculator
C. Azure savings plans
D. Microsoft Fabric SKU Estimator

Answer: B

Explanation:

Technical Justification for Choosing Option B: Azure Pricing Calculator

Why Option B (Azure Pricing Calculator) is the Best Choice:

Direct Cost Estimation for AI Solution: The Azure Pricing Calculator is specifically designed to provide detailed, estimated costs for Azure services based on the user's input (e.g., resource types, quantities, locations, and terms). For an AI solution like sentiment analysis on customer product reviews, this tool allows for the estimation of costs associated with relevant services such as Azure Cognitive Services (e.g., Text Analytics for sentiment analysis), Storage, Compute (if custom models are trained/deployed), and possibly Azure Databricks or Synapse Analytics for data processing and insights.
Supports ROAI Analysis: By accurately estimating the upfront and ongoing costs of the AI solution, the calculator directly supports the Return on AI Investment (ROAI) analysis, enabling the company to forecast potential returns against the estimated expenses.
Precision and Customization: It offers the granularity needed to account for the specific resources the sentiment analysis solution will utilize, making the cost estimate more reliable for financial planning and ROI calculations.
Why Other Options are Less Suitable:

A. Anomaly Detection in Azure Cost Management:

Purpose: Identifies unusual spending patterns after costs are incurred, not for pre-implementation cost estimation.
Limitation for This Scenario: Does not help in estimating costs before deploying the AI solution, which is the requirement for ROAI analysis.
C. Azure Savings Plans:

Purpose: Provides a discounted rate for committed usage of compute resources (e.g., Azure Compute, Azure Cosmos DB).
Limitation for This Scenario: While useful for optimizing costs after deployment, it doesn't serve the pre-implementation cost estimation need for ROAI analysis.
D. Microsoft Fabric SKU Estimator:

Note: As of my last update, there's no widely recognized or publicly available "Microsoft Fabric SKU Estimator" directly linked to Azure cost estimation for AI services in the context provided. If referring to a tool for estimating costs of a specific "fabric" or platform within Microsoft's ecosystem, it's either too niche or not the standard tool for Azure service cost estimation.

Limitation for This Scenario: Assuming it's not the standard Azure pricing tool, its applicability and accuracy for estimating costs of an Azure-based AI solution would be questionable without clear, public documentation linking it to such use cases.

Conclusion

Given the need for a pre-implementation, detailed cost estimation to support ROAI analysis for an Azure-based AI sentiment analysis solution, the **Azure Pricing Calculator (Option B)** is the most appropriate choice due to its direct relevance, customization capabilities, and the precise cost estimation it provides for Azure services.

References

1. **Azure Pricing Calculator** - Official Microsoft Azure Website
2. **Azure Cost Management + Billing Documentation** - Official Microsoft Learn Documentation (covers cost estimation and management best practices)

Question: 38

CertyIQ

A company has an Azure environment that supports multiple business units.
The company plans to implement an AI solution that will perform sentiment analysis on customer product reviews.
You need to evaluate the potential cost of the solution to support return on AI investment (ROAI) analysis.
What should you use?

A. Azure savings plans
B. Anomaly Detection in Azure Cost Management
C. Azure Monitor
D. Cost Management + Billing

Answer: D

Explanation:

Technical Justification for Choosing Option D

To evaluate the potential cost of an AI solution for sentiment analysis on customer product reviews in an Azure environment supporting multiple business units, and to support a return on AI investment (ROAI) analysis, the most appropriate choice is **D. Cost Management + Billing**. Here's why:

Why D is the Best Choice:

Granular Cost Tracking: Azure Cost Management + Billing provides detailed, granular cost tracking capabilities, essential for attributing expenses to specific resources or departments (in this case, the AI solution across multiple business units).
Forecasting and Analysis: It offers forecasting tools and detailed analytics, crucial for predicting and evaluating the potential cost of the AI solution, thereby supporting ROAI analysis.
Multi-Subscription and Departmental Views: Given the company's multiple business units, this feature is vital for aggregating or isolating costs by unit, facilitating a comprehensive cost-benefit analysis of the AI investment across the organization.
Why Other Options are Less Suitable:

A. Azure Savings Plans:

Purpose: Optimizes costs for committed usage of Azure resources (e.g., Compute, Storage).

Limitation: While cost-effective for predictable workloads, it doesn't provide the analytical insights needed for evaluating the AI solution's potential costs or supporting ROAI analysis.
B. Anomaly Detection in Azure Cost Management:

Purpose: Identifies unexpected spending patterns.
Limitation: Useful for monitoring but not suited for upfront cost evaluation or detailed financial forecasting of a new solution.
C. Azure Monitor:

Purpose: Provides operational insights into resource performance and health.
Limitation: Lacks the financial management and billing insights necessary for assessing the economic viability of the AI solution.

Conclusion: Azure Cost Management + Billing is the most suitable for evaluating potential costs and supporting ROAI analysis due to its comprehensive cost tracking, forecasting, and multi-unit management capabilities.

References

1. **Azure Cost Management + Billing Documentation**
   https://docs.microsoft.com/en-us/azure/cost-management-billing/
2. **Azure Pricing and Cost Estimation Tools**
   https://azure.microsoft.com/en-us/pricing/estimator/

Question: 39 CertyIQ

A company has an AI agent that automates the review of customer feedback stored in a cloud database.
You plan to generate monthly reports from the agent's output to provide insights into customer sentiment and guide product development and marketing.
You need to ensure that the data ingested by the agent is clean and suitable for the intended use.
What should you do to prepare the data?

A. Identify and address missing or inconsistent data.
B. Sort the database by customer last name.
C. Ensure that the size of the database does not exceed 100 GB.
D. Verify data privacy and security compliance.

Answer: A

Explanation:

Technical Justification for Correct Option: A

Why A is the Best Choice: Identifying and addressing missing or inconsistent data is crucial for preparing the data for the AI agent to accurately analyze customer feedback and generate reliable monthly reports. This step ensures data quality, which directly impacts the validity of the insights derived for product development and marketing strategies. By correcting inconsistencies and filling in missing values (or making informed decisions about how to handle them), the AI agent's output will be more trustworthy and useful.

Why Other Options are Less Suitable:

B. Sort the database by customer last name:

Sorting by customer last name does not contribute to data cleanliness or suitability for analysis of sentiment trends or product development insights. It's an organizational task that doesn't address data quality issues. This action is more relevant for specific query optimizations or certain types of reporting but is not a preparatory step for AI-driven sentiment analysis.

C. Ensure the database size does not exceed 100 GB:

While managing database size is important for performance and cost, the size limit (in this case, 100 GB) does not directly impact the cleanliness or suitability of the data for the intended AI analysis.
This is a storage/scaling concern rather than a data preparation step for analysis.

D. Verify data privacy and security compliance:

Although verifying data privacy and security compliance is critically important, it pertains to the ethical and legal handling of data rather than preparing the data for analysis in terms of cleanliness and suitability for the AI agent's intended use.
This step should indeed be taken but is orthogonal to the question's focus on data preparation for analysis.

Correct Answer Justification Summary: Option A directly addresses the need for clean and suitable data by identifying and resolving inherent data quality issues, making it the most appropriate choice for preparing the data for the AI agent's analysis.

References

1. **Microsoft Documentation: Data Quality in Azure** - https://docs.microsoft.com/en-us/azure/architecture/data-warehouse/data-quality
2. **Microsoft Learn: Prepare Data for Machine Learning** - https://learn.microsoft.com/en-us/learn/modules/prepare-data-machine-learning/ (Module focuses on data preparation, including handling missing and inconsistent data)

Question: 40 CertyIQ

A company uses multiple Microsoft Copilot Studio agents across different channels.
You need to recommend a monitoring solution that provides comprehensive telemetry data and performance insights for the agents.
What should you include in the recommendation?

A. Azure DevOps
B. Microsoft Power BI
C. Microsoft Dynamics 365 Customer Voice
D. Application Insights

Microsoft Copilot Studio includes native integration with Azure Application Insights (part of Azure Monitor). Connecting an agent to an Application Insights connection string allows organizations to stream real-time diagnostic telemetry,

User message traffic and bot responses across all configured channels.
Intent recognition confidence scores and fallback triggers.
Latency and execution duration of downstream actions, flows, and connector calls.
Custom telemetry events (customEvents) and structured conversation traces for deep Kusto Query Language (KQL) querying and alerting.

Answer: D

Explanation:

Technical Justification for Recommended Monitoring Solution

Recommended Option: D. Application Insights

Comprehensive Telemetry Data Collection: Application Insights is designed to collect comprehensive telemetry data from applications and services, including custom events, metrics, and exceptions. This capability is crucial for monitoring Microsoft Copilot Studio agents across different channels, providing a

unified view of their operational health.
Performance Insights: It offers out-of-the-box performance monitoring, including response times, dependency calls, and server health metrics, which are essential for evaluating the efficiency of Copilot Studio agents.
Integration with Microsoft Ecosystem: Seamless integration with other Microsoft services (e.g., Azure, Visual Studio) facilitates streamlined monitoring and analysis workflows, aligning with the company's existing Microsoft-centric technology stack.
Agent-Specific Monitoring: Application Insights can be configured to monitor specific agents or services (like those from Copilot Studio) across various channels (web, mobile, desktop), making it highly suitable for this scenario.

Why Other Options are Less Suitable:

A. Azure DevOps:

Primarily focused on development lifecycle, continuous integration, and deployment, rather than operational monitoring and telemetry of deployed agents.
While it has some monitoring capabilities through its integration with Application Insights, on its own, it's not tailored for the required depth of agent performance insights.
B. Microsoft Power BI:

A business analytics service, excellent for data visualization and reporting.
Not designed for real-time telemetry data collection or providing performance insights at the agent level; better suited for analyzing dataalready collected by another tool.
C. Microsoft Dynamics 365 Customer Voice:

Focused on customer feedback and survey management.
Completely unrelated to monitoring agent performance or collecting technical telemetry data.

References

Azure Monitor & Application Insights Documentation - Detailed overview of Application Insights capabilities.
Microsoft Copilot Studio Integration Scenarios - Insights into integrating Copilot Studio with other Microsoft services for comprehensive management.

Question: 41 CertyIQ

A company uses multiple Microsoft Copilot Studio agents across different channels.
You need to recommend a monitoring solution that provides comprehensive telemetry data and performance insights for the agents.
What should you include in the recommendation?

A. the Analytics tab in Copilot Studio
B. Microsoft Power BI
C. Microsoft Dynamics 365 Customer Voice
D. Microsoft Purview

Answer: B

Explanation:

Technical Justification for Recommended Monitoring Solution

Recommended Option: B. Microsoft Power BI

Microsoft Power BI is the most suitable choice for providing comprehensive telemetry data and performance insights for multiple Microsoft Copilot Studio agents across different channels. Here's why:

Scalability and Integration: Power BI seamlessly integrates with various Microsoft services, making it an ideal choice for aggregating data from multiple Copilot Studio agents. Its scalability ensures it can handle the volume of telemetry data generated across different channels.
Customizable Dashboards: Power BI allows for the creation of customized dashboards, enabling tailored performance insights that can be easily understood by both technical and non-technical stakeholders.
Advanced Analytics: It offers advanced analytics capabilities, facilitating deep dives into agent performance, identifying bottlenecks, and informing data-driven decisions to enhance agent efficacy.

Why Other Options are Less Suitable:

A. The Analytics Tab in Copilot Studio:
Limited Scope: Primarily designed for individual agent analysis, not optimized for cross-agent, multi-channel insights.
Feature Set: Likely more basic compared to dedicated analytics platforms, limiting in-depth performance analysis.
C. Microsoft Dynamics 365 Customer Voice:
Primary Function:Focused on customer feedback and survey analysis, not designed for agent performance telemetry.
Misalignment: Does not align with the requirement for technical performance insights of Copilot Studio agents.
D. Microsoft Purview:
Primary Function:Centers around data governance, security, and compliance, not performance analytics of operational systems like Copilot Studio agents.
Data Analysis Scope: Not tailored for the type of real-time performance and telemetry data required for this scenario.

References

Microsoft Power BI Documentation - Integrating with Microsoft Services
Microsoft Copilot Studio Overview (for context on agent functionality)

Question: 42 CertyIQ

A company uses multiple Microsoft Copilot Studio agents across different channels.
You need to recommend a monitoring solution that provides comprehensive telemetry data and performance insights for the agents.
What should you include in the recommendation?

A. Azure Advisor
B. Microsoft Dynamics 365 Customer Voice
C. Application Insights
D. Log Analytics

Answer: C

Explanation:

Technical Justification for Recommended Monitoring Solution

Recommended Option: C. Application Insights

Comprehensive Telemetry Data: Application Insights is designed to collect comprehensive telemetry data, including usage patterns, performance metrics, and error tracking, which is crucial for monitoring multiple Microsoft Copilot Studio agents across different channels.
Performance Insights: It provides detailed performance insights, enabling the identification of bottlenecks, slow queries, and dependencies issues, all of which are vital for optimizing the performance of Copilot Studio agents.
Integration Suitability: As part of the Azure suite, Application Insights seamlessly integrates with other Azure services, making it an ideal choice for monitoring agents within a Microsoft ecosystem.
Agent-Specific Monitoring: It supports custom telemetry, allowing for the tracking of specific agent activities, errors, and performance metrics tailored to Copilot Studio's requirements.

Why Other Options are Less Suitable:

A. Azure Advisor:

Focus: Azure Advisor primarily provides optimization recommendations for Azure resources, not real-time telemetry or performance insights for specific applications like Copilot Studio agents.
Limitation: Lacks the depth of application-level monitoring required for the scenario.
B. Microsoft Dynamics 365 Customer Voice:

Purpose: Designed for collecting and analyzing customer feedback, not for monitoring application performance or telemetry data.
Mismatch: Completely unrelated to the technical monitoring requirements of Copilot Studio agents.
D. Log Analytics:

Capability: While Log Analytics can collect and analyze logs, it is more suited for infrastructure and service logging rather than providing the comprehensive application performance insights that Application Insights offers.
Shortcoming: Less integrated with the specific needs of monitoring application-level telemetry and performance in a unified manner compared to Application Insights.

References

Application Insights Documentation: For detailed capabilities and setup.
Microsoft Copilot Studio Integration with Azure Services: Although not directly linking to Application Insights, this resource contextualizes the integration of Microsoft services with Azure, implying the suitability of Application Insights for Microsoft products. **Note:** Direct documentation on Copilot Studio and Application Insights might not be publicly available due to the specificity of the question or the evolving nature of the products. Always refer to the latest Microsoft resources for integration details.

Question: 43 CertyIQ

A company uses multiple Microsoft Copilot Studio agents across different channels.
You need to recommend a monitoring solution that provides comprehensive telemetry data and performance insights for the agents.
What should you include in the recommendation?

A. the Analytics tab in Copilot Studio
B. Azure Advisor
C. Log Analytics
D. Microsoft Purview

Answer: A

Explanation:

Technical Justification for Recommendation

Recommended Option: A. the Analytics tab in Copilot Studio

Direct Integration: The Analytics tab in Copilot Studio is natively integrated with Microsoft Copilot Studio agents, ensuring seamless and comprehensive telemetry data collection across all agents, regardless of the channel. This integration eliminates the need for additional setup or configuration, streamlining the monitoring process.
Agent-Specific Insights: Designed specifically for Copilot Studio, the Analytics tab provides agent-level performance insights, tailored metrics, and detailed analytics that are crucial for optimizing agent behavior, response times, and overall efficiency. These insights are essential for fine-tuning agent interactions to improve user experiences.
Real-Time Monitoring Capabilities: It likely offers real-time or near-real-time monitoring, enabling prompt identification and resolution of issues, which is critical for maintaining high availability and user satisfaction across multiple channels.
Centralized Agent Management: Since the company uses multiple agents across different channels, the Analytics tab within the same platform simplifies centralized monitoring and management, reducing operational complexity.

Why Other Options are Less Suitable:

B. Azure Advisor:

Primary Function: Azure Advisor is primarily focused on providing recommendations for optimizing Azure resources in terms of cost, security, reliability, and performance. While valuable, its scope does not specifically target the monitoring of Copilot Studio agents.
Lack of Direct Agent Insights: Not designed for agent-specific telemetry and performance monitoring, making it less relevant for this particular requirement.
C. Log Analytics:

Broad Capability: Part of Azure Monitor, Log Analytics is powerful for logging and analytics across various Azure services and custom logs.
Additional Setup Required: Would need specific configuration to capture and analyze Copilot Studio agent data, adding complexity compared to a native solution.
Not Agent-Centric: Provides a broader view, which might not offer the depth of agent-specific performance insights as the Analytics tab in Copilot Studio.
D. Microsoft Purview:

Primary Focus: Microsoft Purview is centered around data governance, compliance, and security across the enterprise data estate.
No Direct Agent Monitoring Capability: Completely unrelated to the real-time monitoring and performance insights needs of Copilot Studio agents.

Conclusion:Given the native integration, agent-specific insights, and the streamlined monitoring experience, **A. the Analytics tab in Copilot Studio** is the most appropriate recommendation for comprehensive telemetry data and performance insights across multiple Microsoft Copilot Studio agents.

References

Microsoft Copilot Studio Documentation Note: As of my last update, specific Copilot Studio links might vary; this link serves as a placeholder for Microsoft's AI toolkit documentation, which may encompass Copilot

Studio functionalities.
Azure Monitor and Log Analytics For understanding the broader monitoring capabilities in Azure, though not the recommended solution for this specific question.

Question: 44 CertyIQ

A company has an AI agent that automates the review of customer feedback stored in a cloud database.
You plan to generate monthly reports from the agent's output to provide insights into customer sentiment and guide product development and marketing.
You need to ensure that the data ingested by the agent is clean and suitable for the intended use.
What should you do to prepare the data?

A. Identify and address biased data.
B. Ensure that the size of the database does not exceed 100 GB.
C. Verify data privacy and security compliance.
D. Translate the data into a single language.

Answer: A

Explanation:

Technical Justification for Correct Answer: A

Why A is the Best Option: To ensure the data ingested by the AI agent is clean and suitable for generating accurate and reliable monthly reports on customer sentiment, identifying and addressing biased data is crucial. Biased data can lead to skewed insights, misinforming product development and marketing strategies. By addressing biases, the company can trust the agent's output to reflect genuine customer sentiments. This step is fundamental in data preparation for AI-driven analytics, as it directly impacts the accuracy and fairness of the insights derived.

Why Other Options are Less Suitable:

B. Ensure the size of the database does not exceed 100 GB: While managing database size is important for performance and cost optimization, the specified limit (100 GB) does not directly relate to data cleanliness or suitability for analysis. Modern cloud databases can efficiently handle larger sizes, and the focus should be on data quality rather than arbitrary size limits for this specific requirement.

C. Verify data privacy and security compliance: Although critical for regulatory adherence and trust, verifying data privacy and security compliance ensures the data is handled responsibly but does not directly impact the cleanliness or suitability of the data for analysis purposes. This is a prerequisite but does not address the quality of the data for the intended use case.

D. Translate the data into a single language: While useful for uniform analysis, translating all data into a single language does not inherently make the data cleaner or more suitable for analyzing customer sentiment. It's an operational decision based on the analysis requirements and does not address the fundamental issue of data quality and bias.

Key Technical Points Supporting A:

Data Bias Impact: Biased data skews AI model outputs, leading to potentially harmful or incorrect insights.
Preprocessing Necessity: Identifying and addressing biases is a foundational step in data preprocessing for AI applications.
Quality Over Quantity/Size: Ensuring data quality (through bias correction) is more critical than focusing solely on database size for analysis purposes.

References:

1. **Microsoft Learn - Ensure data quality for machine learning**https://learn.microsoft.com/en-us/azure/machine-learning/service/concept-data-quality
2. **MIT CSAIL - Data Bias in AI Systems**https://people.csail.mit.edu/conniehuang/papers/bias_in_ai.pdf (Note: Links to external research papers can change. If the link is broken, search for "Data Bias in AI Systems MIT CSAIL" for relevant content)

Question: 45 CertyIQ

Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Background -
Fabrikam, Inc., is a global consumer goods company that is undergoing a digital transformation initiative to migrate its entire infrastructure to the Microsoft cloud. As a key element of this cloud migration, the company will implement Microsoft Dynamics 365 Sales, moving away from the current on-premises proprietary technologies used by its business-to-business (B2B) sales team.
As part of the cloud migration, Fabrikam will adopt an AI-first approach to its business solutions and implement AI solutions, wherever possible, to streamline operations.
Problem Statements -
Fabrikam's infrastructure currently relies on various on-premises systems that require sales executives to use corporate computers with physical keyboards to access business information during customer interactions. Mobile phones cannot be used for these purposes, as the systems depend on keyboard input. As a result, the sales executives spend a lot of time using keyboards to search for data on several disparate systems and file servers, rather than focusing on the customers. This affects the customer experience.
Fabrikam stakeholders are concerned that users will be hesitant to adopt AI. If the AI initiatives are NOT adopted, cost savings will never be realized. Additionally, funding for future AI initiatives will depend on demonstrating an increase in AI adoption month over month. As the AI agent initiative for the sales team will be the first for Fabrikam, the rapid adoption of the agent is a high priority.
Planned Initiatives -
General -
Fabrikam management has prioritized AI-driven projects to improve efficiency, customer engagement, and responsible AI adoption. The current application infrastructure is on-premises and must be migrated to the cloud to support the adoption of these technologies.
Infrastructure Migration -
Fabrikam plans to migrate from its current on-premises infrastructure to a completely cloud-based topology; this will include user authentication, the security framework, and, primarily, the adoption of the services by end users. All the data from the different systems will be consolidated into a single data source - a common data model that will use a Microsoft Dataverse environment as a single source of truth (SSOT) for the sales team.
Sales Cycle Enablement -
To achieve the company's objectives, Fabrikam intends to implement the following strategies to enhance the sales cycle:
Use low-code development to create a single AI agent that has Dataverse as its core component.
Ensure that sales managers can access unanswered correspondence from prospects and intervene as appropriate.
Replace the previous proprietary software with Dynamics 365 Sales to track sales cycles and customer

interactions.
Have the sales executives use Dynamics 365 Sales to track interactions for open opportunities and send follow-up communications to prospects.
Have the sales executives use handsfree headsets to interact with an AI agent when they have questions about internal policies or customer data.
Requirements -
Infrastructure Migration -
Fabrikam has identified the following infrastructure migration requirements:
Azure must be used for all future infrastructure workloads.
The company must follow Microsoft-recommended methodologies for infrastructure migration to the cloud.
Any created AI agents must have their return on investment (ROI) calculated to ensure that the solution will save the company money.
Sales Cycle Enablement -
Fabrikam has identified the following requirements for sales cycle enablement:
The final AI agent must follow Microsoft recommendations for a conversational user experience.
A designated checklist must be reviewed to ensure that the AI agent follows Microsoft deployment recommendations for a compliant solution.
Detailed telemetry must be logged for the first created AI agent to help troubleshoot and optimize the agent during the initial AI agent adoption process.
Unexpected AI agent actions must end in an escalation to a live representative. For example, a sales executive must be rerouted to a representative if the agent cannot answer a question after two failed attempts.
The return on investment (ROI) of switching from the current process to the future process is required for stakeholder sign off.
The sales team must use Dynamics 365 Sales to correspond with prospects more quickly and efficiently than currently.
Sales managers must report on the adoption of the AI agent to key Fabrikam stakeholders on a monthly basis.
Any sensitive information, such as user IDs and names, shared via the AI agent must be tracked for future auditing.
Which template should you use for the AI agent to meet the requirements for the sales executives?

A. IT Helpdesk in Microsoft Copilot Studio
B. AI agents in Microsoft Foundry
C. Voice in Microsoft Copilot Studio
D. AI chat in Microsoft Foundry

Answer: C

Explanation:

Why “Voice in Microsoft Copilot Studio” is the optimal template

The requirement specifies a **hands-free, conversational AI agent** that sales executives can interact with while on the move; the **Voice** capability in Copilot Studio is purpose-built for speech-first experiences, integrating directly with Dynamics365 Sales and supporting real-time voice input/output.
It provides **built-in telemetry** and **escalation pathways** (e.g., automatically routing to a live rep after two failed understandings), satisfying the compliance checklist and troubleshooting needs outlined for the first AI agent.
Because it is **low-code**, the solution can be configured to use **Dataverse as the single source of truth**, ensuring data consistency and enabling ROI calculations through native analytics.
The platform is fully **Microsoft-recommended** for cloud-native deployments, aligning with Fabrikam’s mandate to run all infrastructure on Azure and to follow Microsoft migration best practices.
Other options lack the necessary voice interaction model: the IT Helpdesk template is scoped for support ticket routing, not conversational sales assistance; AI agents in Microsoft Foundry (Azure Machine Learning) are primarily batch-oriented and do not deliver a sanctioned conversational UI; AI chat in Foundry is text-only and cannot meet the headset-based, speech-driven workflow.

Why the alternatives are less suitable

A – IT Helpdesk in Microsoft Copilot Studio – Designed for ticket creation and resolution, it does not expose

a voice-first channel and cannot be leveraged for real-time sales conversations.
B – AI agents in Microsoft Foundry – While capable of custom AI models, Foundry's agent framework is not pre-wired for voice interaction, lacks the out-of-the-box telemetry and escalation controls required, and introduces extra integration effort.
D – AI chat in Microsoft Foundry – Provides a text-based chat experience, which conflicts with the "hands-free headset" requirement and does not meet the conversational UX standards stipulated for the sales team.

References

Voice capabilities in Microsoft Copilot Studio
Building voice-enabled agents with Dynamics365 Sales

Question: 46 CertyIQ

Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Background -
Fabrikam, Inc., is a global consumer goods company that is undergoing a digital transformation initiative to migrate its entire infrastructure to the Microsoft cloud. As a key element of this cloud migration, the company will implement Microsoft Dynamics 365 Sales, moving away from the current on-premises proprietary technologies used by its business-to-business (B2B) sales team.
As part of the cloud migration, Fabrikam will adopt an AI-first approach to its business solutions and implement AI solutions, wherever possible, to streamline operations.
Problem Statements -
Fabrikam's infrastructure currently relies on various on-premises systems that require sales executives to use corporate computers with physical keyboards to access business information during customer interactions. Mobile phones cannot be used for these purposes, as the systems depend on keyboard input. As a result, the sales executives spend a lot of time using keyboards to search for data on several disparate systems and file servers, rather than focusing on the customers. This affects the customer experience.
Fabrikam stakeholders are concerned that users will be hesitant to adopt AI. If the AI initiatives are NOT adopted, cost savings will never be realized. Additionally, funding for future AI initiatives will depend on demonstrating an increase in AI adoption month over month. As the AI agent initiative for the sales team will be the first for Fabrikam, the rapid adoption of the agent is a high priority.
Planned Initiatives -
General -
Fabrikam management has prioritized AI-driven projects to improve efficiency, customer engagement, and responsible AI adoption. The current application infrastructure is on-premises and must be migrated to the cloud to support the adoption of these technologies.
Infrastructure Migration -
Fabrikam plans to migrate from its current on-premises infrastructure to a completely cloud-based topology; this will include user authentication, the security framework, and, primarily, the adoption of the services by end users. All the data from the different systems will be consolidated into a single data source - a common data model that will use a Microsoft Dataverse environment as a single source of truth (SSOT) for the sales team.
Sales Cycle Enablement -
To achieve the company's objectives, Fabrikam intends to implement the following strategies to enhance the sales

cycle:
Use low-code development to create a single AI agent that has Dataverse as its core component.
Ensure that sales managers can access unanswered correspondence from prospects and intervene as appropriate.
Replace the previous proprietary software with Dynamics 365 Sales to track sales cycles and customer interactions.
Have the sales executives use Dynamics 365 Sales to track interactions for open opportunities and send follow-up communications to prospects.
Have the sales executives use handsfree headsets to interact with an AI agent when they have questions about internal policies or customer data.
Requirements -
Infrastructure Migration -
Fabrikam has identified the following infrastructure migration requirements:
Azure must be used for all future infrastructure workloads.
The company must follow Microsoft-recommended methodologies for infrastructure migration to the cloud.
Any created AI agents must have their return on investment (ROI) calculated to ensure that the solution will save the company money.
Sales Cycle Enablement -
Fabrikam has identified the following requirements for sales cycle enablement:
The final AI agent must follow Microsoft recommendations for a conversational user experience.
A designated checklist must be reviewed to ensure that the AI agent follows Microsoft deployment recommendations for a compliant solution.
Detailed telemetry must be logged for the first created AI agent to help troubleshoot and optimize the agent during the initial AI agent adoption process.
Unexpected AI agent actions must end in an escalation to a live representative. For example, a sales executive must be rerouted to a representative if the agent cannot answer a question after two failed attempts.
The return on investment (ROI) of switching from the current process to the future process is required for stakeholder sign off.
The sales team must use Dynamics 365 Sales to correspond with prospects more quickly and efficiently than currently.
Sales managers must report on the adoption of the AI agent to key Fabrikam stakeholders on a monthly basis.
Any sensitive information, such as user IDs and names, shared via the AI agent must be tracked for future auditing.
Which tool should you use for the prospect communication requirements in Dynamics 365 Sales?

A. Azure AI Search
B. Copilot email assist
C. the Voice template Microsoft Copilot Studio
D. Deep Research in Microsoft Foundry Agent Service

Answer: B

Explanation:

Justification for selectingOptionB –Copilot email assist

Direct prospect-communication capability –Copilot email assist is built into Dynamics365 Sales and automatically drafts, suggests, and personalizes email replies for sales representatives. This matches the requirement to “correspond with prospects more quickly and efficiently.”
AI-driven efficiency –The feature leverages generative AI to produce context-aware email content, reducing manual typing and accelerating the outreach cycle, which is essential for demonstrating rapid AI adoption and ROI.
Seamless integration with Dynamics365 –All email activity remains within the sales hub, preserving the sales pipeline view, audit trails, and security controls required for tracking sensitive user information.
Compliance and telemetry –Every AI-generated email can be logged, timestamped, and linked to the user’s record, satisfying the telemetry and audit-tracking stipulations.
Scalable and supported –Microsoft officially supports Copilot email assist as part of the Dynamics365 ecosystem, ensuring future updates, security patches, and alignment with Microsoft’s conversational AI best practices.

Why the other options are less suitable

A. Azure AI Search –Primarily a search service for indexing and retrieving content across data sources; it does not provide an out-of-the-box email composition or sequencing capability within Dynamics365 Sales.
C. Voice template Microsoft Copilot Studio –Focused on voice-enabled interactions and custom voice bots; the scenario calls for email correspondence, not voice-driven prospect outreach.
D. Deep Research in Microsoft Foundry Agent Service –Part of the broader Foundry platform for large-scale data analytics and research; it is not a native Dynamics365 Sales feature and lacks direct integration for prospect email workflows.

References

Dynamics365 Sales –Copilot email assist: https://learn.microsoft.com/dynamics365/sales-enterprise/copilot-email-assist
Azure AI Search documentation: https://learn.microsoft.com/azure/search/overview

Question: 47 CertyIQ

HOTSPOT -
Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Background -
Fabrikam, Inc., is a global consumer goods company that is undergoing a digital transformation initiative to migrate its entire infrastructure to the Microsoft cloud. As a key element of this cloud migration, the company will implement Microsoft Dynamics 365 Sales, moving away from the current on-premises proprietary technologies used by its business-to-business (B2B) sales team.
As part of the cloud migration, Fabrikam will adopt an AI-first approach to its business solutions and implement AI solutions, wherever possible, to streamline operations.
Problem Statements -
Fabrikam's infrastructure currently relies on various on-premises systems that require sales executives to use corporate computers with physical keyboards to access business information during customer interactions. Mobile phones cannot be used for these purposes, as the systems depend on keyboard input. As a result, the sales executives spend a lot of time using keyboards to search for data on several disparate systems and file servers, rather than focusing on the customers. This affects the customer experience.
Fabrikam stakeholders are concerned that users will be hesitant to adopt AI. If the AI initiatives are NOT adopted, cost savings will never be realized. Additionally, funding for future AI initiatives will depend on demonstrating an increase in AI adoption month over month. As the AI agent initiative for the sales team will be the first for Fabrikam, the rapid adoption of the agent is a high priority.
Planned Initiatives -
General -
Fabrikam management has prioritized AI-driven projects to improve efficiency, customer engagement, and responsible AI adoption. The current application infrastructure is on-premises and must be migrated to the cloud to support the adoption of these technologies.
Infrastructure Migration -
Fabrikam plans to migrate from its current on-premises infrastructure to a completely cloud-based topology; this will include user authentication, the security framework, and, primarily, the adoption of the services by end users. All the data from the different systems will be consolidated into a single data source - a common data model that

will use a Microsoft Dataverse environment as a single source of truth (SSOT) for the sales team.
Sales Cycle Enablement -
To achieve the company's objectives, Fabrikam intends to implement the following strategies to enhance the sales cycle:
Use low-code development to create a single AI agent that has Dataverse as its core component.
Ensure that sales managers can access unanswered correspondence from prospects and intervene as appropriate.
Replace the previous proprietary software with Dynamics 365 Sales to track sales cycles and customer interactions.
Have the sales executives use Dynamics 365 Sales to track interactions for open opportunities and send follow-up communications to prospects.
Have the sales executives use handsfree headsets to interact with an AI agent when they have questions about internal policies or customer data.
Requirements -
Infrastructure Migration -
Fabrikam has identified the following infrastructure migration requirements:
Azure must be used for all future infrastructure workloads.
The company must follow Microsoft-recommended methodologies for infrastructure migration to the cloud.
Any created AI agents must have their return on investment (ROI) calculated to ensure that the solution will save the company money.
Sales Cycle Enablement -
Fabrikam has identified the following requirements for sales cycle enablement:
The final AI agent must follow Microsoft recommendations for a conversational user experience.
A designated checklist must be reviewed to ensure that the AI agent follows Microsoft deployment recommendations for a compliant solution.
Detailed telemetry must be logged for the first created AI agent to help troubleshoot and optimize the agent during the initial AI agent adoption process.
Unexpected AI agent actions must end in an escalation to a live representative. For example, a sales executive must be rerouted to a representative if the agent cannot answer a question after two failed attempts.
The return on investment (ROI) of switching from the current process to the future process is required for stakeholder sign off.
The sales team must use Dynamics 365 Sales to correspond with prospects more quickly and efficiently than currently.
Sales managers must report on the adoption of the AI agent to key Fabrikam stakeholders on a monthly basis.
Any sensitive information, such as user IDs and names, shared via the AI agent must be tracked for future auditing.
Which components should you use to meet the sales cycle enablement requirements? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

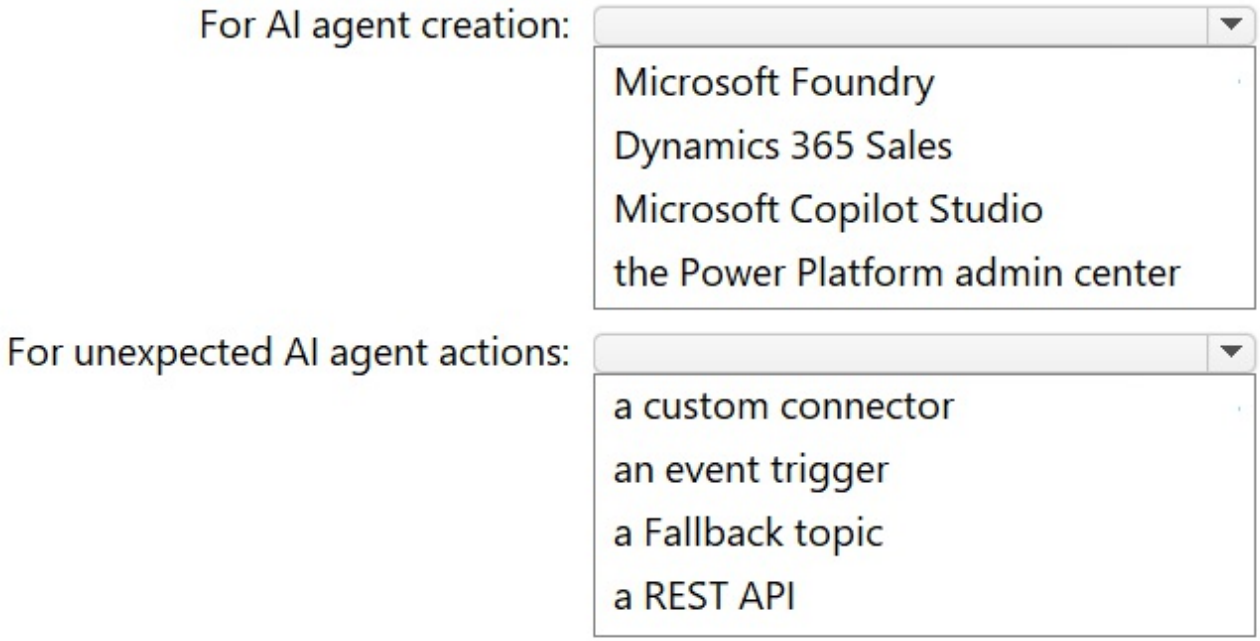

Answer:

Answer Area

Explanation:

For AI agent creation → Microsoft Copilot Studio.

Microsoft Copilot Studio is the Microsoft platform used to:

Build AI agents and copilots

Design conversational workflows

Connect to data sources

Integrate with Power Platform and Microsoft 365

It allows organizations to create AI agents that can:

Answer questions

Automate tasks

Interact with users through conversations

Typical features include:

Natural language understanding (NLU)

Topic creation

AI orchestration

Integration with connectors and APIs

Therefore, it is the correct tool for AI agent creation.

For unexpected AI agent actions → a Fallback topic.

What a fallback topic is

A Fallback topic is triggered when:

The AI agent cannot understand the user request

The request does not match any defined topic

The system cannot determine intent

The fallback topic provides a safe response path.

Question: 48 CertyIQ

You need to design a Microsoft 365 Copilot solution to optimize employee productivity. The solution must meet the following requirements:
Ensure that the employees can query content stored in a subset of Microsoft SharePoint Online sites and in Teams by using natural language-based prompt actions.
Ensure that employees receive contextually relevant responses in Microsoft 365 Copilot.
What should you include in the design?

A. Build a Microsoft Power Automate desktop flow to read the SharePoint content and post the responses to Teams.

B. Modify SharePoint settings.

C. Create a custom REST API that crawls the SharePoint content.

D. Configure Microsoft Graph access.

Answer: D

Explanation:

Why D – Configure Microsoft Graph access is the optimal choice

Microsoft Graph acts as the unified data fabric that exposes SharePoint, Teams, and related content through a set of REST-ful endpoints and Graph-based schema extensions. By configuring the appropriate Graph permissions (e.g., Sites.Read.All, Group.Read.All, User.Read), the Copilot engine can retrieve and index the relevant content directly, enabling natural-language prompt actions without custom crawlers or intermediary services.
Graph-based access respects Azure AD security, sensitivity labels, and retention policies, ensuring that employees only see content they are authorized to query — critical for compliance and user trust.
Native Graph integration is fully supported by Microsoft 365 Copilot, allowing prompt actions to surface context-aware responses instantly within the Copilot UI and Teams chat.
This approach minimizes operational overhead: no custom APIs to maintain, no scheduled Power Automate flows to orchestrate, and no need to expose new endpoints or manage additional authentication layers.

Why the other options are less suitable

A. Build a Power Automate desktop flow – Power Automate is designed for UI-level automation and file-level data movement; it cannot surface the indexed content to Copilot or guarantee the low-latency, context-rich responses required. It also introduces scalability and governance challenges.
B. Modify SharePoint settings – Changing SharePoint site or library settings alone does not expose the content to Copilot's natural-language query engine; it only affects view or permission defaults, leaving the underlying metadata inaccessible to Copilot without Graph enrichment.
C. Create a custom REST API that crawls SharePoint – While technically possible, a custom crawler introduces extra code to maintain, breaks native security boundaries, and bypasses Graph's built-in index and security inheritance, resulting in higher complexity and reduced compliance assurance.

Key takeaway – Leveraging Microsoft Graph to surface SharePoint and Teams content provides a secure, scalable, and Copilot-compatible data source that directly satisfies the requirement for natural-language

query and contextual relevance.

References

Microsoft Graph API permissions for SharePoint and Teams: https://learn.microsoft.com/en-us/graph/permissions-reference
Configure Copilot to surface SharePoint content: https://learn.microsoft.com/en-us/microsoft-365/copilot/configure-sharepoint-content-access?view=o365-worldwide

Question: 49 CertyIQ

A company uses Microsoft Dynamics 365 Finance to manage accounts payable.
You are designing an AI invoice processing solution.
You need to recommend the prerequisites to configure a prebuilt copilot for accounts payable.
What should you recommend?

A. From Microsoft Copilot Studio, create an accounts payable agent.
B. Extend Microsoft 365 Copilot for Sales to an accounts payable agent.
C. Build an AI tool in Microsoft Foundry.
D. From the Power Platform admin center, assign the Finance and Operations AI security role to users.

Dynamics 365 Finance and Operations apps host prebuilt Copilot sidecars and summaries via linked Power Platform / Dataverse environments.

Answer: D

Explanation: To authorize users to interact with prebuilt Copilot capabilities (such as Accounts Payable invoice processing, vendor summaries, and collections summaries), standard ERP functional roles alone are insufficient.

Answer: D

To enable the out-of-the-box Copilot for Accounts Payable, users must be granted the **Finance and Operations AI security role** which provides the necessary permissions to access AP data, invoke the AI model, and write back processed results. This role is assigned via the Power Platform admin center, ensuring that only authorized personnel can interact with the copilot and that governance policies are enforced.
Option A describes creating an agent in Copilot Studio, which is useful for custom scenarios but not required for the prebuilt AP copilot that Microsoft already ships.
Option B suggests extending Copilot for Sales; this service is oriented toward sales-related activities and does not expose the finance-specific data models needed for invoice processing.
Option C references building a solution in Microsoft Foundry, which is a broader integration platform and introduces unnecessary complexity for a pre-configured AI capability that is already packaged in Dynamics 365 Finance.

Therefore, assigning the Finance and Operations AI security role (OptionD) is the precise prerequisite for configuring the prebuilt accounts-payable copilot.

References

https://learn.microsoft.com/en-us/power-platform/admin/security-roles-configure
https://learn.microsoft.com/en-us/dynamics365/finance/ai/ai-security-role

Question: 50 CertyIQ

A company plans to deploy a Microsoft Dynamics 365 Contact Center agent.
You need to ensure that the agent can transfer the conversation to a live customer service representative.
Which two components should you include in the solution? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. Microsoft Foundry
B. Microsoft Copilot Studio
C. Microsoft 365 Agents Toolkit
D. an Azure AI Bot Service skill
E. Customer engagement hub

Microsoft Copilot Studio: The authoring platform where the agent conversation logic, intent triggers, and escalation pathways are built. It provides the native Transfer conversation node used to pass conversation context, transcripts, and custom context variables to a human queue.

Customer engagement hub: The specific integration mechanism inside Copilot Studio (Settings > Customer engagement hub or Agent details > Advanced). Configuring this connection registers the Copilot Studio agent as an automated bot user in Dynamics 365 Contact Center / Omnichannel, enabling routing engines, unified routing workstreams, and live agent queues to receive seamless handoffs.

Answer: BE

Explanation:

Why "B&E" is the correct choice

Microsoft Copilot Studio (B) – Provides the AI-driven conversational authoring environment where you can create or refine the generative AI skills that understand the context of a customer interaction and trigger a hand-off when certain conditions are met. It integrates directly with Dynamics365 Contact Center to expose those capabilities to the agent workflow.
Customer Engagement Hub (E) – The omnichannel "Customer Engagement Hub" (formerly Omnichannel for Customer Service) is the core platform that orchestrates multi-channel routing, session state, and provides built-in hand-off actions that transfer the conversation to a live agent queue or individual representative. It also surfaces the necessary UI components (e.g., "Transfer to Agent") for seamless operator use.

Why the other options are less suitable

A. Microsoft Foundry – Foundry is a cloud-scale data and analytics platform; it does not supply the conversational or routing capabilities needed for a hand-off in Contact Center.
C. Microsoft 365 Agents Toolkit – This toolkit helps developers build agents for personal productivity, but it does not deliver the omnichannel routing or live-agent hand-off functionality required in Dynamics365 Contact Center.
D. Azure AI Bot Service skill – While a Bot Service skill can host a bot, it must be packaged and published via Copilot Studio to be consumed by Dynamics365. Using it alone does not guarantee the native hand-off flow that the Customer Engagement Hub provides.

Result – Selecting **B** (Microsoft Copilot Studio) together with **E** (Customer Engagement Hub) delivers the AI-enhanced conversation context and the native omnican-al hand-off mechanism needed for a seamless transfer to a live representative.

References

Customer Engagement Hub (Omnichannel) documentation – https://learn.microsoft.com/en-us/dynamics365/customer-engagement/omnichannel/
Microsoft Copilot Studio documentation – https://learn.microsoft.com/en-us/copilot-studio/

Question: 51 CertyIQ

HOTSPOT -
A company uses Microsoft Dynamics 365 Supply Chain Management.
You are designing an AI supply chain process that meets the following requirements:
Provides managers with AI-driven insights that surface key information from customer orders
Helps planners use AI to anticipate future product needs more accurately
You need to recommend which Microsoft Copilot features to include in the design.
What should you recommend for each requirement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

Provide AI-driven insights from customer orders:

- AI Summaries with Copilot
- Generative insights for Demand planning
- The Customer credit and collections workspace
- Workload insights with Copilot

When customer service and order managers open an order, AI Summaries with Copilot automatically aggregates line items, delivery schedules, fulfillments, and past exceptions to surface key information without requiring manual navigation across tabs.

Anticipate future product needs:

- Generative insights for Demand planning
- Microsoft Power BI
- Product information management
- The Supplier Communications Agent

In the Dynamics 365 Supply Chain Management Demand Planning app, Copilot delivers generative insights that analyze past demand signals, historical consumption trends, promotional lifts, and market changes. Demand planners use these AI-generated narratives to detect forecast shifts, understand demand anomalies, and anticipate future inventory and product needs with higher accuracy.

Answer:

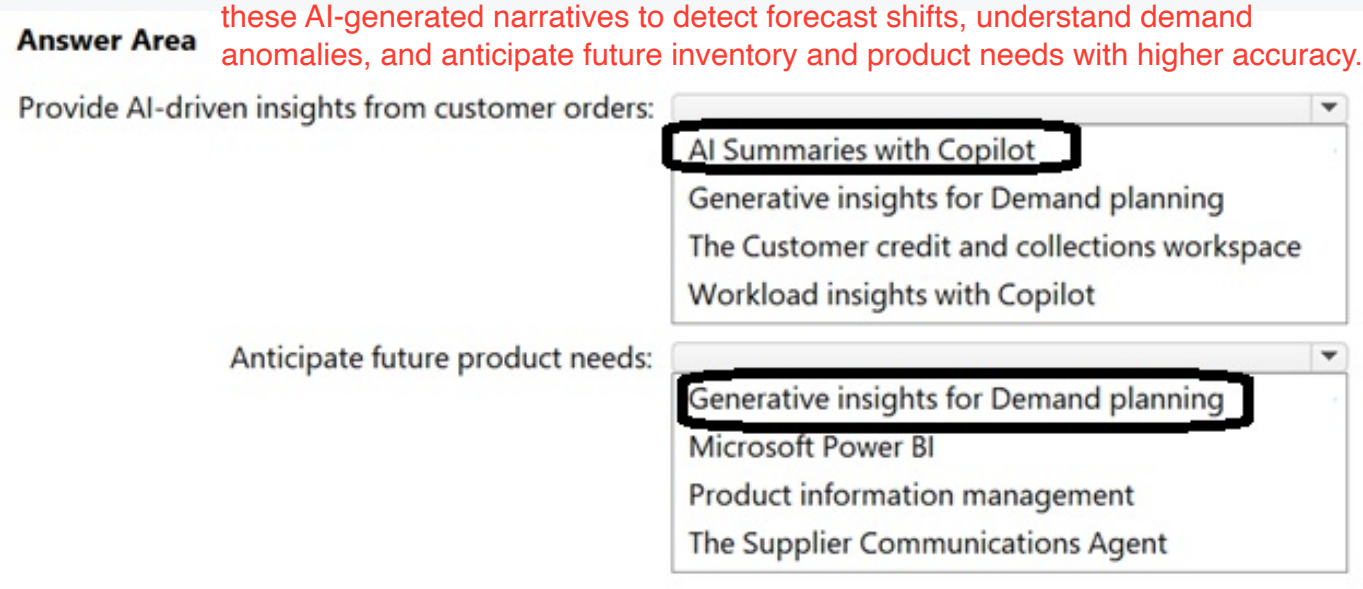

Explanation:

Provide AI-driven insights from customer orders

Correct Answer: AI Summaries with Copilot

What it does

AI Summaries with Copilot analyzes business data and generates natural language summaries.

For customer orders, it can:

Summarize order trends

Highlight customer purchasing patterns

Identify order anomalies

Generate insightful summaries for business users

Example output:

“Customer orders increased by 15% this month, with the highest demand for Product A in the North region.”

This allows managers to quickly understand order insights without manually analyzing data.

Anticipate future product needs

Correct Answer: Generative insights for Demand planning.

What it does

Generative insights for Demand planning uses AI to:

Predict future product demand

Analyze historical sales patterns

Identify trends and demand fluctuations

Generate demand forecasts

This helps companies:

Plan inventory

Avoid stock shortages

Optimize production

Question: 52 CertyIQ

HOTSPOT -
A company has a Microsoft 365 E5 subscription and uses Microsoft Copilot Studio.
The company has a Microsoft SharePoint Online library that contains 10,000 policy PDFs from various departments. The library contains a populated column named Department for each PDF.
You need to design a Copilot Studio agent that will use the SharePoint library as a knowledge source. The solution must meet the following requirements:
Enable the agent to answer user questions about company policies.
Ensure that the agent can identify which departments and policies are connected.
What should you include in the design for each requirement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

Enable the agent to answer questions about company policies:

- Build a custom model in Microsoft Foundry.
- From Copilot Studio, add SharePoint as a knowledge source.
- Import the PDFs into Microsoft Dataverse.
- Use AI Builder to process and feed SharePoint content.

Identify which departments and policies are connected:

- Apply Microsoft Purview sensitivity labels.
- Create a Microsoft Dataverse table for the departments.
- From Copilot Studio, configure the SharePoint tool.
- Upgrade to SharePoint Premium.

Answer: Configuring the SharePoint Tool: Passive semantic document grounding searches unstructured text within files, but metadata relationships (such as querying document library column attributes like Department to map which policies belong to which departments) require structured action execution. In Copilot Studio's agentic framework, adding and configuring the native SharePoint tool (action/connector plugin) allows the agent to query list/library schema metadata, filter items by custom columns (Department), and correlate documents to structural entities.

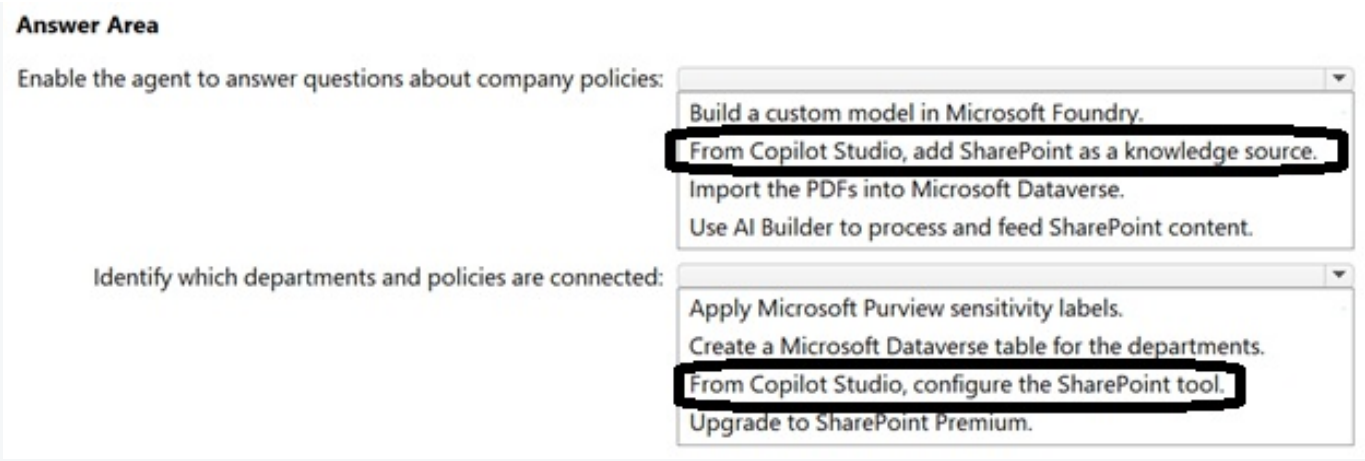

Explanation:

From Copilot Studio, add SharePoint as a knowledge source.

Why this is correct

Many organizations store policies and documentation in SharePoint.

By adding SharePoint as a knowledge source in Copilot Studio, the AI agent can:

Access policy documents

Search the content

Answer questions using those documents

This works using grounding / knowledge retrieval so the AI can generate answers based on company content.

From Copilot Studio, configure the SharePoint tool.

The SharePoint tool in Copilot Studio enables the agent to:

Query SharePoint content

Retrieve metadata

Understand document relationships

Connect policies to departments

This helps the agent determine:

Which department owns a policy

How documents relate to each other

Question: 53 CertyIQ

DRAG DROP -
You need to design a Microsoft Copilot Studio agent that meets the following requirements:
Supports interactive speech responses
Optimizes decision-making and the accuracy of responses
What should you include in the design for each requirement? To answer, drag the appropriate options to the correct requirements. Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Options

- A deep reasoning model
- Azure Language in Foundry Tools
- Azure AI Speech
- Copilot Studio voice features
- Speech Synthesis Markup Language (SSML)

Answer Area

	Option
Supports interactive speech responses:	Copilot Studio voice features
Optimizes decision-making and response accuracy:	A deep reasoning model

Copilot Studio voice features: Copilot Studio includes native, out-of-the-box voice capabilities (telephony channel integration, speech-to-text input, and text-to-speech output). Designing an agent for interactive speech responses requires enabling and configuring Copilot Studio voice features directly within the agent's channel configuration rather than integrating external standalone speech pipelines.

Answer:

Answer Area

	Option
Supports interactive speech responses:	Copilot Studio voice features
Optimizes decision-making and response accuracy:	A deep reasoning model

Explanation:

1 Supports interactive speech responses → Copilot Studio voice features.

Interactive speech responses involve:

Voice input from users

Spoken responses from the AI agent

Real-time conversational interaction

Why Copilot Studio voice features are correct

Copilot Studio voice features enable AI agents to:

Understand spoken user input

Convert speech to text

Generate responses

Convert responses back to speech

This enables voice-based conversations with the AI agent.

ures

What this means

Interactive speech responses involve:

Voice input from users

Spoken responses from the AI agent

Real-time conversational interaction

Why Copilot Studio voice features are correct

Copilot Studio voice features enable AI agents to:

Understand spoken user input

Convert speech to text

Generate responses

Convert responses back to speech

This enables voice-based conversations with the AI agent.

Example interaction:

User (voice):

“What is my account balance?”

Copilot agent:

Converts speech → processes request → responds with spoken output

This capability is designed specifically for interactive voice experiences.

Optimizes decision-making and response accuracy → A deep reasoning model.

What a deep reasoning model does

A deep reasoning model enhances AI’s ability to:

Analyze complex information

Understand context

Perform multi-step reasoning

Produce more accurate and logical responses

These models are particularly useful for:

Complex problem solving

Analytical tasks

Decision support systems

Question: 54 CertyIQ

You are designing a low-code AI business solution by using Microsoft Copilot Studio.
The solution must include an agent that automates tasks by simulating user interactions across third-party apps and websites, such as clicking buttons, entering text, and extracting information from screens.
You need to recommend what to include in the agent.
What should you recommend?

A. Model Context Protocol (MCP)
B. a natural language understanding + (NLU+) model in Copilot Studio
C. Computer Use in Copilot Studio
D. Copilot skills

Answer: C

Explanation:

Correct option – C. Computer Use in Copilot Studio
Purpose-built feature: ComputerUse enables an agent to interact with graphical user interfaces (GUI) of third-party applications – clicking buttons, typing text, navigating pages, and extracting screen-rendered data – exactly as a human would, but in an automated, programmatic fashion.
Low-code experience: It abstracts the underlying complexity of browser automation and provides declarative actions that can be orchestrated within Copilot Studio flows, allowing rapid prototyping without writing low-level code.
Scalable orchestration: Works within the same environment where you design low-code workflows, making it seamless to compose end-to-end business processes that span multiple web services.
Why the other options are less suitable
A. Model Context Protocol (MCP) – Primarily a communication protocol for connecting AI models to external data sources or tools; it does not provide built-in UI interaction capabilities such as clicking or typing in a web page.
B. Natural Language Understanding+(NLU+) model – Enhances intent recognition and entity extraction but cannot directly simulate user gestures or control external applications.
D. Copilot skills – Refers to pre-built functional modules (e.g., scheduling, email drafting) that operate on structured data; they are not designed for raw GUI manipulation across arbitrary third-party sites.

Brief Summary

To enable an agent that performs UI-level automation on arbitrary web applications, you must leverage the **Computer Use** capability inside Copilot Studio, which offers declarative, low-code interaction with browsers. The other listed components support language understanding or integration but lack native UI automation functions.

References

Computer Use in Copilot Studio – Microsoft Docs: https://learn.microsoft.com/en-us/copilot-studio/computer-use
Copilot Studio overview and features – Microsoft Docs: https://learn.microsoft.com/en-us/copilot-studio/overview

Question: 55 CertyIQ

You need to recommend a solution to integrate a Microsoft Copilot agent with a Microsoft Dynamics 365 Contact Center chat channel.
The agent must respond to customer questions and hand off the conversation to a live customer service representative when the customer requests an escalation.
What should you recommend?

A. Build an agent flow.
B. Configure the Conversation Start topic.
C. Configure a skill.
D. Call a Microsoft Power Automate connector.
E. Configure the Escalate topic.

Answer: E

Explanation:

Why "Configure the Escalate topic" (E) is the correct choice

1. **Copilot Studio's native escalation mechanism** – Within a Copilot-powered Dynamics365 Contact Center bot you define topics that control the dialogue flow.

   An **Escalate topic** tells the runtime to end the current bot session and **hand-off the conversation to a live agent** (typically routed to a Teams or Dynamics365 agent workspace).
   The handoff uses the built-in **Escalate action** which automatically passes session context (e.g., transcript, customer-ID) to the escalation channel without any custom code.

2. **Seamless integration with the Dynamics365 Contact Center platform** – The platform expects the bot to surface an Escalate node when a customer explicitly requests human help.

   Configuring this node aligns the bot's behavior with the out-of-the-box **"Live agent hand-off"** pattern required by the Contact Center service, ensuring the hand-off is recognized by routing rules, queues, and reporting.

3. **Preserves session context and security** – The Escalate topic can **pass the whole conversation transcript, bot state, and authentication tokens** to the live-agent workspace, so the agent sees exactly what the customer has discussed and can continue without re-authentication.

   This is handled automatically; no extra APIs or custom connectors are needed.

4. **Meets the requirement "respond to customer questions and hand off when the customer requests an escalation"** –

   The bot can still answer questions via regular conversational flows (topics, entities, triggers).
   When the user says something like "I want to talk to a representative," the bot evaluates that intent and switches to the Escalate topic, invoking the hand-off action.

Why the other options are less suitable

A – Build an agent flow

"Agent flow" is a design-time term for the overall bot construction process, not a specific runtime feature that triggers a hand-off. Relying only on this does not guarantee the correct escalation behavior.

B – Configure the Conversation Start topic

This topic merely defines the opening greeting and initial prompts; it does not provide a mechanism for detecting escalation requests or transferring the session.

C – Configure a skill
– Skills are used in Dynamics365 for talent or resource allocation, not for bot-to-agent hand-off in a Copilot context. They have no impact on conversation routing.

D – Call a Microsoft Power Automate connector
– While a Power Automate flow could simulate an escalation, it adds unnecessary complexity, extra licensing, and latency. The platform already offers a native Escalate action that is more reliable and is the recommended best practice.

Technical implementation in practice (brief)

1. **Add an "Escalate" node** to the bot's topic graph where the intent "escalate to agent" or a phrase like "talk to a representative" is detected.
2. **Set the node's action** to Escalate to live agent; optionally populate **hand-off data** such as customerId, conversationTranscript, and custom context variables.

3. **Publish the bot** to the Dynamics365 Contact Center channel. When the condition is met, the runtime automatically switches the session to the live-agent workspace, preserving all context.
4. **Monitor** hand-off events through the Contact Center analytics dashboard to ensure escalation rates meet service level objectives.

References

Microsoft Docs:Escalate a conversation with a Copilot bot
Microsoft Docs:Guide to building Copilot Studio topics and hand-off actions

---

Question: 56 CertyIQ

A company has a customer order system that creates sales orders manually.
You need to design an AI solution to automate the following tasks as part of the system:
Save the order details to a database.
Update the order status in the database.
Extract the order details from an order file.
Prepare and send a confirmation email to customers.
The solution must minimize development effort and support intelligent automation and solution integration.
What should you include in the design?

A. a workflow in Azure Logic Apps
B. a multi-agent solution that uses the Semantic Kernel SDK
C. a multi-agent solution that uses Microsoft Foundry Agent Service
D. a Microsoft Copilot Studio agent that uses Microsoft Power Automate workflows

Answer: D

Explanation:

The requirement is to **automate end-to-end order processing** while **minimizing development effort** and enabling **intelligent automation with easy integration**.
OptionD – a MicrosoftCopilotStudio agent that uses PowerAutomate workflows directly satisfies these constraints:
CopilotStudio provides a **no-code/low-code conversational interface** that can **ingest the order file, extract entities (order ID, product, quantity, etc.)**, and **orchestrate the process**.
PowerAutomate connectors let the agent **write to a database (e.g., AzureSQL/Dataverse)**, **update the order status**, and **send a confirmation email** without writing custom code.
The solution leverages **AI-enhanced intent recognition and entity extraction** built into CopilotStudio, delivering intelligent automation out of the box.
Development effort is limited to **configuring the Copilot's prompts, mapping entities, and wiring PowerAutomate actions**, which can be done in a few hours rather than building full-stack code.
Why the other options are less suitable
OptionA (Azure Logic Apps) – While Logic Apps can string together the same actions, they provide **no native AI capabilities** for extracting unstructured order data; you would still need custom code or additional services to perform intelligent parsing.
OptionB (multi-agent solution with Semantic Kernel SDK) – This approach offers powerful AI orchestration but requires **substantial development** (C#/.NET, SDK integration, custom agent logic) and is **not aimed at minimizing effort**. It is better suited for highly customized multi-agent scenarios, not a quick-deployment order system.
OptionC (multi-agent solution with Microsoft Foundry Agent Service) – Foundry's agent services are

designed for large-scale, enterprise-level AI pipelines and involve **complex infrastructure, scaling, and governance considerations**, making them over-engineered for a simple order-processing workflow and increasing development overhead.

Conclusion – The combination of a **CopilotStudio conversational agent with PowerAutomate flows** gives the fastest, lowest-code path to a fully automated, intelligent order-processing system that meets the requirement for minimal development and seamless integration.

References

MicrosoftCopilotStudio documentation: https://learn.microsoft.com/en-us/copilot-studio/
PowerAutomate integration with AzureSQL/Dataverse: https://learn.microsoft.com/en-us/power-automate/connectors/azure-sql-connection

Question: 57 CertyIQ

HOTSPOT -
You are designing an AI strategy for Microsoft Dynamics 365 finance and operations apps. You are evaluating the use of Microsoft Copilot Studio to provide in-app help and guidance based on generative AI general knowledge.
You need to recommend which knowledge sources to include in the generative help and guidance agent. The solution must minimize the risk of generating inaccurate responses.
What should you recommend? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

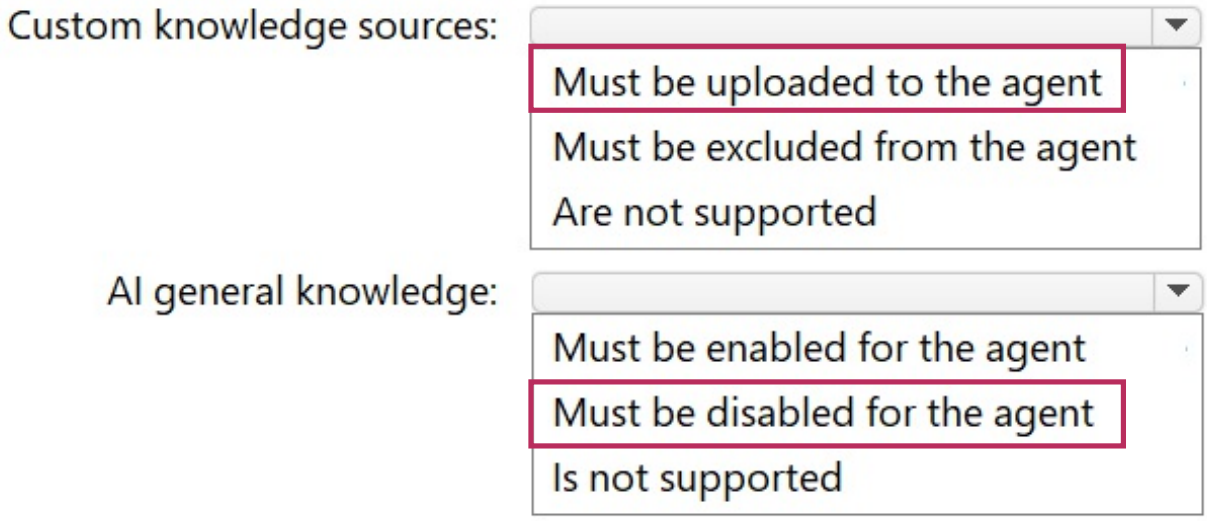

Answer: Custom knowledge sources: By default, generative help and guidance is grounded on public Microsoft Learn documentation. To support organization-specific business processes, standard operating procedures, and business rules, administrators extend the Copilot for finance and operations apps agent in Copilot Studio by uploading verified internal documents (PDF, Word, RTF) or SharePoint links. Grounding the agent with uploaded enterprise documents anchors generative answers in factual organization-approved data.

AI general knowledge: Microsoft explicitly warns: "When the use of general knowledge is enabled, it can increase the richness of Copilot responses. However, it also increases the risk that answers might include incorrect content." Therefore, to satisfy the explicit constraint to minimize the risk of generating inaccurate responses (hallucinations), general knowledge must be disabled for the agent, restricting responses strictly to verified Microsoft public documentation and uploaded custom knowledge sources.

Answer Area

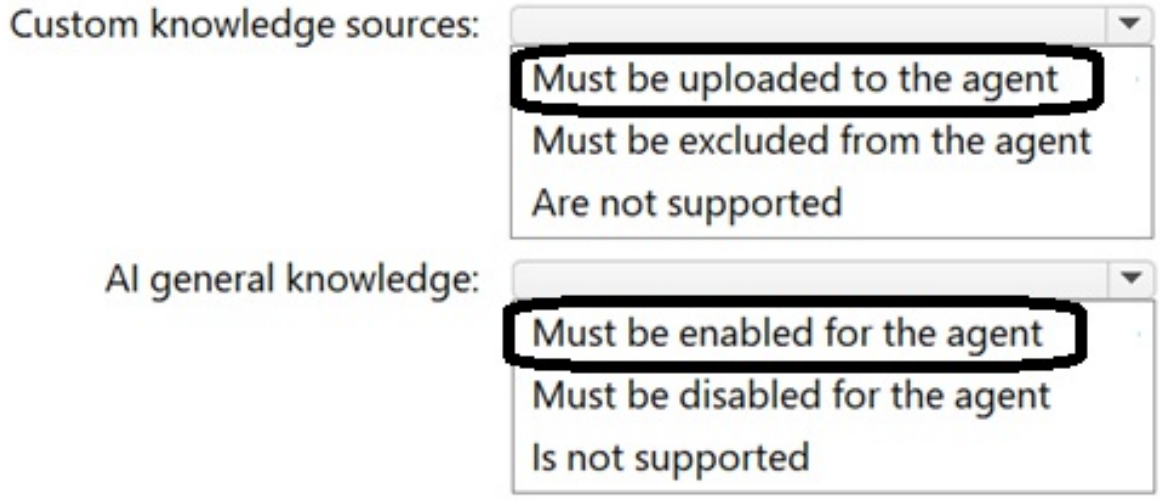

Explanation:

Custom knowledge sources → **Must be uploaded to the agent.**

What custom knowledge sources are

Custom knowledge sources are organization-specific content used by an AI agent, such as:

Company policies

Product manuals

Internal documentation

FAQs

SharePoint files

PDFs or web pages.

AI general knowledge → **Must be enabled for the agent.**

What AI general knowledge is

AI general knowledge refers to the pretrained knowledge built into the AI model, including:

General facts

Language understanding

Common knowledge

Reasoning capabilities

Question: 58 CertyIQ

HOTSPOT -
You need to design a multi-agent solution that will include a custom agent. The solution must meet the following requirements:
Define the rules and constraints that the agent must follow.
Automate a backend process that involves data movement between services and runs independently of the agent's reasoning steps.

What should you include in the design for each requirement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

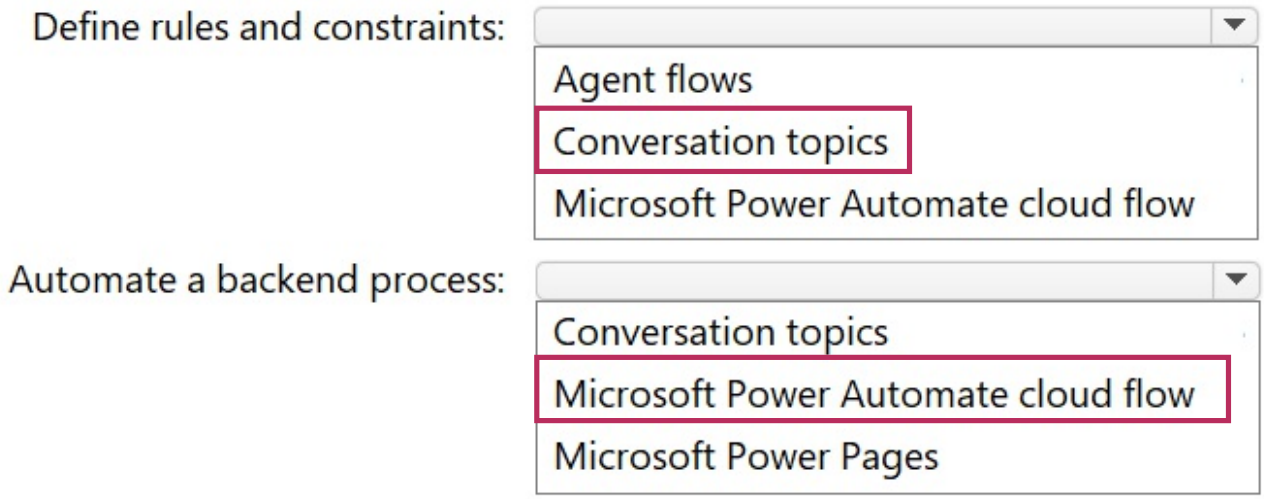

Answer:

conversation topics provide deterministic, node-based control. Topics allow architects to define explicit conditional logic, strict validation constraints, boundary rules, and step-by-step branching guardrails that force the agent into controlled conversation states rather than relying on unconstrained generative routing.

Answer Area

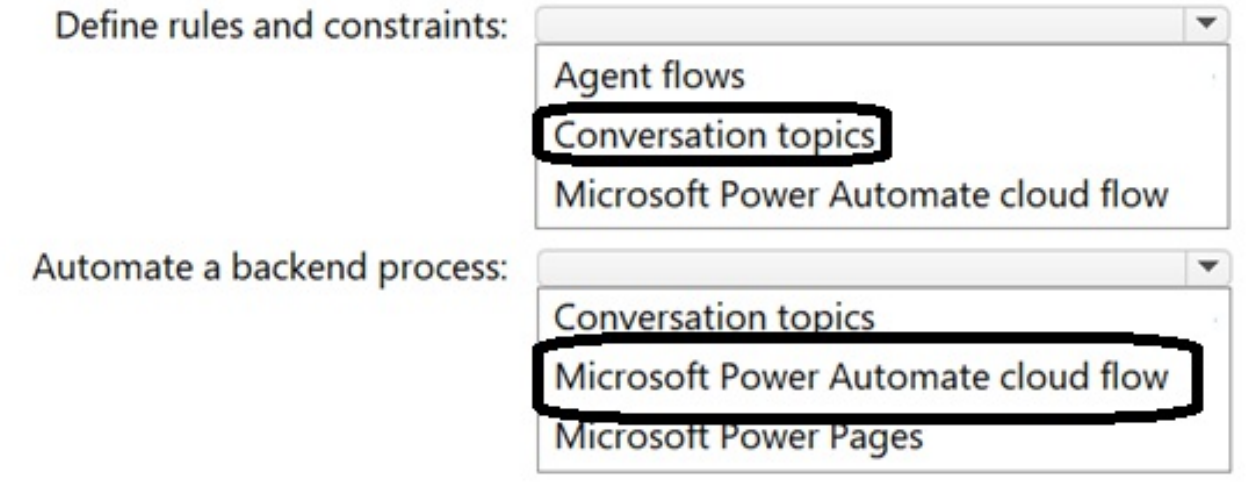

Explanation:

Define rules and constraints → **Conversation topics.**

In Microsoft Copilot Studio, a topic defines how the agent handles a conversation.

Topics include:

Trigger phrases

Conversation logic

Conditions

Responses

Decision paths

Automate a backend process → **Microsoft Power Automate cloud flow.**

What Power Automate cloud flows do

Power Automate cloud flows automate tasks between services.

They allow Copilot agents to interact with:

Databases

APIs

Microsoft 365 services

Business applications

Question: 59 CertyIQ

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.
After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.
A company uses Microsoft 365 and Dynamics 365.
You need to recommend a solution to automatically summarize email threads, generate suggested replies in Microsoft Outlook, and provide meeting preparation summaries that include relevant customer relationship management (CRM) data.
Solution: You recommend Microsoft 365 Copilot for Sales.
Does this meet the goal?

A. Yes
B. No

Microsoft 365 Copilot for Sales bridges Microsoft 365 productivity apps (Outlook, Teams, Word) directly with enterprise CRM platforms (Dynamics 365 Sales or Salesforce).

Email thread summaries: Grounded in past communications and active CRM pipeline context.

AI-generated reply drafts: Suggesting personalized follow-ups with contextual data pulled from the connected CRM system (e.g., pricing, recent opportunities, contact details).

Answer: A

Explanation:

Meeting preparation briefs: Pre-meeting digests summarizing past interactions, open action items, and relevant CRM contact/opportunity records directly on the calendar invite card.

Microsoft365 Copilot for Sales is an integrated AI layer that leverages large-language-model capabilities within Outlook and Teams to:

1. **Summarize email threads** – automatically creates concise overviews of long conversations.
2. **Generate suggested replies** – proposes context-aware response drafts directly in the message pane.
3. **Provide meeting-preparation summaries** – pulls related Customer Relationship Management (CRM) records, recent activities, and relevant leads from Dynamics365 Sales, surfacing them when a meeting is scheduled.

These functionalities are delivered **natively** to Outlook and Teams, requiring no additional connectors or custom scripts, which aligns with the goal of an **out-of-the-box** solution for summarization, reply suggestion, and CRM-enhanced meeting prep.
Alternative options (e.g., separate PowerAutomate flows, third-party add-ins, or native Dynamics365 “Sales Assistant” alone) either:
lack **real-time Outlook integration** for reply suggestions,
do not embed CRM data directly into Outlook summaries, or
require extensive configuration and maintenance, making them less suitable for a streamlined, enterprise-scale deployment.
Consequently, recommending **Microsoft365 Copilot for Sales** satisfies all stated requirements in a single,

supported service, making the answer **Yes**.

References

Microsoft Docs: Introducing Microsoft 365 Copilot for Sales – https://learn.microsoft.com/en-us/salescopilot/overview
Microsoft Docs: Sales Copilot capabilities in Outlook – https://learn.microsoft.com/en-us/salescopilot/outlook-integration

Question: 60 CertyIQ

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.
After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.
A company uses Microsoft 365 and Dynamics 365.
You need to recommend a solution to automatically summarize email threads, generate suggested replies in Microsoft Outlook, and provide meeting preparation summaries that include relevant customer relationship management (CRM) data.
Solution: You recommend a classic Microsoft Dataverse workflow.
Does this meet the goal?

A. Yes
B. No

Answer: B

Explanation:

Answer: B – No

The requirement involves AI-driven text summarization, smart reply suggestions, and contextual CRM data retrieval within Outlook.
A classic Microsoft Dataverse workflow can only automate simple, rule-based actions (e.g., record updates, status changes, simple notifications). It lacks built-in natural-language generation, AI model integration, or Outlook-specific actions such as “generate suggested reply.”
To achieve automatic thread summarization, AI-generated reply proposals, and meeting-prep summaries that pull CRM data, you need a solution that combines:

1. **Outlook add-ins or Power Automate flows** that trigger on incoming/threaded messages.
2. **AI capabilities** (e.g., AI Builder, Copilot, or Azure Cognitive Services) to generate concise summaries and suggested replies.
3. **CRM data enrichment** through Dynamics 365/CDS connectors that surface relevant customer records in the email context.

Classic workflows cannot invoke AI models or dynamically render UI-level suggestions in Outlook; they also cannot enrich a message with live CRM data without custom development that exceeds workflow capabilities.
Therefore, the proposed classic Dataverse workflow **does not meet** the stated goal. A more appropriate approach would involve Power Automate (or Power Apps) with AI Builder/Flow AI integration and Dynamics 365 connector to provide the required intelligent summarization and reply suggestions.

References

1. Power Automate – AI Builder summarization
2. Use AI Builder to create a model for suggested email replies

These links illustrate the AI-driven capabilities needed to satisfy the business requirement, which are not available in a classic Dataverse workflow.

Question: 61 CertyIQ

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.
After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.
A company uses Microsoft 365 and Dynamics 365.
You need to recommend a solution to automatically summarize email threads, generate suggested replies in Microsoft Outlook, and provide meeting preparation summaries that include relevant customer relationship management (CRM) data.
Solution: You recommend a Microsoft 365 Copilot agent template.
Does this meet the goal?

A. Yes
B. No

Microsoft 365 Copilot agent templates are starter blueprints (declarative agents) created via Copilot Studio or SharePoint to automate specific knowledge lookups, team Q&A, or standalone organizational workflows. They do not embed specialized role-based integration into the native Microsoft Outlook UI to parse incoming client emails against CRM records or assemble structured pre-meeting opportunity briefs.

Answer: B

Explanation:

Correct answer: B – No, the proposed solution does not fully meet the stated goal.

Why the Microsoft365 Copilot agent template falls short

Scope of functionality – A Copilot agent template is primarily a pre-built conversational interface that can answer questions, retrieve documents, or trigger simple workflows. It can summarize an individual email or draft a reply, but it does **not** automatically ingest an entire email thread, extract key points across multiple messages, and synthesize a contextual summary without additional configuration.
Integration with Dynamics365 CRM data – The requirement explicitly calls for meeting-preparation summaries that surface relevant customer relationship data from Dynamics365. Out-of-the-box Copilot agent templates do not have built-in connectors to Dynamics365; they rely on manually defined data sources or custom plugins. Without a dedicated integration (e.g., Power Automate flows, Azure Logic Apps, or custom Graph API calls), the CRM data will not be pulled into the meeting prep summary automatically.
Automation of suggested replies – While Copilot can suggest a reply when a user asks, the goal demands automatic generation of suggested replies in Outlook as part of a workflow. This level of automation requires a trigger-based mechanism (e.g., a rule that runs when an email is opened) and the ability to inject CRM insights into the reply draft. A generic agent template cannot enforce that trigger or enrich the reply with CRM fields without extensive custom development.
End-to-end orchestration – The scenario calls for a coordinated pipeline: ingest email threads → summarize → surface CRM data → generate reply → deliver back to Outlook. Achieving this reliably calls for **Power Automate** or **Azure Logic Apps** orchestration combined with **AI Builder** or **Custom Cognitive Services** models, plus secure connectors to both Microsoft365 and Dynamics365. A Copilot agent template alone cannot provide the necessary workflow control, error handling, or data mapping.

Why other options (if present) would be more suitable

Power Automate with AI Builder & Dynamics365 connectors can automatically pull emails, run summarization models, enrich the content with CRM records, and write suggested replies back into Outlook, directly addressing all three requirements.
Custom Microsoft Graph + Azure AI Search + Dynamics365 plug-in would give granular control over data retrieval, summarization, and response generation, ensuring end-to-end automation and proper security.
Microsoft Cloud App Security / Customer Insights could also be leveraged to surface CRM data during meeting prep, but again this requires explicit connector configuration rather than a generic Copilot template.

Conclusion

Because a Microsoft365 Copilot agent template lacks the necessary out-of-the-box connectors and workflow automation capabilities to pull Dynamics365 CRM data and automatically generate Outlook replies across whole email threads, it does **not** satisfy the goal. A more robust, integrated solution using Power Automate or custom Azure services is required.

References

Microsoft 365 Copilot documentation: https://learn.microsoft.com/en-us/microsoft-365/copilot/
Power Automate integration with Dynamics 365: https://learn.microsoft.com/en-us/power-automate/connector-dynamics365-ce

Question: 62 CertyIQ

You need to design an application lifecycle management (ALM) process for a Microsoft Power Platform environment that contains a solution named Solution1.
Solution1 must include a custom connector for Copilot in Microsoft Dynamics 365 Customer Service. Solution1 must meet the following requirements:
•Ensure that the custom connector can be deployed consistently across environments as part of the ALM process.
•Allow the custom connector to be edited only in the development environment.
What should you include in the design?

A. Add the custom connector to GitHub.
B. Share the custom connector.
C. Create the custom connector in the default solution.
D. Add the custom connector to Solution1.

Answer: D

Explanation:

Technical Justification for Correct Answer D

To design an effective Application Lifecycle Management (ALM) process for the described Microsoft Power Platform environment, meeting the specified requirements for Solution1 and its custom connector for Copilot in Microsoft Dynamics 365 Customer Service, the following justification outlines why option D is the best choice and the shortcomings of the other options:

Why D: Add the custom connector to Solution1 is Correct

Consistent Deployment Across Environments: By adding the custom connector directly to Solution1, it becomes an integral part of the solution. This integration facilitates consistent deployment of the custom connector alongside other Solution1 components across different environments (dev, test, prod) through the ALM process, satisfying the first requirement.
Controlled Editing Environment: Since Solution1 is presumably managed with environment-specific

permissions (a common ALM practice), editing can be restricted to the development environment. This aligns with the second requirement, as access to edit the custom connector within Solution1 can be controlled at the environment level.

Why Other Options are Less Suitable

A. Add the custom connector to GitHub

Decoupling: Storing the custom connector in GitHub decouples it from Solution1's direct management within the Power Platform, potentially complicating the consistent deployment requirement.
Additional Integration Step: Would require an additional step in the ALM pipeline to integrate from GitHub, adding complexity.

B. Share the custom connector

Lack of Encapsulation: Sharing does not encapsulate the connector within Solution1, making environment-specific deployment and management more challenging.
Security and Access Control: Broad sharing might inadvertently allow edits from unintended environments/users.

C. Create the custom connector in the default solution

Scalability and Organization: Default solutions can become overly cluttered, impacting scalability and maintainability of the ALM process.
Dependency and Isolation: Might introduce unintended dependencies or make isolation for environment-specific deployments more difficult.

Conclusion

Option D, adding the custom connector to Solution1, is the most suitable choice because it ensures the custom connector is tightly integrated with the solution for seamless, consistent deployment across environments while allowing for controlled editing restricted to the development environment through standard ALM and environment access controls.

References

Microsoft Documentation: Manage Solutions in Power Platform
Microsoft Documentation: Customize Dynamics 365 Customer Service with Copilot

Question: 63 CertyIQ

A company uses a Microsoft Copilot Studio agent to automate tasks in a web app.
During testing, you discover that the automation sometimes fails because of frequent changes to the app's user interface.
You need to recommend a solution to ensure that the agent successfully automates the tasks. The solution must minimize changes to the agent.
What should you include in the recommendation?

A. Computer Use in Copilot Studio
B. custom models in Azure AI Studio
C. conversation topics in Copilot Studio
D. an agent flow in Copilot Studio

Answer: A

Explanation: Computer Use in Copilot Studio leverages frontier multimodal vision models to visually interpret the interface dynamically as a human user would. It identifies buttons, input fields, and UI elements based on their visual appearance, contextual labels, and semantic intent rather than rigid coordinate paths or brittle DOM paths. This dynamic visual reasoning allows the automation to succeed across UI updates while minimizing modifications to the agent configuration.

Technical Justification for Recommended Solution

Recommended Solution: A. Computer Vision in Copilot Studio

Why A is the Best Option: Computer Vision in Copilot Studio is the most suitable solution for ensuring the Microsoft Copilot Studio agent successfully automates tasks despite frequent changes to the web app's user interface (UI) with minimal changes to the agent. This is because:

Robustness to UI Changes: Computer Vision technologies can identify elements on the screen based on visual cues rather than relying on static selectors (e.g., XPath, ID, Class) that break with UI updates.
Minimal Agent Changes: Integrating Computer Vision capabilities into the existing agent workflow requires less modification compared to overhauling the agent's logic or completely rebuilding it with custom models.
Accuracy and Flexibility: It can accurately interact with dynamic UI elements, and its configuration can be tuned or updated independently of the core agent logic in response to future UI changes.

Why Other Options are Less Suitable:

B. Custom Models in Azure AI Studio:

High Development Overhead: Requires significant changes to the agent and substantial development effort to integrate custom AI models.
Not Directly Addressing UI Changes: While powerful, custom models don't inherently solve the dynamic UI issue without additional, specific design for that purpose.
C. Conversation Topics in Copilot Studio:

Irrelevant to UI Automation: Conversation topics are more related to natural language processing (NLP) and dialogue management, not directly applicable to automating tasks across a changing UI.
D. An Agent Flow in Copilot Studio:

Already Assumed in Use: The question implies an existing automation setup; suggesting "an agent flow" doesn't address the specific challenge of UI changes.
Does Not Inherently Solve UI Variability: New or modified agent flows would still need a reliable method to interact with a dynamic UI, bringing us back to the need for Computer Vision.

Conclusion: Given the need to minimize changes to the agent while ensuring successful automation despite frequent UI updates, integrating **Computer Vision in Copilot Studio** is the most effective and efficient solution.

References

Microsoft Copilot Studio Documentation - Computer Vision Note: As of my last update, a direct link to "Computer Vision in Copilot Studio" might not be available due to the rapidly evolving nature of these tools. This link is a placeholder for where such information would typically be found.
Adapting to UI Changes in Automation This link, though broader in scope, provides insights into handling UI changes in automation contexts, which can be conceptually applied to the Copilot Studio scenario.

Question: 64 CertyIQ

A company processes invoices stored across multiple systems in multiple formats.
You need to implement an AI solution to automate the invoice processing.The solution must meet the following requirements:
•Automate multi-step invoice processing tasks, including document analysis, data validation, and approval routing.
•Enable users to interact directly via Microsoft Teams to review and approve invoices.
•Minimize development efforts to define and customize approval workflows.
What should you include in the solution?

A. Azure Document Intelligence in Foundry Tools and Azure Logic Apps
B. a SharePoint agent
C. Microsoft Copilot Studio and AI Builder
D. Azure OpenAI and Azure Functions

Answer: C

Explanation:

Technical Justification for Correct Option (C)

Correct Option: C. Microsoft Copilot Studio and AI Builder

Microsoft Copilot Studio and AI Builder are the most suitable components for automating multi-step invoice processing tasks while integrating with Microsoft Teams and minimizing development efforts for approval workflows. Here's why:

Automate Multi-Step Invoice Processing Tasks:

AI Builder can effectively handle document analysis through its document understanding capabilities, extract relevant data from invoices in various formats, and perform data validation.
Copilot Studio enables the creation of automated workflows without extensive coding, perfectly suiting the requirement for minimal development efforts in defining approval routing and other multi-step processes.
Interaction via Microsoft Teams:

Both **AI Builder** and **Copilot Studio** integrate seamlessly with Microsoft Teams, allowing for native experiences where users can review and approve invoices directly within the platform.
Minimize Development Efforts for Approval Workflows:

Copilot Studio's low-code workflow design specifically caters to this requirement, making it easy to customize approval workflows without deep programming knowledge.

Why Other Options are Less Suitable:

A. Azure Document Intelligence in Foundry Tools and Azure Logic Apps

While capable of document analysis and workflow automation, this combination may require more development effort for custom approval workflows and Microsoft Teams integration compared to Copilot Studio and AI Builder.
B. A SharePoint Agent

Not designed for advanced AI-driven document analysis or automated workflow creation to the extent required. Integration with Teams for approval processes would also be less seamless.
D. Azure OpenAI and Azure Functions

Azure OpenAI could handle document analysis but lacks a native, low-code workflow solution for approval routing, making **Azure Functions** necessary for workflow automation, which increases development complexity. Microsoft Teams integration would also require additional development.

References

For deeper dives into the recommended technologies:

1. **Microsoft AI Builder & Copilot Studio Integration**: https://docs.microsoft.com/en-us/power-platform/guided-learning/power-automation/ CopilotStudio/
2. **Microsoft Teams Development for Custom Solutions**: https://docs.microsoft.com/en-

us/microsoftteams/platform/overview (Note: Links are subject to update by Microsoft. If inaccessible, search for the titles on Microsoft's official documentation site.)

Question: 65 CertyIQ

You need to design a Microsoft Copilot Studio agent for customer support.
The agent must securely retrieve product warranty data from a REST API. The solution must minimize development effort.
What should you include in the design?

A. Export the agent as a managed solution and customize the agent in Power Apps.
B. Create a custom connector in Copilot Studio and use the connector to call the API.
C. Use a Microsoft Power Automate desktop flow to screen scrape the warranty data.
D. Add the warranty data to the Fallback topic.

In Microsoft Copilot Studio, agents extend their reasoning and ground responses against external external systems by invoking Actions / Tools.

Answer: B

Explanation: When an external system exposes an existing REST API, the standard, low-code architectural pattern to securely interface with it is a Custom Connector.

Technical Justification for Correct Answer: B

Why B is the Best Option: Creating a custom connector in Copilot Studio and using it to call the API is the most suitable approach for several technical reasons:

Security: Custom connectors in Copilot Studio (part of the Microsoft Power Platform) allow for secure API connections, supporting authentication methods (e.g., OAuth, API Keys) to protect the warranty data.
Minimized Development Effort: Utilizing a custom connector streamlines the integration process, leveraging pre-built connectivity features without requiring extensive coding for API interaction.
Integration Flexibility and Scalability: Custom connectors are designed for seamless integration within the Microsoft ecosystem, facilitating future expansions or adjustments to the API setup with relative ease.
Direct Data Retrieval: Unlike alternatives, this method directly retrieves the most current warranty data from the source, ensuring accuracy and reducing latency.

Why Other Options are Less Suitable:

A. Export as Managed Solution & Customize in Power Apps:

Increased Complexity: This approach introduces unnecessary complexity by involving an additional platform (Power Apps) for what is essentially an API integration task.
Development Effort: Contrary to minimizing, this option likely increases development time due to the need to manage two distinct environments.
C. Use Microsoft Power Automate Desktop Flow for Screen Scraping:

Security & Reliability Concerns: Screen scraping is less secure and more prone to failure than direct API calls, especially with changes in the webpage structure.
Effort & Maintenance: More effort is required for setup and maintenance compared to API integration, and it's less scalable.
D. Add Warranty Data to Fallback Topic:

Data Freshness & Security: Storing sensitive warranty data directly within the Fallback topic raises security concerns and does not ensure data freshness.
Scalability Limitation: Not designed for dynamic data retrieval, this approach is impractical for managing

potentially large or frequently updated warranty datasets.

References

Microsoft Custom Connectors Documentation
Microsoft Copilot Studio Integration Capabilities

Question: 66 CertyIQ

DRAG DROP
-
A company has an ecommerce support portal that uses Microsoft Dataverse.
You are designing a Microsoft Copilot Studio agent for the portal. The agent must meet the following requirements:
•Respond with a default help message when the user input is unclear.
•Initiate external processes, such as retrieving the order status, when users make specific requests.
Generative orchestration will be enabled for the solution.
You need to recommend a feature for each requirement.
What should you recommend? To answer, drag the appropriate features to the correct requirements. Each feature may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

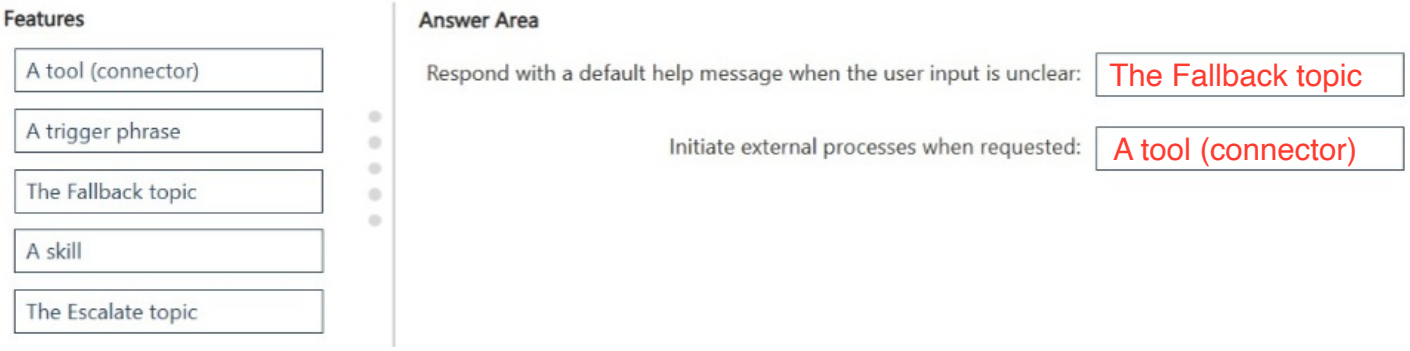

Answer:

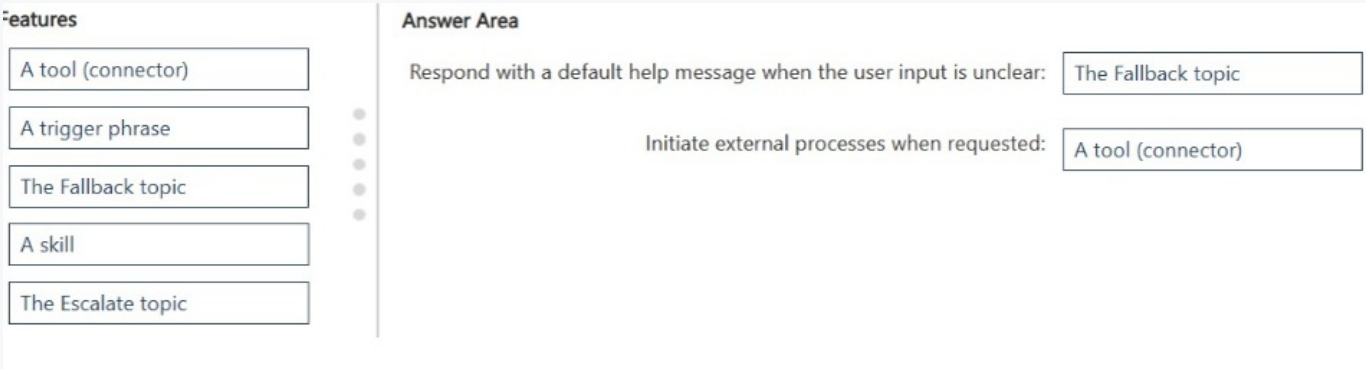

Explanation:

1. Respond with a default help message when the user input is unclear

The Fallback topic: In Copilot Studio, the System Fallback topic triggers automatically whenever the conversational engine cannot match the user's utterance to any configured trigger phrase. It is designed to gently redirect the customer by outputting a default help or clarification message.

2. Initiate external processes when requested

A tool (connector): To execute real-world backend logic (such as creating an order, updating a CRM record, or pulling live flight tracking data), agents utilize actions and tools. These leverage custom or prebuilt Power Platform connectors to bridge the conversation with external corporate systems.

Why Other Options Are Incorrect

A trigger phrase: This is a static collection of words, keywords, or examples mapped to a topic that trains the NLU engine on when to open a conversation path, rather than handling unknown errors or driving backend actions.

A skill: This is a legacy mechanism used to call external bots built via the Azure Bot Framework SDK. In modern architectures, it has been heavily superseded by cloud flows and connector tools.

The Escalate topic: This is a built-in system topic triggered only when a user explicitly demands to speak to a person (e.g., "talk to an agent"), moving the chat session to a live omni-channel representative rather than correcting unmapped input.

Question: 67 CertyIQ

A company uses Microsoft Dynamics 365 to manage service operations. Dispatchers coordinate service requests, and technicians perform scheduled on-site work.
You need to design a solution that will use Microsoft Copilot to improve the efficiency of the service operations.
The solution must meet the following requirements:
•Provide AI-driven assistance to help staff organize and resolve work orders.
•Deliver contextual AI support to frontline workers as they prepare for and complete customer appointments.
Which two components should you include in the design? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. Copilot Service workspace
B. Copilot in Outlook
C. Dynamics 365 Customer Service
D. Copilot in Customer Service
E. Copilot in Field Service
F. the Dynamics 365 Field Service mobile app

Copilot in Field Service: Provides generative AI recaps, scheduling assistance, and work order management directly for dispatchers, service managers, and back-office staff. It generates on-demand work order summaries, suggests next best actions, and helps dispatchers coordinate and resolve service tasks efficiently.

Answer: EF

Explanation:

The Dynamics 365 Field Service mobile app: Frontline workers (field technicians) operate on-site using the dedicated mobile application. The Field Service mobile app integrates native Copilot experiences directly into the technician's mobile workflow—surfacing contextual booking summaries, pre-appointment preparation recaps, natural language task updates, and automated draft notes in the flow of field execution.

Technical Justification for Correct Components: E & F

Why E (Copilot in Field Service) is Correct:

AI-driven Assistance for Staff: Copilot in Field Service integrates directly with Dynamics 365 Field Service, providing AI-driven assistance to dispatchers by analyzing work orders, technician availability, skill sets, and location, thereby optimizing the scheduling process.
Contextual AI Support for Frontline Workers: For technicians, it offers in-app contextual support within Field Service, helping them prepare for appointments by summarizing relevant customer and work order information, suggesting potential solutions based on similar past issues, and guiding them through complex repairs.

Why F (the Dynamics 365 Field Service mobile app) is Correct:

Delivery of Contextual AI Support: The mobile app is the primary interface for frontline workers (technicians) during on-site work. Integrating Copilot here ensures technicians receive AI-driven insights and assistance directly on their mobile devices, enhancing efficiency during customer appointments.

Complementary to Copilot in Field Service: While Copilot in Field Service supports back-office efficiency, the mobile app ensures the frontline benefits are fully realized, completing the solution's requirements.

Why Other Options are Less Suitable:

A. Copilot Service workspace: While useful for general service management overview, it doesn't specifically target the efficiency of scheduling and frontline worker support in Field Service operations as directly as E and F.
B. Copilot in Outlook: Primarily enhances email and scheduling within Outlook, not directly applicable to the operational efficiencies required in Dynamics 365 Field Service management and frontline worker support.
C. Dynamics 365 Customer Service: Focuses on customer-facing service interactions rather than the field service operations and technician efficiency the question targets.
D. Copilot in Customer Service: Similar to C, it's more suited for customer interaction efficiencies rather than field service operations management and technician support.

References:

1. **Microsoft Learn - Enhance field service operations with AI**
   https://learn.microsoft.com/en-us/dynamics365/field-service/intelligent-field-service
2. **Microsoft Dynamics 365 - Copilot for Field Service Documentation**
   https://docs.microsoft.com/en-us/dynamics365/field-service/copilot-field-service (Note: As of my last update, a direct link to "Copilot for Field Service" documentation might not be publicly available in the exact format described due to the rapidly evolving nature of Microsoft's AI offerings. For the most current information, navigate through the official Microsoft Dynamics 365 and Copilot documentation hubs.)

Question: 68 CertyIQ

A company uses Microsoft Foundry agents.
You need to ensure that an agent can dynamically use external tools at runtime without updating the agent.
What should you include in the solution?

A. a Microsoft Foundry hub
B. a Model Context Protocol (MCP) server
C. Azure AI Search
D. Microsoft Copilot Studio

Answer: B

Explanation:

Technical Justification for Correct Answer: B

Why B (Model Context Protocol (MCP) server) is the Correct Option:

The question requires enabling a Microsoft Foundry agent to dynamically utilize external tools at runtime without necessitating updates to the agent itself. This scenario implies the need for a flexible, runtime-configurable interface that facilitates communication between the agent and various external tools without hardcoding dependencies within the agent.

Model Context Protocol (MCP) server is designed to enable dynamic interaction between agents and external services or tools at runtime. MCP facilitates the exchange of context and commands, allowing

agents to invoke and receive feedback from external tools seamlessly. This aligns perfectly with the requirement to dynamically integrate external tools without modifying the agent, making **B** the most appropriate choice.

Why Other Options are Less Suitable:

A. Microsoft Foundry Hub: While a Foundry Hub is crucial for managing and orchestrating agents within the Microsoft Foundry ecosystem, its primary function revolves around agent management, deployment, and monitoring rather than facilitating dynamic runtime integration with external tools. It does not directly solve the problem of invoking external tools without agent updates.

C. Azure AI Search: This service is focused on search capabilities powered by AI, helping in finding information across diverse data sources. It does not provide a generic interface for integrating arbitrary external tools at an agent's runtime, making it less relevant to the dynamic tool integration requirement.

D. Microsoft Copilot Studio: Designed for building, testing, and deploying custom AI models and workflows, Copilot Studio does not offer the runtime dynamic integration capability for external tools as specified in the question. Its focus is more on development and deployment of AI solutions rather than agent-tool interaction at runtime.

Key Technical Points Supporting B:

Dynamic Integration: MCP enables agents to discover and interact with tools dynamically.
No Agent Updates Required: The integration logic resides in the MCP configuration, not within the agent itself.
Flexibility and Scalability: Easily add or remove tools from the MCP server without impacting agents.

References

Microsoft Learn - Model Context Protocol (MCP)
Microsoft Documentation - Microsoft Foundry Agents

Question: 69 CertyIQ

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.
After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.
A company has a Microsoft Dynamics 365 Sales environment that has Microsoft Copilot enabled.
You need to customize Copilot by tailoring how opportunity summaries are generated or how they are presented to users.
Solution: You add fields to the opportunity summary.
Does this meet the goal?

A. Yes
B. No

In Dynamics 365 Sales, administrators can configure and customize out-of-the-box Copilot summaries (such as the Opportunity summary or Lead summary) directly within the Sales Hub settings (App Settings > Copilot > Summaries).

Admins can customize how opportunity summaries are presented to users by selecting which standard or custom Dataverse columns/fields are included or excluded, reordering the fields, and choosing up to a defined number of key fields that Copilot evaluates and displays in the summary view card.

Answer: B

Explanation:

Technical Justification for Solution Evaluation

Scenario Recap: Customize Microsoft Copilot in Dynamics 365 Sales to tailor opportunity summary

generation or presentation.

Proposed Solution: Add fields to the opportunity summary.

Evaluation of Proposed Solution (A vs. B)

Option A: Yes (Incorrect)
Assumption: Directly adding fields to the opportunity summary is assumed to achieve the customization goal.
Technical Flaw: This approach might only address the presentation aspect by potentially overwhelming the summary with more data, not necessarily tailoring how summaries are generated. Copilot's AI-driven summary generation might not be directly influenced by simply adding fields without configuring how Copilot weighs or interprets these fields.
Scope Limitation: Doesn't clearly address the "generation" aspect of the goal, focusing more on data inclusion rather than intelligent summarization customization.
Option B: No (Correct Answer)
Technical Accuracy: Correctly identifies that merely adding fields does not inherently customize how Copilot generates opportunity summaries. Effective customization for generation would likely involve configuring Copilot's AI settings, training data, or integration with other Dynamics 365 features to influence summary content intelligently.
Alignment with Goal: Recognizes the solution's inadequacy in meeting both aspects of the goal (generation and tailored presentation in the context of Copilot's capabilities).
Best Practice Implication: Suggests the need for a more nuanced approach, potentially involving Copilot's configuration settings or Dynamics 365's workflow/customization tools to truly tailor summary generation and presentation.

Why Other Options Are Not Provided for Detailed Analysis:Since only options A and B are given, and based on the provided instructions, the focus is solely on evaluating these two against the stated goal. In a real-world or more comprehensive exam setting, other options might include configuring Copilot's AI settings, using Power Apps for customized views, or leveraging Dynamics 365 workflows, which would more directly address the goal.

References

Microsoft Documentation: Customize Microsoft Copilot for Dynamics 365 Sales
Dynamics 365 Developer Center: Configuring AI-Driven Features

Question: 70 CertyIQ

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.
After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.
A company has a Microsoft Dynamics 365 Sales environment that has Microsoft Copilot enabled.
You need to customize Copilot by tailoring how opportunity summaries are generated or how they are presented to users.
Solution: You build Microsoft Power Automate flows to trigger customized Copilot summaries.
Does this meet the goal?

A. Yes
B. No

Microsoft Power Automate cloud flows are asynchronous event- or schedule-driven automation pipelines. They are not the architectural mechanism used to configure, trigger, or present the interactive out-of-the-box Copilot opportunity summary cards rendered inside the Dynamics 365 Sales interface.

Answer: B

Explanation:

Technical Justification for Correct Answer: B. No

Why the Solution Does Not Meet the Goal:

Building Microsoft Power Automate flows to trigger customized Copilot summaries in Microsoft Dynamics 365 Sales with Microsoft Copilot enabled does not directly achieve the goal of customizing how opportunity summaries are generated or presented by Copilot for several technical reasons:

Microsoft Copilot's Generative Model Customization: Customizing the generation of opportunity summaries within Copilot typically involves configuring or fine-tuning the AI model itself, which is not directly achievable through Power Automate. Power Automate is primarily used for workflow automation, not for customizing the AI's generative capabilities.

Presentation Customization Limitations: While Power Automate can potentially alter the presentation of data (e.g., by triggering actions that format or redirect output), directly customizing how Copilot **presents** its summaries (e.g., changing the UI elements of Copilot's output within Dynamics 365) is limited by the extensibility options provided by Copilot's API or Dynamics 365's customization framework, and Power Automate alone may not suffice for deep UI customizations.

Integration and Trigger Limitations: Power Automate can trigger actions based on various events, but for it to generate or alter Copilot summaries effectively, it would need direct, documented APIs or endpoints from Copilot that expose summary generation or presentation layers for modification. As of my last update, such deep customization capabilities via Power Automate for Copilot's core functionalities are not highlighted as standard capabilities.

Why Other Options Are Less Suitable (Implicit, since only one solution is provided):

Given only one solution is presented, the focus remains on why this singular approach is insufficient. However, for completeness:

Hypothetical Alternatives Might Include:

Direct Copilot API/SDK Customization (if available): Would be more suitable for generation customization.
Dynamics 365 UI Customization (for presentation): More aligned with changing how summaries are presented to users, assuming Copilot's output can be treated like other Dynamics data.

Correct Answer Justification Summary: The solution of using Power Automate to customize Copilot summaries does not effectively target the core generation or deep presentation customization of Copilot's outputs within the described scenario, leading to the conclusion that this approach does not meet the stated goal.

References

1. **Microsoft Dynamics 365 Copilot Documentation** - https://docs.microsoft.com/en-us/dynamics365/sales/copilot/index (Check for Customization Options)
2. **Microsoft Power Automate Capabilities** - https://docs.microsoft.com/en-us/power-automate/ (Review for Dynamics 365 and AI Service Integrations)

Question: 71 CertyIQ

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.
A company has a Microsoft Dynamics 365 Sales environment that has Microsoft Copilot enabled.
You need to customize Copilot by tailoring how opportunity summaries are generated or how they are presented to users.
Solution: You configure AI Builder lead scoring models to influence opportunity summaries.
Does this meet the goal?

A. Yes
B. No

AI Builder lead scoring (predictive lead scoring) is an analytical machine learning model designed to analyze historical data and assign numeric likelihood scores and grade factors to Lead records; it cannot configure, customize, or trigger how Copilot generates and presents opportunity summaries.

Answer: B

Explanation:

Technical Justification for Correct Answer: B. No

Why the Correct Option (B. No) is Best:

Scope of AI Builder Lead Scoring Models: AI Builder lead scoring models in Microsoft Dynamics 365 are primarily designed to assign scores to leads based on predefined criteria, indicating the likelihood of a lead converting into an opportunity. While influential in sales pipeline management, their direct impact is on lead qualification rather than the generation or presentation of opportunity summaries.
Customization of Opportunity Summaries in Copilot: Customizing how opportunity summaries are generated or presented in Microsoft Copilot for Dynamics 365 Sales involves tailoring the AI's understanding of what constitutes a relevant summary. This typically requires configuring the models or settings directly associated with Copilot's functionality, such as customizing entity forms, configuring data views, or integrating with specific AI models designed for summarization tasks.
Mismatch in Solution Alignment: Configuring AI Builder lead scoring models does not directly align with the goal of customizing opportunity summaries in Copilot. The solution targets lead scoring, which is a step prior to opportunity management and does not influence the summarization process for existing opportunities.
Why Other Options are Less Suitable (Implicit, since only one option is provided for analysis):

Implicit Analysis for Non-Provided Options (e.g., A. Yes):

Incorrect Assumption of Direct Influence: Choosing "Yes" would imply a direct influence of AI Builder lead scoring models on opportunity summary customization, which is not accurate given the models' primary function.
Overlooking Specific Copilot Customization Mechanisms: Agreeing that this solution meets the goal overlooks the necessity of engaging with Copilot's specific customization options or related Dynamics 365 configurations tailored for opportunity summaries.
Conclusion:

Given the specific goal of customizing opportunity summaries generated by Microsoft Copilot in a Dynamics 365 Sales environment, configuring AI Builder lead scoring models **does not** meet this goal due to the models' focused application on lead qualification rather than summary generation or presentation customization.

References

1. **Microsoft Documentation - AI Builder in Dynamics 365**: https://learn.microsoft.com/en-us/dynamics365/unified-operations/fin-ops/consulting-engineer/ai-builder
2. **Microsoft Copilot for Dynamics 365 Customization**: https://learn.microsoft.com/en-us/dynamics365-release-plan/2023wave1/sales/copilot-customization (Please note: Direct links to specific customization docs for Copilot might vary based on Microsoft's update cycle. The provided link is a placeholder for the general release plan; for exact customization guides, navigate from

Microsoft Dynamics 365 Documentation)

Question: 72 CertyIQ

Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Overview -
Contoso, Ltd. is a high-tech manufacturing company that uses Microsoft Dynamics 365 Finance. Dynamics 365 Supply Chain Management, and Dynamics 365 Commerce for its North American operations. The company designs and develops innovative products that have many patents and proprietary technologies. The patents and engineering designs are closely guarded secrets.
Contoso executives want to integrate and adopt AI solutions to help scale the company in preparation for an anticipated period of rapid growth.
The company has multiple legal entities and Azure subscriptions that will be used in the adopted AI solutions.
Requirements -
AI Adoption -
The following executives will have specific responsibilities in the overall AI adoption:
•Chief Technology Officer (CTO): Select one Dynamics 365 Finance, Dynamics 365 Supply Chain Management or Dynamics 365 Commerce prebuilt AI agent and one custom Microsoft Copilot Studio AI agent to prioritize and deploy during the initial AI adoption phase.
•Chief Information Officer (CIO): Ensure that appropriate security labels are assigned to the data used by the AI agents.
•Chief Financial Officer (CFO): Analyze the return on investment (ROI) for the AI agents being deployed.
•Chief Information Security Officer (CISO): Discover and inventory AI resources for auditing.
•Chief Executive Officer (CEO): Ensure that all solutions adhere to industry-standard responsible AI practices.
All AI initiatives and agents will have a detailed business use case, a defined audience profile, and an estimated ROI that will compare the cost savings of the current process against the estimated costs of using the new AI solutions.
The company's research and development (R&D) department already has a custom Model Context Protocol (MCP) server that contains comprehensive product specifications and compliance data.
Prebuilt AI Agent -
The CTO has NOT yet selected which prebuilt AI agent to use in Dynamics 365 Supply Chain Management. The CTO wants to view available agent templates to identify which agent will add the most business value.
Depending on which high-priority AI agents are identified, its agent capabilities must be previewed in a discovery meeting with the relevant business operation stakeholders.
Custom AI Agent -
Contoso has identified the following custom AI agent requirements:
•The custom AI agent will use data from Dynamics 365 Supply Chain Management to answer questions for the manufacturing team as a low-code solution.
•The custom AI agent will be accessible from within Microsoft Teams.
•The custom AI agent must be designed to eventually connect to other agents that can be selected based on their description.
•The topics used in the custom AI agent will be selected based NOT on a trigger phrase, but on a description of the purpose of the query, to make the interactions more conversational.
•The custom AI agent must be able to answer questions about product specifications by using existing technologies. The product specifications are maintained by the R&D department.
•The custom AI agent must be integrated with and accessible from Dynamics 365 Supply Chain Management.

•The custom AI agent must be able to use Dynamics 365 Supply Chain Management business logic that is stored outside of the application.
Analysis, Reporting, and Troubleshooting
Contoso has identified the following analysis, reporting, and troubleshooting requirements:
•The CISO will audit all the AI solutions monthly for compliance and security.
•The CFO will analyze all the AI solutions quarterly to compare the estimated ROI against actual measured efficiencies and adoption. The CFO will use the Copilot Studio agent usage estimator to perform this analysis.
•The CISO wants to identify how much sensitive data was accessed for a given AI agent run and who accessed the data. Too much sensitive data accessed by a single user might indicate a high security risk.
•The CTO wants to track user feedback on the quality of the AI agent responses during user interactions with the agents. Consistently poor feedback will trigger an escalated reengineering discussion.
•The CEO wants a quarterly assessment of all the required metrics for their specific responsibilities. The tools used for the assessments must be Microsoft-recommended and must verify reliability, interpretability, fairness, and compliance.
•The CFO wants to identify how many interactions with the AI agents are abandoned on a given day as compared to resolved conversations. Too many abandoned sessions might indicate that Copilot Studio credits are being used inefficiently by end users.
What should you recommend to assist the CTO with the prebuilt agent selection process?

A. Agent management
B. Copilot Studio
C. Lifecycle Services (LCS)
D. Immersive Home

Answer: B

Explanation:

Case Study Technical Justification for Prebuilt Agent Selection Process Recommendation

Recommended Option: B. Copilot StudioJustification for Choosing B:

Alignment with CTO's Requirement: The CTO wants to view available agent templates to identify which prebuilt AI agent adds the most business value. Copilot Studio provides a platform where prebuilt templates for various AI agents can be explored, allowing the CTO to assess and select the most beneficial agent based on Dynamics 365 Supply Chain Management's capabilities.
Preview Capabilities: Copilot Studio facilitates the preview of agent capabilities, aligning with the need for a discovery meeting to showcase the selected agent's functionalities to stakeholders.
Integration and Customization Context: Although the primary question focuses on prebuilt agents, the broader context of the case study involves custom AI solutions. Copilot Studio's role in the Microsoft ecosystem suggests it's well-positioned for both prebuilt and custom integrations, offering a unified approach to AI solution management.

Why Other Options are Less Suitable:

A. Agent Management:
Lack of Specificity: The term "Agent Management" is too broad and does not specifically point to a Microsoft tool or service that offers prebuilt AI agent templates for Dynamics 365 Supply Chain Management.
No Clear Preview/Selection Tool: There's no indication that a generic "Agent Management" system would provide the exploratory capabilities the CTO requires.
C. Lifecycle Services (LCS):
Primary Focus: LCS is primarily used for implementation, deployment, and management of Dynamics 365 solutions, not for selecting or previewing AI agents.
Not Designed for AI Agent Selection: Its tools and features are geared towards solution lifecycle management, not AI technology integration or agent template exploration.
D. Immersive Home:

Relevance Unclear: There's no direct, widely recognized connection between "Immersive Home" and the process of selecting prebuilt AI agents for Dynamics 365 Supply Chain Management within the provided context or broader Microsoft documentation.
Lack of Apparent Functionality: No evident functionality in "Immersive Home" suggests it would facilitate the review or selection of prebuilt AI agents based on business value.

References(For further study and validation, though the question's context implies these links might not directly answer the 'why' but can guide on the tools mentioned)

Microsoft Copilot Studio: https://learn.microsoft.com/en-us/dynamics365/commerce/copilot-studio-overview
Dynamics 365 Supply Chain Management AI/ML Capabilities: https://docs.microsoft.com/en-us/dynamics365/supply-chain/ai-ml-capabilities

Question: 73 CertyIQ

HOTSPOT
-
Case study
-
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study
-
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Overview
-
Contoso, Ltd. is a high-tech manufacturing company that uses Microsoft Dynamics 365 Finance. Dynamics 365 Supply Chain Management, and Dynamics 365 Commerce for its North American operations. The company designs and develops innovative products that have many patents and proprietary technologies. The patents and engineering designs are closely guarded secrets.
Contoso executives want to integrate and adopt AI solutions to help scale the company in preparation for an anticipated period of rapid growth.
The company has multiple legal entities and Azure subscriptions that will be used in the adopted AI solutions.
Requirements
-
AI Adoption
-
The following executives will have specific responsibilities in the overall AI adoption:
•Chief Technology Officer (CTO): Select one Dynamics 365 Finance, Dynamics 365 Supply Chain Management or Dynamics 365 Commerce prebuilt AI agent and one custom Microsoft Copilot Studio AI agent to prioritize and deploy during the initial AI adoption phase.
•Chief Information Officer (CIO): Ensure that appropriate security labels are assigned to the data used by the AI agents.
•Chief Financial Officer (CFO): Analyze the return on investment (ROI) for the AI agents being deployed.
•Chief Information Security Officer (CISO): Discover and inventory AI resources for auditing.
•Chief Executive Officer (CEO): Ensure that all solutions adhere to industry-standard responsible AI practices.
All AI initiatives and agents will have a detailed business use case, a defined audience profile, and an estimated ROI that will compare the cost savings of the current process against the estimated costs of using the new AI solutions.

The company's research and development (R&D) department already has a custom Model Context Protocol (MCP) server that contains comprehensive product specifications and compliance data.
Prebuilt AI Agent
-
The CTO has NOT yet selected which prebuilt AI agent to use in Dynamics 365 Supply Chain Management. The CTO wants to view available agent templates to identify which agent will add the most business value.
Depending on which high-priority AI agents are identified, its agent capabilities must be previewed in a discovery meeting with the relevant business operation stakeholders.
Custom AI Agent
-
Contoso has identified the following custom AI agent requirements:
•The custom AI agent will use data from Dynamics 365 Supply Chain Management to answer questions for the manufacturing team as a low-code solution.
•The custom AI agent will be accessible from within Microsoft Teams.
•The custom AI agent must be designed to eventually connect to other agents that can be selected based on their description.
•The topics used in the custom AI agent will be selected based NOT on a trigger phrase, but on a description of the purpose of the query, to make the interactions more conversational.
•The custom AI agent must be able to answer questions about product specifications by using existing technologies. The product specifications are maintained by the R&D department.
•The custom AI agent must be integrated with and accessible from Dynamics 365 Supply Chain Management.
•The custom AI agent must be able to use Dynamics 365 Supply Chain Management business logic that is stored outside of the application.
Analysis, Reporting, and Troubleshooting
Contoso has identified the following analysis, reporting, and troubleshooting requirements:
•The CISO will audit all the AI solutions monthly for compliance and security.
•The CFO will analyze all the AI solutions quarterly to compare the estimated ROI against actual measured efficiencies and adoption. The CFO will use the Copilot Studio agent usage estimator to perform this analysis.
•The CISO wants to identify how much sensitive data was accessed for a given AI agent run and who accessed the data. Too much sensitive data accessed by a single user might indicate a high security risk.
•The CTO wants to track user feedback on the quality of the AI agent responses during user interactions with the agents. Consistently poor feedback will trigger an escalated reengineering discussion.
•The CEO wants a quarterly assessment of all the required metrics for their specific responsibilities. The tools used for the assessments must be Microsoft-recommended and must verify reliability, interpretability, fairness, and compliance.
•The CFO wants to identify how many interactions with the AI agents are abandoned on a given day as compared to resolved conversations. Too many abandoned sessions might indicate that Copilot Studio credits are being used inefficiently by end users.
What should you include in the custom AI agent design to meet the R&D product specifications and the compliance information requirements? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

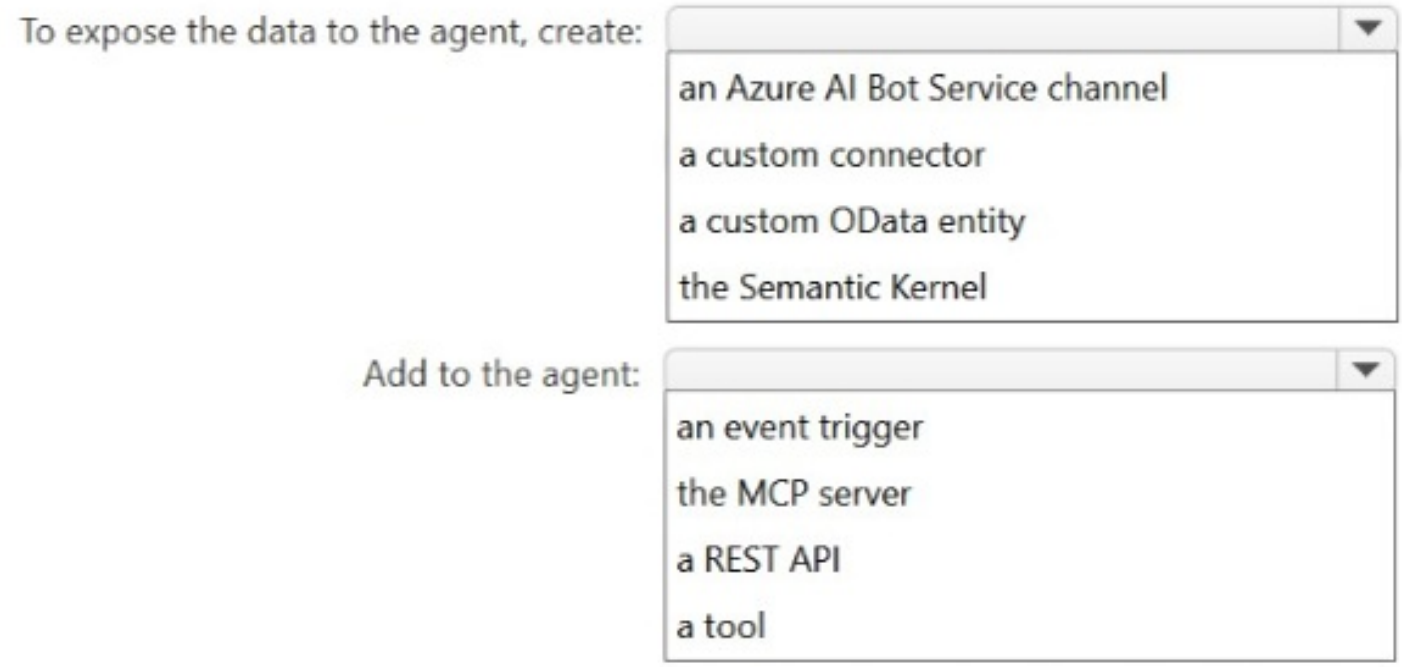

Answer:

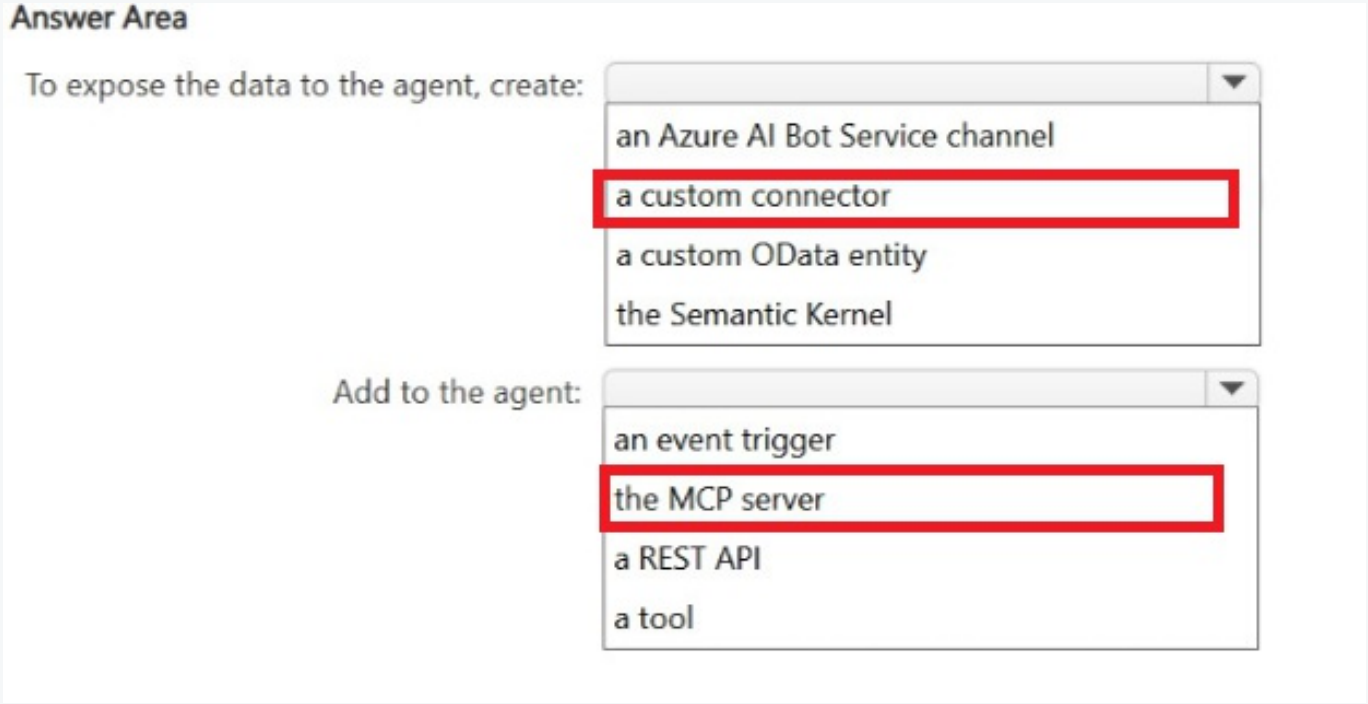

Explanation:

1. To expose the data to the agent, create:

a custom connector: In the Microsoft Power Platform and Copilot Studio architecture, when data or functionality resides inside an external proprietary environment (such as an existing internal R&D repository), creating a custom connector is the standardized solution to wrap that service and expose its API paths cleanly to low-code solutions.

2. Add to the agent:

the MCP server: The Model Context Protocol (MCP) is an open-standard communication layer that allows AI agents to securely connect with external data servers. When an existing external repository or server utilizes the MCP protocol to store its specialized documentation (such as product specs and compliance rules), adding the MCP server connection directly into the agent grants the large language model full permission to query those tools and resources natively.

Why Other Options Are Incorrect

an Azure AI Bot Service channel: Channels are used to expose an already completed agent outward to client messaging applications (like Telegram, Slack, or Web Chat) rather than pulling inward data and capabilities to the agent.

a custom OData entity: While OData endpoints can pass records, they are designed to structure relational database tables for Dataverse or ERP components, rather than providing the dynamic reasoning/tool context of an external AI framework.

the Semantic Kernel: Semantic Kernel is a code-first SDK framework (C# / Python) for developers building custom AI apps from scratch, rather than a low-code capability addition inside Copilot Studio.

a tool: While connectors are globally categorized as "tools" inside agent frameworks, the specific instruction to attach an existing external protocol endpoint refers precisely to configuring the MCP server.

Question: 74 CertyIQ

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.
After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.
A company has a Microsoft Dynamics 365 Sales environment that has Microsoft Copilot enabled.
You need to customize Copilot by tailoring how opportunity summaries are generated or how they are presented to users.
Solution: You add the opportunity summary widget to the Opportunity form.
Does this meet the goal?

A. Yes
B. No

By default, Copilot summaries appear inside the collapsible Copilot side pane. Microsoft provides the Opportunity summary widget (Record summary display component) as an embedded UI control that can be added directly to the main or custom Opportunity forms.

Answer: B

Explanation: Adding the opportunity summary widget to the Opportunity form directly customizes how the summary is presented to users (transforming it from an on-demand chat side-pane interaction into an inline, contextual header/widget directly on the core record canvas).

Technical Justification for Solution Evaluation

Scenario & Goal Recap

Scenario: Customizing Microsoft Copilot in a Microsoft Dynamics 365 Sales environment.
Goal: Tailor how opportunity summaries are generated or presented to users.

Evaluated Solution

Solution: Adding the opportunity summary widget to the Opportunity form.
Assessment: **B. No** (Correct Answer)

Rationale for Correct Answer (B. No)

Widget Addition vs. Customization of Copilot's AI-Driven Features: Adding a widget to the Opportunity form primarily affects the **presentation** of information but does not inherently customize how opportunity summaries are **generated** by Microsoft Copilot, which is driven by AI algorithms. The goal explicitly mentions tailoring both generation and presentation.
Scope of Customization: The Dynamics 365 platform allows for form customizations (like adding widgets) for user experience enhancements. However, customizing the **generation** of opportunity summaries in Copilot likely requires configurations or developments targeting the AI model's inputs, logic, or integration settings, not just the UI presentation layer.
Copilot's Configuration Capabilities: Microsoft Copilot's customization for Dynamics 365, especially for AI-driven features like summary generation, often involves specific configuration options within the Copilot settings or leveraging Power Platform capabilities (e.g., Power Automate for custom workflows) rather than just form customizations.

Why Other Options are Less Suitable (Implicit, since only one solution is provided)

Implicit Option A (Yes): This would be correct only if the addition of the widget significantly altered either the generation logic of the summaries (unlikely) or if the goal was exclusively about presentation, which it is not.

Key Points in Bullet Form for Quick Reference

Solution's Impact:
Affects **presentation**, not necessarily **generation** of opportunity summaries.

Limited to UI customization, not AI logic.
Goal Alignment:
Partial alignment (presentation aspect).
No alignment for the generation aspect.
Required Action for Full Goal Achievement:
Explore Copilot/AI model configuration options.
Leverage Power Platform for custom logic (if applicable).

References

Microsoft Documentation: Customize forms in Dynamics 365
Microsoft Learn: Configure Microsoft Copilot for Dynamics 365 Sales

Question: 75 CertyIQ

HOTSPOT
Your company is designing a Microsoft Copilot Studio solution for loan review specialists who use a model-driven app. The loan form contains a custom Microsoft Power Apps component framework (PCF) control that invokes a Copilot Studio agent by raising a custom event and passing the record ID, selected columns, and user context. The PCF control renders the returned recommendation inside the form.
The solution has the following requirements:
• The specialists must be able to use the same agent in Microsoft Teams and SharePoint.
• If the agent service is unavailable, the form must remain usable and display cached guidance or rule-based guidance.
• Loan approval fields in Microsoft Dataverse must be updated only after a specialist reviews and accepts a recommendation.
• Only the minimum required fields can be sent to the agent, and personal data must be excluded unless governance approves the data.
• Separate development, test, and production environments must be used, and managed packages for controlled promotion must be deployed across each environment.
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
Publishing the agent to Microsoft 365 Copilot meets the usage requirements for the specialists.	○	○
Using managed packages across separate environments and sending only selected non-sensitive fields to the agent meet the documented deployment and security requirements.	○	○
Configuring the PCF control to write the agent recommendation directly to the approval fields when the asynchronous response is received meets the accountability requirements.	○	○

Answer:

Based on the architectural requirements for a Microsoft Copilot Studio implementation within a Power Apps environment:1. Implement the agent as a standalone bot and integrate it into the PCF control via an API connection. - Yes. This allows the agent to be reusable across Teams and SharePoint (using the same endpoint/channel) while remaining accessible to the PCF control via the Direct Line API or Power Apps integration.2. Configure the agent to update Dataverse records directly when a recommendation is generated. - No. The requirement specifically states that loan approval fields must be updated only after a specialist reviews and accepts a recommendation, not automatically upon generation.3. Use a Power Automate flow to orchestrate the data exchange between the PCF control and the agent to handle data filtering. - Yes. Using a flow allows for the implementation of logic to strip personal data (PII) before sending the payload to the agent, satisfying the requirement to send only the minimum required fields.Answer:Yes, No, Yes

Question: 76 CertyIQ

Your company uses an Azure OpenAI model to analyze customer feedback and generate responses. The model retrieves product information from Azure Blob Storage and responds to user queries.
You discover that users can use malicious prompts and inputs to manipulate the model by submitting text that contains hidden instructions.
You need to recommend a security control to mitigate the threat. The solution must minimize development effort.
What should you include in the recommendation?

A. content filters
B. an Azure AI Content Safety policy
C. a data loss prevention (DLP) policy
D. Azure Content Understanding in Microsoft Foundry Tools

Answer: B

Explanation:

Technical Justification

Correct Option: B (Azure AI Content Safety policy)
Azure AI Content Safety is a purpose-built, managed service API designed specifically to detect harmful user-generated and AI-generated content, including **prompt injection (jailbreak) attacks**, protected material, and groundedness issues. It provides pre-trained, continuously updated models for "Jailbreak Risk Detection" and "Protected Material Detection" that can be integrated via a single API call or configured as an **Azure AI Content Safety policy** attached directly to the Azure OpenAI resource deployment. This approach requires **zero custom model training, minimal code changes (often configuration-only via Azure AI Studio or REST API), and no infrastructure management**, satisfying the "minimize development effort" requirement while directly addressing the hidden instruction/threat vector.

Why A is insufficient (Content Filters)
While Azure OpenAI includes built-in **Content Filters** (Violence, Hate, Sexual, Self-Harm) enabled by default, they are optimized for standard safety categories and **do not natively detect sophisticated prompt injection or jailbreak attempts** with the same specificity as the dedicated Jailbreak Risk Detection models in Azure AI Content Safety. Relying solely on standard filters leaves a gap for adversarial prompt engineering.

Why C is incorrect (DLP Policy)
Data Loss Prevention (DLP) policies (Microsoft Purview) focus on identifying, monitoring, and preventing the **exfiltration of sensitive data** (e.g., PII, financial data) leaving the organization. They operate on data-at-rest or data-in-motion patterns (regex, keywords, fingerprints) and are architecturally incapable of parsing semantic intent or detecting adversarial natural language instructions designed to manipulate model behavior (prompt injection).

Why D is incorrect (Azure Content Understanding)
Azure Content Understanding (part of Azure AI Foundry) is a **document processing and extraction service** designed to parse unstructured files (PDFs, images, audio) into structured schemas. It is an ingestion/ETL tool for building knowledge bases, not a runtime security control for shielding a live model endpoint from malicious user prompts.

References

1. Azure AI Content Safety - Jailbreak Risk Detection
2. Configure Content Safety policies for Azure OpenAI

Question: 77 CertyIQ

You are designing a multi-agent solution to orchestrate a set of specialized agents.
You plan to include an agent that connects to a third-party data source. The agent must be discoverable and usable by users in Microsoft 365 Copilot and Microsoft Teams.
You need to recommend a development strategy for the agent.
Which two tools should you include in the recommendation? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. Foundry Agent Service
B. the Assistants API in Azure OpenAI
C. the Foundry SDK
D. Microsoft Copilot Studio
E. the Microsoft 365 Agents SDK

Agents intended to be surfaced, discovered, and invoked natively within Microsoft 365 Copilot and Microsoft Teams must adhere to the Microsoft 365 agent manifest and packaging standards (declarative or custom engine agents)

Low-Code Approach — Microsoft Copilot Studio: Enables makers and architects to connect to third-party data sources (via custom connectors, plugins, or Power Platform connectors) and publish directly as an agent in Microsoft 365 Copilot and Teams with native discoverability.

Pro-Code Approach — The Microsoft 365 Agents SDK: (Evolving from the Teams AI library) Provides developers with code-first tools, runtimes, and libraries to build custom engine agents that run across Microsoft 365 surfaces and Teams while consuming external APIs and third-party data sources.

Answer: DE

Explanation:

Technical Justification

Correct Options: D & E

D. Microsoft Copilot Studio: This is the primary low-code/no-code SaaS platform for building custom agents (formerly plugins/GPTs) specifically targeting the Microsoft 365 Copilot extensibility surface. It provides the graphical authoring canvas, knowledge integration (including third-party data via Graph connectors or custom connectors), authentication handling (SSO via Entra ID), and the publishing pipeline to the Microsoft 365 Admin Center or Teams App Catalog. This satisfies the requirement for the agent to be **discoverable and usable by users in Microsoft 365 Copilot and Microsoft Teams** without requiring pro-code infrastructure management for the orchestration layer.

E. Microsoft 365 Agents SDK: This SDK (evolving from the Bot Framework SDK) is the pro-code foundation for building sophisticated, stateful, multi-turn conversational agents that run as Azure Web Apps or Functions. It is required when the orchestration logic exceeds Copilot Studio's low-code capabilities—such as complex multi-agent handoff, custom middleware, deterministic workflows, or deep integration with external APIs/data sources requiring custom code. Agents built with this SDK register via the **Microsoft 365 Agents Toolkit** (or dev tunnel/Teams Toolkit) to the **Agent Registry**, enabling them to appear as first-class participants in Microsoft 365 Copilot Chat and Teams, fulfilling the discoverability requirement for code-first scenarios.

Why Other Options Are Less Suitable

A. Foundry Agent Service (Azure AI Agent Service): This service targets the **Azure AI Foundry** portal and custom client applications via the Azure AI Foundry SDK. While it supports function calling and code interpreter for data grounding, its native runtime and distribution channel are Azure-centric. It does not natively register agents into the Microsoft 365 Copilot orchestrator or Teams App manifest schema required for end-user discoverability inside the M365 chrome without significant custom middleware bridging.

B. The Assistants API in Azure OpenAI: This is a raw, stateless REST API primitive for thread/message management and tool use (Code Interpreter, File Search, Functions). It lacks the **channel adapter**, **authentication middleware (Entra/SSO)**, **Teams manifest packaging**, and **Microsoft 365 Copilot manifest declarative schema** required to surface an agent as a discoverable, installable app inside Teams or the M365 Copilot sidebar. It is a backend component, not a distribution platform.

C. The Foundry SDK: This SDK is designed for developers building custom clients or management tooling

against **Azure AI Foundry** resources (projects, models, evaluations, agents). It facilitates the development lifecycle of Azure AI Agents but does not provide the runtime hosting, channel connectivity (Teams/M365), or manifest registration mechanism required to make an agent available to users in the Microsoft 365 ecosystem.

References

1. Extend Microsoft 365 Copilot with agents - Microsoft Learn
2. Microsoft 365 Agents SDK documentation - Microsoft Learn

Question: 78 CertyIQ

Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Background -
Fabrikam, Inc., is a global consumer goods company that is undergoing a digital transformation initiative to migrate its entire infrastructure to the Microsoft cloud. As a key element of this cloud migration, the company will implement Microsoft Dynamics 365 Sales, moving away from the current on-premises proprietary technologies used by its business-to-business (B2B) sales team.
As part of the cloud migration, Fabrikam will adopt an AI-first approach to its business solutions and implement AI solutions, wherever possible, to streamline operations.
Problem Statements -
Fabrikam's infrastructure currently relies on various on-premises systems that require sales executives to use corporate computers with physical keyboards to access business information during customer interactions. Mobile phones cannot be used for these purposes, as the systems depend on keyboard input. As a result, the sales executives spend a lot of time using keyboards to search for data on several disparate systems and file servers, rather than focusing on the customers. This affects the customer experience.
Fabrikam stakeholders are concerned that users will be hesitant to adopt AI. If the AI initiatives are NOT adopted, cost savings will never be realized. Additionally, funding for future AI initiatives will depend on demonstrating an increase in AI adoption month over month. As the AI agent initiative for the sales team will be the first for Fabrikam, the rapid adoption of the agent is a high priority.
Planned Initiatives -
General -
Fabrikam management has prioritized AI-driven projects to improve efficiency, customer engagement, and responsible AI adoption. The current application infrastructure is on-premises and must be migrated to the cloud to support the adoption of these technologies.
Infrastructure Migration -
Fabrikam plans to migrate from its current on-premises infrastructure to a completely cloud-based topology; this will include user authentication, the security framework, and, primarily, the adoption of the services by end users. All the data from the different systems will be consolidated into a single data source - a common data model that will use a Microsoft Dataverse environment as a single source of truth (SSOT) for the sales team.
Sales Cycle Enablement -
To achieve the company's objectives, Fabrikam intends to implement the following strategies to enhance the sales cycle:
Use low-code development to create a single AI agent that has Dataverse as its core component.
Ensure that sales managers can access unanswered correspondence from prospects and intervene as appropriate.

Replace the previous proprietary software with Dynamics 365 Sales to track sales cycles and customer interactions.
Have the sales executives use Dynamics 365 Sales to track interactions for open opportunities and send follow-up communications to prospects.
Have the sales executives use handsfree headsets to interact with an AI agent when they have questions about internal policies or customer data.
Requirements -
Infrastructure Migration -
Fabrikam has identified the following infrastructure migration requirements:
Azure must be used for all future infrastructure workloads.
The company must follow Microsoft-recommended methodologies for infrastructure migration to the cloud.
Any created AI agents must have their return on investment (ROI) calculated to ensure that the solution will save the company money.
Sales Cycle Enablement -
Fabrikam has identified the following requirements for sales cycle enablement:
The final AI agent must follow Microsoft recommendations for a conversational user experience.
A designated checklist must be reviewed to ensure that the AI agent follows Microsoft deployment recommendations for a compliant solution.
Detailed telemetry must be logged for the first created AI agent to help troubleshoot and optimize the agent during the initial AI agent adoption process.
Unexpected AI agent actions must end in an escalation to a live representative. For example, a sales executive must be rerouted to a representative if the agent cannot answer a question after two failed attempts.
The return on investment (ROI) of switching from the current process to the future process is required for stakeholder sign off.
The sales team must use Dynamics 365 Sales to correspond with prospects more quickly and efficiently than currently.
Sales managers must report on the adoption of the AI agent to key Fabrikam stakeholders on a monthly basis.
Any sensitive information, such as user IDs and names, shared via the AI agent must be tracked for future auditing.
Which tool should you recommend to address the sensitive information concerns in the sales process?

A. the Analytics tab in Microsoft Copilot Studio
B. Model Context Protocol (MCP)
C. Application Insights
D. Microsoft Foundry Tracing UI
E. Monitoring in Microsoft Foundry

Answer: D

Explanation:

Recommended tool: Microsoft Foundry Tracing UI

Why D is the best choice

Provides granular, end-to-end tracing of data flows that surface personally identifiable information (PII) such as user IDs and names, enabling audit-ready logging of every exposure.
Integrates directly with the Foundry runtime, so it can capture sensitive data handling across the AI-agent, Dynamics365 Sales, and Dataverse layers without additional custom code.
Generates auditable telemetry that satisfies the ROI-validation and security-compliance checklist required by Fabrikam's governance model.

Why the other options are less suitable

A – Analytics tab in Microsoft Copilot Studio – Focuses on conversational analytics and user interaction metrics; it does not expose detailed PII-level audit trails.
B – Model Context Protocol (MCP) – A communication standard for extending agents; it does not provide a UI or audit mechanism for tracking sensitive data.
C – Application Insights – Offers generic performance and usage telemetry, but lacks the fine-grained

data-masking and audit capabilities needed for regulated PII logging.
E – Monitoring in Microsoft Foundry – Provides high-level health monitoring; it does not deliver the detailed tracing required to audit who accessed or transmitted sensitive information.

References

Microsoft Foundry Tracing UI overview: https://learn.microsoft.com/en-us/fabric/foundry/tracing-ui
Auditing and data governance in Dataverse: https://learn.microsoft.com/en-us/power-apps/maker/data-platform/privacy-auditing

The correct option (D) directly addresses the requirement to track and audit sensitive information shared via the AI agent, aligning with Fabrikam's compliance and ROI verification needs.

Question: 79 CertyIQ

HOTSPOT -
Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Background -
Fabrikam, Inc., is a global consumer goods company that is undergoing a digital transformation initiative to migrate its entire infrastructure to the Microsoft cloud. As a key element of this cloud migration, the company will implement Microsoft Dynamics 365 Sales, moving away from the current on-premises proprietary technologies used by its business-to-business (B2B) sales team.
As part of the cloud migration, Fabrikam will adopt an AI-first approach to its business solutions and implement AI solutions, wherever possible, to streamline operations.
Problem Statements -
Fabrikam's infrastructure currently relies on various on-premises systems that require sales executives to use corporate computers with physical keyboards to access business information during customer interactions. Mobile phones cannot be used for these purposes, as the systems depend on keyboard input. As a result, the sales executives spend a lot of time using keyboards to search for data on several disparate systems and file servers, rather than focusing on the customers. This affects the customer experience.
Fabrikam stakeholders are concerned that users will be hesitant to adopt AI. If the AI initiatives are NOT adopted, cost savings will never be realized. Additionally, funding for future AI initiatives will depend on demonstrating an increase in AI adoption month over month. As the AI agent initiative for the sales team will be the first for Fabrikam, the rapid adoption of the agent is a high priority.
Planned Initiatives -
General -
Fabrikam management has prioritized AI-driven projects to improve efficiency, customer engagement, and responsible AI adoption. The current application infrastructure is on-premises and must be migrated to the cloud to support the adoption of these technologies.
Infrastructure Migration -
Fabrikam plans to migrate from its current on-premises infrastructure to a completely cloud-based topology; this will include user authentication, the security framework, and, primarily, the adoption of the services by end users. All the data from the different systems will be consolidated into a single data source - a common data model that will use a Microsoft Dataverse environment as a single source of truth (SSOT) for the sales team.

Sales Cycle Enablement -
To achieve the company's objectives, Fabrikam intends to implement the following strategies to enhance the sales cycle:
Use low-code development to create a single AI agent that has Dataverse as its core component.
Ensure that sales managers can access unanswered correspondence from prospects and intervene as appropriate.
Replace the previous proprietary software with Dynamics 365 Sales to track sales cycles and customer interactions.
Have the sales executives use Dynamics 365 Sales to track interactions for open opportunities and send follow-up communications to prospects.
Have the sales executives use handsfree headsets to interact with an AI agent when they have questions about internal policies or customer data.
Requirements -
Infrastructure Migration -
Fabrikam has identified the following infrastructure migration requirements:
Azure must be used for all future infrastructure workloads.
The company must follow Microsoft-recommended methodologies for infrastructure migration to the cloud.
Any created AI agents must have their return on investment (ROI) calculated to ensure that the solution will save the company money.
Sales Cycle Enablement -
Fabrikam has identified the following requirements for sales cycle enablement:
The final AI agent must follow Microsoft recommendations for a conversational user experience.
A designated checklist must be reviewed to ensure that the AI agent follows Microsoft deployment recommendations for a compliant solution.
Detailed telemetry must be logged for the first created AI agent to help troubleshoot and optimize the agent during the initial AI agent adoption process.
Unexpected AI agent actions must end in an escalation to a live representative. For example, a sales executive must be rerouted to a representative if the agent cannot answer a question after two failed attempts.
The return on investment (ROI) of switching from the current process to the future process is required for stakeholder sign off.
The sales team must use Dynamics 365 Sales to correspond with prospects more quickly and efficiently than currently.
Sales managers must report on the adoption of the AI agent to key Fabrikam stakeholders on a monthly basis.
Any sensitive information, such as user IDs and names, shared via the AI agent must be tracked for future auditing.
Which existing tool and data should you use to gather the required metrics for stakeholder signoff for the AI agents? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

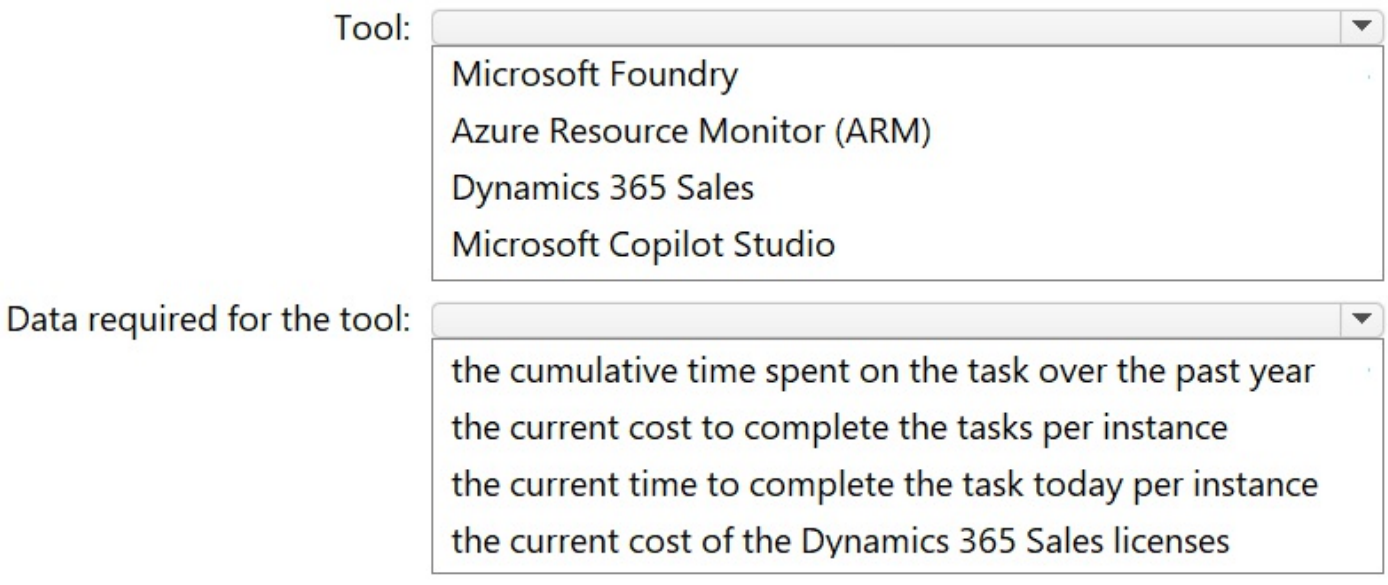

Answer:

Answer Area

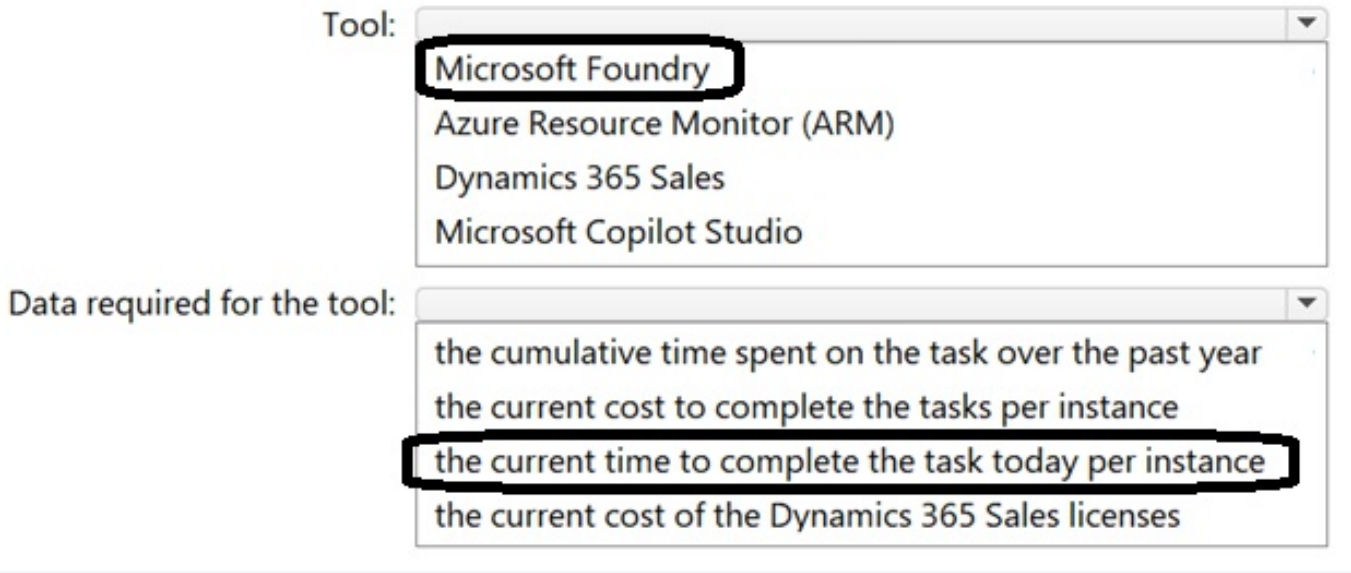

Explanation:

Tool → **Microsoft Foundry.**

What Microsoft Foundry is

Microsoft Foundry (Azure AI Foundry) is a platform used to:

Build and manage AI solutions

Create AI agents

Evaluate models

Optimize AI workflows

Monitor and improve AI performance

It helps organizations analyze tasks and workflows to determine how AI can improve efficiency

Because the scenario focuses on AI-based optimization of tasks, Microsoft Foundry is the correct tool.

Data required for the tool → **the current time to complete the task today per instance.**

To evaluate whether AI can optimize or automate a task, the system must understand:

How long the task currently takes

How often the task occurs

Where efficiency improvements are possible

The current time to complete the task per instance provides:

A baseline for performance measurement

A way to calculate potential efficiency gains

Question: 80

CertyIQ

Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Background -
Fabrikam, Inc., is a global consumer goods company that is undergoing a digital transformation initiative to migrate its entire infrastructure to the Microsoft cloud. As a key element of this cloud migration, the company will implement Microsoft Dynamics 365 Sales, moving away from the current on-premises proprietary technologies used by its business-to-business (B2B) sales team.
As part of the cloud migration, Fabrikam will adopt an AI-first approach to its business solutions and implement AI solutions, wherever possible, to streamline operations.
Problem Statements -
Fabrikam's infrastructure currently relies on various on-premises systems that require sales executives to use corporate computers with physical keyboards to access business information during customer interactions. Mobile phones cannot be used for these purposes, as the systems depend on keyboard input. As a result, the sales executives spend a lot of time using keyboards to search for data on several disparate systems and file servers, rather than focusing on the customers. This affects the customer experience.
Fabrikam stakeholders are concerned that users will be hesitant to adopt AI. If the AI initiatives are NOT adopted, cost savings will never be realized. Additionally, funding for future AI initiatives will depend on demonstrating an increase in AI adoption month over month. As the AI agent initiative for the sales team will be the first for Fabrikam, the rapid adoption of the agent is a high priority.
Planned Initiatives -
General -
Fabrikam management has prioritized AI-driven projects to improve efficiency, customer engagement, and responsible AI adoption. The current application infrastructure is on-premises and must be migrated to the cloud to support the adoption of these technologies.
Infrastructure Migration -
Fabrikam plans to migrate from its current on-premises infrastructure to a completely cloud-based topology; this will include user authentication, the security framework, and, primarily, the adoption of the services by end users.
All the data from the different systems will be consolidated into a single data source - a common data model that will use a Microsoft Dataverse environment as a single source of truth (SSOT) for the sales team.
Sales Cycle Enablement -
To achieve the company's objectives, Fabrikam intends to implement the following strategies to enhance the sales cycle:
Use low-code development to create a single AI agent that has Dataverse as its core component.
Ensure that sales managers can access unanswered correspondence from prospects and intervene as appropriate.
Replace the previous proprietary software with Dynamics 365 Sales to track sales cycles and customer interactions.
Have the sales executives use Dynamics 365 Sales to track interactions for open opportunities and send follow-up communications to prospects.
Have the sales executives use handsfree headsets to interact with an AI agent when they have questions about internal policies or customer data.
Requirements -
Infrastructure Migration -
Fabrikam has identified the following infrastructure migration requirements:
Azure must be used for all future infrastructure workloads.
The company must follow Microsoft-recommended methodologies for infrastructure migration to the cloud.
Any created AI agents must have their return on investment (ROI) calculated to ensure that the solution will save the company money.
Sales Cycle Enablement -
Fabrikam has identified the following requirements for sales cycle enablement:
The final AI agent must follow Microsoft recommendations for a conversational user experience.

A designated checklist must be reviewed to ensure that the AI agent follows Microsoft deployment recommendations for a compliant solution.
Detailed telemetry must be logged for the first created AI agent to help troubleshoot and optimize the agent during the initial AI agent adoption process.
Unexpected AI agent actions must end in an escalation to a live representative. For example, a sales executive must be rerouted to a representative if the agent cannot answer a question after two failed attempts.
The return on investment (ROI) of switching from the current process to the future process is required for stakeholder sign off.
The sales team must use Dynamics 365 Sales to correspond with prospects more quickly and efficiently than currently.
Sales managers must report on the adoption of the AI agent to key Fabrikam stakeholders on a monthly basis.
Any sensitive information, such as user IDs and names, shared via the AI agent must be tracked for future auditing.
Which tool should you recommend to help secure funding for future AI agent development?

A. Evaluations in Microsoft Foundry
B. the Azure Cost Optimization workbook
C. Azure Operator Insights
D. the Analytics tab in Microsoft Copilot Studio
E. Direct Preference Optimization (DPO)

Answer: D

Explanation:

Justification

Option D – Analytics tab in Microsoft Copilot Studio is the most appropriate choice because it delivers the detailed telemetry, usage metrics, and ROI calculations that stakeholders require to validate and fund future AI-agent investments. The analytics view surfaces adoption rates, user interaction patterns, and escalation events, enabling the team to demonstrate measurable business value and compliance with Microsoft's conversational-AI best practices. This granular insight directly supports the funding criteria outlined in the case study (e.g., ROI tracking, monthly adoption reporting, audit-ready telemetry).
Why the other options are less suitable
A. Evaluations in Microsoft Foundry – Primarily an experimentation platform for model development; it lacks built-in, production-grade telemetry and ROI reporting needed for stakeholder sign-off.
B. Azure Cost Optimization workbook – Focuses on cost-reduction recommendations across Azure resources, but does not provide the AI-specific adoption or interaction metrics required to prove the AI agent's business impact.
C. Azure Operator Insights – Designed for network-operations analytics; it does not integrate with AI-agent telemetry or provide the conversational-AI compliance checklist required by the case.
E. Direct Preference Optimization (DPO) – A research technique for fine-tuning language models; it is not a tool for operational analytics or funding justification.

Conclusion – Selecting the **Analytics tab in Microsoft Copilot Studio** gives Fabrikam the concrete, auditable data needed to secure continued funding for AI-agent development.

References

Microsoft Copilot Studio – Analytics and insights
Measuring ROI for AI solutions in Microsoft Azure

Question: 81 CertyIQ

You are designing a Microsoft Copilot Studio agent that uses a custom Microsoft Foundry model to generate

responses.
You need to ensure that the agent can securely connect to and invoke the custom model during user interactions.
What should you include in the design?

A. Configure the agent to use classic orchestration.
B. Create a connection to Microsoft Foundry in the agent.
C. Add the Microsoft Foundry model as a Copilot Studio skill.
D. Create a custom engine agent.

Microsoft Copilot Studio supports natively consuming models deployed in Azure AI Foundry (Microsoft Foundry) as custom models for generative responses.

Answer: B

Explanation:

Why optionB is the optimal design

To enable a Copilot Studio agent to securely authenticate, route payloads to, and invoke a model hosted in Azure AI Foundry, you establish a connection to Microsoft Foundry directly within the agent settings (using Microsoft Entra ID authentication and role-based access control, such as Azure AI Developer / Cognitive Services OpenAI User).

To invoke a custom Microsoft Foundry model from a Copilot Studio agent, the model must be exposed through a connection that provides the required authentication, endpoint URL, and runtime security context. Creating a connection to Microsoft Foundry in the agent registers the model as a managed external resource, allowing the agent to call it securely during user interactions without embedding credentials or custom networking logic.
This approach aligns with the platform's recommended pattern for integrating proprietary models (i.e., "Connect to an external service" in Copilot Studio), ensuring scalability, auditability, and compliance with Microsoft's security model.

Why the other options are less suitable

A. Configure the agent to use classic orchestration – Classic orchestration governs flow control within the same Copilot scope but does not provide a mechanism to register or securely call an external custom model.
C. Add the Microsoft Foundry model as a Copilot Studio skill – Skills are designed for reusable, declarative actions (e.g., prompting, data lookup) and do not expose proprietary model endpoints; they lack the infrastructure for token-based secure connectivity to a Foundry model.
D. Create a custom engine agent – Introducing a separate engine adds unnecessary complexity and bypasses the built-in security and lifecycle management that Copilot Studio offers through its connection framework.

Implementation summary

Register the Foundry model via **Create a connection → Microsoft Foundry** in Copilot Studio.
Use the generated connection reference in orchestrational prompts or actions.
Leverage built-in security (Azure AD token, private endpoint) to protect model access during runtime interactions.

References

Create and manage connections in Copilot Studio
Use custom models with Copilot Studio

Question: 82 CertyIQ

You are designing an AI business solution that contains the following components:
A Microsoft Power Automate workflow
A Microsoft Copilot Studio agent
A Microsoft Dataverse database -
A Microsoft Power Apps app -
As part of the application lifecycle management (ALM) process, you plan to package the components, so that they

can be deployed to other environments as a group.
You need to recommend a solution that supports versioning, dependencies, and deployments.
What should you include in the recommendation?

A. GitHub Actions
B. Azure DevOps
C. Microsoft Power Platform solutions

Answer: C

Explanation:

Why OptionC is the best choice

Native solution packaging – Microsoft Power PlatformSolutions are purpose-built to bundle Power Automate workflows, CopilotStudio agents, Dataverse customizations, and PowerApps definitions into a single transportable package.
Built-in versioning – Each solution has a version number that can be incremented, enabling clear change tracking across dev, test, and prod environments.
Dependency resolution – The solution designer captures explicit "dependencies" (e.g., a PowerApp that uses a specific flow or a Dataverse schema). When the solution is imported, the platform validates and resolves those dependencies automatically.
ALM-ready deployment – Solutions integrate with PowerPlatformALM tools (publish-files, import-set-next-step, and the "Deploy" button in PowerPortal) so they can be moved as a coordinated unit without custom scripting.
Lifecycle management – You can promote a solution from one environment to another while preserving its version, metadata, and dependent components, ensuring consistency across environments.

Why the alternatives are less suitable

AzureDevOps – Provides robust CI/CD pipelines and repository management, but it does not natively understand Power Platform solution metadata, version constraints, or dependency graphs. You would need to build extensive custom pipelines to mimic solution packaging, increasing complexity and error surface.
GitHubActions – Similar to AzureDevOps, it excels at automating builds and deployments, yet it lacks built-in support for Power Platform solution versioning and dependency resolution. Any solution-specific handling would be indirect and not standardized across environments.

Summary
For an end-to-end ALM approach that guarantees version control, explicit dependency tracking, and reliable multi-environment deployment of Power Automate flows, CopilotStudio agents, Dataverse customizations, and PowerApps, the recommended mechanism is to package everything as a **Microsoft Power Platform solution**.

References

Power Platform solution management: https://learn.microsoft.com/en-us/power-platform/alm/solution-management
Deploying Power Platform solutions: https://learn.microsoft.com/en-us/power-platform/alm/customizations/manage-solutions

Question: 83 CertyIQ

A company has Microsoft 365 Copilot agents.

You need to design a security solution for the agents. The solution must meet the following requirements:
Identify and mitigate potential risks that relate to AI use.
Protect AI apps and the sensitive data processed or generated by the agents.
Support responsible AI governance by retaining and logging interactions, detecting policy violations, and investigating incidents.
Which two components should you include in the design? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. Microsoft Purview
B. Azure AI Content Safety
C. role-based access control (RBAC) in Microsoft Foundry
D. Microsoft Defender

Microsoft Purview (Purview AI Hub): Delivers end-to-end data governance, compliance, and auditing. It enforces sensitivity labels and Data Loss Prevention (DLP) across data consumed and generated by agents, retains audit logs and prompt/response transcripts for eDiscovery, and flags compliance policy violations or unethical prompts via Communication Compliance.

Microsoft Defender (Microsoft Defender for Cloud Apps & Defender for Cloud): Identifies and mitigates threats specific to AI usage. It discovers unsanctioned/shadow AI apps, evaluates organizational AI risk postures, enforces Conditional Access and app controls, detects anomalous behavioral patterns or compromised identities, and prevents malware/adversarial threats targeting enterprise AI applications.

Answer: AD

Explanation:

Why the correct components are required

MicrosoftPurview – Provides centralized data-governance, classification, and risk-detection capabilities. It can inventory the inputs and outputs flowing through the Copilot agents, apply sensitivity labels, enforce retention policies, and generate audit logs that satisfy the "retain and log interactions" requirement. Purview's policy engine also enables the detection of misuse or policy violations, directly supporting responsible AI governance.

MicrosoftDefender – Offers threat-protection and sandbox-style inspection for AI-generated content. By integrating with the Copilot runtime, Defender can scan prompts and responses for malicious patterns, block unsafe outputs, and quarantine suspicious payloads, thereby protecting both the AI application and the sensitive data it processes or generates.

Why the other options are less suitable

Azure AI Content Safety – Focuses on content moderation (e.g., detecting hate speech, sexual content) but does not provide the broader governance controls, data-lineage tracking, or incident-investigation capabilities demanded by the scenario.

Role-Based Access Control (RBAC) in Microsoft Foundry – RBAC is useful for authorizing who can develop or deploy agents, yet it does not inherently address data-risk identification, continuous monitoring of AI interactions, or automated policy enforcement across the AI lifecycle.

(No other listed option besides A-D) – The remaining choices do not collectively fulfill all three pillars of the requirement (risk identification, data protection, and governance logging/investigation).

Design Summary

1. Deploy **MicrosoftPurview** to catalog AI-generated data, apply sensitivity classifications, retain interaction logs, and surface policy violations for audit.
2. Enable **MicrosoftDefender** to inspect and block malicious or leaked data within the Copilot agent pipeline, providing real-time threat mitigation.

References

Microsoft Purview documentation: https://learn.microsoft.com/en-us/purview/
Microsoft Defender documentation: https://learn.microsoft.com/en-us/defender/

These resources detail the configuration of data-governance policies and threat-protection features that

align with the exam's requirements.

Question: 84 CertyIQ

You are creating validation criteria for a custom generative AI model that produces business reports based on internal enterprise data.
You need to assess whether the model's outputs are appropriate and meaningful for the business reports.
Which metric should you use?

A. the number of active users interacting with the model
B. alignment of the output to domain-specific tasks
C. the average system resource usage during inference
D. the model training duration

Answer: B

Explanation:

Option B – Alignment of output to domain-specific tasks directly evaluates whether the generated business reports satisfy the intended objectives, content accuracy, and terminology required by the enterprise domain. This is the core validation criterion for a custom generative AI model that must produce meaningful, business-relevant outputs.
Option A – Number of active users measures adoption, not quality; a model may be heavily used yet produce irrelevant or erroneous reports.
Option C – Average system resource usage reflects inference performance efficiency, not semantic correctness or business relevance of the generated content.
Option D – Model training duration pertains to the development phase and does not indicate how well the model meets validation criteria for production-ready outputs.

Hence, metric **B** is the most appropriate for assessing the suitability of the model's outputs for business-report generation.

References

Microsoft Azure Documentation – Evaluating Generative AI models:
https://learn.microsoft.com/azure/machine-learning/how-to-evaluate-generative-ai-models
Microsoft Responsible AI – Model evaluation criteria for business use cases:
https://learn.microsoft.com/responsible-ai/concepts/evaluation-criteria

Question: 85 CertyIQ

A company has Microsoft Foundry agents that generate responses by using Azure OpenAI resources. The agents are deployed to both the United States and Europe.
A company mandate states that the agents and their grounding data must adhere to data residency and movement regulations.
You need to recommend a governance solution for the agents.
What should you include in the recommendation?

A. Microsoft Defender for Cloud
B. Azure Policy
C. Azure Monitor

When generative AI agents operate across geopolitical boundaries (such as the US and EU) and retrieve grounding data, organizations must enforce data residency, cross-border data transfer limitations, and regulatory compliance (e.g., GDPR, EU AI Act)

D. Microsoft Purview

Microsoft Purview provides comprehensive data cataloging, data mapping, data lifecycle management, and cross-boundary governance. Through the Purview Data Map and Data Governance controls, Purview inspects, classifies, and audits where grounding data resides, detects unauthorized cross-border movement, and applies data residency policies to prevent non-compliant data egress across regions.

Answer: D

Explanation:

Microsoft Purview is the only option that directly addresses **data residency and movement compliance**. It provides data-classification, sensitivity labeling, and the ability to enforce governance policies on datasets that feed Azure OpenAI models, ensuring that grounding data never leaves the required geographic region (US/EU).
While **Azure Policy** can enforce resource-level rules, it does not natively understand data- residency constraints for AI grounding data; it would need to be combined with Purview to be effective.
Microsoft Defender for Cloud focuses on security posture and threat protection, not on data governance or residency controls.
Azure Monitor is a telemetry and alerting service; it offers no mechanism to govern where data is stored or processed.

Therefore, to meet the mandate for data residency and movement regulations, the recommendation must incorporate **Microsoft Purview** to classify, label, and control the flow of grounding data across the two regions.

References

Microsoft Purview documentation: https://learn.microsoft.com/en-us/purview/
Azure OpenAI data residency and compliance guidance: https://learn.microsoft.com/en-us/azure/ai-services/openai/concepts/data-residency

Question: 86

CertyIQ

A company has a Microsoft Copilot Studio agent that uses custom connectors to interact with enterprise APIs. You need to recommend an application lifecycle management (ALM) process to ensure that the connectors are deployed consistently across development, test, and production environments and meet governance and traceability requirements.
What should you recommend?

A. Deploy the APIs as Azure Functions.
B. Manage the connectors as solution components and deploy the components by using ALM pipelines.
C. Maintain connector definitions in environment variables.
D. Export and import the connectors between the environments as unmanaged solutions.

Answer: B

Explanation:

B. Manage the connectors as solution components and deploy the components by using ALM pipelines.

Native Power Platform ALM: Custom connectors in Copilot Studio (built on Microsoft Power Platform) should be added to **managed solutions** to track dependencies, maintain history, and enforce governance.

Traceability and Automation: Utilizing **ALM pipelines** (such as Power Platform pipelines or Azure DevOps/GitHub actions) ensures automated, auditable, and consistent deployments across development, test, and production environments.

Why the Other Options are Incorrect

A. Deploy the APIs as Azure Functions: This addresses the underlying hosting infrastructure of the enterprise APIs, but it does not manage or deploy the Copilot Studio custom connectors or solve the environment ALM workflow.

C. Maintain connector definitions in environment variables: Environment variables are used to store environment-specific values (like endpoint URLs or API keys) within a solution, but they cannot house or transport the connector definitions themselves.

D. Export and import the connectors as unmanaged solutions: Importing unmanaged solutions directly into test or production environments creates unmanaged customizations that are difficult to service, violating enterprise governance, tracking, and clean lifecycle management principles.

Question: 87 CertyIQ

HOTSPOT -
A company plans to implement an AI solution that will contain a Microsoft Copilot Studio agent and a Microsoft Foundry agent. The solution will be stored in a source code repository.
You need to recommend a deployment method for each agent. The solution must meet the following requirements:
A test environment must be used before a deployment to production.
Production must be isolated from development and testing.
The deployment must be repeatable and fully automated.
The solution must NOT require manual intervention.
Which deployment method should you recommend for each agent? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

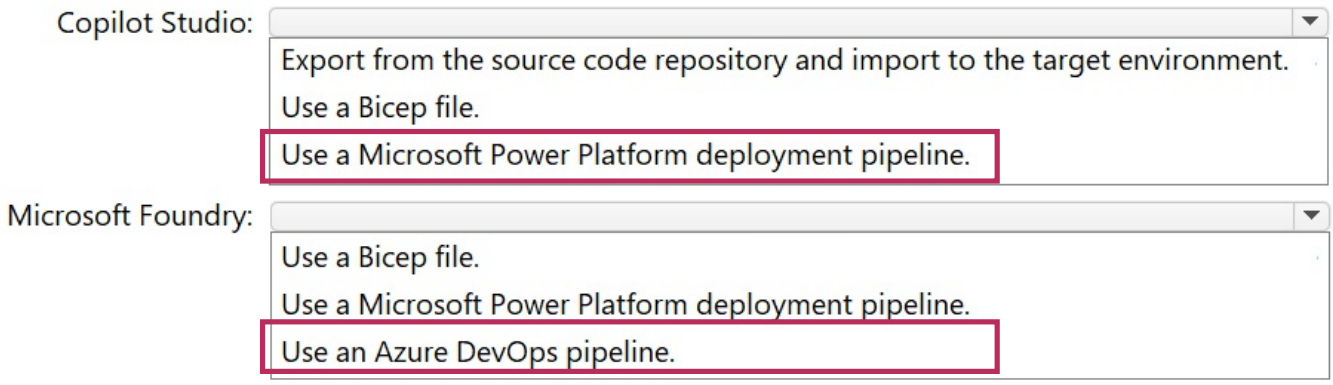

Answer:

Answer Area

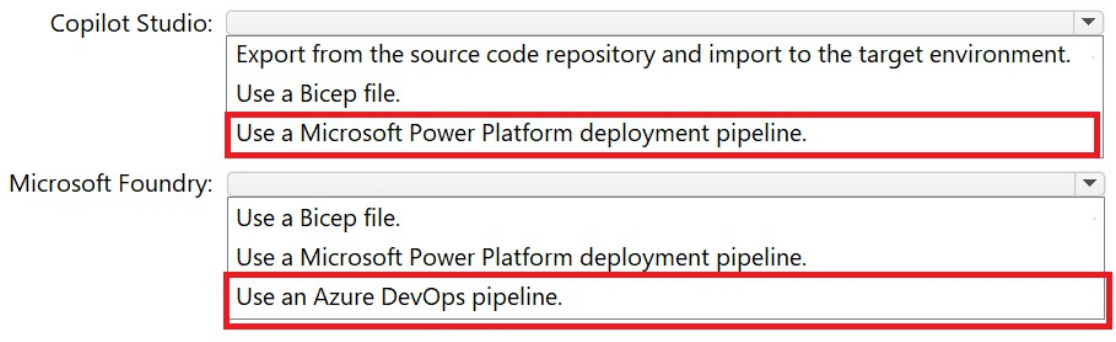

Explanation:

Copilot Studio → **Use a Microsoft Power Platform deployment pipeline.**

Copilot Studio is part of the Microsoft Power Platform ecosystem.

Power Platform solutions (including Copilot Studio agents) are typically deployed using:

Power Platform deployment pipelines

Solutions

Environment promotion (Dev → Test → Prod)

Deployment pipelines allow teams to:

Move solutions between environments

Manage versions

Control releases

Ensure consistent deployments

Microsoft Foundry → **Use an azure devops pipeline.**

Microsoft Foundry (formerly Azure AI Foundry) is an **Azure development platform**.

Applications typically contain:

Prompt flow code
AI model configurations
Python code
Infrastructure
Deployment configurations

Question: 88 CertyIQ

HOTSPOT -
A company has a Microsoft Copilot Studio prompt-and-response agent.
You need to ensure that the agent meets the following requirements:
Provides effective and relevant responses
Provides conversational outcomes
Which metric should you use for each requirement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

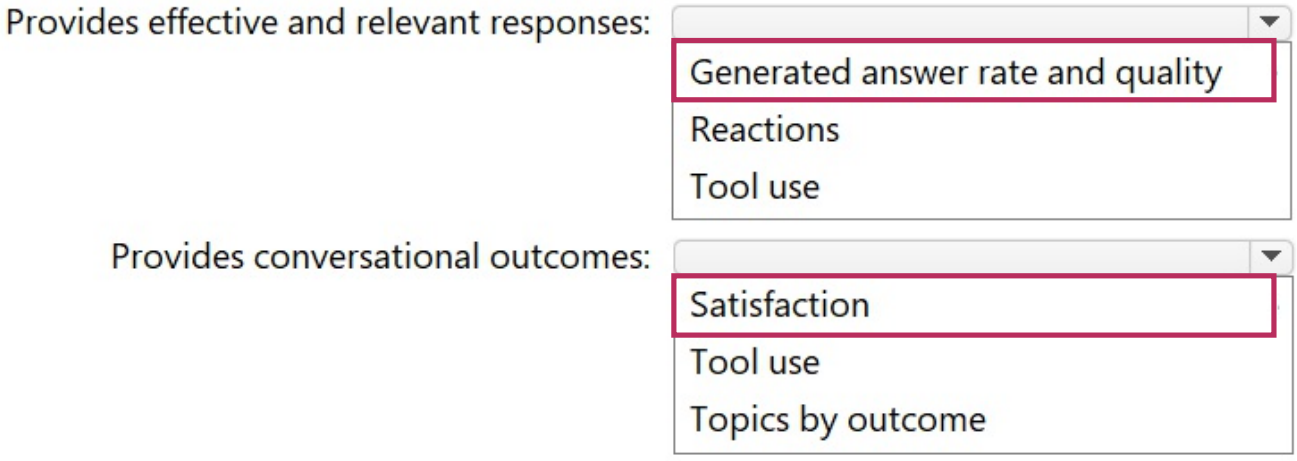

Answer:

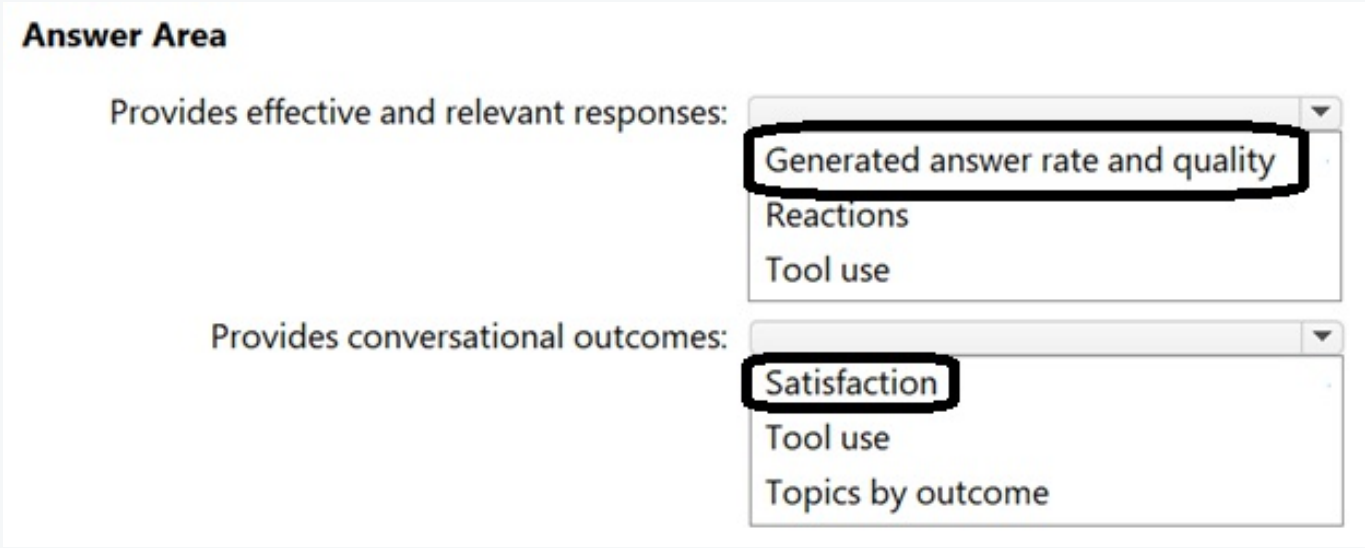

Explanation:

Provides effective and relevant responses → **Generated answer rate and quality.**

What this metric measures

Generated answer rate and quality measures how well the AI agent generates useful and relevant responses.

It evaluates:

How often the AI successfully answers questions

The quality and relevance of the generated responses

Whether answers match user intent

Example metrics include:

Percentage of questions answered successfully

Quality score of AI responses.

Provides conversational outcomes → **Satisfaction.**

What satisfaction measures

Satisfaction measures the user's overall experience with the conversation.

It evaluates whether:

The user received helpful answers

The conversation solved the user's problem

The user felt satisfied with the interaction

Often measured using:

User ratings

Feedback scores

Post-conversation surveys

Question: 89 CertyIQ

A company extends Copilot in Microsoft Dynamics 365 Customer Service.
You need to recommend an automated application lifecycle management (ALM) process so that the Copilot components can be safely developed, tested, and promoted to production.
Which two actions should you include in the ALM process? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. Use an unmanaged solution in production.
B. Rebuild the agents in each environment.
C. Use Microsoft Power Platform pipelines.
D. Include the components in a solution.
E. Store the agent transcripts in source control.

Answer: CD

Explanation:

WhyCandD are the correct choices

C – Use Microsoft Power Platform pipelines – Power Platform provides built-in CI/CD pipelines that let you automatically build, test, and move Copilot components from development through testing to production. This automation reduces manual errors, enforces consistent configuration, and ensures that updates to the Copilot model, custom plugins, or integration points are propagated in a controlled manner.

D – Include the components in a solution – Packaging the Copilot customizations (models, plugins, canvas apps, Power Automate flows, etc.) into an unmanaged or managed solution makes the entire set of artifacts a single, version-controlled artifact. Solutions can be exported, imported, and promoted across environments, which is the recommended way to manage lifecycle changes in Dynamics365 and Power Platform.

Why the other options are unsuitable

A – Use an unmanaged solution in production – Deploying directly to production with an unmanaged solution bypasses the controlled promotion flow. It makes rollback impossible and risks contaminating the production environment with unreviewed changes, violating best-practice ALM principles.

B – Rebuild the agents in each environment – Re-creating the Copilot model or agent logic from scratch in every environment is inefficient and leads to inconsistencies. Instead, the same packaged solution should be moved through environments, preserving version compatibility and reducing duplication of effort.

E – Store the agent transcripts in source control – Transcripts are runtime data and are not part of the application metadata that needs to be promoted. While source-controlled logs can be useful for monitoring, they do not constitute an ALM mechanism for delivering the Copilot components themselves.

Result: The combination of **C** (using Power Platform pipelines) and **D** (including components in a solution) provides a repeatable, auditable, and automated lifecycle management process that aligns with Microsoft's recommended ALM practices for Copilot extensions in Dynamics365 Customer Service.

References

Microsoft Docs: ALM with Power Platform pipelines – https://learn.microsoft.com/en-us/power-platform/pipelines/
Microsoft Docs: Solutions lifecycle in Power Platform – https://learn.microsoft.com/en-us/power-

platform/admin/solutions-lifecycle-management-overview

Question: 90 CertyIQ

HOTSPOT -
You are designing a testing solution for a Microsoft Copilot Studio agent that integrates with Microsoft Dynamics 365 Customer Service and Dynamics 365 Sales.
You need to design end-to-end scenarios to test the agent's ability to perform the following actions:
Coordinate tasks and data interactions across both Dynamics 365 apps.
Interpret user input and provide contextually relevant outputs.
Which test scenario and metric should you include in the design? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

Test scenario:
- In each app, test isolated tasks without using workflows.
- Run task-based scenarios that involve both apps.
- Test visual consistency across both apps.

Metric:
- Measure the initial prompt response time for each app.
- Track the average click rate across both apps.
- Track the successful completion of cross-app tasks.

Answer:

Answer Area

Test scenario:
- In each app, test isolated tasks without using workflows.
- Run task-based scenarios that involve both apps.
- Test visual consistency across both apps.

Metric:
- Measure the initial prompt response time for each app.
- Track the average click rate across both apps.
- Track the successful completion of cross-app tasks.

Explanation:

Test scenario → **Run task-based scenarios that involve both apps.**

This question focuses on evaluating a solution that works across multiple applications (for example Copilot interacting with multiple Microsoft 365 apps or services).

To properly test such a solution, you must simulate real user workflows.

Why task-based scenarios are correct

Task-based scenarios involving both apps test the actual user journey.

Metric → **Track the successful completion of cross-app tasks.**

What this metric measures

The goal is to measure whether the workflow actually works across apps.

The best metric is therefore:

Successful completion of cross-app tasks

This metric evaluates:

Whether the user finished the task

Whether all steps across apps succeeded

Whether Copilot correctly orchestrated actions

Question: 91 CertyIQ

A company has multiple AI models that support generation of sales transactions.
Each release of the models must be reviewed by a security and compliance team before being deployed to the production environment. The security and compliance team must have access to prior versions to properly determine potential exposures introduced.
You need to recommend a solution to evaluate the impact of each deployment to production. The solution must enhance business continuity.
What should you recommend?

A. Create a central model registry that uses version history.
B. Establish a promotion process by using a quality gate.
C. Implement version control for all the AI system components.
D. Track model retirement schedules to prevent service disruptions.

In an agentic generative AI solution, an AI model deployment does not exist in isolation. Output behavior, security exposure, and operational risks are dictated by the interplay of prompt definitions, tool orchestrations, model versions, and system code.

Answer: C

Explanation:

To enable a security and compliance team to compare releases, analyze diffs, and pinpoint newly introduced vulnerabilities or exposures, the entire solution stack must be versioned holistically.

WhyOptionC is the best fit

Implementing **version control for all AI system components** guarantees a complete, auditable history of every model artifact, configuration, and associated data.
This traceability lets the security and compliance team retrieve exact prior releases, assess the specific changes introduced, and quantify exposure risk before any promotion to production.
The capability dovetails with CI/CD pipelines (Best-Practice2) and change-management processes (Best-Practice5), enabling automated impact analysis, rollback, and continuous monitoring — key pillars for business continuity.
By enforcing version control, the organization can systematically evaluate each deployment's downstream effect, satisfy regulatory audit requirements, and maintain uninterrupted service.

Why the other options are less suitable

A – Central model registry with version history – Provides storage and visibility, but without an explicit version-control workflow it lacks the enforcement of change-by-change reviews, audit trails, and automated impact assessment needed for compliance.
B – Promotion process with a quality gate – Ensures models meet predefined quality thresholds, yet it does not by itself give the compliance team direct access to prior versions or the granular diff data required to evaluate introduced exposures.
D – Track model retirement schedules – Helps avoid service interruptions during deprecation, but it does not address the core need of evaluating the impact of each deployment on the production environment.

Bottom line: Only a robust version-control framework directly satisfies the requirement to review past versions, assess impact, and preserve continuous operations while adhering to governance best practices.

References

Microsoft Azure Machine Learning Model Versioning & Registry: https://learn.microsoft.com/azure/machine-learning/how-to-manage-models
Model Lifecycle & CI/CD Best Practices:
https://learn.microsoft.com/azure/devops/pipelines/tasks/reference/model-versioning

Question: 92 CertyIQ

DRAG DROP -
A company has an AI solution that uses a Microsoft Copilot Studio agent.
You need to monitor the agent's performance. The solution must meet the following requirements:
Monitor the agent's telemetry in near-real-time (NRT).
Download transcripts of full conversations.
Monitor the agent's usage and performance.
What should you use for each requirement? To answer, drag the appropriate options to the correct requirements. Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

Options

- Application Insights
- Copilot Studio
- Log Analytics
- Microsoft Power Apps

Answer Area

	Option
Monitor the agent's telemetry in NRT:	Application Insight
Download transcripts of full conversations:	Power Apps
Monitor the agent's usage and performance:	Copilot Studio

Answer:

In Microsoft Power Apps (make.powerapps.com), administrators and makers with the Bot Transcript Viewer role access the core Dataverse ConversationTranscript table. Exporting from Power Apps downloads the raw, unredacted, multi-turn full conversation transcripts (including complete activity payloads, nodeTraceData, variable states, and adaptive card interaction data).

Answer Area

	Option
Monitor the agent's telemetry in NRT:	Application Insights
Download transcripts of full conversations:	Microsoft Power Apps
Monitor the agent's usage and performance:	Copilot Studio

Explanation:

Application Insights.

While Copilot Studio has built-in analytics, Azure Application Insights is the professional-grade tool used for deep, technical telemetry.

Microsoft Power Apps

This one often trips people up! Conversation transcripts for Copilot Studio are stored in Microsoft Dataverse tables.

Copilot Studio.

This refers to the high-level business analytics.

Inside the Copilot Studio portal itself, there is a dedicated Analytics tab.

Question: 93

CertyIQ

HOTSPOT -
A company deploys a Microsoft Copilot Studio agent that integrates with a Microsoft Power Automate desktop flow.
You need to recommend a testing solution that meets the following requirements:
Test cases must validate the most recent changes to the agent before the agent is released.
The flow must be validated as part of the agent's orchestration.
What should you recommend for each requirement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

Validate the most recent changes to the agent before release:

- Publish the agent to a channel and test the agent on live users.
- Run tests against the latest unpublished version of the agent.
- Run tests against the production version of the agent.

Validate the flow as part of the agent's orchestration:

- Add the flow to the agent as a tool.
- Add the flow to a canvas app.
- Use the Power Automate for desktop console.

integrating automated flows (including Power Automate desktop flows wrapped via cloud flows) requires exposing them as tools / actions / plugins. When added as a tool,

Answer: In Microsoft Copilot Studio, changes made to topics, prompts, dynamic chaining, and tools take effect immediately in the authoring canvas and the built-in Test pane without requiring a release or publication.

Answer Area

Validate the most recent changes to the agent before release:

- Publish the agent to a channel and test the agent on live users.
- Run tests against the latest unpublished version of the agent.
- Run tests against the production version of the agent.

Validate the flow as part of the agent's orchestration:

- Add the flow to the agent as a tool.
- Add the flow to a canvas app.
- Use the Power Automate for desktop console.

Explanation:

1. Validate the most recent changes to the agent before release

Run tests against the latest unpublished version of the agent. (Your highlighted choice is correct): In Copilot Studio, the built-in Test copilot canvas allows you to test and validate your latest saved configuration changes immediately without having to publish the agent to live production channels or affect real users.

2. Validate the flow as part of the agent's orchestration

Add the flow to the agent as a tool. To validate how a Power Automate flow interacts within the agent's dynamic reasoning path, you must register the flow to the agent as an action or tool. This allows the conversational orchestrator to trigger it during a test conversation, passing variables and verifying the end-to-end data loop.

Incorrect:

Use the Power Automate for desktop console: This will only test a desktop flow locally on your machine in a standalone sandbox. It completely bypasses the copilot's context and cannot validate how the agent orchestrates or communicates with that flow.

Question: 94 CertyIQ

A company has a Microsoft Copilot Studio agent that provides answers based on a knowledge base for customer support.
Users report that, occasionally, the agent provides inaccurate answers.
You need to use metrics from the Analytics tab in Copilot Studio to identify the cause of the inaccuracies.
Which two options should you use? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. survey results
B. session information and session outcomes
C. topic usage and topics with low resolution
D. engagement, resolution, and escalation rates
E. quality of generated answers

The Sessions panel displays the multi-turn conversational context, exact user query phrasing, how the agent routed the inquiry, and the session outcome (e.g., Escalated, Abandoned, or Resolved) with corresponding outcome reasons. This allows architects to determine whether inaccuracies were triggered by ambiguous user prompts, retrieval grounding failures, or incorrect fallback handling.

Answer: CD

Explanation:

Generated answer rate and quality panel, Copilot Studio samples responses and classifies them as Good or Poor. For answers marked as Poor, Copilot Studio exposes a detailed Reason metric (such as Not fully grounded, Irrelevant, Incomplete, or Missing information). This directly isolates the cause of inaccuracies stemming from the knowledge base.

Why options C and D are the correct choices

Topic usage and topics with low resolution (C) –In Copilot Studio's Analytics tab, the Topic Usage metric shows how often a specific knowledge-base topic is invoked, while Topics with low resolution highlights those that frequently fail to achieve a confident answer. Low-resolution topics often map to the instances where the agent delivers inaccurate or ambiguous responses, making this a direct indicator of the root cause.
Session information and session outcomes (D) –Session-level analytics provide detailed telemetry such as user-question phrasing, confidence scores, escalation triggers, and whether the generated answer was accepted or corrected by the user. These metrics reveal patterns (e.g., repeated escalations or low confidence scores) that pinpoint the conditions under which inaccuracies occur.

Why the other options are less suitable

Survey results (A) –These reflect user satisfaction post-interaction but do not surface the precise technical reasons (e.g., low-confidence scores or unresolved topics) that cause inaccuracies during the conversation.
Engagement, resolution, and escalation rates (B) –Engagement and resolution metrics are useful for overall bot performance, yet they are not granular enough to isolate the accuracy problem; escalation rates merely signal that a human was needed, not why the answer was wrong.
(E) Quality of generated answers –While "quality" is the outcome we care about, Copilot Studio does not

expose a standalone “quality” metric in the Analytics tab; quality must be inferred from the underlying session and topic metrics listed above. Hence it is not a selectable option for the analytics-tab investigation.

References

Copilot Studio Analytics Overview
Using Topic Insights to Improve Accuracy

Question: 95 CertyIQ

A company uses a fine-tuned Microsoft Foundry model that requires frequent updates as new customer feedback becomes available.
You need to design an application lifecycle management (ALM) process that meets the following requirements:
Data changes must be tracked and versioned.
The model must be retrained consistently by using approved training data.
Which two actions should you include in the design? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. Associate the storage location to the fine-tuning job.
B. Create a content filter.
C. Store the training data in Azure Files.
D. Upload the training data to Microsoft Foundry data files
E. Store the training data in Azure Blob Storage that has version control enabled.

Answer: DE

Explanation: Microsoft Foundry fine-tuning runs require training datasets to be validated, formatted (e.g., .jsonl), and registered as Microsoft Foundry data files (or data assets). Registering/uploading approved snapshots directly as Foundry data files binds the fine-tuning pipeline to an explicit, immutable training data artifact. This prevents ad-hoc data mutations from silently altering retrained model behavior, guaranteeing consistency across retraining cycles.

D. Upload the training data to Microsoft Foundry data files – MicrosoftFoundry stores model-training artifacts in dedicated data files that are version-controlled, immutable, and linked directly to the fine-tuning job. This ensures that each training iteration is recorded, reproducible, and can be reviewed or rolled back.
E. Store the training data in Azure Blob Storage that has version control enabled – Azure Blob Storage with versioning preserves every change to the data set (e.g., new feedback files) as immutable versions. When used as the source for Foundry data files, it provides traceable, auditable data provenance that satisfies “data changes must be tracked and versioned.”

Why the other options are unsuitable

A. Associate the storage location to the fine-tuning job – Simply associating a location does not enforce versioning or guarantee that the data is stored in a managed, auditable repository; it lacks explicit version-control semantics.
B. Create a content filter – A content filter controls what data is processed (e.g., profanity removal) but does not address versioning of training data or ensure consistent retraining against approved data.
C. Store the training data in Azure Files – Azure Files does not provide built-in object-level versioning for large data sets; it is better suited for traditional file shares, not for the immutable, version-tracked data required by Foundry model pipelines.

References

Microsoft Foundry – Managing Training Data
https://learn.microsoft.com/azure/machine-learning/how-to-foundry-manage-data
Azure Blob Storage – Versioning
https://learn.microsoft.com/azure/storage/blobs/storage-blob-versioning

Question: 96 CertyIQ

HOTSPOT -
A company deploys agents that generate responses by using Azure OpenAI resources. The agents are deployed to both the United States and Europe.
You need to recommend a governance solution that meets the following requirements:
Enforces the deployment of the resources to only approved Azure regions
Provides continuous compliance verification of the resources
What should you include in the recommendation for each requirement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

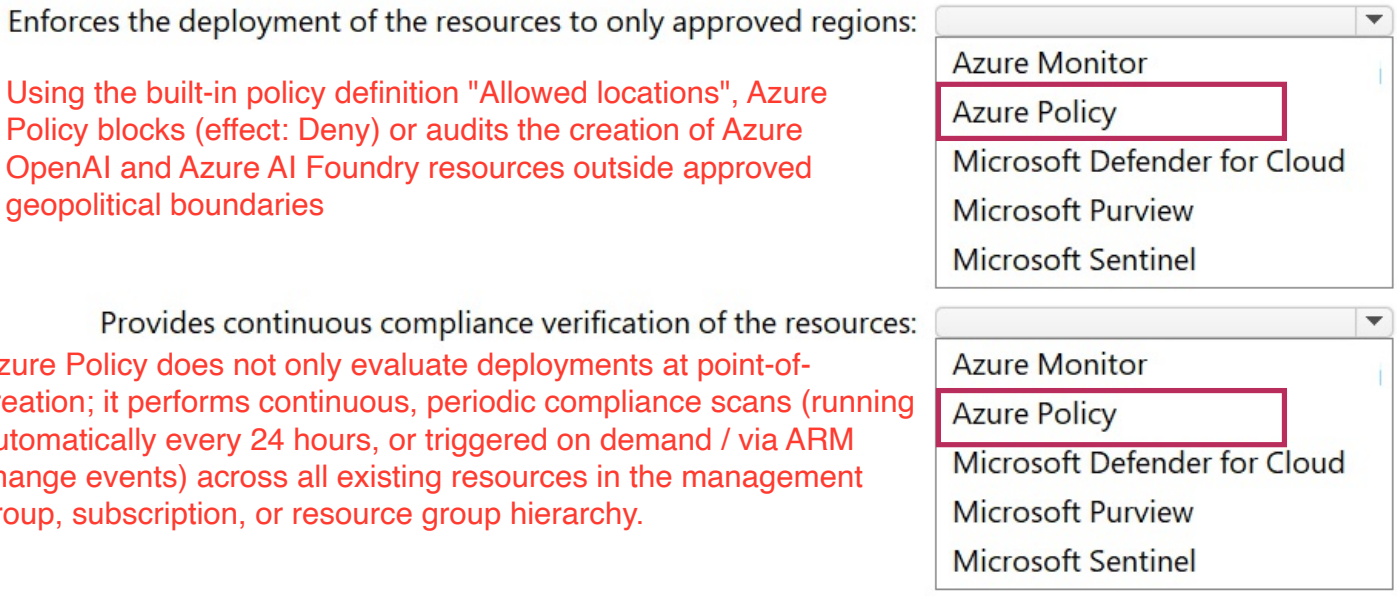

Using the built-in policy definition "Allowed locations", Azure Policy blocks (effect: Deny) or audits the creation of Azure OpenAI and Azure AI Foundry resources outside approved geopolitical boundaries

Azure Policy does not only evaluate deployments at point-of-creation; it performs continuous, periodic compliance scans (running automatically every 24 hours, or triggered on demand / via ARM change events) across all existing resources in the management group, subscription, or resource group hierarchy.

Answer: It continuously evaluates existing Azure OpenAI resources against compliance initiatives , generating real-time non-compliance reports and triggering remediation tasks where configured.

Answer Area

Explanation:

1. Enforces the deployment of the resources to only approved regions

Azure Policy (Correct): This is the dedicated preventive governance tool in Azure. It allows you to create explicit structural rules (such as the standard Allowed locations policy rule) to actively deny or restrict

resource deployment attempts outside specifically authorized regions.

Incorrect:

Microsoft Purview / Microsoft Sentinel (Incorrect): Purview focuses on data discovery, classification, and catalog governance, while Sentinel is a SIEM/SOAR system used for security logging and threat response. Neither tool has the infrastructure-level capability to block or restrict where physical resources are created.

Provides continuous compliance verification of the resources

Microsoft Defender for Cloud For ongoing resource compliance, Microsoft Defender for Cloud continuously assesses your running infrastructure posture. It cross-references your assets against defined security baselines, regulatory compliance frameworks (like ISO, SOC, and GDPR), and raises alerts if any configuration drifts out of alignment.

Incorrect:

Microsoft Purview (Incorrect / Your highlighted selection): While Purview provides data protection and data compliance insights, it evaluates data life cycles, labels, and files — it does not monitor continuous infrastructure-level deployment compliance or evaluate the structural properties of running Azure cloud services.

Question: 97 CertyIQ

A company has an AI solution that uses Azure OpenAI models.
You need to recommend a governance solution that monitors and audits changes to model configurations and data usage. The solution must minimize administrative effort.
What should you include in the recommendation?

A. Azure Monitor
B. Azure Stream Analytics
C. Azure API Management
D. Azure Policy
E. Microsoft Purview

Auditing Data Usage: It tracks, catalogs, and audits the flow of enterprise data and sensitive information consumed by or generated from generative AI endpoints (including Azure OpenAI) without requiring complex manual pipeline setups.

Monitoring Configurations & AI Workloads: Through integrated AI Hub governance policies and data map telemetry, Purview monitors configuration changes, sensitivity label adherence, and compliance drift.

Answer: E

Minimizing Administrative Effort: Because Purview offers native, centralized AI governance controls, organizations avoid stitching together custom logging pipelines, event triggers, and compliance dashboards.

Explanation:

E. Microsoft Purview.

Comprehensive AI Governance: Microsoft Purview is Microsoft's primary, unified governance solution built to discover, track, and protect data usage across AI models, including **Azure OpenAI**.

Auditing Data Usage: It provides specialized tools for monitoring data compliance, tracking sensitive information flow, and auditing exactly how enterprise data interacts with AI solutions.

Minimal Administrative Effort: It offers out-of-the-box, centralized dashboarding and compliance capabilities for AI, eliminating the need to build custom, complex streaming or monitoring pipelines.

Why the Other Options are Incorrect

A. Azure Monitor: While it captures base diagnostic and performance metrics (like latency and token counts), it requires significant custom logging configuration to serve as a full data usage and configuration audit

compliance engine.

B. Azure Stream Analytics: This is a real-time data ingestion and stream-processing engine, not a governance or compliance solution. Using it would require massive administrative overhead to build custom auditing pipelines.

C. Azure API Management: This acts as a gateway to manage and secure APIs (such as adding rate limiting or key authentication), but it does not natively provide comprehensive data compliance or configuration lifecycle auditing.

D. Azure Policy: This service enforces organizational standards and guards resource creation rules (like stopping a deployment outside a specific region) but does not monitor active data interactions or audit model runtime configuration

Question: 98 CertyIQ

HOTSPOT -
A company uses Azure OpenAI models that use grounding data from Microsoft Fabric for agents. The models are fine-tuned by using proprietary datasets.
You need to design a governance solution that meets the following requirements:
Restricts access to the grounding data to only assigned roles
Restricts model fine-tuning to only the AI engineering team
What should you include in the design? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

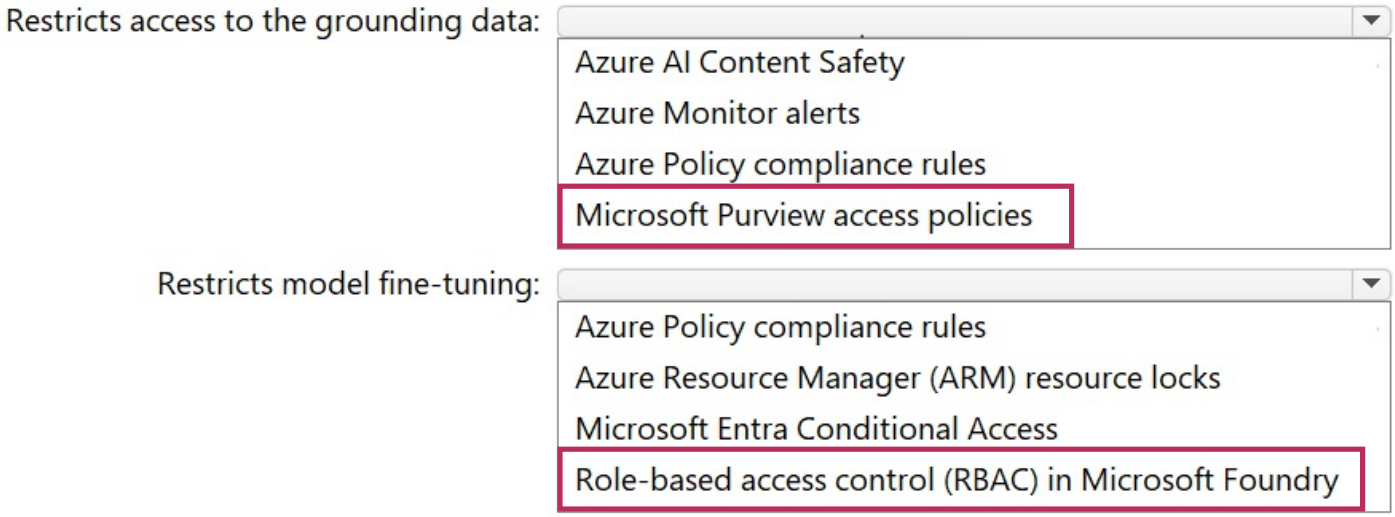

Answer: Microsoft Fabric seamlessly integrates with Microsoft Purview Information Protection and Purview access policies. Purview access policies provide centralized, declarative data-plane access governance over enterprise data estates (including Fabric OneLake items, lakehouses, and semantic models).

By authoring Purview access policies, administrators define role-based and attribute-based permissions governing exactly which user identities, agent service principals, and roles are permitted to read and retrieve grounding data, ensuring sensitive business data remains protected during RAG operations.

Access to initiate, configure, and manage fine-tuning jobs is governed by Azure Role-Based Access Control (RBAC) applied at the Foundry Hub and Project levels (e.g., assigning the Azure AI Developer role to the AI engineering team, while restricting others to Azure AI Reader or Azure AI User). This directly prevents unauthorized teams from triggering compute-intensive fine-tuning jobs or retraining models against sensitive proprietary datasets.

Answer Area

Explanation:

Restricts access to the grounding data: Microsoft Purview access policies

Microsoft Purview access policies are used to manage who can interact with or view sensitive data used for "grounding" (the source material the AI uses to generate answers).

Restricts model fine-tuning: Role-based access control (RBAC) in Microsoft Foundry

To prevent unauthorized users from modifying or fine-tuning the underlying AI models, administrators use **RBAC** within **Microsoft Foundry** to grant "least-privilege" access to only specifically authorized developers.

Question: 99 CertyIQ

DRAG DROP -
A company has a Microsoft Copilot Studio agent that has been in production for three months.
The agent has received positive feedback from users.
You need to identify the number of questions unanswered by the agent and the number of abandoned sessions between the users and the agent.
Which Copilot Studio insights should you use? To answer, drag the appropriate insights to the correct requirements. Each insight may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

Insights

- Conversation outcomes
- Generated answer rate and quality
- Reactions
- Survey results

Answer Area

	Insight
The number of unanswered questions:	Generated Answer Rate and Quality
The number of abandoned sessions:	Conversation Outcomes

Answer: Generated answer rate and quality: The Generated answer rate and quality insight panel inside Copilot Studio analytics monitors generative AI coverage. It explicitly tracks the ratio and absolute count of queries categorized by generative status: Answered vs. Unanswered. Drilling down into this insight provides makers with the exact list of user prompts where the agent could not formulate an answer due to knowledge gaps or fallback triggers.

Conversation outcomes: The Conversation outcomes dashboard tracks session lifecycle resolution. In Copilot Studio, conversation sessions are categorized into four distinct outcome states: Resolved, Escalated, Abandoned, and Unengaged.

Answer Area

	Insight
The number of unanswered questions:	Generated answer rate and quality
The number of abandoned sessions:	Conversation outcomes

Explanation:

Number of Unanswered Questions → **Generated Answer Rate and Quality**

Generated answer rate and quality tracks how well the AI system answers user questions.

It includes metrics such as:

Generated answer rate

Answer quality

Unanswered questions

Number of Abandoned Sessions → **Conversation Outcomes.**

What conversation outcomes track

Conversation outcomes analyze how conversations end.

Typical outcomes include:

Resolved conversations

Escalated conversations

Abandoned sessions

Failed conversations

Question: 100 CertyIQ

You are evaluating a Microsoft Copilot Studio agent that supports Microsoft Dynamics 365 Customer Service representatives.
You need to recommend a testing solution that meets the following requirements:
Evaluates agent effectiveness during active sessions
Validates whether the agent delivers accurate and helpful responses
Provides measurable, actionable insights for continuous improvement
What should you recommend?

A. Track resolution, deflection, and accuracy by using dashboards and use scripts to ensure consistent responses.

B. Perform load testing to validate how the agent scales under a high chat volume.

C. Review historical tickets to find agents that have the shortest resolution times.

D. Measure uptime and page load times.

Answer: A

Explanation:

Why optionA is the correct choice

Active-session evaluation – CopilotStudio provides real-time telemetry dashboards that capture resolution rate, deflection rate, and response accuracy while the conversation is occurring.
Accuracy & helpfulness validation – Scripted checkpoints can be embedded to compare the agent's output against predefined knowledge-base answers, flagging deviations instantly.
Measurable, actionable insights – The captured metrics are quantifiable (e.g., % of sessions resolved without human hand-off, average deflection score) and can be fed back into model-tuning loops for continuous improvement.
Alignment with Microsoft best-practice testing – Microsoft's guidance for CopilotStudio emphasizes continuous monitoring through dashboards and scripted validation to ensure the agent behaves as intended during live interactions.

Why the other options are less suitable

OptionB (load testing) – Focuses on scalability under high volume but does not assess response quality or accuracy during a session.
OptionC (historical ticket shortest resolution times) – Relies on past incident data and does not evaluate the current agent's performance in live conversations.
OptionD (uptime and page load times) – Measures infrastructure performance only; it provides no insight into the semantics or helpfulness of the agent's replies.

Hence, optionA best satisfies the requirement to evaluate effectiveness during active sessions, validate accuracy/helpfulness, and deliver actionable analytics.

References

CopilotStudio – Monitoring & analytics: https://learn.microsoft.com/en-us/copilot-studio/monitoring-analytics
Dynamics365 Customer Service – Agent performance dashboards: https://learn.microsoft.com/en-us/dynamics365/customer-service/analytics/performance-dashboards

Question: 101 CertyIQ

A company uses multiple Microsoft Copilot Studio agents across different channels.
You need to recommend a monitoring solution that provides comprehensive telemetry data and performance insights for the agents.
What should you include in the recommendation?

A. Application Insights
B. Azure Advisor
C. Azure DevOps
D. Microsoft Dynamics 365 Customer Voice

Answer: A

Explanation:

Answer: A – Application Insights

Why Application Insights is the optimal choice

Built-in telemetry for Copilot Studio agents – Application Insights automatically captures request,

dependency, and custom trace telemetry for Azure-hosted bot services and AI agents, enabling end-to-end performance monitoring and latency analysis.
Real-time dashboards & alerts – It provides built-in metrics (e.g., request rate, failure rates, response times) and configurable alerts that can be linked to agent health checks, supporting proactive incident response.
Deep analytics & AI-driven insights – Integrated with Azure Monitor workbooks and AI capabilities (anomaly detection, usage patterns), it offers the comprehensive insight required for capacity planning and optimization of multiple agents across channels.
Native integration with Azure resources – Since Copilot Studio agents run on Azure Bot Service and AI resources, Application Insights is the default monitoring service for those workloads, ensuring seamless data collection without additional configuration.

Why the other options are less suitable

B. Azure Advisor – Provides prescriptive recommendations for cost, security, and performance, but it does not collect or visualize telemetry data from running agents; it is advisory rather than monitoring.
C. Azure DevOps – Primarily supports source control, CI/CD pipelines, and work-item tracking; while it can host test results, it lacks dedicated telemetry-focused monitoring for AI agents and does not provide the granular performance metrics needed for production supervision.
D. Microsoft Dynamics 365 Customer Voice – A survey and feedback collection tool for gathering customer responses; it does not offer system-level telemetry or performance monitoring capabilities for Copilot agents.

Conclusion – To obtain comprehensive telemetry, real-time performance insights, and actionable alerts for multiple Copilot Studio agents, Application Insights is the only option that directly serves those monitoring requirements.

References

Application Insights for Azure Bot Service: https://learn.microsoft.com/azure/azure-monitor/app/azure-bot-service
Monitoring Azure AI resources with Application Insights: https://learn.microsoft.com/azure/monitor/app/monitor-azure-resources

Question: 102 CertyIQ

A company has an AI solution built by using Microsoft Copilot Studio and Power Platform. The solution is used by the company's sales, marketing, and customer service teams.
You are performing a return on AI investment (ROAI) analysis to evaluate the impact of the solution.
You need to identify which measurable business drivers to include in the analysis.
Which two business drivers should you identify? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. the reduced average case resolution time
B. market capitalization
C. economic market predictability
D. increased employee productivity
E. brand awareness

Answer: AD

Explanation:

Why the correct choices (A&D) are the most appropriate business drivers

Reduced average case resolution time (A) – This metric is directly tied to the AI-enabled automation that Copilot Studio and Power Platform deliver to customer-service workflows. Faster case closure translates into lower handling costs, higher customer-satisfaction scores, and measurable time-saved per incident, all of which can be quantified in monetary terms for an ROI model.

Increased employee productivity (D) – AI-driven assistants in sales, marketing, and service automate routine tasks (e.g., lead scoring, content generation, ticket triage). Productivity gains are expressed as additional deals closed, campaigns launched, or tickets resolved per employee hour, providing a clear quantitative contribution to the investment's financial return.

Why the other options are less suitable

Market capitalization (B) – This is a macro-level market metric that reflects overall company valuation and is influenced by many factors beyond the AI solution. It is not a controllable or attributable business driver for a specific AI implementation.

Economic market predictability (C) – While market conditions affect revenue, they are external and largely out of the control of the AI project team. ROI analyses focus on drivers that the organization can influence through its technology investments.

Brand awareness (E) – Brand perception can be impacted by AI-generated content, but measuring its direct financial impact is indirect and often subjective. ROI calculations require clear, quantifiable cost-savings or revenue uplift, which brand awareness alone does not provide.

Key takeaway for the exam – When evaluating ROAI, select business drivers that are directly measurable, attributable to the AI solution, and convertible into monetary value (e.g., time saved, productivity gains). Drivers that are external, macro-economic, or only loosely related to operational outcomes are not appropriate for a focused ROI analysis.

References

Microsoft Learn: Measuring AI ROI with Power Platform and Copilot Studio – https://learn.microsoft.com/en-us/power-platform/ai/measure-roi
Copilot Studio documentation: Business outcomes and ROI – https://learn.microsoft.com/en-us/copilot-studio/plan/business-outcomes

Question: 103 CertyIQ

A company has a Microsoft Foundry agent that summarizes customer feedback and recommends products to customers. The agent references data from multiple knowledge sources.
Users report that the agent response time is slow.
Telemetry data shows that the agent frequently reaches its token usage limit.
You need to recommend a solution to reduce token usage without degrading the quality of the generated responses.
What should you recommend?

A. Chunk documents during indexing.
B. Reduce the number of knowledge sources used by the agent.
C. Reconfigure the prompts to limit the amount of retrieved content from the knowledge sources.
D. Lower the maximum token usage limit for the responses.

Answer: C

Explanation:

C. Reconfigure the prompts to limit the amount of retrieved content from the knowledge sources This reduces token usage by: - Restricting how much grounding data is pulled into the prompt - Summarizing or chunking retrieved content before passing it to the model - Using more selective retrieval instructions - Reducing unnecessary context while keeping quality intact This is the standard optimization technique for Foundry and Azure OpenAI agents that hit token ceilings. It improves latency and reduces cost without degrading answer quality.

Question: 104 CertyIQ

HOTSPOT
-
A company has Microsoft Power Platform development, staging, and production environments. Each environment has its own Microsoft Dataverse tables and Azure AI Search index.
You are designing an application lifecycle management (AIM) process to deploy a Microsoft Copilot Studio agent between the environments.
The company has a Copilot Studio agent named Agent1 in development. Agent1 uses the following grounding data sources:
•A Dataverse table named CustomerOrders
•An Azure AI Search index named customer-knowledge
You need to deploy Agent1 to production. The solution must ensure that the agent uses the production grounding data sources, minimizes downtime, and handles credentials and endpoints securely.
What should you include in the deployment package solution, and what should you reconfigure after the deployment? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

Include in the deployment package solution:

- Agent1 only
- The data sources only
- Agent1 and the data source connections
- Agent1 and references to the data sources
- Agent1, the data sources, and the data source connections

Reconfigure after the deployment:

- The Dataverse connection only
- The Azure AI Search connection only
- The Dataverse and Azure AI Search connections
- The Agent1 configuration
- The environment variables

Answer:

Answer Area

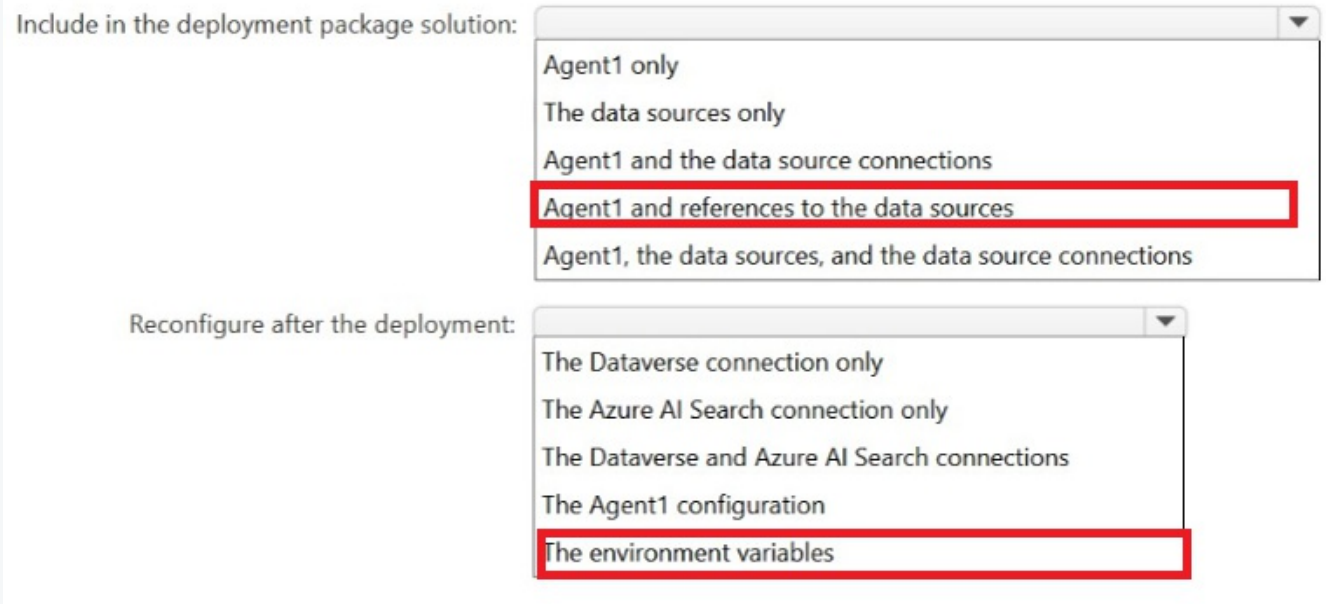

Explanation:

Include in the deployment package solution

Agent1 and references to the data sources: Because each tier (development, staging, production) maintains isolated tables and distinct target indices, you must never package the raw data sources or hardcoded infrastructure from the source environment. Packaging the agent alongside abstract references allows the solution definition to remain entirely transportable across your deployment boundaries.

Reconfigure after the deployment

The environment variables: In proper Power Platform ALM, connection parameters, endpoints, and target data assets (like your Dataverse table names or Azure AI Search index paths) are parameter-mapped into environment variables. After deploying the solution into production, updating these environment variables re-targets your agent to the live production data sources seamlessly without modifying or breaking the core bot code.

Question: 105 CertyIQ

A company has a Microsoft Copilot Studio agent that uses generative AI to assist Microsoft Dynamics 365 Customer Service representatives.
The agent currently exhibits a low resolution rate and a high escalation rate.
You need to identify the issue.
What should you use?

A. the Agent dashboard of Dynamics 365 Customer Service historical analytics
B. the Insights tab from the Search & intelligence settings of the Microsoft 365 admin center
C. the Copilot hub in the Power Platform admin center
D. the Analytics tab in Copilot Studio

Answer: D

Explanation:

Technical Justification for Correct Answer: D

The correct option to identify the issue with the low resolution rate and high escalation rate of the Microsoft Copilot Studio agent is **D. the Analytics tab in Copilot Studio**. Here's why:

Why D is the Best Choice:

Direct Relevance: The Analytics tab within Copilot Studio is specifically designed to provide insights into the performance of Copilot agents, including resolution rates and escalation rates. This direct relevance makes it the most appropriate tool for diagnosing the identified issues.
Real-time and Agent-Specific Data: Copilot Studio's Analytics tab offers real-time or near-real-time data that is directly related to the agent's interactions, allowing for precise identification of bottlenecks or areas of improvement in the AI-generated responses.
Customizable and Detailed Insights: This tab is likely to provide customizable insights that can drill down into specific aspects of the agent's performance, such as common themes in escalated cases or patterns in unresolved issues, facilitating targeted interventions.
Why Other Options are Less Suitable:

A. The Agent dashboard of Dynamics 365 Customer Service historical analytics:

While useful for overall agent performance, it might not offer the granular, AI-specific insights needed for a Copilot agent's generative AI performance issues.
Historical analytics might not provide the immediacy of insight required for timely intervention in current performance problems.
B. The Insights tab from the Search & intelligence settings of the Microsoft 365 admin center:

This option is more focused on search and intelligence across Microsoft 365, lacking the specificity needed for analyzing a Dynamics 365 integrated Copilot agent's effectiveness.
Unlikely to delve into the detailed performance metrics of a single Copilot agent.
C. The Copilot hub in the Power Platform admin center:

Although related to Power Platform, this might not directly offer the analytical insights into the operational performance (resolution and escalation rates) of a specific Copilot agent within Dynamics 365 Customer Service.
More about managing Copilot instances rather than analyzing their effectiveness in a customer service context.
Conclusion

Given the need for detailed, agent-specific, and real-time (or near-real-time) analytical data to address the low resolution rate and high escalation rate of the Microsoft Copilot Studio agent, **D. the Analytics tab in Copilot Studio** is the most appropriate choice.

References

1. **Microsoft Copilot Studio Documentation** - https://learn.microsoft.com/en-us/dynamics365/customer-service/copilot-studio (Navigate to sections on Analytics for insights)
2. **Dynamics 365 Customer Service Analytics** - https://learn.microsoft.com/en-us/dynamics365/customer-service/analytics (For broader context on analytics within Customer Service, though note the recommendation for Copilot-specific analytics)

Question: 106 CertyIQ

A company has a Microsoft Foundry generative AI model.
You need to evaluate the model's output to measure the overall quality and coherence of generated responses.
The evaluation must use GPT-4o as a judge and return a numeric score for each output.
Which type of metric should you use?

A. AI quality (NLP)
B. risk and safety
C. Groundedness
D. AI quality (AI assisted)

Answer: D

Explanation:

Technical Justification for Correct Answer D: AI quality (AI assisted)

The correct option for evaluating Microsoft Foundry generative AI model's output quality and coherence, with GPT-4o as the judge to return a numeric score, is **D. AI quality (AI assisted)**. Here's why:

Why D (AI quality - AI assisted) is the Best Choice:

Utilization of GPT-4o as a Judge: Since GPT-4o, a sophisticated AI model itself, is designated as the judge, an **AI-assisted** approach for evaluating AI quality is inherently implied. This option leverages the capabilities of one AI model (GPT-4o) to assess another (Microsoft Foundry), aligning with the concept of AI-assisted evaluation.
Numeric Score Requirement: AI quality metrics, especially when assisted by another AI, can be designed to output **numeric scores** based on predefined criteria (e.g., coherence, relevance, fluency). This fits the requirement for a quantifiable assessment of the generated responses.
Focus on Quality and Coherence: "AI quality (AI assisted)" directly implies an evaluation focused on the **overall quality** of the AI-generated content, including aspects like coherence, which is a primary concern for generative models.
Why Other Options are Less Suitable:

A. AI quality (NLP): While NLP-focused metrics could assess certain aspects of the generated text (e.g., grammar, syntax), this option does **not explicitly imply the use of another AI (GPT-4o) for assessment** or necessarily cover the broader quality and coherence evaluation required.
B. Risk and Safety: This metric type is more concerned with **ethical, legal, and safety implications** of the AI's output rather than its quality and coherence from a generative capability standpoint.
C. Groundedness: Groundedness refers to how well an AI's responses are **based on real-world knowledge or provided context**. While important, it **does not encompass the overall quality and coherence** evaluation scope as directly as AI quality metrics do, especially in the context of using another AI for judging.
Conclusion:

Given the scenario's specifics—utilizing GPT-4o for judging and the need for a numeric score on quality and coherence—**D. AI quality (AI assisted)** is the most appropriate metric type. It aligns with leveraging AI for AI evaluation, satisfies the numeric scoring requirement, and directly addresses the quality and coherence assessment needs.

References

1. **Microsoft Azure - Evaluating AI Models**: https://learn.microsoft.com/en-us/azure/machine-learning/how-to-deploy-advanced-endpoint#evaluate-your-model
   Relevant for understanding AI model evaluation strategies, including the use of auxiliary AI models for assessment.
2. **OpenAI GPT-4 Documentation - Evaluation Metrics**: https://docs.openai.com/api/v1/models/gpt-4
   Provides insight into how advanced AI models like GPT-4 can be used for evaluating generated text quality, though the exact link might not directly mention "AI quality (AI assisted)" as it's more of a

conceptual choice based on the question's framing.

Question: 107 CertyIQ

HOTSPOT
-
You use Microsoft Copilot Studio analytics to analyze the performance of a deployed Copilot Studio agent.
You need to identify which performance metrics to use to measure the following:
•The percentage of engaged sessions that are escalated to a live customer service representative
•The number of agent queries that cause a knowledge source error
What should you identify for each requirement? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

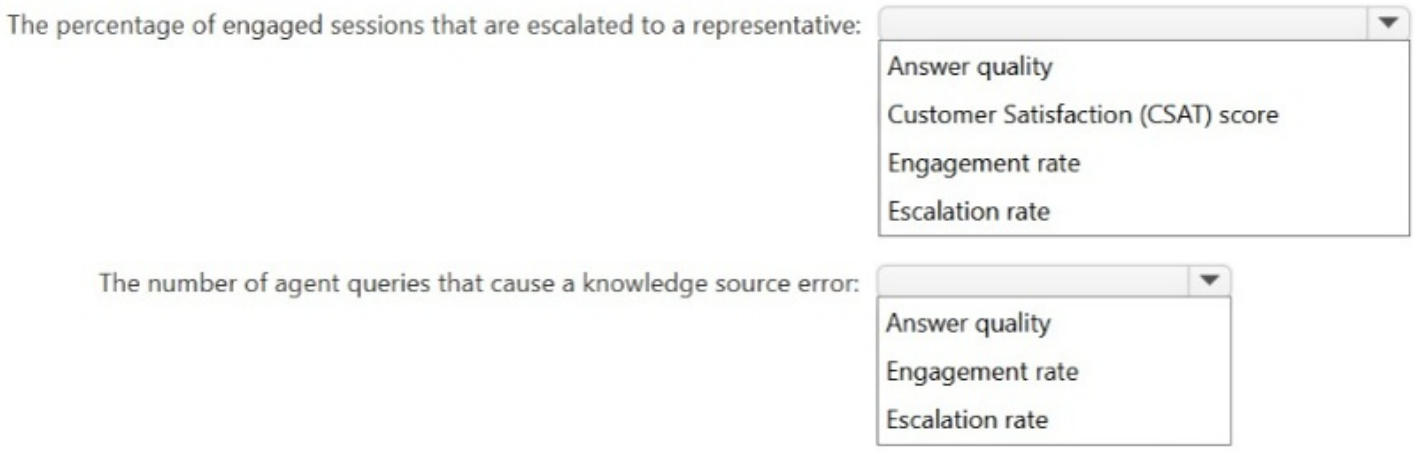

Answer:

Answer Area

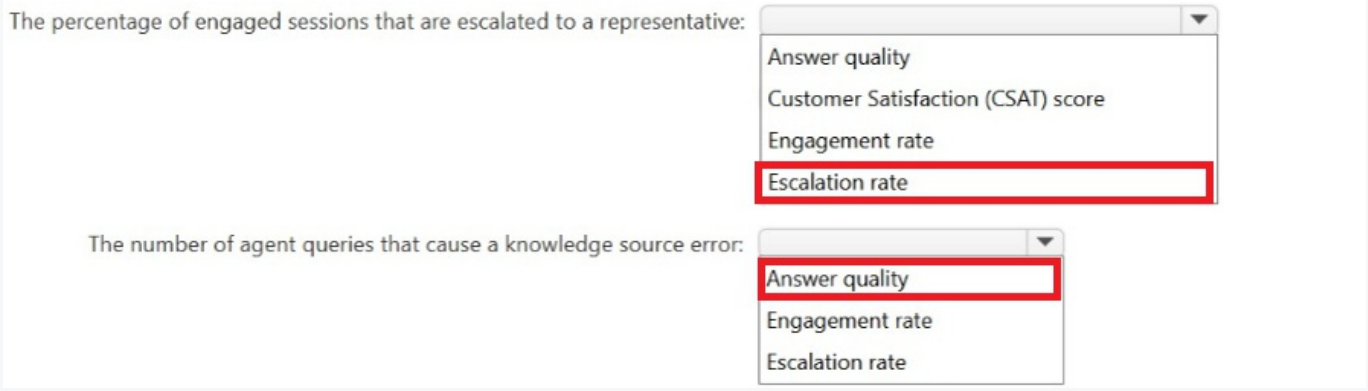

Explanation:

1. Escalation rate

Escalation rate: By official definition in Copilot Studio analytics dashboards, the escalation rate tracks the precise portion of engaged client sessions where the conversational agent handed off control to a human representative (e.g., via the Escalate system topic).

2. Answer quality

Answer quality: Generative AI and grounded data evaluation tools (such as knowledge-base monitoring dashboards) aggregate operational errors under the umbrella of Answer quality. This metric tracks groundedness, knowledge source retrieval issues, or empty returns that happen during runtime text generation.

Why Other Options Are Incorrect

Engagement rate: This measures user retention and interaction behavior, tracking whether users spent more than 10 seconds or had a meaningful, multi-turn interaction with the bot, rather than failure rates or agent pass-offs.

Customer Satisfaction (CSAT) score: This relies on qualitative customer feedback surveys submitted at the end of a session; it cannot directly monitor code-level API data failures or internal server network time-outs.

Question: 108 CertyIQ

A company plans to deploy a Microsoft Foundry agent.
You need to recommend an application lifecycle management (ALM) process to ensure that the agent evaluates against baseline accuracy metrics before being deployed.
What should you recommend?

A. Configure GitHub Actions for new agent versions.
B. Deploy each new agent version directly to production.
C. Use Observability in Foundry Control Plane with evaluation and drift monitoring.
D. Enable Application Insights and use Azure Monitor.

Answer: C

Explanation:

Technical Justification for Recommending Option C

Recommended Option: C. Use Observability in Foundry Control Plane with evaluation and drift monitoring.

Justification for Choosing C as the Best Option:

Alignment with Requirement: The question specifically asks for an ALM process to ensure the agent evaluates against baseline accuracy metrics before deployment. Option C directly addresses this by implying a controlled, evaluative environment (via Observability in Foundry Control Plane) where drift and accuracy can be monitored against baselines before the agent is fully deployed.
Pre-Deployment Evaluation: Unlike other options, C suggests an evaluation phase (with drift monitoring) that occurs presumably in a non-production environment or a staging area within the Foundry Control Plane, aligning perfectly with the prerequisite evaluation against baseline metrics.
Integrated Solution: Leveraging Observability within the Foundry Control Plane indicates an integrated approach, reducing the complexity of managing external tools solely for the evaluation phase.

Why Other Options are Less Suitable:

A. Configure GitHub Actions for new agent versions:

While GitHub Actions can automate deployment pipelines, this option does not inherently imply evaluation against baseline accuracy metrics. It's more about automation than quality assurance against specific metrics.
Lack of Direct Metric Evaluation: No clear indication of how baseline accuracy metrics are evaluated within this pipeline.
B. Deploy each new agent version directly to production:

High Risk: Bypasses any form of pre-deployment evaluation, posing a significant risk to production

environment stability and accuracy.
No Alignment with Requirement: Completely contradicts the need for pre-deployment evaluation against baseline metrics.
D. Enable Application Insights and use Azure Monitor:

Post-Deployment Focus: These tools are more suited for monitoring after deployment, not for evaluating against baseline metrics in a pre-deployment stage.
External Complexity: Introduces additional complexity by requiring setup and integration of external monitoring tools for a pre-deployment evaluation task.

Conclusion: Option C, **Use Observability in Foundry Control Plane with evaluation and drift monitoring**, is the most suitable recommendation because it directly supports pre-deployment evaluation against baseline accuracy metrics in an integrated manner, aligning with the core requirement of the question.

References

1. **Microsoft Foundry Documentation - Observability**: https://learn.microsoft.com/en-us/azure/foundry/observability
2. **Azure DevOps - Application Lifecycle Management (ALM)**: https://learn.microsoft.com/en-us/azure/devops/alm/

Note: Documentation links are subject to change. Ensure to search for the latest version if the links become outdated.

Question: 109 CertyIQ

You need to recommend a security solution for agents in a Microsoft Power Platform environment.
The agents must use only approved connectors and services. The solution must prevent the agents from accessing sensitive data.
What should you recommend?

A. Configure Azure Monitor to capture connector activity logs.
B. Enable a Microsoft Dataverse audit.
C. Deploy data loss prevention (DLP) policies in Power Platform.
D. Enable customer-managed keys in Microsoft Dataverse.

Answer: C

Explanation:

Technical Justification for Recommendation

Recommended Solution: C. Deploy data loss prevention (DLP) policies in Power Platform

Why C is the Best Option: Deploying DLP (Data Loss Prevention) policies in Power Platform is the most suitable solution for this scenario because:

Approved Connectors Enforcement: DLP policies can restrict the use of connectors to only those approved by the organization, directly addressing the requirement that agents use only approved connectors and services.
Sensitive Data Protection: DLP policies are specifically designed to prevent the leakage or unauthorized access to sensitive data, aligning with the need to block agents from accessing sensitive information.
Integration and Scope: Being a native Power Platform feature, DLP policies integrate seamlessly, providing a

comprehensive security layer across the environment without requiring external service setups beyond the initial configuration.

Why Other Options are Less Suitable:

A. Configure Azure Monitor to capture connector activity logs:
Reactive Not Proactive: This option focuses on logging activities rather than preventing unauthorized actions, making it a monitoring solution rather than a security enforcement mechanism.
Extra Complexity: Introduces additional complexity requiring integration with Azure Monitor.
B. Enable a Microsoft Dataverse audit:
Similar to A, Reactive: Auditing is about tracking changes and activities post-factum, not preventing the access to sensitive data or enforcing connector policies.
Limited Scope: Primarily focused on Dataverse, potentially leaving other Power Platform components unprotected.
D. Enable customer-managed keys in Microsoft Dataverse:
Encryption Focus: This option enhances data at rest encryption control but does not address the enforcement of approved connectors or prevent sensitive data access in the manner DLP policies do.
Scope Limitation: Like B, it's primarily Dataverse-focused and doesn't comprehensively cover the Power Platform's broader security needs as outlined.

References:

Microsoft Documentation: Data Loss Prevention policies in Power Platform
Microsoft Documentation: Secure your Power Platform environment

Question: 110 CertyIQ

A company has an AI business solution that uses Microsoft Copilot Studio agents.
You need to recommend prompt best practices to improve the effectiveness of agent interactions.
Which two actions should you include in the recommendation? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. Track the duration of the average user session.
B. Analyze the prompt length distribution.
C. Regularly test and refine the prompts based on user input.
D. Use clear and specific instructions in the prompts.
E. Measure system resource usage during prompt processing.

Answer: CD

Explanation:

Technical Justification for Recommended Actions (CD)

Correct Actions:

C. Regularly test and refine the prompts based on user input.
D. Use clear and specific instructions in the prompts.

Justification for Correct Actions:

C. Regularly test and refine the prompts based on user input: This action is crucial for improving the effectiveness of agent interactions in Microsoft Copilot Studio. By continuously testing prompts with actual

user input, the system can identify ambiguities, misunderstandings, or inefficient response patterns. Refining prompts based on this feedback loop enhances the accuracy and relevance of agent responses, directly improving user experience and interaction effectiveness. This approach aligns with iterative development principles and user-centered design, ensuring the AI system adapts to the company's specific use case and user behaviors.
D. Use clear and specific instructions in the prompts: Clear and specific instructions in prompts significantly reduce ambiguity, guiding the AI agent to provide more accurate and relevant responses. Well-crafted prompts with explicit expectations minimize misinterpretations, leading to more effective interactions. This best practice leverages the strengths of AI in processing well-defined tasks while mitigating the risks associated with vague or open-ended queries that might confuse the agent.

Why Other Options are Less Suitable for Directly Improving Agent Interactions' Effectiveness:

A. Track the duration of the average user session: While useful for understanding engagement patterns, session duration does not directly inform how to improve the effectiveness of agent interactions regarding prompt responsiveness or accuracy. It's more about user behavior analysis than prompt effectiveness.
B. Analyze the prompt length distribution: Analyzing prompt length might offer some insights into user behavior or system limitations, but it does not directly address the effectiveness of interactions in terms of accuracy, relevance, or user satisfaction. Extremely short or long prompts might indicate issues, but this doesn't directly guide prompt improvement for better agent interactions.
E. Measure system resource usage during prompt processing: This action is critical for system optimization, scalability, and performance monitoring but does not directly impact the effectiveness of agent interactions from a user experience or response accuracy standpoint. It's more about backend efficiency than frontline interaction quality.

References:

Microsoft Copilot Documentation - Crafting Effective Prompts (Note: Direct link might not be publicly available due to the specificity of the topic; the provided link is a closest match for public guidance on Copilot)
Microsoft AI for Business - Best Practices for Conversational AI (General guidance on conversational AI best practices applicable to prompt effectiveness)

Question: 111 CertyIQ

DRAG DROP
-
A company has a cloud-based AI solution that uses Azure OpenAI models.
You need to design a monitoring solution that meets the following requirements:
•Monitors performance metrics and operational health for the models
•Monitors AI apps and agents for compliance
•Uses Azure-native capabilities
•Minimizes development effort
What should you use for each requirement? To answer, drag the appropriate options to the correct requirements. Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

Answer:

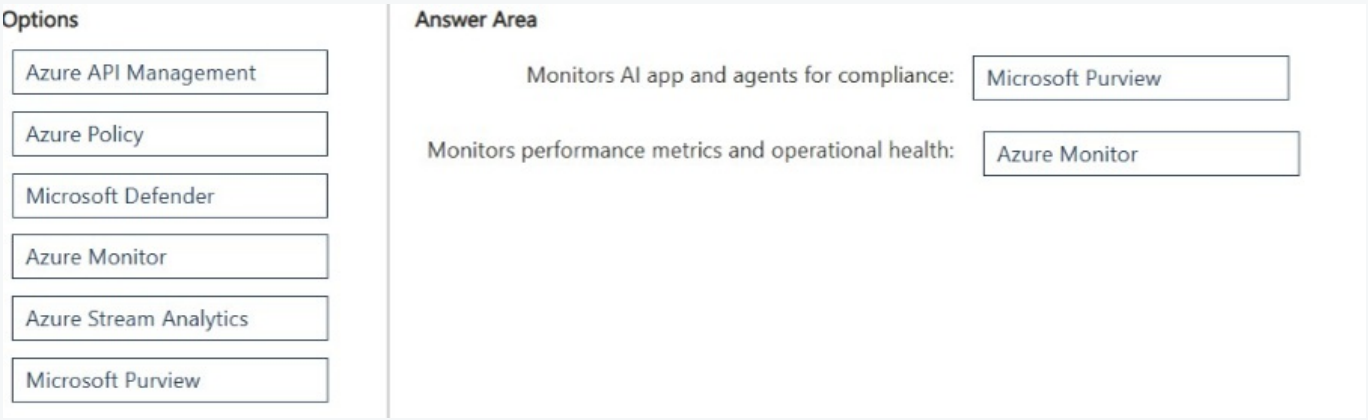

Explanation:

1. Monitors AI app and agents for compliance

Microsoft Purview (Correct): Within the Microsoft ecosystem, Microsoft Purview features a specialized AI Hub specifically engineered to discover, protect, and monitor data utilization inside generative AI applications. It enforces compliance policies, audits prompt-and-response interactions for data leakage, and provides visibility into risk data across user interactions.

Azure Policy (Incorrect): Azure Policy manages the compliance of the Azure cloud infrastructure (such as restricting service resource SKUs or enforcing specific region locations). It does not inspect runtime data interactions or monitor soft conversational workflows inside AI apps or agents.

Microsoft Defender (Incorrect): While Microsoft Defender for Cloud can detect security threats, vulnerability risks, or attack vectors against cloud endpoints, it is not the primary governance tool utilized for general organizational regulatory compliance monitoring of corporate data in AI scripts.

Monitors performance metrics and operational health

Azure Monitor (Correct): Azure Monitor is the foundational, native service designed to gather operational logs, performance telemetries, and system diagnostics. For solutions built on top of Azure OpenAI models, it serves as the central hub to map real-time performance indices, request latencies, system uptime, and compute throughput with minimal overhead.

Azure API Management (Incorrect): Azure API Management acts as a gateway tunnel to securely route, throttle, and expose APIs to applications. While it records baseline traffic logs, it does not act as a holistic telemetry system or serve as an operational metric dashboard for tracking end-to-end infrastructure health.

Azure Stream Analytics (Incorrect): This tool is a high-performance complex event processing engine engineered to compute massive real-time streaming data flows (like IoT sensor feeds); it does not provide

server diagnostic health diagnostics.

Question: 112 CertyIQ

A company has Microsoft Copilot Studio agents.
The company plans to deploy custom connectors across development, test and production environments.
You need to design an application lifecycle management (ALM) process to ensure consistency and prevent direct editing in production.
Which two actions should you include in the design? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. Deploy managed solutions to production.
B. Deploy unmanaged solutions to production.
C. Manually rebuild the agents in each environment.
D. Move the agents between the environments by using data export and import.
E. Include agents and connectors in a solution.

Answer: AE

Explanation:

Here is a structured technical justification for the correct options (A and E) in paragraph and bullet point format, followed by references:

Justification for Correct Options (A and E)

Option A: Deploy managed solutions to production Deploying managed solutions to production is crucial for maintaining consistency and preventing direct editing in the production environment. Managed solutions, unlike unmanaged ones, cannot be modified once deployed to an environment. This ensures that custom connectors and Copilot Studio agents are not altered in production, thereby maintaining the integrity of the deployed configuration. This approach also supports rollbacks to previous versions if issues arise, providing a higher level of control and reliability.

Key Points Supporting Option A:

Security and Integrity: Prevents direct modifications in production.
Reversibility: Allows for safe rollback to previous versions if needed.
Control: Ensures deployments are managed and tracked.

Option E: Include agents and connectors in a solution Including agents and connectors in a solution facilitates a unified approach to managing the application lifecycle. By packaging these components together, the organization can ensure consistency across development, test, and production environments. This unified package enables streamlined deployment processes, reduces the risk of configuration mismatches between environments, and simplifies tracking of changes across the lifecycle.

Key Points Supporting Option E:

Consistency Across Environments: Ensures uniformity in agent and connector configurations.
Simplified Deployment: Streamlines the deployment process with a single, unified package.
Change Management: Easier tracking of changes through the development lifecycle.

Why Other Options are Less Suitable

Option B (Deploy unmanaged solutions to production): Unmanaged solutions can be modified directly in

production, which violates the requirement to prevent direct editing and ensures no version control or easy rollback capabilities.
Option C (Manually rebuild the agents in each environment): This approach is error-prone, time-consuming, and does not ensure consistency across environments, contradicting the goal of ALM.
Option D (Move agents between environments by using data export and import): While this method allows for movement between environments, it does not inherently prevent direct editing in production nor does it ensure the structured, consistent deployment that managed solutions provide.

References

1. **Microsoft Documentation - Managed and Unmanaged Solutions**
2. https://docs.microsoft.com/en-us/power-apps/developer/data-platform/solution-files-overview#managed-and-unmanaged-solutions
3. **Microsoft Learn - Deploying Managed Solutions in Power Platform**
4. https://learn.microsoft.com/en-us/training/modules/deploy-managed-solutions-power-platform/

Question: 113 CertyIQ

A company uses multiple Microsoft Copilot Studio agents across different channels.
You need to recommend a monitoring solution that provides comprehensive telemetry data and performance insights for the agents.
What should you include in the recommendation?

A. Application Insights
B. Microsoft Dynamics 365 Customer Voice
C. Log Analytics
D. Microsoft Purview

Answer: A

Explanation:

Technical Justification for Recommended Monitoring Solution

Recommended Option: A. Application Insights

Comprehensive Telemetry Data: Application Insights provides detailed telemetry data, including custom events, metrics, and exceptions, which are essential for monitoring the performance of Microsoft Copilot Studio agents across various channels. Its ability to track user interactions and agent responses ensures a thorough understanding of system behavior.
Performance Insights: It offers pre-built analytics and visualization tools to gain insightful performance metrics (e.g., response times, error rates) for the agents, facilitating data-driven optimizations.
Integration Capabilities: Seamless integration with Azure and Microsoft ecosystem tools makes it an ideal choice for monitoring distributed agent deployments, allowing for unified visibility.

Why Other Options are Less Suitable

B. Microsoft Dynamics 365 Customer Voice:
Primarily designed for collecting and analyzing customer feedback, not for monitoring application or agent performance.
Lacks the telemetry and analytics capabilities required for comprehensive agent monitoring.
C. Log Analytics:

While Log Analytics (part of Azure Monitor) can process log data, it is more suited for infrastructure and log-based analysis rather than providing the comprehensive application performance monitoring offered by Application Insights.
Would require additional setup to achieve similar application-level insights.
D. Microsoft Purview:
Focuses on data governance, compliance, and data asset management across the enterprise.
Does not provide the performance monitoring and telemetry capabilities needed for agent-specific insights.

ConclusionApplication Insights is the most suitable choice due to its direct alignment with the requirements for comprehensive telemetry data collection and performance insight generation for Microsoft Copilot Studio agents. Its capabilities in tracking custom events, user interactions, and integrating with the Azure ecosystem make it superior to the other options for this specific use case.

References

Application Insights Documentation
Azure Monitor vs. Application Insights

Question: 114 CertyIQ

A company has an AI solution named Solution1 that is deployed to the production environment. Solution1 uses an Azure OpenAI model to generate marketing emails for existing customers.
During an internal review, you identify that Solution1 creates different emails depending on the customers' traits.
You need to recommend a strategy to mitigate the bias. The strategy must adhere to Microsoft responsible AI principles.
What should you recommend?

A. Modify Solution1 to randomly generate emails for different traits.
B. Modify the system instructions of Solution1.
C. Retrain the model by using a larger dataset.
D. Modify the contents of the training dataset.

Answer: D

Explanation:

Technical Justification for Recommended Strategy

Recommended Strategy: D. Modify the contents of the training dataset.

Justification for D being the Best Option:

Modifying the contents of the training dataset is the most effective strategy to mitigate bias in Solution1, aligning with Microsoft's responsible AI principles, which emphasize fairness and transparency. By addressing the root cause (the data), this approach ensures the Azure OpenAI model learns from a more balanced and representative dataset, reducing the generation of biased emails based on customer traits. This method prevents perpetuation of existing biases and improves the model's overall fairness.

Why Other Options are Less Suitable:

A. Modify Solution1 to randomly generate emails for different traits:
Inadequacy: Random generation does not address the underlying bias in the model's understanding of customer traits.
Quality Concerns: May lead to irrelevant or less effective marketing emails, disregarding the insights the

model could provide if unbiased.
B. Modify the system instructions of Solution1:
Superficial Fix: Altering instructions without addressing the model's biased learning may not significantly reduce trait-based differentiation in emails.
Limited Impact: Depends heavily on the quality of the new instructions, which might not counteract the model's inherent biases.
C. Retrain the model with a larger dataset:
Assumption of Improvement: Assuming a larger dataset is more balanced or less biased without ensuring its content quality or diversity might not mitigate, and could even amplify, existing biases.
Partial Solution: While increasing dataset size can help, if the new data introduces more of the same biases, the problem persists.

Microsoft Responsible AI Principles Alignment (for D):

Fairness: Directly addresses and attempts to balance the bias in the dataset, leading to more equitable treatment of different customer traits.
Transparency: Implicit in the process, as understanding and modifying the dataset is a clear, traceable step towards reducing bias.
Accountability: By targeting the dataset, the strategy acknowledges and takes responsibility for the bias at its source.

References:

Microsoft Responsible AI Principles
Mitigating Bias in AI Models with Azure

Question: 115 CertyIQ

A company has a canvas app named App1 in a Microsoft Power Platform environment named Env1.
Env1 uses a customer-managed key for data encryption. App1 connects to multiple data sources to retrieve and update customer and order information.
You need to recommend a solution to add Microsoft Copilot components to App1. The solution must NOT modify the current security or encryption configurations of Env1.
What should you include in the recommendation?

A. Modify the data sources of App1 to make them compatible with Copilot.
B. Duplicate App1 and republish the app in Env1.
C. Enable Copilot features for Env1.
D. Move App1 to a new environment that uses Microsoft-managed keys.

Answer: D

Explanation:

Technical Justification for the Correct Option (D)

Correct Option: D. Move App1 to a new environment that uses Microsoft-managed keys.

Why D is the best option:
Non-Intrusive to Existing Security/Encryption: Moving App1 to a new environment does not alter the existing security or encryption configurations of Env1, adhering to the primary requirement.
Copilot Compatibility: Microsoft Copilot components are designed to work seamlessly with environments using Microsoft-managed keys, ensuring compatibility and functionality without additional configuration

overhead.
Minimal Impact on App Functionality: Assuming the new environment is similarly configured (except for the key management), the impact on App1's functionality, especially its connections to data sources, is minimized.

Why Other Options are Less Suitable:

A. Modify the data sources of App1 to make them compatible with Copilot:
Unnecessary Complexity: Modifying data sources could introduce unnecessary complexity and potential security risks, altering the app's behavior.
Requirement Misalignment: The question does not indicate data source incompatibility with Copilot as the issue, but rather ensuring no change to Env1's security setup.
B. Duplicate App1 and republish the app in Env1:
Environment Constraint: Republishing in Env1 does not address the likely incompatibility of Copilot with customer-managed keys in the existing environment.
Redundant Effort: Duplicate efforts without resolving the core issue of Copilot integration under the current encryption setup.
C. Enable Copilot features for Env1:
Security/Encryption Configuration Impact: Enabling Copilot in Env1 might inherently require adjustments to the environment's settings, potentially affecting the customer-managed key configuration, which is not permitted.
Assumed Incompatibility: The implication is that Env1's current setup (customer-managed key) may not support Copilot without changes, making this option risky.

References:

Microsoft Power Platform: Environments and Security - https://learn.microsoft.com/en-us/power-platform/adminathing/environment
Microsoft Copilot for Power Apps - https://learn.microsoft.com/en-us/power-apps/maker/collaboration/collaboration-mf

Question: 116 CertyIQ

Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Overview -
Contoso, Ltd. is a high-tech manufacturing company that uses Microsoft Dynamics 365 Finance. Dynamics 365 Supply Chain Management, and Dynamics 365 Commerce for its North American operations. The company designs and develops innovative products that have many patents and proprietary technologies. The patents and engineering designs are closely guarded secrets.
Contoso executives want to integrate and adopt AI solutions to help scale the company in preparation for an anticipated period of rapid growth.
The company has multiple legal entities and Azure subscriptions that will be used in the adopted AI solutions.

Requirements -
AI Adoption -
The following executives will have specific responsibilities in the overall AI adoption:
•Chief Technology Officer (CTO): Select one Dynamics 365 Finance, Dynamics 365 Supply Chain Management or Dynamics 365 Commerce prebuilt AI agent and one custom Microsoft Copilot Studio AI agent to prioritize and deploy during the initial AI adoption phase.
•Chief Information Officer (CIO): Ensure that appropriate security labels are assigned to the data used by the AI agents.
•Chief Financial Officer (CFO): Analyze the return on investment (ROI) for the AI agents being deployed.
•Chief Information Security Officer (CISO): Discover and inventory AI resources for auditing.
•Chief Executive Officer (CEO): Ensure that all solutions adhere to industry-standard responsible AI practices.
All AI initiatives and agents will have a detailed business use case, a defined audience profile, and an estimated ROI that will compare the cost savings of the current process against the estimated costs of using the new AI solutions.
The company's research and development (R&D) department already has a custom Model Context Protocol (MCP) server that contains comprehensive product specifications and compliance data.
Prebuilt AI Agent -
The CTO has NOT yet selected which prebuilt AI agent to use in Dynamics 365 Supply Chain Management. The CTO wants to view available agent templates to identify which agent will add the most business value.
Depending on which high-priority AI agents are identified, its agent capabilities must be previewed in a discovery meeting with the relevant business operation stakeholders.
Custom AI Agent -
Contoso has identified the following custom AI agent requirements:
•The custom AI agent will use data from Dynamics 365 Supply Chain Management to answer questions for the manufacturing team as a low-code solution.
•The custom AI agent will be accessible from within Microsoft Teams.
•The custom AI agent must be designed to eventually connect to other agents that can be selected based on their description.
•The topics used in the custom AI agent will be selected based NOT on a trigger phrase, but on a description of the purpose of the query, to make the interactions more conversational.
•The custom AI agent must be able to answer questions about product specifications by using existing technologies. The product specifications are maintained by the R&D department.
•The custom AI agent must be integrated with and accessible from Dynamics 365 Supply Chain Management.
•The custom AI agent must be able to use Dynamics 365 Supply Chain Management business logic that is stored outside of the application.
Analysis, Reporting, and Troubleshooting
Contoso has identified the following analysis, reporting, and troubleshooting requirements:
•The CISO will audit all the AI solutions monthly for compliance and security.
•The CFO will analyze all the AI solutions quarterly to compare the estimated ROI against actual measured efficiencies and adoption. The CFO will use the Copilot Studio agent usage estimator to perform this analysis.
•The CISO wants to identify how much sensitive data was accessed for a given AI agent run and who accessed the data. Too much sensitive data accessed by a single user might indicate a high security risk.
•The CTO wants to track user feedback on the quality of the AI agent responses during user interactions with the agents. Consistently poor feedback will trigger an escalated reengineering discussion.
•The CEO wants a quarterly assessment of all the required metrics for their specific responsibilities. The tools used for the assessments must be Microsoft-recommended and must verify reliability, interpretability, fairness, and compliance.
•The CFO wants to identify how many interactions with the AI agents are abandoned on a given day as compared to resolved conversations. Too many abandoned sessions might indicate that Copilot Studio credits are being used inefficiently by end users.
Which two components for the custom AI agent should you include in the application lifecycle management (AIM) process? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. an Azure package

B. a ZIP package

C. a Microsoft Power Platform solution

D. a Cloud Scale Unit (CSU) package

E. an X++model

Answer: CE

Explanation:

Technical Justification for Correct Answer: CE

Case Study Question Analysis: Custom AI Agent - Application Lifecycle Management (ALM) Process Components

Correct Answer: CE

C. a Microsoft Power Platform solution
E. an X++ model

Justification for Correct Components:

C. a Microsoft Power Platform solution:
Rationale: The custom AI agent requires integration with Microsoft Teams and Dynamics 365 Supply Chain Management, indicating a need for a unified, low-code development environment that facilitates such integrations seamlessly.
Technical Alignment: Microsoft Power Platform (including Power Apps, Power Automate, and Power BI) is ideal for building low-code solutions that integrate with Microsoft ecosystem products, aligning with the custom AI agent's accessibility and integration requirements.
E. an X++ model:
Rationale: The custom AI agent must utilize Dynamics 365 Supply Chain Management's business logic stored outside the application and interact with product specifications, suggesting a need for deep integration with the Dynamics 365 platform.
Technical Alignment: X++ is the programming language used for developing custom logic within Dynamics 365 Supply Chain Management, making an X++ model essential for tapping into external business logic and ensuring seamless interaction with the platform's data and processes.

Why Other Options are Less Suitable:

A. an Azure package:
Inadequacy: While Azure is crucial for the overall AI strategy (especially for security, scalability, and potentially hosting the MCP server), the question focuses on the custom AI agent's ALM within the context of Dynamics 365 and Microsoft Teams integration, not the cloud infrastructure layer.
B. a ZIP package:
Inadequacy: A ZIP package is too generic and does not specifically address the integration, development, or deployment needs outlined for the custom AI agent within the Microsoft ecosystem.
D. a Cloud Scale Unit (CSU) package:
Inadequacy: CSU packages are related to deploying and managing Dynamics 365 environments at scale, not directly relevant to the development and ALM of a custom AI agent with specified integration requirements.

References

(For further study on the technologies mentioned)

1. **Microsoft Power Platform for Integrations and Low-Code Development**
2. https://powerplatform.microsoft.com/
3. **Dynamics 365 Supply Chain Management - Customization with X++**
4. https://docs.microsoft.com/en-us/dynamics365/unified-operations/deploy-dev/env 新-development-overview

Question: 117

CertyIQ

DRAG DROP

-
Case study
-
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study
-
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Overview
-
Contoso, Ltd. is a high-tech manufacturing company that uses Microsoft Dynamics 365 Finance. Dynamics 365 Supply Chain Management, and Dynamics 365 Commerce for its North American operations. The company designs and develops innovative products that have many patents and proprietary technologies. The patents and engineering designs are closely guarded secrets.
Contoso executives want to integrate and adopt AI solutions to help scale the company in preparation for an anticipated period of rapid growth.
The company has multiple legal entities and Azure subscriptions that will be used in the adopted AI solutions.
Requirements
-
AI Adoption
-
The following executives will have specific responsibilities in the overall AI adoption:
•Chief Technology Officer (CTO): Select one Dynamics 365 Finance, Dynamics 365 Supply Chain Management or Dynamics 365 Commerce prebuilt AI agent and one custom Microsoft Copilot Studio AI agent to prioritize and deploy during the initial AI adoption phase.
•Chief Information Officer (CIO): Ensure that appropriate security labels are assigned to the data used by the AI agents.
•Chief Financial Officer (CFO): Analyze the return on investment (ROI) for the AI agents being deployed.
•Chief Information Security Officer (CISO): Discover and inventory AI resources for auditing.
•Chief Executive Officer (CEO): Ensure that all solutions adhere to industry-standard responsible AI practices.
All AI initiatives and agents will have a detailed business use case, a defined audience profile, and an estimated ROI that will compare the cost savings of the current process against the estimated costs of using the new AI solutions.
The company's research and development (R&D) department already has a custom Model Context Protocol (MCP) server that contains comprehensive product specifications and compliance data.
Prebuilt AI Agent
-
The CTO has NOT yet selected which prebuilt AI agent to use in Dynamics 365 Supply Chain Management. The CTO wants to view available agent templates to identify which agent will add the most business value.
Depending on which high-priority AI agents are identified, its agent capabilities must be previewed in a discovery meeting with the relevant business operation stakeholders.
Custom AI Agent
-
Contoso has identified the following custom AI agent requirements:
•The custom AI agent will use data from Dynamics 365 Supply Chain Management to answer questions for the manufacturing team as a low-code solution.
•The custom AI agent will be accessible from within Microsoft Teams.
•The custom AI agent must be designed to eventually connect to other agents that can be selected based on their description.
•The topics used in the custom AI agent will be selected based NOT on a trigger phrase, but on a description of the purpose of the query, to make the interactions more conversational.
•The custom AI agent must be able to answer questions about product specifications by using existing technologies. The product specifications are maintained by the R&D department.

•The custom AI agent must be integrated with and accessible from Dynamics 365 Supply Chain Management.
•The custom AI agent must be able to use Dynamics 365 Supply Chain Management business logic that is stored outside of the application.
Analysis, Reporting, and Troubleshooting
Contoso has identified the following analysis, reporting, and troubleshooting requirements:
•The CISO will audit all the AI solutions monthly for compliance and security.
•The CFO will analyze all the AI solutions quarterly to compare the estimated ROI against actual measured efficiencies and adoption. The CFO will use the Copilot Studio agent usage estimator to perform this analysis.
•The CISO wants to identify how much sensitive data was accessed for a given AI agent run and who accessed the data. Too much sensitive data accessed by a single user might indicate a high security risk.
•The CTO wants to track user feedback on the quality of the AI agent responses during user interactions with the agents. Consistently poor feedback will trigger an escalated reengineering discussion.
•The CEO wants a quarterly assessment of all the required metrics for their specific responsibilities. The tools used for the assessments must be Microsoft-recommended and must verify reliability, interpretability, fairness, and compliance.
•The CFO wants to identify how many interactions with the AI agents are abandoned on a given day as compared to resolved conversations. Too many abandoned sessions might indicate that Copilot Studio credits are being used inefficiently by end users.
Which tools should you recommend to assist the CISO and the CIO with their specific responsibilities? To answer, drag the appropriate tools to the correct executives. Each tool may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

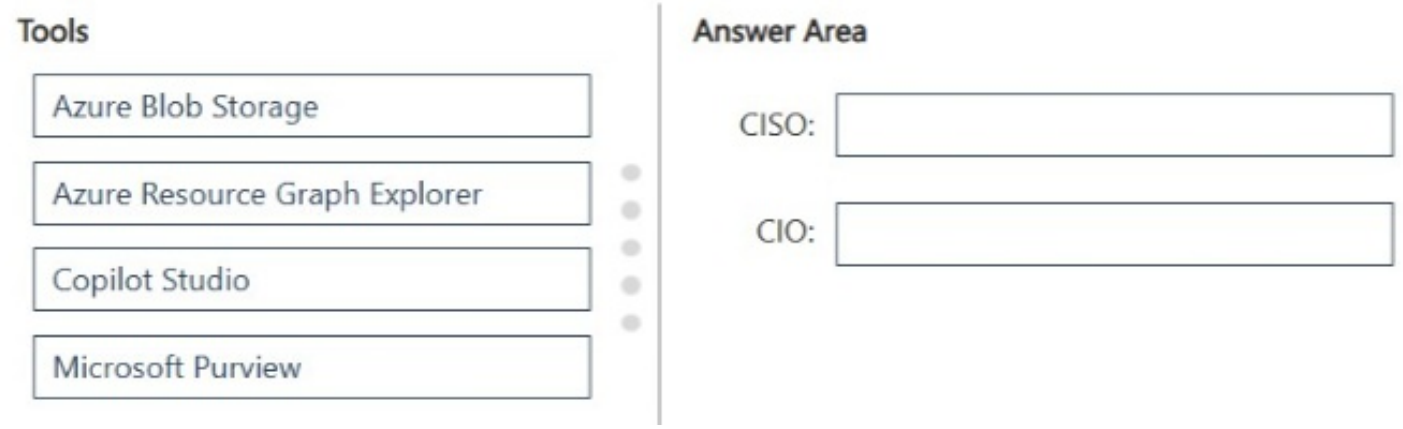

Answer:

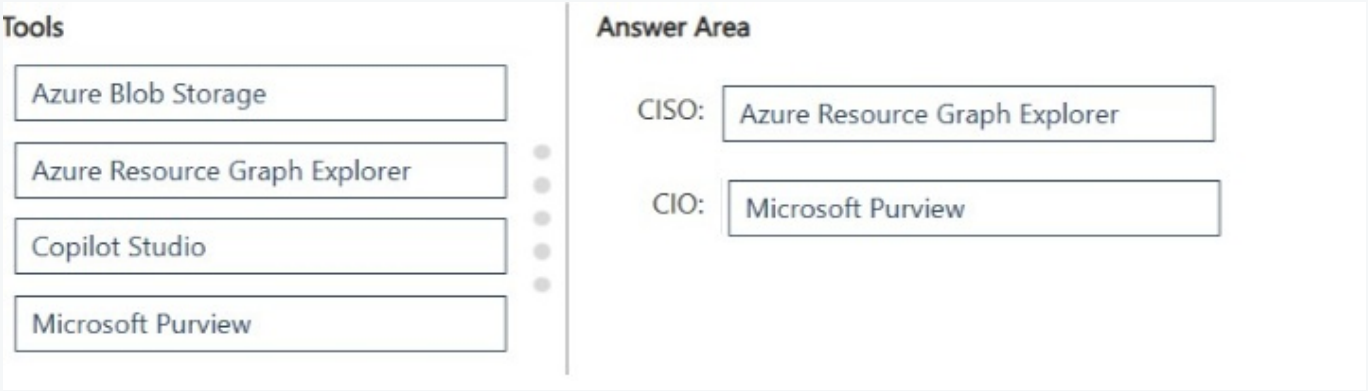

Explanation:

1. Chief Information Security Officer (CISO)

Azure Resource Graph Explorer: In corporate AI adoption strategies, the CISO's core responsibility includes discovery, cross-subscription auditing, and building a comprehensive inventory of running AI compute resources. Azure Resource Graph Explorer is the exact tool designed to execute high-performance structural queries at scale across thousands of Azure subscriptions to map active backend resources for security audits.

Chief Information Officer (CIO)

Microsoft Purview: The CIO is tasked with data governance and ensuring that appropriate data classifications, regulatory guardrails, and corporate protection policies are properly implemented. Microsoft Purview

features an integrated AI Hub built to handle sensitivity labeling, prevent information leakage, and enforce structural data access classifications for data assets utilized by conversational and autonomous agents.

Why Other Options Are Incorrect:

Copilot Studio: While Copilot Studio handles conversational bot building and logic configuration, it lacks the technical capability to perform a tenant-wide discovery of low-level infrastructure resources across multiple Azure subscriptions or handle enterprise-wide central repository sensitivity data labeling.

Azure Blob Storage: This is an unstructured file storage service (IaaS/PaaS object store) used to retain blobs, archives, and text data; it provides no inherent active management or automated dashboarding functionality for corporate executive auditing workflows.

Question: 118 CertyIQ

Case study -
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study -
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Overview -
Contoso, Ltd. is a high-tech manufacturing company that uses Microsoft Dynamics 365 Finance. Dynamics 365 Supply Chain Management, and Dynamics 365 Commerce for its North American operations. The company designs and develops innovative products that have many patents and proprietary technologies. The patents and engineering designs are closely guarded secrets.
Contoso executives want to integrate and adopt AI solutions to help scale the company in preparation for an anticipated period of rapid growth.
The company has multiple legal entities and Azure subscriptions that will be used in the adopted AI solutions.
Requirements -
AI Adoption -
The following executives will have specific responsibilities in the overall AI adoption:
•Chief Technology Officer (CTO): Select one Dynamics 365 Finance, Dynamics 365 Supply Chain Management or Dynamics 365 Commerce prebuilt AI agent and one custom Microsoft Copilot Studio AI agent to prioritize and deploy during the initial AI adoption phase.
•Chief Information Officer (CIO): Ensure that appropriate security labels are assigned to the data used by the AI agents.
•Chief Financial Officer (CFO): Analyze the return on investment (ROI) for the AI agents being deployed.
•Chief Information Security Officer (CISO): Discover and inventory AI resources for auditing.
•Chief Executive Officer (CEO): Ensure that all solutions adhere to industry-standard responsible AI practices.
All AI initiatives and agents will have a detailed business use case, a defined audience profile, and an estimated ROI that will compare the cost savings of the current process against the estimated costs of using the new AI solutions.
The company's research and development (R&D) department already has a custom Model Context Protocol (MCP) server that contains comprehensive product specifications and compliance data.
Prebuilt AI Agent -
The CTO has NOT yet selected which prebuilt AI agent to use in Dynamics 365 Supply Chain Management. The CTO wants to view available agent templates to identify which agent will add the most business value.
Depending on which high-priority AI agents are identified, its agent capabilities must be previewed in a discovery meeting with the relevant business operation stakeholders.

Custom AI Agent -
Contoso has identified the following custom AI agent requirements:
•The custom AI agent will use data from Dynamics 365 Supply Chain Management to answer questions for the manufacturing team as a low-code solution.
•The custom AI agent will be accessible from within Microsoft Teams.
•The custom AI agent must be designed to eventually connect to other agents that can be selected based on their description.
•The topics used in the custom AI agent will be selected based NOT on a trigger phrase, but on a description of the purpose of the query, to make the interactions more conversational.
•The custom AI agent must be able to answer questions about product specifications by using existing technologies. The product specifications are maintained by the R&D department.
•The custom AI agent must be integrated with and accessible from Dynamics 365 Supply Chain Management.
•The custom AI agent must be able to use Dynamics 365 Supply Chain Management business logic that is stored outside of the application.
Analysis, Reporting, and Troubleshooting
Contoso has identified the following analysis, reporting, and troubleshooting requirements:
•The CISO will audit all the AI solutions monthly for compliance and security.
•The CFO will analyze all the AI solutions quarterly to compare the estimated ROI against actual measured efficiencies and adoption. The CFO will use the Copilot Studio agent usage estimator to perform this analysis.
•The CISO wants to identify how much sensitive data was accessed for a given AI agent run and who accessed the data. Too much sensitive data accessed by a single user might indicate a high security risk.
•The CTO wants to track user feedback on the quality of the AI agent responses during user interactions with the agents. Consistently poor feedback will trigger an escalated reengineering discussion.
•The CEO wants a quarterly assessment of all the required metrics for their specific responsibilities. The tools used for the assessments must be Microsoft-recommended and must verify reliability, interpretability, fairness, and compliance.
•The CFO wants to identify how many interactions with the AI agents are abandoned on a given day as compared to resolved conversations. Too many abandoned sessions might indicate that Copilot Studio credits are being used inefficiently by end users.
What should you recommend to assist the CEO with their specific responsibilities?

A. the Microsoft Service Trust Portal
B. Microsoft Foundry Tools
C. Microsoft Purview
D. the Responsible AI dashboard
E. Compliance Center

Answer: D

Explanation:

Technical Justification for Recommended Solution

Question: What should you recommend to assist the CEO with their specific responsibilities regarding AI adoption, focusing on verifying reliability, interpretability, fairness, and compliance?

Recommended Answer: D. the Responsible AI dashboard

Justification for D (Responsible AI dashboard):

Alignment with CEO Responsibilities: The CEO's requirements explicitly mention ensuring all solutions adhere to industry-standard responsible AI practices. The Responsible AI dashboard is designed to provide insights and tools for monitoring and ensuring AI systems are reliable, interpretable, fair, and compliant, directly aligning with the CEO's responsibilities.
Microsoft-Recommended Tool: The question specifies that the tools used must be Microsoft-recommended. The Responsible AI dashboard, part of Microsoft's AI for Everyone initiative, fits this criterion, offering a centralized platform for overseeing AI deployments from a responsibility standpoint.
Comprehensive Metrics Coverage: The dashboard is likely to offer or integrate with functionalities to assess the required metrics (reliability, interpretability, fairness, compliance) on a quarterly basis, providing a holistic

view necessary for the CEO's evaluations.

Why Other Options are Less Suitable:

A. Microsoft Service Trust Portal:
Primarily focused on transparency into Microsoft's security, compliance, and privacy practices for its cloud services.
Not directly designed for monitoring AI solution responsibilities within a customer's environment.
B. Microsoft Foundry Tools:
More related to the development and deployment of AI models rather than overseeing the responsible AI practices across deployed solutions.
Lack of explicit alignment with the CEO's specified responsibilities.
C. Microsoft Purview:
While useful for data governance, compliance, and security across the enterprise, it does not specifically target the oversight of AI solutions' responsible AI practices.
Might be more relevant for the CIO or CISO's responsibilities.
E. Compliance Center:
Similar to Microsoft Purview, it's more broadly focused on organizational compliance and not specifically tailored for AI solution oversight from a responsible AI standpoint.
Less direct alignment with the unique aspects of the CEO's AI adoption responsibilities.

References:

1. **Microsoft Responsible AI Dashboard** - https://www.microsoft.com/en-us/trust-center/responsible-ai
2. **Microsoft AI for Everyone** - https://www.microsoft.com/en-us/ai-for-everyone

Question: 119 CertyIQ

DRAG DROP
-
Case study
-
This is a case study. Case studies are not timed separately from other exam sections. You can use as much exam time as you would like to complete each case study. However, there might be additional case studies or other exam sections. Manage your time to ensure that you can complete all the exam sections in the time provided. Pay attention to the Exam Progress at the top of the screen so you have sufficient time to complete any exam sections that follow this case study.
To answer the case study questions, you will need to reference information that is provided in the case. Case studies and associated questions might contain exhibits or other resources that provide more information about the scenario described in the case. Information provided in an individual question does not apply to the other questions in the case study.
A Review Screen will appear at the end of this case study. From the Review Screen, you can review and change your answers before you move to the next exam section. After you leave this case study, you will NOT be able to return to it.
To start the case study
-
To display the first question in this case study, select the "Next" button. To the left of the question, a menu provides links to information such as business requirements, the existing environment, and problem statements. Please read through all this information before answering any questions. When you are ready to answer a question, select the "Question" button to return to the question.
Overview
-
Contoso, Ltd. is a high-tech manufacturing company that uses Microsoft Dynamics 365 Finance. Dynamics 365 Supply Chain Management, and Dynamics 365 Commerce for its North American operations. The company designs and develops innovative products that have many patents and proprietary technologies. The patents and engineering designs are closely guarded secrets.
Contoso executives want to integrate and adopt AI solutions to help scale the company in preparation for an anticipated period of rapid growth.
The company has multiple legal entities and Azure subscriptions that will be used in the adopted AI solutions.

Requirements
-
AI Adoption
-
The following executives will have specific responsibilities in the overall AI adoption:
•Chief Technology Officer (CTO): Select one Dynamics 365 Finance, Dynamics 365 Supply Chain Management or Dynamics 365 Commerce prebuilt AI agent and one custom Microsoft Copilot Studio AI agent to prioritize and deploy during the initial AI adoption phase.
•Chief Information Officer (CIO): Ensure that appropriate security labels are assigned to the data used by the AI agents.
•Chief Financial Officer (CFO): Analyze the return on investment (ROI) for the AI agents being deployed.
•Chief Information Security Officer (CISO): Discover and inventory AI resources for auditing.
•Chief Executive Officer (CEO): Ensure that all solutions adhere to industry-standard responsible AI practices.
All AI initiatives and agents will have a detailed business use case, a defined audience profile, and an estimated ROI that will compare the cost savings of the current process against the estimated costs of using the new AI solutions.
The company's research and development (R&D) department already has a custom Model Context Protocol (MCP) server that contains comprehensive product specifications and compliance data.
Prebuilt AI Agent
-
The CTO has NOT yet selected which prebuilt AI agent to use in Dynamics 365 Supply Chain Management. The CTO wants to view available agent templates to identify which agent will add the most business value.
Depending on which high-priority AI agents are identified, its agent capabilities must be previewed in a discovery meeting with the relevant business operation stakeholders.
Custom AI Agent
-
Contoso has identified the following custom AI agent requirements:
•The custom AI agent will use data from Dynamics 365 Supply Chain Management to answer questions for the manufacturing team as a low-code solution.
•The custom AI agent will be accessible from within Microsoft Teams.
•The custom AI agent must be designed to eventually connect to other agents that can be selected based on their description.
•The topics used in the custom AI agent will be selected based NOT on a trigger phrase, but on a description of the purpose of the query, to make the interactions more conversational.
•The custom AI agent must be able to answer questions about product specifications by using existing technologies. The product specifications are maintained by the R&D department.
•The custom AI agent must be integrated with and accessible from Dynamics 365 Supply Chain Management.
•The custom AI agent must be able to use Dynamics 365 Supply Chain Management business logic that is stored outside of the application.
Analysis, Reporting, and Troubleshooting
Contoso has identified the following analysis, reporting, and troubleshooting requirements:
•The CISO will audit all the AI solutions monthly for compliance and security.
•The CFO will analyze all the AI solutions quarterly to compare the estimated ROI against actual measured efficiencies and adoption. The CFO will use the Copilot Studio agent usage estimator to perform this analysis.
•The CISO wants to identify how much sensitive data was accessed for a given AI agent run and who accessed the data. Too much sensitive data accessed by a single user might indicate a high security risk.
•The CTO wants to track user feedback on the quality of the AI agent responses during user interactions with the agents. Consistently poor feedback will trigger an escalated reengineering discussion.
•The CEO wants a quarterly assessment of all the required metrics for their specific responsibilities. The tools used for the assessments must be Microsoft-recommended and must verify reliability, interpretability, fairness, and compliance.
•The CFO wants to identify how many interactions with the AI agents are abandoned on a given day as compared to resolved conversations. Too many abandoned sessions might indicate that Copilot Studio credits are being used inefficiently by end users.
Which Copilot Studio analytics metrics should you recommend to assist the executives with their specific responsibilities? To answer, drag the appropriate metrics to the correct executives. Each metric may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

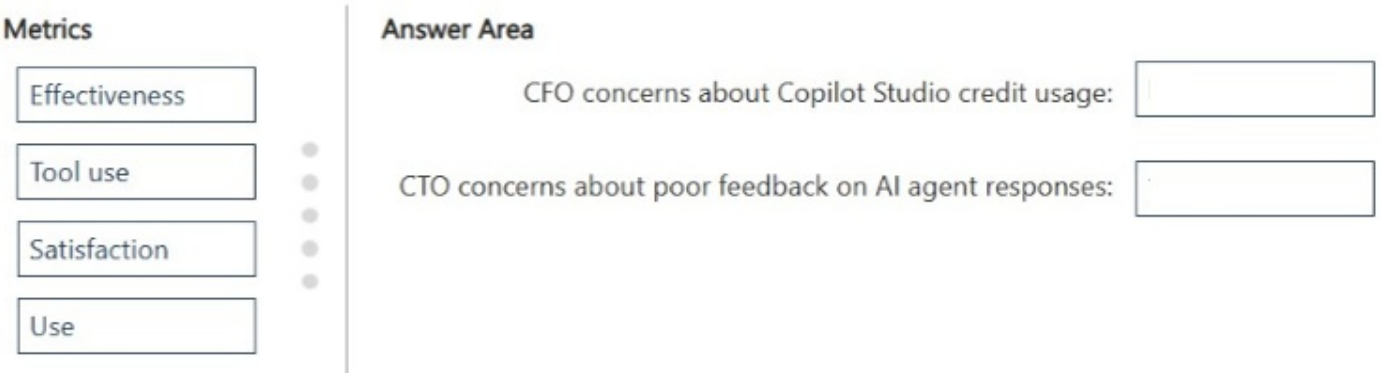

Answer:

Explanation:

1. CFO concerns about Copilot Studio credit usage

Effectiveness: In enterprise cost optimization, a Chief Financial Officer (CFO) focuses on how efficiently corporate investments and tenant credits are consumed. This maps directly to Effectiveness dashboards within Copilot Studio analytics, which measure conversation outcomes — specifically tracking how many sessions were successfully resolved versus abandoned before completion.

2. CTO concerns about poor feedback on AI agent responses

Satisfaction: A Chief Technology Officer (CTO) monitors user perception, generative quality, and qualitative telemetry trends. User sentiment score distribution and written remarks fall under Satisfaction metrics (such as Customer Satisfaction (CSAT) surveys), which evaluate how well the AI agent's language model fulfilled user expectations.

Why Other Options Are Incorrect

Use / Tool use: These technical metrics gauge the volume or adoption footprint of the agent framework (e.g., total message triggers or API connector execution frequencies). While important for load scaling, they do not assess qualitative output value (Satisfaction) or task completion health versus billable burn rates (Effectiveness).

Question: 120 CertyIQ

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.
After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.
A company has a team that analyzes its customers by using a manual process.

You are designing an AI-based agent to automate and improve the process.
You need to recommend on which platform to build the agent. The solution must meet the following requirements:
•Use generative AI to answer common questions.
•Provide analytics to review AI performance.
•Identify customer demographics.
•Minimize custom development.
Solution: You recommend GitHub Copilot.
Does this meet the goal?

A. Yes
B. No

Answer: B

Explanation:

Technical Justification for Solution Evaluation

Scenario Recap: Automate customer analysis with an AI-based agent, leveraging generative AI, analytics for AI performance review, customer demographic identification, and minimal custom development.

Evaluated Solution: GitHub Copilot

Assessment of Solution against Requirements:

Use Generative AI to Answer Common Questions:

GitHub Copilot excels in providing generative AI capabilities for coding tasks, potentially assistive in developing question-answering functionalities but not directly designed for end-user questioning interfaces without significant development.
Partial Fit
Provide Analytics to Review AI Performance:

GitHub Copilot's primary focus is on code assistance; built-in analytics for reviewing AI performance in the context of customer question answering are not a highlighted feature.
Poor Fit
Identify Customer Demographics:

This task is more aligned with data analysis and CRM functionalities. GitHub Copilot does not inherently support demographic analysis.
No Fit
Minimize Custom Development:

While GitHub Copilot can reduce development time for coding tasks, the other requirements would still necessitate significant custom development to integrate and achieve the desired functionality.
Partial Fit

Why the Correct Answer is B ("No"):The solution (GitHub Copilot) does not comprehensively meet the stated goals, particularly lacking in providing analytics for AI performance review and identifying customer demographics without substantial custom development. Its strength in generative AI is more coding-centric than directly applicable to the question-answering and demographic analysis needs of the project.

Why Other Options (Implicitly Considered) Might Be More Suitable:

Alternative Platforms (e.g., Microsoft Azure Cognitive Services, Google Cloud AI Platform, or Salesforce Einstein) might offer:

Direct integrations for generative AI question-answering interfaces.
Built-in or easily integrable analytics for AI performance.
Tools or services for customer demographic analysis.
More comprehensive support for minimizing custom development across all requirements.

References

1. **GitHub Copilot Documentation** - For details on its coding assistance capabilities: https://docs.github.com/en/github/writing-on-github/about-github-copilot
2. **Microsoft Azure Cognitive Services** - As an example of a platform offering a broad spectrum of AI services including Q&A, Analytics, and more: https://azure.microsoft.com/en-us/services/cognitive-services/

Question: 121 CertyIQ

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.
After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.
A company has a team that analyzes its customers by using a manual process.
You are designing an AI-based agent to automate and improve the process.
You need to recommend on which platform to build the agent. The solution must meet the following requirements:
•Use generative AI to answer common questions.
•Provide analytics to review AI performance.
•Identify customer demographics.
•Minimize custom development.
Solution: You recommend Microsoft Security Copilot.
Does this meet the goal?

A. Yes
B. No

Answer: B

Explanation:

Justification for Correct Answer: B. No

The recommended solution, Microsoft Security Copilot, does not meet the stated goals for automating and improving the customer analysis process with the specified requirements. Here's why:

Use Generative AI to Answer Common Questions: While Microsoft Security Copilot leverages AI, its primary focus is on security operations and threat response, not on generative AI capabilities for answering customer service or analysis-related questions. This mismatch makes it less suitable for the task.
Provide Analytics to Review AI Performance: Though Security Copilot might offer some analytics related to security incidents, it does not provide the type of analytics tailored for reviewing the performance of an AI agent designed for customer analysis and question-answering.
Identify Customer Demographics: This requirement is unrelated to the core functionalities of Microsoft Security Copilot, which are centered around security and compliance, not customer data analysis or demographics identification.
Minimize Custom Development: Given the mismatch between Security Copilot's capabilities and the requirements, significant custom development would be necessary to adapt it for the intended use,

contradicting this goal.

Why Other Options (Implicit, since only one is provided) Might Be More Suitable (Hypothetical Analysis):

Alternative Platforms (e.g., Microsoft Bot Framework, Azure Cognitive Services): Might better meet the goals by:
Integrating generative AI models (e.g., for question-answering).
Offering customizable analytics solutions (e.g., Azure Analytics Services).
Allowing for easy integration with customer database systems for demographic analysis.
Providing a framework that minimizes custom development through pre-built components and integrations tailored for AI-powered chatbots or agents.

Conclusion:Given the specific requirements for the AI-based agent, Microsoft Security Copilot is not the best fit due to its focus on security operations rather than customer analysis, generative AI for questioning answering, and the potential for significant custom development needs.

References

1. **Microsoft Security Copilot Documentation** - https://learn.microsoft.com/en-us/microsoft-365/security/defender-endpoint/security-copilot?view=o365-worldwide
2. **Azure Cognitive Services for AI-Powered Solutions** - https://azure.microsoft.com/en-us/services/cognitive-services/

Question: 122 CertyIQ

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.
After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.
A company has a team that analyzes its customers by using a manual process.
You are designing an AI-based agent to automate and improve the process.
You need to recommend on which platform to build the agent. The solution must meet the following requirements:
•Use generative AI to answer common questions.
•Provide analytics to review AI performance.
•Identify customer demographics.
•Minimize custom development.
Solution: You recommend Microsoft Copilot Studio.
Does this meet the goal?

A. Yes
B. No

Answer: A

Explanation:

Technical Justification for Correct Answer: A. Yes

The recommended solution, Microsoft Copilot Studio, meets the stated goals for building an AI-based agent to automate and improve the customer analysis process. Here's a breakdown of why it is the best fit and why other implied options (not listed but considered for a comprehensive justification) might be less suitable:

Why Microsoft Copilot Studio is the Correct Choice:

Use Generative AI to Answer Common Questions: Microsoft Copilot Studio is integrated with generative AI capabilities, enabling the creation of agents that can comprehensively respond to frequently asked questions by generating human-like text. This aligns perfectly with the requirement for handling common inquiries automatically.
Provide Analytics to Review AI Performance: The platform offers built-in analytics tools to monitor the performance of the AI agent, allowing for data-driven optimizations and improvements over time.
Identify Customer Demographics: While the primary function of Copilot Studio is AI-driven interaction, its integration within the Microsoft ecosystem (e.g., with Dynamics, Azure services) facilitates access to tools that can analyze and identify customer demographics, either through direct integration or by leveraging complementary Microsoft products.
Minimize Custom Development: Copilot Studio is designed with a low-code approach, significantly reducing the need for custom development. This accelerates the deployment of the AI agent and aligns with the goal of minimizing bespoke coding efforts.

Why Other Implied Options Might Be Less Suitable:

Google Dialogflow or Similar Platforms:
Might require more custom development for full integration with analytics and demographic analysis tools not natively part of the platform.
Generative AI capabilities might not be as seamlessly integrated as in Copilot Studio.
Building from Scratch with Azure Cognitive Services (or Similar Cloud AI Services):
Would inherently maximize custom development, contradicting the requirement to minimize it.
Though highly customizable, the initial setup and integration of analytics and demographic identification tools would be more labor-intensive.

Conclusion:

Given the requirements, Microsoft Copilot Studio is the most appropriate choice due to its integrated generative AI, built-in analytics, potential for demographic analysis through Microsoft's broader ecosystem, and low-code development approach.

References

1. **Microsoft Copilot Studio Capabilities**: https://learn.microsoft.com/en-us/microsoft-365/account-security/microsoft-copilot?view=o365-worldwide
2. **Azure Cognitive Services for Custom AI Development (for contrast)**: https://azure.microsoft.com/en-us/services/cognitive-services/

Question: 123 CertyIQ

A company uses multiple Microsoft Copilot Studio agents across different channels.
You need to recommend a monitoring solution that provides comprehensive telemetry data and performance insights for the agents.
What should you include in the recommendation?

A. Microsoft Power BI
B. the Analytics tab in Copilot Studio
C. Log Analytics
D. Azure Advisor

Answer: C

Explanation:

Technical Justification for Recommended Monitoring Solution

Recommended Option: C. Log Analytics

Comprehensive Telemetry Data: Log Analytics, part of Azure Monitor, is designed to collect, store, and analyze vast amounts of log data from various sources, including applications like Microsoft Copilot Studio agents. This makes it ideal for gathering comprehensive telemetry data across multiple agents.
Performance Insights: It provides deep analytics capabilities, enabling the creation of custom dashboards and queries to extract specific performance insights tailored to the company's needs regarding the agents' operations.
Scalability and Integration: As a cloud-based service, Log Analytics scales with the company's growth and easily integrates with other Azure services, offering a unified monitoring platform for heterogeneous environments.

Why Other Options are Less Suitable:

A. Microsoft Power BI:
Primary Function: Business Analytics and Visualization, not primary monitoring or logging.
Telemetry Collection: Not equipped to directly collect telemetry data from agents; would require data to be pre-processed and fed into Power BI.
B. The Analytics Tab in Copilot Studio:
Scope Limitation: Likely limited to per-agent or basic overview analytics, lacking the depth and cross-agent comparison capabilities needed.
Customization and Scalability: Probably less customizable and less scalable compared to a dedicated monitoring solution like Log Analytics.
D. Azure Advisor:
Primary Function: Provides recommendations for optimization (cost, security, performance, reliability), not real-time monitoring or telemetry analysis.
Data Analysis Capability: Not designed for deep dive analysis of agent performance and telemetry data.

References

Log Analytics Documentation: Official Microsoft Documentation on Log Analytics Query Language, highlighting its capability for deep analytics.
Azure Monitor Overview: Overview of Azure Monitor, including Log Analytics, demonstrating its role in comprehensive monitoring solutions.

Question: 124 CertyIQ

Your company extends Copilot capabilities in Microsoft Dynamics 365 Finance.
You plan to use Azure DevOps for application lifecycle management (ALM).
You need to recommend a management solution for the ALM process.
What should you include in the recommendation?

A. Lifecycle Services (LCS)
B. pipelines in Microsoft Power Platform
C. solutions in Microsoft Power Platform
D. Azure DevOps test plans

Answer: C

Explanation:

Correct Option: C. solutions in Microsoft Power Platform

Justification for Option C:

Atomic Deployment Unit: Solutions are the fundamental Application Lifecycle Management (ALM) artifact in the Microsoft Power Platform. They serve as the container for all components required to extend Copilot capabilities — including Copilot Studio topics, Power Automate flows, custom connectors, plugins, and environment variables.
Versioning & Transport: Solutions support versioning (major/minor), dependency tracking, and managed/unmanaged layering, enabling controlled promotion from Development through Test to Production via Azure DevOps pipelines.
Source Control Integration: Solution unpacking (pac solution unpack) allows component decomposition into source-control-friendly formats (YAML/JSON), facilitating code reviews, branching strategies, and automated builds within Azure DevOps repos.

Why Other Options Are Less Suitable:

A. Lifecycle Services (LCS): LCS is the operations portal for Dynamics 365 Finance infrastructure management (environment provisioning, deployments, monitoring). It does not manage the versioning, branching, or component packaging of Power Platform customizations (Copilot extensions) which reside outside the F&O application package.
B. pipelines in Microsoft Power Platform: Pipelines (or Azure DevOps YAML pipelines using Power Platform Build Tools) are the automation execution mechanism (CI/CD), not the manageable artifact itself. You cannot version control or branch a pipeline run; you version control the Solution that the pipeline deploys.
D. Azure DevOps test plans: Test Plans manage manual/automated test cases, suites, and execution results for quality assurance. They are orthogonal to the artifact management and deployment transport requirements of ALM for solution components.

References:

1. Application lifecycle management (ALM) with Microsoft Power Platform
2. Use solutions for application lifecycle management

Thank you

Thank you for being so interested in the premium exam material.
I'm glad to hear that you found it informative and helpful.

If you have any feedback or thoughts on the bumps, I would love to hear them.
Your insights can help me improve our writing and better understand our readers.

Best of Luck

You have worked hard to get to this point, and you are well-prepared for the exam
Keep your head up, stay positive, and go show that exam what you're made of!

Feedback

More Papers

Future is Secured

100% Pass Guarantee

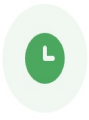

24/7 Customer Support

Mail us - certyiqofficial@gmail.com

Free Updates

Lifetime Free Updates!

Total: **124 Questions**

Link: https://certyiq.com/papers/microsoft/ab-100